

MaiAgent User Guide

Table of Contents

1. Registration & Login

- 1.1 Apply for an Account
- 1.2 How to Log In to the Platform
- 1.3 Forgot Password

2. MaiGPT Overview

- 2.1 What is MaiGPT?
- 2.2 Create Documents and Spreadsheets from Templates
- 2.3 AI Chat
- 2.4 Canvas
- 2.5 Deep Research
- 2.6 AI Image Generation
- 2.7 Connectors
- 2.8 Using on Different Devices
- 2.9 Web Version
- 2.10 iOS App
- 2.11 Using on LINE
- 2.12 Using on Microsoft Teams
- 2.13 FAQ
- 2.14 Usage & Quotas
- 2.15 How to Choose Between Models
- 2.16 Security & Privacy

3. Agent Marketplace

- 3.1 What Are Official Built-in Agents?
- 3.2 Use Official Built-in Agents
- 3.3 Enable and Manage

4. Agent Builder Overview

- 4.1 What is Agent Builder?
- 4.2 Agent (AI Assistant)
- 4.3 Knowledge Base
- 4.4 Database
- 4.5 Agent Schedule
- 4.6 How to Set Up Agent Schedule
- 4.7 Tools
- 4.8 Code Interpreter (Code Execution)
- 4.9 Browser Tool (Browser Operations)
- 4.10 Voice Agent
- 4.11 Skills
- 4.12 Crawler
- 4.13 I Want to Do X, What Should I Use?

5. AI Assistant

- 5.1 AI Assistant Features
- 5.2 Create an AI Assistant
- 5.3 Role Prompt Design Guide
- 5.4 Add a Role Prompt to an AI Assistant
- 5.5 Multi-Language Support
- 5.6 AI Customer Service Quality Management
- 5.7 Cross-Conversation Search
- 5.8 Custom Maximum Output Tokens

6. Knowledge Base

- 6.1 Knowledge Base Overview
- 6.2 How to Create a Knowledge Base: Basic Settings
- 6.3 How to Use the Crawler (Web Scraping) Feature
- 6.4 How to Create FAQ
- 6.5 Document Management: Tags & Metadata
- 6.6 Move Knowledge Base Files
- 6.7 PDF Document Translation
- 6.8 Search Testing

7. Database

- 7.1 Text to SQL Feature
- 7.2 Create & Configure a Database
- 7.3 Text to SQL with MaiAgent Database
- 7.4 Link AI Assistants & Group Permissions

8. Tools

- 8.1 Tools Feature Overview
- 8.2 Create an MCP Tool
- 8.3 Customize Tool Message Display
- 8.4 Contacts MCP Credential Setup
- 8.5 Create an API Tool
- 8.6 Built-in AI Image Generation Tool
- 8.7 Gemini 3 Image Generation Tool
- 8.8 Text to SQL with Supabase
- 8.9 Configure Tools for an AI Assistant
- 8.10 Agent UI

9. Skills

- 9.1 Skills Feature Overview
- 9.2 Skills Management

10. Agent Teams

- 10.1 What is Agent Teams?
- 10.2 Create and orchestrate teams
- 10.3 Team execution traces

11. AgentOps

- 11.1 AgentOps Overview
- 11.2 Test Dataset Management
- 11.3 Automated Evaluation & AI Assistant Monitoring
- 11.4 Evaluation Insights Report
- 11.5 Tool Execution Records
- 11.6 Add Conversations to FAQ

12. Hook Message Interception

- 12.1 Hook Message Interception

12.2 Create and Configure Hooks

12.3 Hook Execution Logs

13. Meeting Records

13.1 Meeting Records Feature Overview

13.2 Create a Meeting Record

13.3 Live Meeting Transcription

13.4 Meeting Records Quick Access

13.5 Meeting Record Backup Audio

13.6 Offline Meeting Transcription

13.7 Meeting Record Details

13.8 Meeting Access Control

13.9 AI Meeting Summary & Q&A

13.10 Download & Knowledge Base Integration

14. Customer Service Conversations & Platform Integration

14.1 Choose an Integration Platform

14.2 Internal Q&A Features

14.3 Web Chat Overview

14.4 Integrate with Platform: Website

14.5 Passwordless Conversations with Contact Source ID

14.6 Citation Source Display

14.7 Markdown Rendering

14.8 MaiGPT Window Mode Embed

14.9 Reply Loading Style

14.10 Chat Platform Settings Overview

14.11 AI Awareness Settings: Time, Name & Location

14.12 Integrate with Platform: LINE

14.13 Integrate with Platform: FB Messenger

14.14 Integrate with Platform: Telegram

14.15 Integrate with Platform: Microsoft Teams

14.16 Integrate with Platform: WhatsApp

14.17 Integrate with Platform: Email

14.18 Integrate with Platform: Slack

14.19 Integrate with Platform: LINE WORKS

- 14.20 Integrate with Platform: Team+
- 14.21 All Conversations Features
- 14.22 Conversation Viewer
- 14.23 Human Handoff Settings
- 14.24 Taking Over and Replying as a Human Agent
- 14.25 Handoff Notification
- 14.26 Inbox Search
- 14.27 Conversation Platform Analytics
- 14.28 Webhook

15. Organization Settings

- 15.1 Organization Overview
- 15.2 Organization & Member Management
- 15.3 Organization Management
- 15.4 Delete Organization
- 15.5 Credits Billing System
- 15.6 Member Credits Quota Management
- 15.7 Conversation Platform Credits Quota Management
- 15.8 Roles and Permissions
- 15.9 Role Permission Management
- 15.10 MaiGPT Role Usage Permissions
- 15.11 Role Resource Access Permissions
- 15.12 Role Member Filtering
- 15.13 Differences Between Roles and Contacts
- 15.14 Role Permission Planning Guide
- 15.15 Member Management
- 15.16 Model Access Permissions
- 15.17 Contacts
- 15.18 Knowledge Management Permissions: Query Metadata
- 15.19 Getting Started - Using the Query Builder
- 15.20 Third-Party Login (SSO)
- 15.21 SAML Integration
- 15.22 ADFS Integration
- 15.23 EIP Integration
- 15.24 Keycloak Integration

- 15.25 Identity Sync (LDAP)
- 15.26 Usage Analytics
- 15.27 Profile: Custom Attributes

16. Developer Tools

- 16.1 API Call Logs
- 16.2 MaiAgent MCP

17. Use Cases

- 17.1 Text Customer Service
- 17.2 Regulation Lookup Assistant
- 17.3 Product Lookup Assistant
- 17.4 Internal Knowledge Management
- 17.5 Professional Knowledge Management
- 17.6 Meeting Requirement Management
- 17.7 Financial Analysis Assistant
- 17.8 Billing & Expense Assistant
- 17.9 Credit Analysis Assistant
- 17.10 Voice Customer Service
- 17.11 IVR Customer Service Intent Recognition
- 17.12 Voice Call Summary
- 17.13 Voice Call Quality Inspection
- 17.14 Image & Vision
- 17.15 Invoice Recognition
- 17.16 Official Document Scanning
- 17.17 Business Card Scanning
- 17.18 ID Card Recognition
- 17.19 External Data Integration
- 17.20 Agricultural Data Query
- 17.21 AI Assistant Optimization Consultant

18. MaiAgent AI Assistant

- 18.1 Scan QR Code for Real-Time Interaction

19. Others

- 19.1 Trial & Subscription Plans

Registration & Login

Apply for an Account

Please contact us via email and provide the following information:

```
Company Information:
Tax ID Number:
Name:
Phone Number:
LINE ID (for easier communication, not required):
```

Send the information to sales@maiagent.ai 😊

How to Log In to the Platform

If you already have an account and password, visit the MaiAgent AI Assistant Platform login page:

<https://admin.maiagent.ai/login>


Enter your account and password to log in and start using the platform.




Forgot Password

If you forgot your login password, click ["Forgot Password?"](#) to reset it.

1. Click the Forgot Password Button



2. Enter Your Registered Email



3. Check Your Email for the Password Reset Email and Click the "Reset Password" Button



4. Set Your New Password



MaiGPT Overview

What is MaiGPT?

MaiGPT is your company's dedicated AI chat assistant. You can talk to it just like ChatGPT, but it can also query your company's internal data, connect to your work tools, and all conversations stay within the company — nothing leaks out.

Overview: MaiGPT Feature Modules



MaiGPT six major feature modules overview

MaiGPT six major feature modules overview

AI chat is the core experience — choose a model, upload files, and ask questions directly. Other features extend the chat capabilities: Canvas lets you edit AI-generated documents in real time, Deep Research has AI automatically search the web to produce comprehensive reports, AI Image Generation lets you describe visuals with text, and Connectors let AI access your work tools like Notion and Slack.

Six Features at a Glance

Feature	What It Does	Typical Use Cases
<u>AI Chat</u>	Multi-model conversation with file uploads, knowledge base Q&A, and web search	Daily Q&A, document analysis, translation
<u>Canvas</u>	Edit AI-generated long-form content in real time	Writing reports, emails, proposals, creating charts
<u>Deep Research</u>	AI automatically searches the web to produce structured research reports	Market research, competitive analysis, topic research
<u>AI Image Generation</u>	Generate images from text descriptions with iterative editing	Marketing materials, presentation illustrations, concept visualization
<u>Connectors</u>	Connect external work tools so AI can query them directly	Search Notion notes, find Slack messages, locate Google Drive files
<u>Multi-Device Access</u>	Web, iOS App, LINE, Microsoft Teams	Use MaiGPT anytime, anywhere

How Is It Different from Agent Builder?



MaiGPT vs Agent Builder use case comparison

MaiGPT vs Agent Builder: positioning, audience, use cases, and operation at a glance

	MaiGPT	Agent Builder
Audience	All company employees	IT admins Developers
Purpose	Chat directly with AI to get daily work done	Build AI assistants for customers or internal use
How to Use	Just type, like ChatGPT	Configure role instructions, knowledge bases, tools, and other modules

Quick rule of thumb: Want to use AI yourself → MaiGPT; Want to build AI assistants for others → Agent Builder.

Recommended Getting Started Order

If this is your first time using MaiGPT, follow this order to get started:

```
Step 1 → Start a conversation – pick an AI model and ask a question
Step 2 → Try uploading a file – let AI analyze the document for you
Step 3 → Use Canvas – ask AI to write a report or email for you
Step 4 → Expand as needed: connect Notion
        Slack, use Deep Research
```


You don't need to learn everything at once. Start with chat + file uploads, then gradually explore advanced features.

Still Have Questions?

FAQ — Answers about usage limits, model selection, security, and privacy

Create Documents and Spreadsheets from Templates

The [Document](#) and [Spreadsheet](#) shortcuts on the MaiGPT home page let you select a template and data sources before returning to the conversation and asking AI to generate content. Documents work well for reports, summaries, and proposals, while spreadsheets are ideal for organizing data into fields and rows.



Filter document templates by source and category, then add a knowledge base, skill, or database below



The spreadsheet shortcut provides built-in templates such as Gantt charts, dashboards, registers, and spreadsheets

Use Case

Project manager Mai needs to prepare a weekly progress report. She selects the weekly report template under **Document**, adds the project knowledge base, and applies the template. Maigpt brings the template and selected data sources into the conversation, where she can specify the scope for the week and generate a draft.

Create a Document or Spreadsheet

Choose a Document or Spreadsheet

In the **Productivity** section of the Maigpt home page, select **Document** or **Spreadsheet**.

Find a Template

Use search, source, or category filters to find a suitable template. You can also select the star icon to add frequently used templates to **Favorites**.

Organization templates are provided by your company, while **Built-in** templates are included with MaiGPT. The templates displayed depend on your organization's settings.

Add Data Sources

At the bottom of the panel, select the **knowledge bases**, **skills**, or **databases** to reference. Add only the sources needed for the content to keep the subsequent conversation focused.

Apply the Template and Describe Your Request

Select a template, then click **Apply**. The selected template appears above the input box, which is also prefilled with prompt text. Close the template panel, add the time period, audience, output tone, or required fields to your message, and submit the generation request.

When generation is complete, MaiGPT lists a summary of the results and the generated files in its response. Review the content, then click **Download** to get the document or spreadsheet.

已完成「星河茶坊」2026 年 9 月第 2 週週報，並套用週進度報告範本。

內容包含：

- 本週完成：官網改版
- 下週計畫：進行會員測試
- 主要風險：測試樣本不足

下載週進度報告 DOCX

已產出 1 個檔案

	xinghe_tea_weekly_report_2026-09-week2.docx	 下載
815.8 KB · 已依照「週進度報告」範本建立星河茶坊 2026 年 9 月第 2 週週報。		

MaiGPT displays the weekly report summary, generated file name, and download button

The workflow is the same for spreadsheets. The following example uses the **Project Schedule Gantt Chart** template to generate an Excel file from three fictitious work items. When it is ready, review the item summary, file name, and file size in the response, then download the **.xlsx** file.



The spreadsheet result lists three work items and provides an Excel file for download

Manage Your Own Templates

Select **Import Document Template** or **Import Spreadsheet Template** in the upper-right corner of the panel to add custom templates permitted by your organization. Imported templates appear in the template list and can be reused later.

If the data source list is empty, first confirm that you have access to the relevant knowledge base, skill, or database, then contact your organization administrator for assistance.

AI Chat

AI Chat is the core experience of MaiGPT. Select an AI model, type your question, and get an answer. You can also upload files for AI analysis, query your company knowledge base, search the web in real time, and enable deep thinking to tackle complex problems.

Chat Capabilities Overview

Capability	What It Does	Best For
Select AI Model	Switch between different models (Claude, GPT, Gemini)	Choosing the best AI for your task
Upload Files & Images	Upload documents for AI to read, summarize, and analyze	Contract review, report analysis, image recognition
Knowledge Base Q&A	Query your company's internal knowledge base	Company policies, product manuals, SOPs
Web Search	Search the web in real time for the latest information	News, market trends, competitor intelligence
Deep Thinking	AI reasons before answering, ideal for complex questions	Multi-step analysis, logical reasoning, decision-making
Chat Memory	Remembers conversation context and your preferences	Continuing discussions, long-term collaboration
Share Conversations	Share an entire conversation with members of your organization so they can view it or copy it and continue asking questions	Sharing results, continuing work

Quick Start

Type your question — Enter your question in the input box and press Enter to send

Upload a file — Click the paperclip icon or drag and drop a file directly

Switch models — Click the model name at the top and select a different AI

Not sure where to start? Just type your question. The AI will automatically respond based on your input — no setup required.

Select AI Model

MaiGPT includes multiple built-in AI models that you can freely switch between based on your task. Each model has its own strengths, and choosing the right one helps you get better answers.

How to Switch Models

Click the model name at the top of the chat area, then select your desired model from the dropdown menu.



Click the model name at the top to open the dropdown and switch AI models

After switching, all subsequent responses will use the new model. Your previous chat history remains unaffected.

How to Choose a Model

My Task	Recommended Model	Reason
Writing reports, analyzing documents	Claude	Strong Chinese comprehension, excels at long-form writing
Code, technical questions	Claude GPT	Both excel here — try both and compare
Quick Q&A, translation	Gemini	Faster response times
Creative ideation, brainstorming	Claude GPT	Better creative quality
Need the latest information	Pair with web search	Any model works

Not sure which to pick? Start with the default model. If you're not satisfied with the response quality, try switching to another model and asking the same question to compare which answer better meets your needs.

FAQ

Will my previous conversation disappear after switching models?

No. Your chat history is fully preserved — only subsequent responses will be generated by the new model.

Why can't I see certain models?

Available models are configured by your company administrator. If you need access to a specific model, contact your IT admin.

Do different models cost the same?

Different models have different usage calculations, but this is typically managed by your company, so you don't need to worry about costs.

Upload Files & Images

You can send files or images to MaiGPT and let the AI read, summarize, analyze, or translate them for you.

How to Upload

There are two ways:

Click the paperclip icon — Located next to the input box, click it to select files

Drag and drop — Drag files from your desktop into the chat area

After uploading, the file will appear above the input box. Then type your question and press Enter to send.

Supported Formats

Type	Formats
Documents	PDF, Word (.docx), Excel (.xlsx), PowerPoint (.pptx), TXT
Images	PNG, JPG, GIF
Other	CSV, JSON, Markdown

What You Can Do

Uploaded Content	What You Can Ask AI
A contract PDF	"What clauses in this contract are unfavorable to us?"
An Excel report	"Summarize the key trends in this report"
A product image	"What are the product specifications in this image?"
A meeting presentation	"List the three key takeaways from this presentation"
An English document	"Translate this into Traditional Chinese"

Tips

Upload multiple files at once You can select and upload multiple files at the same time, and the AI will analyze them together. For example, upload three quotes and ask "Compare the differences between these three quotes."

File size limits The size limits for documents and audio files in a chat depend on your company's settings; the current default is 10 MB. Images have a separate limit. If a file exceeds the limit, MaiGPT displays the next steps instead of uploading it directly.

When File Content Cannot Be Read

For example, a project manager wants to summarize a PDF sent by a partner, but the file was corrupted during transmission. She clicks **Add** → **Upload File**, selects the file, and enters a request for a summary. If the system cannot read the content, the response may explain that it cannot provide a summary and ask her to upload the file again. When she sees this type of message, she can first confirm that the original file opens normally, then obtain the complete file and upload it again.




In this verification, MaiGPT explained that it could not summarize the file and asked the user to upload it again.

If parsing fails, first open the original file on your computer to confirm that it is not corrupted and that its content is complete. Then download it again or ask the sender to resend it. Do not use the file name as the basis for a summary. After uploading the file again, verify that the response is actually based on the file's content.

What to Do If a File Exceeds the Limit

For example, project manager Mai wants MaiGPT to summarize an 11 MB fictional company report. After selecting the file, a message indicates that it exceeds the 10 MB limit. Mai can then upload and attach the file through a knowledge base before returning to the chat to continue asking questions.



When a file exceeds the chat limit, follow the on-screen instructions to process it through a knowledge base

Confirm the file name and current size limit in the notification window.

If you have permission to use knowledge bases, click [Go to Knowledge Base](#). Create or open a knowledge base, upload the file, and wait for processing to finish.

In the knowledge base's **Linked AI Assistant** settings, link the knowledge base to the MaiGPT assistant you want to use. Return to **AI Chat**, click **Select Knowledge Base**, select the knowledge base you just used, and then enter the question you want MaiGPT to process.

The instructions above apply to the MaiGPT admin web interface. Web Chat end users do not have access to knowledge bases. If [Go to Knowledge Base](#) is not displayed, contact your organization administrator and ask them to add the file to the assistant's knowledge base. Alternatively, paste the sections you need analyzed, or split the file into smaller files and upload them again. Audio files are more likely to reach the limit, so try compressing them first.

Knowledge Base Q&A

MaiGPT can connect to your company's knowledge base, allowing you to query internal company data through natural conversation. AI responses include source references that you can click to view the original documents.


How to Use

Just ask your question. If your MaiGPT is connected to a company knowledge base, the AI will automatically search for relevant information to answer your question.

"What is the company's leave request process?"

"What are the differences between Product A and Product B?"

"What should new employees do on their first day?"



After selecting a knowledge base and asking a question, the AI uses the query_files tool to search the knowledge base and return results

Answers Include Sources

AI responses indicate data sources, typically displayed as citation markers or links. Click them to view the relevant sections of the original document.

This allows you to:

- **Verify answer accuracy** — See which document the AI based its answer on
- **Dive deeper** — If the summary isn't enough, click through to read the full original text

Can't Find an Answer?

If the AI responds with "I couldn't find relevant information in the knowledge base," it may be because:

Situation	What to Do
The data hasn't been uploaded to the knowledge base yet	Contact your admin to upload the relevant documents
Your question is too vague	Try asking with more specific keywords
Knowledge base access restrictions	You may not have access permissions for that data — contact your admin

"Not found" doesn't mean the company doesn't have the data Knowledge base content is uploaded and maintained by administrators. If you're sure certain data exists but the AI can't find it, notify your admin to confirm whether it has been imported into the knowledge base.

Web Search

MaiGPT can search the web in real time to get the latest information for answering your questions. It's ideal for queries about news, market trends, competitor intelligence, and other scenarios that require up-to-date data.

How to Use

Click the web search button next to the input box, or mention that you need the latest information in your question:

- "Search for today's Taiwan stock market closing price"
- "What are the recent AI-related regulatory updates?"
- "Look up what new products competitor X has recently launched"



After using the Web Search Tool to search web sources, the AI compiles and responds with results and source links

The AI automatically searches multiple websites, compiles the findings into a comprehensive answer, and includes source links.

When to Use Web Search

Good For	Not Needed
Latest news and current events	Company internal data (use the knowledge base)
Product prices and specifications	Writing reports, translation (no external data needed)
Regulatory and policy updates	Analyzing uploaded files
Market trend research	Casual conversation, brainstorming

Web Search vs Deep Research Web search is great for quickly looking up a single question — you get answers in seconds. If you need to deeply research a topic and produce a comprehensive report, use Deep Research.

Deep Thinking

Deep Thinking lets the AI "think it through" before answering, making it ideal for complex questions that require reasoning, analysis, or multi-step judgment. When enabled, you can see the AI's thinking process.

How to Use

Click the Deep Thinking button next to the input box to enable it, then ask your question as usual.

The AI response will have two parts:

Thinking process (expandable/collapsible) — Shows how the AI analyzed your question

Final answer — The complete response after reasoning

When to Enable

Good to Enable	Not Needed
Complex analysis and comparisons	Simple Q&A ("What's the weather today")
Logical reasoning required	Translation
Math calculations, data analysis	Casual conversation
Multi-faceted decision-making	Text polishing, rewriting

Deep Thinking takes longer Because the AI needs to think before answering, response times are slower than normal mode. If your question isn't complex, stick with the standard mode.

Chat Memory

MaiGPT remembers your conversation content, allowing you to continue previous discussions without re-explaining the background every time.

Two Types of Memory

Conversation Memory (Short-Term)

Within the same conversation, the AI remembers everything discussed earlier. You can continue the discussion seamlessly:

You: "Summarize the key Q1 performance highlights for me"

AI: (responds with the summary)

You: "Now add a Q2 comparison" ← The AI knows you're talking about performance

Starting a new conversation clears the short-term memory, and the AI starts fresh.

Long-Term Memory

MaiGPT can remember important information across conversations, such as your role, preferences, and frequently used formats.

You: "Remember that I'm in the marketing department, and I always want reports in bullet-point format"

AI: Got it, I'll remember that.

Even in new conversations afterward, the AI will remember these preferences.

Tips

Make the most of a single conversation If you're working on a topic (such as writing a report), try to complete it within the same conversation. This way the AI retains all the context, resulting in better responses.

Need a fresh start? If you want to switch to a completely different topic, start a new conversation to prevent content from the previous topic from influencing responses on the new one.

Organize Conversations: Pin Frequently Used Chats

As your conversations pile up, important or frequently used ones get pushed down by newer chats and become hard to find. You can "pin" them to keep them fixed at the top of your "Chat History" for one-click access anytime.

How to Pin

In the "Chat History" panel on the left, hover your mouse over the conversation you want to pin
Click the  menu that appears on the right

Select **Pin**



Hover over a conversation and click the  menu to see the "Pin" option

Once pinned, the conversation moves to the **Pinned** section at the top of "Chat History," with a  icon next to its title.



Pinned conversations are grouped at the top in the "Pinned" section, while regular conversations are sorted by time below

Unpin

Click the  menu on a pinned conversation and select **Unpin**. The conversation will return to the regular time-sorted list below.

Sorting & Rules

Item	Description
Sort order	Both the Pinned section and the regular list are sorted by the time of the last message, from newest to oldest — regardless of when pinning occurred
Pin limit	There is no limit — you can pin multiple conversations at the same time
Persistence	Pin status is saved and persists after refreshing or logging in again

Item	Description
Collapse section	Click the "Pinned" header to expand or collapse the entire section
Scope	You can only pin your own conversations

Good Scenarios for Pinning

Finance

Banking: Pin reconciliation chats or financial plan comparison conversations you regularly track

Tech

Electronics Manufacturing: Pin frequently referenced product spec or technical troubleshooting conversations

Educational Institutions: Pin recurring admissions Q&A or curriculum planning discussions

Service

Retail: Pin ongoing campaign planning or membership benefits inquiry conversations

Or any workflow you return to and continue every day

Pinning only affects your own view Pinning is a personal organization method — it only changes the sort order you see in "Chat History." It does not affect other team members, nor does it alter the conversation content itself.

Share Conversations

You can share an entire conversation with members of your organization. Anyone with access can open the link to view the full conversation. If you allow copying, they can also copy it to their own conversation list with one click and continue asking questions.

Use Case

A knowledge base administrator uses MaiGPT to compile a "Frequently Asked Questions for New Employees" conversation and wants everyone in the organization to use it. At the top of the conversation, they click [Share](#) and select [Quick Share with Entire Organization](#). The link is immediately copied to the clipboard and can be pasted into an announcement. Members can open the link to view the entire conversation. To ask follow-up questions about their own situation, they can click [Continue Conversation](#) to copy it to their conversation list and keep chatting without affecting the original conversation.

How to Share

The share option appears at the top of the conversation page only after the current conversation contains at least one message.

1. Click "Share"

Click **Share** at the top of the conversation page to display two options:

- Quick Share with Entire Organization**: Immediately generates a link and copies it to the clipboard. Everyone in the organization can view and copy the conversation, with no additional settings required.
- Customize Sharing Scope**: Choose the recipients and permissions yourself, as described in the next step.



The Share button provides two options: Quick Share with Entire Organization and Customize Sharing Scope

2. Set the Sharing Scope and Permissions

After selecting **Customize Sharing Scope**, configure the following in the dialog:

Field	Description
Sharing Scope	Choose Everyone in Organization , Specific Groups , or Specific Members . For groups or members, use the selection field below to select one or more recipients


Field	Description
Share Title	Optional. The title recipients see on the shared page and in the Browse Organization Shares list. If left blank, the conversation title is used
Description	Optional. Briefly describe the conversation's key points or purpose
Sharing Permission	Can Copy (default) lets recipients copy the conversation to their own list and continue asking questions. Read Only lets recipients view it only



Set the sharing scope, title, description, and permissions in the same dialog

3. Generate and Send the Link

Click **Generate Share Link**. The link is automatically copied to the clipboard, and you can click **Copy Link** to copy it again. Send the link to the recipients. They must sign in to MaiAgent to open it.



The link is copied automatically when generated, and you can click Copy Link to copy it again

What Recipients See

After signing in and opening the link, recipients arrive at the shared page. The title, sharer, sharing time, and message count appear at the top, with the entire conversation below. They can view the content but cannot type on this page.

When permission is set to **Can Copy**, a **Continue Conversation** button appears at the bottom. Clicking it copies the conversation into a new conversation owned by the recipient, displays "Copied to your conversation list!", and automatically opens it so they can continue asking questions.

When permission is set to **Read Only**, the page displays "This conversation is shared as read-only" and does not show a copy button.

Members outside the sharing scope, or accounts from another organization, see "Unable to load this shared conversation. The link may no longer be valid."



A shared page with Can Copy permission has a Continue Conversation button at the bottom; click it to copy the conversation to your own list


The copy is a new conversation After clicking **Continue Conversation**, the recipient continues in a new conversation under their own account. Your original conversation is unchanged, and you cannot see their subsequent messages. The two conversations remain independent.

Browse Organization Shares and Manage My Shares

Browse Organization Shares at the top of the conversation page lists all shares in the organization that you are allowed to view. Search by title or description, then click **View** to open a shared page.

On this page, click **My Shares** to manage every conversation you have shared:

Action	Description
Edit	Change the share title, description, and sharing permission. The sharing scope and recipients cannot be changed after creation. To change recipients, cancel the share and create a new one
Copy Link	Copy the same share link again
Cancel Sharing	After cancellation, others who open the link see that it cannot be loaded. Conversations already copied to their own lists are unaffected



Use My Shares to edit a share, copy its link, or cancel sharing

Sharing Rules

Item	Description
Shared content	The share captures the content at the moment you click Share. Messages subsequently added to the original conversation do not appear on the shared page. Create a new share to share newer content
Excluded content	The AI's thinking process and tool execution records do not appear on the shared page. Only the conversation text, code, and Canvas content are included
Who can open it	Only signed-in members of the same organization. When the scope is limited to groups or members, anyone not selected cannot open it
Expiration	There is no expiration date. The share remains available until you cancel it
After the original conversation is deleted	The shared page remains available to view and copy, preserving the version captured when it was shared
Who can manage it	Only the person who created the share can edit or cancel it

Review Earlier Content While a Response Is Being Generated

While MaiGPT is generating a long response, you can scroll up to review earlier content. The page stays where you scrolled instead of being pulled back down as new text appears. When you scroll **above the latest question**, a white **Back to Latest** pill appears above the input box. Click it to return to the latest question and resume automatic scrolling.

Use Case

A customer service manager asks MaiGPT to compile a lengthy training checklist. While the AI is still writing, she scrolls up to verify the requirements she provided in the previous exchange. The page remains at the section she wants to read instead of being pushed down by newly generated text, and a **Back to Latest** pill appears above the input box. After checking the requirements, she clicks the pill to return to the current question and watches MaiGPT finish the response.


Steps

1. Scroll Up

While MaiGPT is generating a response, scroll up with the mouse wheel, a touch gesture, or the scrollbar.

2. Read in Place

The page stays at your current position, and newly generated content does not move it. When you scroll above the latest question, a **Back to Latest** pill appears above the input box.



After you scroll above the latest question, Back to Latest appears above the input box. Click it to return to that question

3. Return to the Latest Content

Click **Back to Latest** to return to the latest question and resume automatically following new content. You can also scroll down to the bottom yourself to resume following.

Back to Latest appears only when you are above the latest question. If you scroll down within the same long response to read content that has already been generated, the page still stays in place, but the pill does not appear. Scroll to the bottom to resume following the latest content.

Next Steps

- [Write Documents with Canvas](#)
- [Meeting Records Quick Access](#)
- [Back to MaiGPT Overview](#)

Canvas

Canvas is MaiGPT's built-in editor. When you ask AI to write a report, email, or create a chart, the content appears in the Canvas area on the right side, where you can directly edit and modify it, then ask AI to continue refining.



The left side is the chat area, and the right side is the Canvas editing area

When to Use Canvas

Scenario	Example Prompt
Writing a report	"Write a Q1 performance summary for me"
Writing an email	"Help me reply to this customer inquiry"
Organizing meeting notes	"Format the following content as meeting minutes"

Scenario	Example Prompt
Drawing a flowchart	"Draw a leave request process flowchart"
Creating a table	"Create a comparison table for Plan A and Plan B"

How to Use

Opening Canvas

AI automatically opens Canvas when generating longer document content. You can also ask for it explicitly:

"Use Canvas to write a..."

Editing Content

Once Canvas is open, you can:

Edit text directly on the Canvas — Click anywhere to start editing

Ask AI to make changes — Say "Make the second paragraph more concise" in the chat box

Select specific content for AI to adjust — Select text, then tell AI how you want it changed

Exporting

When finished, click the export button in the top right corner:

Copy — Copy all content to clipboard

Download — Download as a file

What Canvas Can Produce

Type	Good For	Example
Documents	Reports, emails, proposals, announcements	"Write a new product launch announcement"
Tables	Comparative analysis, data organization	"Compare quotes from three suppliers"
Charts	Flowcharts, org charts, sequence diagrams	"Draw a customer complaint handling process"

Tips

Let AI Iterate for You You don't need to get it perfect on the first try. Let AI generate a first draft, then refine step by step: 1. "Write a rough draft first" 2. "Make the tone more formal" 3. "Add supporting data" 4. "Change the last paragraph to a call to action"

Multi-Round Edits Are Never Lost Canvas preserves all your changes. Feel free to go back and forth with adjustments — you don't need to worry about losing previous versions.

Writing Documents

The most common use of Canvas is writing documents. You can ask AI to write reports, emails, proposals, announcements, and meeting minutes, then edit and refine them directly on the Canvas.

Quick Start

Tell AI what you want to write:

"Write a quote reply email for a client"

"Use Canvas to organize today's meeting minutes"

"Write an internal announcement for a new feature launch"

AI will produce the complete document in Canvas, and you can start editing immediately.

Editing and Modifying

Direct Editing

Click anywhere on the Canvas to edit text directly, just like editing in Word.

Ask AI to Edit

Type modification instructions in the chat box, and AI will update the content on Canvas directly:

Document Type	Example Prompt
Email	"Write an email politely but firmly declining a supplier's quote"
Meeting minutes	"Format the following content as meeting minutes: (paste notes)"
Proposal	"Write a proposal outline for implementing AI customer service, including a benefits analysis"
Announcement	"Write a system maintenance announcement, scheduled for Saturday 2-6 AM"

Exporting Documents

Click export in the top right corner when finished:

Copy — Paste into Email, Slack, Word, or anywhere else

Download — Save as a file

Creating Charts

You can use natural language descriptions to have AI create flowcharts, org charts, sequence diagrams, and other charts in Canvas. No need to learn any syntax — just describe what you want.

Quick Start

Tell AI what chart you want:

"Draw a leave request process flowchart"

"Draw our department's organizational chart"

"Draw a process from user placing an order to shipping"



Visualized charts generated by AI in Canvas

Available Chart Types

Chart Type	Example Prompt	Use Case
Flowchart	"Draw the reimbursement process"	Steps, decision branches
Org chart	"Draw the department structure"	Hierarchical relationships
Sequence diagram	"Draw a sequence diagram for API calls"	System interaction order
Gantt chart	"Draw a project timeline Gantt chart"	Project scheduling
Mind map	"Use a mind map to organize marketing strategies"	Brainstorming and categorization

Modifying Charts

Want to adjust after generation? Just say:

You Say	AI Does
"Add a review step after step two"	Adds a node
"Change the flowchart to left-to-right"	Adjusts direction
"Remove the last step"	Deletes a node
"Change the colors to a blue theme"	Adjusts styling

Exporting Charts

When finished, you can:

Copy — Paste into presentations or documents

Screenshot — Capture the chart directly from Canvas

Next Steps

[Deep Research](#)

[Back to MaiGPT Overview](#)

Deep Research

Deep Research is MaiGPT's automated research feature. Give it a topic, and AI will automatically search large volumes of web data, analyze and organize it, and produce a comprehensive research report.

It's ideal when you need to deeply understand a topic but don't want to spend hours searching and organizing information yourself.

How Is It Different from Regular Chat

	Regular Chat	Deep Research
Speed	Seconds	Minutes
Search Scope	Single search	Automatically searches 15-25 sources
Response Length	A few paragraphs	Full report (4,000+ words)
Source Citations	Sometimes	Always, with citations for every section
Best For	Quick Q&A	In-depth research, market analysis, technical investigation

When to Use It


Scenario	Example
Market research	"Research the current state and trends of the AI customer service market in Taiwan"
Competitive analysis	"Analyze the product differences between Company X and Company Y"
Technical investigation	"Research the latest developments and best practices in RAG technology"
Policies and regulations	"Compile the latest progress on AI-related regulations in Taiwan"
Industry reports	"Research global SaaS industry trends for 2025"

Starting a Research Session

In just three steps, AI will automatically produce a complete research report.

1. Click the Deep Research Button

Find and click the Deep Research button in the chat interface.



Enter a research topic, then click the blue "Deep Research" button to start

2. Enter Your Research Topic

Describe your research topic in one or two sentences. The more specific, the better the report quality.

General	Better
"Research AI"	"Research the current state, challenges, and success stories of generative AI adoption in Taiwanese enterprises"
"Analyze competitors"	"Compare MaiAgent, Dify, and Coze on their enterprise RAG capabilities"
"Look up regulations"	"Compile 2025 AI-related regulations and regulatory guidelines in Taiwan"

3. Wait for AI to Finish

AI will begin researching automatically. You can see:

- Planning phase** — AI breaks the topic into sub-questions
- Search phase** — Searches the web for each sub-question
- Analysis phase** — Organizes and cross-references data
- Report phase** — Produces the full report and saves it to Canvas automatically

The entire process typically takes 2-5 minutes, depending on topic complexity. You can work on other things — you'll be notified when it's done.

Give AI More Context Beyond the topic, you can also tell AI: - Your role: "I'm a marketing director" - Report purpose: "This is for a weekly meeting presentation" - Special focus: "Focus on the Taiwan market" - Format requirements: "Use bullet points, don't exceed 3,000 words"

Reading and Exporting Reports

After Deep Research completes, the report appears automatically in Canvas. You can read, edit, or export it directly.

Reading the Report

The report includes the following structure:

Section	Content
Summary	Key findings and conclusions
Body	Detailed analysis organized by topic in chapters
Source Citations	Each section cites its data sources, with clickable links to original text



After research completes, the report appears automatically in Canvas

Modifying the Report

After the report is generated, you can continue refining it in Canvas:

You Say	AI Does
"The summary is too long, condense it to 3 points"	Shortens the summary
"Add data charts"	Adds charts
"Translate this section to English"	Translates the specified section
"Reformat as a presentation outline"	Reformats the layout

Exporting

Click the export button in the top right corner of Canvas:

Copy — Paste into Email, Word, Notion, etc.

Download — Save as a file

Paste Directly into Presentations After copying, you can paste directly into PowerPoint or Google Slides with formatting preserved. This is perfect for quickly turning research results into weekly meeting reports.

Next Steps

[Continue editing with Canvas](#)

[Back to MaiGPT Overview](#)

AI Image Generation

MaiGPT has built-in AI image generation. Describe the visual you want in text, and AI will generate an image for you. Chinese descriptions are supported — no need to write in English. After generation, you can iteratively edit the image until you're satisfied.

Generate Images from Text Descriptions

Describe the image you want directly in the conversation:

"Draw an illustration of a modern office with bright tones"

"Generate a product branding image for marketing with a gradient blue background"

"Draw a cute mascot — a cat wearing a business suit"



Enter a text description and AI generates the corresponding image in real time

Description Tips

The more specific your description, the closer the image matches your expectations:

Generic Description	Better Description
"Draw a cat"	"An orange cat sitting on a windowsill, sunset in the background, watercolor style"
"Make a poster"	"An A4 portrait poster with the title '2025 Annual Conference,' tech style, blue-purple tones"
"Product image"	"A cup of coffee on a white background, overhead angle, a few coffee beans beside it, commercial photography style"

You can specify:

- **Style** — Watercolor, flat design, realistic, cartoon, commercial photography
- **Color tone** — Bright, warm, cool, monochrome
- **Composition** — Overhead, front-facing, close-up, panoramic
- **Purpose** — Social media post, presentation graphic, marketing material

Downloading Images

After an image is generated, click it to view in full size. Right-click or long-press to save it to your device.

Editing and Iterating on Images

Not satisfied with the AI-generated image? No need to start over — you can keep refining it in the same conversation until you're happy with the result.

How to Edit

After an image is generated, tell AI what you want to change directly in the conversation:

- "Change the background to white"
- "Change the person's expression to a smile"
- "Add the company logo in the bottom right corner"
- "Make the overall tone warmer"

AI will generate a new version based on your instructions.



Give modification instructions in the conversation and AI generates a new version of the image

Generating Based on an Existing Image

You can also upload an image and have AI use it as a reference to generate a new one:

(Upload a product photo)

"Using this image's style as a reference, generate similar marketing material for me"

"Convert this image to cartoon style"

Tips

Edit Step by Step Don't request too many changes at once. Modify one thing at a time for the best results: 1. First confirm the composition and subject are OK 2. Then adjust the color tone and style 3. Finally add details (text, logo, etc.)

Not Happy? Start Over If modifications keep drifting off course, just re-describe from scratch. AI generates independently each time, so starting fresh is sometimes faster than continuing to edit.

Next Steps

[Connectors](#)

[Back to MaiGPT Overview](#)

Connectors

Connectors let MaiGPT access your third-party tools such as Notion, Slack, Google Drive, and GitHub. Once connected, you can query and interact with data in these tools directly from the conversation.

Available Tools


Tool	What You Can Do After Connecting
Notion	Search pages, query databases, create new pages
Slack	Search messages, view channel content
Google Drive	Search and read documents
GitHub	View repositories, issues, pull requests

Available connectors are configured by your company administrator. If the tool you need is not listed, contact your admin.

How to Connect

1. Open Connector Settings

Click the tools icon in the chat interface to find the connectors section.



Click the "Connectors" button to see the list of available connectors

2. Authorize Your Account

Select the tool you want to connect, click "Connect," and you'll be redirected to that tool's authorization page. Follow the prompts to complete authorization.

3. Start Using

Once connected, ask questions directly in the conversation:

"Search Notion for pages about the Q1 plan"

"Look up the quote document in Google Drive"

Managing Connected Tools

You can do the following at any time in the connector settings:

View — Currently connected tools

Disconnect — Remove a tool's connection

Reauthorize — Reauthorize if a connection has expired

Notion

After connecting Notion, you can search Notion pages, query databases, and even create new pages directly in MaiGPT.

Connection Steps

- Open connector settings and find Notion
- Click "Connect" to go to the Notion authorization page
- Select the pages or workspace you want to grant MaiGPT access to
- Click "Allow" to complete authorization

What You Can Do

Once connected, say the following in the conversation:

You Say	MaiGPT Does
"Search Notion for pages about product specifications"	Searches and lists relevant pages
"Check the Notion project database for incomplete tasks"	Queries the database and filters results
"Create a new page in Notion titled Meeting Minutes"	Creates a new page
"Save this summary to Notion"	Writes content to Notion

Important Notes

Authorization Scope When authorizing Notion, you can choose which pages to share with MaiGPT. Only pages you've authorized can be searched. If a page can't be found, verify that you've authorized access to it.

Slack

After connecting Slack, you can search Slack messages and channel content in MaiGPT to quickly find previous discussions.

Connection Steps

- Open connector settings and find Slack
- Click "Connect" to go to the Slack authorization page
- Confirm the authorization scope and click "Allow"

What You Can Do

You Say	MaiGPT Does
"Search Slack for discussions about server migration"	Searches relevant messages
"Check what's been discussed recently in #general"	Summarizes channel content
"Find the link John mentioned last week"	Searches a specific person's messages

Google Drive

After connecting Google Drive, you can search and read your Google Drive documents in MaiGPT.

Connection Steps

- Open connector settings and find Google Drive
- Click "Connect" to go to the Google authorization page
- Select your Google account and allow access

What You Can Do

You Say	MaiGPT Does
"Search Google Drive for the quote document"	Searches and lists relevant files
"Open the last meeting minutes"	Finds and reads the document content
"Summarize that market analysis report in Drive"	Reads the document and produces a summary

Important Notes

Only Google Native Formats Are Supported Currently, only Google native formats such as Google Docs, Sheets, and Slides can be read. Office formats like .pptx, .docx, and .xlsx cannot be read directly. Please convert them to Google format in Google Drive before referencing them.

Shared Drive and External Sharing Limitations Files stored in a Shared Drive or existing as shortcuts may still fail to be read even after authorization. Ensure the authorized account has access permissions to the Shared Drive, and use the original file link instead of shortcuts.

GitHub

After connecting GitHub, you can query repositories, issues, pull requests, and other development-related information in MaiGPT.

Connection Steps

- Open connector settings and find GitHub
- Click "Connect" to go to the GitHub authorization page
- Select the repositories to authorize and confirm

What You Can Do

You Say	MaiGPT Does
"Check what open issues main-repo has recently"	Lists issues
"Show me a summary of changes in PR #123"	Reads the PR and summarizes
"Search the repo for code related to authentication"	Searches code

Next Steps

- [Manage access permissions](#)
- [Back to MaiGPT Overview](#)

Using on Different Devices

MaiGPT supports multiple access methods so you can chat with AI on any device. Conversation history syncs automatically, so you can continue discussions regardless of which device you log in from.

Supported Devices and Platforms

Platform	Description	Best For
Web	Use directly in your browser, no installation required	Office computers, laptops
iOS App	Native app for iPhone iPad	On the go, commuting, mobile work
LINE	Chat through the LINE official account	Users who prefer LINE
Microsoft Teams	Use directly within Teams	Teams already using Microsoft Teams

Which One Should You Choose?

Most complete experience → Web (supports all features)

Most convenient for anytime use → iOS App or LINE

Integrate with your workflow → Microsoft Teams

Conversation history is shared across all platforms. You can start a conversation on your computer and continue on your phone when you're out.

Next Steps

[Web](#)

[iOS App](#)

[Using on LINE](#)

[Using on Microsoft Teams](#)

[Back to MaiGPT Overview](#)

Web Version

The MaiGPT web version works in any browser with no software installation required.

How to Use

Open your browser and go to the MaiGPT URL provided by your company
Log in with your account and password (or use SSO)
Start chatting

Recommended Browsers

Browser	Support
Google Chrome	Recommended
Microsoft Edge	Supported
Safari	Supported
Firefox	Supported

Use the latest version of your browser for the best experience.

Next Steps

[iOS App](#)
[Using on LINE](#)
[Back to MaiGPT Overview](#)

iOS App

MaiGPT offers an iOS App so you can use it anytime on iPhone and iPad.

Installation

Open the App Store

Search for "MaiGPT" or scan the QR code provided by your company

Download and install

Login

After opening the app, log in with the same account and password as the web version. Conversation history syncs across devices.

Features

The app supports the same core features as the web version:

AI Chat

Upload photos (take a photo directly or select from your photo library)

Knowledge Base Q&A

Web Search

Some advanced features (Canvas, Deep Research) are recommended to be used on the web version for a better experience.

Next Steps

[Using on LINE](#)

[Using on Microsoft Teams](#)

[Back to MaiGPT Overview](#)

Using on LINE

If your company has connected MaiGPT to a LINE official account, you can chat with AI directly in LINE without opening a separate web page.

Getting Started

Scan the LINE QR code provided by your company, or search for the official account name

Add as a friend

Send a message to start chatting

What You Can Do on LINE

Feature	Supported
Text conversation	Yes
Send images for AI recognition	Yes
Knowledge Base Q&A	Yes
Upload files	Depends on configuration
Canvas	Not supported (use the web version)
Deep Research	Not supported (use the web version)

LINE is suited for quick Q&A and simple tasks. For advanced features like Canvas or Deep Research, switch to the web version.

Next Steps

[Using on Microsoft Teams](#)

[Web](#)

[Back to MaiGPT Overview](#)

Using on Microsoft Teams

If your company has integrated MaiGPT into Microsoft Teams, you can chat with AI directly within Teams.

Getting Started

Find the MaiGPT app in the left sidebar of Teams (or have your IT admin install it)

Click to open

Send a message to start chatting

What You Can Do on Teams

Feature	Supported
Text conversation	Yes
Send images	Yes
Knowledge Base Q&A	Yes
Upload files	Depends on configuration
Canvas	Not supported (use the web version)
Deep Research	Not supported (use the web version)

The Teams integration lets you use AI without leaving your work environment. It's ideal for quick Q&A and daily work assistance.

Next Steps

[Web](#)

[Using on LINE](#)

[Back to MaiGPT Overview](#)

FAQ

The most common questions when using MaiGPT.

Question Categories

Category	What It Covers
Usage and Quotas	What operations consume usage, how to save usage
How to Choose Between Models	Which scenarios Claude, GPT, and Gemini are best for
Security and Privacy	How data is protected, whether conversations can be leaked

Quick Answers

How is MaiGPT different from ChatGPT?

MaiGPT is your company's dedicated AI assistant. Data stays within the company, it can query the company knowledge base, it supports multiple AI models, and it's managed centrally by your company's IT team.

Do I need to pay?

No. Usage is managed and paid for by your company.

Are my conversations secure?

All conversations are processed within the company environment and are not used for AI model training. See [Security and Privacy](#) for details.

Next Steps

- [Usage and Quotas](#)
- [How to Choose Between Models](#)
- [Security and Privacy](#)
- [Back to MaiGPT Overview](#)

Usage & Quotas

MaiGPT uses pay-as-you-go billing — you pay for what you use. Below is an overview of usage basics.

What Is Usage

Each time you chat with AI, the system calculates usage based on the length and complexity of the messages. Usage is typically managed and paid for by your company, so you don't need to pay anything yourself.

What Operations Consume Usage

Operation	Usage
Regular chat	Low
Uploading files for analysis	Medium (depends on file size)
Knowledge Base Q&A	Medium
Deep Research	High (performs multiple searches and analyses)
AI Image Generation	Medium
Deep Thinking	Higher (more than regular chat)

How to Check Usage

How to check usage varies by company configuration. If your interface shows usage information, you can view it directly. Otherwise, contact your admin to learn about your usage status.

Tips for Saving Usage

More precise questions use less usage — Specific questions produce shorter responses and consume less usage

Make good use of the same conversation — Continuing a conversation is more efficient than starting a new one, since you don't need to repeat background information

Choose the right feature — Simple questions don't need Deep Research or Deep Thinking

Next Steps

[How to Choose Between Models](#)

[Security and Privacy](#)

[Back to MaiGPT Overview](#)

How to Choose Between Models

MaiGPT offers multiple AI models. Here's a quick guide to help you understand the differences.

Quick Comparison

Model	Strengths	Speed	Best For
Claude	Chinese comprehension, long-form writing, analytical reasoning	Medium	Writing reports, analyzing documents, complex reasoning
GPT	Well-rounded general ability, code	Medium	Code, translation, general Q&A
Gemini	Fast response speed, multimodal	Fast	Quick Q&A, image understanding
DeepSeek	Strong reasoning ability	Medium	Math, logical reasoning
Gemma	Google open-source, lightweight and efficient	Fast	Lightweight deployment, edge computing
Grok	Real-time information, humorous style	Medium	Current events Q&A, creative writing
Llama	Open source, privately deployable	Varies	On-premises deployment needs

Not Sure How to Choose?

Use this decision flow:

General work (writing documents, research, translation) → Use the default model

Not satisfied with quality → Try switching to Claude or GPT

Need maximum speed → Use Gemini

Complex reasoning or calculations → Use Claude or DeepSeek

There is no "best" model Every model has strengths and weaknesses. The most practical approach is to try the same question with different models and find the one you prefer.

Next Steps

[Select an AI Model \(How-To\)](#)

[Security and Privacy](#)

[Back to MaiGPT Overview](#)

Security & Privacy

MaiGPT is designed from the ground up to protect enterprise data security. Below are answers to your most frequently asked security questions.

FAQ

Can my conversation content be seen by external parties?

No. Your conversation data is restricted to your company's environment and cannot be accessed by other companies or external parties.

Will AI use my conversations to train models?

No. The AI models used by MaiGPT do not use your conversation data for training.

Can company administrators see my conversations?

This depends on your company's permission settings. Consult your IT admin for the specific privacy policy.

Are uploaded files secure?

Uploaded files are stored in your company's secure environment, and all transfers are encrypted end-to-end.

Are connectors (Notion, Slack, etc.) secure?

Connectors use standard OAuth authorization. MaiGPT does not store your account passwords. You can revoke authorization at any time.

Usage Recommendations

Avoid entering highly confidential information Although MaiGPT has comprehensive security protections, avoid entering passwords, credit card numbers, national ID numbers, or other highly sensitive information directly in conversations.

Next Steps

[Usage and Quotas](#)

[Back to MaiGPT Overview](#)

Agent Marketplace

What Are Official Built-in Agents?

Official built-in Agents are AI assistants designed, tested, and published by MaiAgent. Their role instructions, models, and tools are already configured. Once an organization administrator enables one, members can start chatting with it right away—without building an Agent, fine-tuning role instructions, or preparing a knowledge base first.

In the [Agent Marketplace](#), official built-in Agents and Agents created by your organization appear on separate tabs: [Official Built-in](#) and [Organization-built](#).

How They Differ from Organization-built Agents

	Official Built-in Agent	Organization-built Agent
Creator	Designed, tested, and maintained by MaiAgent	Created using Agent Builder
Setup required	Ready to use	Must be built manually from scratch
Role instructions tools model	Cannot be modified	Fully customizable
Knowledge base	Import your own documents	Fully customizable
Version updates	Automatically updated to the latest official version	Maintained by your organization
Quantity limit	Concurrent enablement limit depends on your plan	Depends on your plan (see Trial and Subscription Plans)
Who can use it	Administrators specify the Authorized Users	Determined by role permissions

In short: **use an official built-in Agent for speed; build your own for customization.** Both can coexist. Many organizations use official built-in Agents for general tasks and build their own Agents for company-specific workflows.

What You Can and Cannot Change After Enabling

When you enable an official built-in Agent, the system creates a **locked assistant** in your organization based on the official blueprint.

You can:

Import your organization's documents into the Agent's dedicated knowledge base

Specify which roles/groups can use it

Complete any connector authorization it requires

You cannot:

Modify the role instructions (System Prompt)

Add or remove tools or skills

Change the language model

It is normal not to see an edit button for an official built-in Agent in the **AI Assistants** list. Its settings are centrally maintained by MaiAgent to ensure consistent quality across organizations and allow automatic updates when a new official version is released. If you need customized role instructions or a custom combination of tools, use [Agent Builder](#) to create your own Agent.

Versions Update Automatically

Each official built-in Agent card displays a **Version**. When MaiAgent releases a new version—with improved prompts, behavior fixes, or expanded capabilities—**organizations that have enabled the Agent are automatically updated to the latest version**. You neither need nor have the option to upgrade manually or pin a version.

If an Agent's response style or output format seems different from before, MaiAgent has usually released a new version. Content you imported into its knowledge base is not affected.

Concurrent Enablement Limits Depend on Your Plan

Official built-in Agents have a **concurrent enablement limit** and a **cooldown period after enablement**. Allowances vary by plan. See [Trial and Subscription Plans • Feature Differences](#).

If your organization's plan does not include official built-in Agents, the **Official Built-in** tab will not appear in the **Agent Marketplace**. Contact the sales team to upgrade your plan.

For details about enabling Agents, cooldown calculation, and what to do when you reach the limit, see [Enable and Manage](#).

Currently Available Scenarios

Official built-in Agents are grouped into two categories—**Functions** and **Industries**—with 14 scenario categories in total:

Functions

Scenario Category	Example Agents
Human Resources	Interview question generation, résumé evaluation, job description writing
Legal	Contract risk review, nondisclosure agreement review, regulatory compliance checks
General Affairs	Administrative announcement writing, procurement price comparison, document classification and archiving
Accounting	Expense receipt organization, financial statement summary and analysis, business tax Form 401 analysis
Project Management	PRD writing, requirements organization and analysis, competitor analysis
Personal Productivity	Meeting summaries, presentation design, email response writing

Industries

Scenario Category	Example Agents
Healthcare	ICD-10 classification code suggestions, discharge summary preparation, nursing handoff reports
Manufacturing	Quality incident 8D reports, SPC process control analysis, equipment knowledge Q&A
Banking	Credit report writing, KYC due diligence summaries, financial regulation interpretation lookup
Securities	Equity research intelligence, internal compliance policy Q&A, pre-meeting preparation for financial advisors
Retail	Daily sales summaries, product listing copy, customer complaint and review analysis

Scenario Category	Example Agents
General Administration	Public petition classification and assignment, press release drafting, preliminary grant application review
Government Accounting	Budget and final account execution analysis
Legal Affairs and Procurement	Regulation interpretation summaries, tender and procurement document checks

The table above lists only representative examples for each category. New official built-in Agents are added continuously. For the current list, see the [Agent Marketplace](#) → [Official Built-in](#) tab.

Next Steps

Ready to get started → [Use Official Built-in Agents](#)

Organization administrator enabling Agents for your team → [Enable and Manage](#)

Use Official Built-in Agents

This page is for **all members**. It explains how to find the right official built-in Agent in the **Agent Marketplace**, how to work with it, and who to contact if you run into an issue.

If you are an organization administrator deciding which Agents to make available to your team, see [Enable and Manage](#).

How Official Built-in Agents Work

Every official built-in Agent follows the same workflow: **you provide source material and requirements, it asks follow-up questions until the details are clear, and then it produces the result.**



The four-step workflow for official built-in Agents

1. Follow the Greeting to Enter Your Task or Upload Files

When you start a conversation, the Agent first sends a **greeting** explaining what it can help with and what you need to provide. Follow its instructions: type a description of your request or upload files, such as contracts, résumés, reports, or meeting recordings, depending on the Agent.

2. The Agent Offers Options and Recommendations to Clarify the Output

The Agent does not immediately start working as soon as it receives your material. It asks about key decisions—for example, “How many years of experience should these interview questions target?” or

“Should the report use bullets or prose?” —and **provides options directly**, so you can click one or add a short response.

At this stage, some Agents also **retrieve external information through connectors**, such as data from your calendar, cloud drive, or a third-party system. The first time a connector is used, you will be asked to authorize access.

Do not skip this step. The extra 30 seconds you spend answering clarification questions directly improves the accuracy of the output.

3. Generate a Draft

Once the details are confirmed, the Agent generates the result. You can continue asking for changes in the conversation—adjust the tone, add sections, or change the structure—without starting over.

4. Choose Whether to Export to a Specific File Format

When you are satisfied with the draft, ask the Agent to export it as a file. Supported formats vary by Agent and commonly include Word, Excel, PDF, and presentation files. The generated file appears in the conversation for you to download.

Choose an Agent from the Agent Marketplace

Open the Agent Marketplace

After signing in, select **Agent Marketplace** from the left sidebar.

Switch to the “Official Built-in” Tab

At the top of the marketplace, switch from the **Organization-built** tab to **Official Built-in**.

Search or Filter by Scenario Category

Search: Use the search bar at the top to search by Agent name, description, or scenario keyword

Scenario categories: Narrow the list using category tags such as Human Resources, Legal, Healthcare, or Manufacturing, or select **All Scenarios** to see everything



Agent Marketplace → Official Built-in: tabs and scenario categories appear at the top, with the enabled count in the upper-right corner

Cards are divided into two sections:

Section	Meaning	Available on the Card
Enabled	Your organization has enabled the Agent, so you can use it immediately	Preview and Start Conversation
Not Enabled	Your organization has not enabled the Agent	Preview only; no Start Conversation option

The upper-right corner of the page shows the number of Agents currently Enabled by your organization and its limit. This allowance is managed by the organization administrator.

Preview Before You Decide

Not sure what an Agent actually does? Select Preview (the eye icon) in the upper-right corner of its card. You can preview it without enabling it.

The window is titled Conversation × Guide Sync and plays two synchronized columns:

Left column: Replays one Example Scenario, from the greeting and the user's request through each step completed by the Agent

Right column—Guide • What This Step Does: As the conversation advances, the corresponding STEP 01, STEP 02, and so on is highlighted, explaining what happens and which capabilities are used

After playback finishes, the right column changes to **Guide • Select Each Step for Details**. Select any step to expand it, or choose **Replay** to watch again.

Two more details appear at the bottom of the window:

Connectors: The third-party services used by this Agent, such as Microsoft 365 or Google Workspace

Enablement status: Agents that are not enabled display **Must Be Enabled by an Organization Administrator**; enabled Agents offer **Start Conversation**



Preview window: the left column replays an example conversation while the right column highlights the corresponding guide step

Start a Conversation

Select **Start Conversation** on a card in the **Enabled** section to open a new conversation, then follow the four steps under How Official Built-in Agents Work.



Step 3: After confirming the details, the Agent generates a structured draft

Return to Recently Used Agents

After you have **had a conversation with at least one official built-in Agent**, a **Recently Used** row appears at the top of the marketplace, ordered from most to least recent conversation. Select **Continue Conversation** to resume your previous work without browsing through the categories again.

This row includes only Agents that are enabled and that you have personally used. It is normal for the row not to appear when you have just joined an organization and have not started using any Agents.

When an Official Built-in Agent Is Unavailable

These three situations look similar but have different causes:

What You See	What It Means	What to Do
The Agent is in the Not Enabled section and has Preview but no Start Conversation option	Your organization has not enabled this Agent	Ask your organization administrator to enable it
The Agent is in the Enabled section but displays No Permission to Use	Your organization enabled it, but you are not in this Agent's Authorized Users list	Ask your organization administrator to add your role/group to the authorized users

What You See	What It Means	What to Do
The Official Built-in tab does not appear in the marketplace	Your plan does not include official built-in Agents	Contact the sales team to upgrade your plan. See Trial and Subscription Plans

If a search returns no results, select **Clear Search** to view every Agent in the category, or select **Clear Selected** to remove the category filter.

If a category has no enabled Agents, it displays **No Enabled Agents in This Category** and a **Request Enablement** button. This button takes you to [Built-in Agent Management](#), but an organization administrator with the required permission must still enable the Agent.

Enable and Manage


This page is for organization administrators. For instructions for general members, see [Use Official Built-in Agents](#).

Open Built-in Agent Management

Go to [Organization Settings](#) → [Built-in Agent Management](#).

You can also select [Go to Built-in Agent Management](#) in the [Agent Marketplace](#) to open it directly.

If [Built-in Agent Management](#) is not in the menu, your plan does not include official built-in Agents. The screen will display [Built-in Agents cannot be enabled under the current plan. Please upgrade your plan](#). Contact the sales team to upgrade. If you can view the page but the switches are locked, you do not have permission to manage enablement. Contact your organization administrator.



Built-in Agent Management: enablement count and filters appear at the top; action icons and an enablement switch appear on the right of each row

Each row shows the Agent's name, scenario category, version, and current [Authorized: {Role Name}](#) setting. Deactivated Agents display [Deactivated, Knowledge Base Retained](#), while Agents in cooldown display [Cooldown: N Days Remaining](#).

Three action icons, followed by the enablement switch, appear from left to right on each row:

Icon	Purpose
Configure Authorized Users	Specify which roles/groups can use the Agent
Clear Knowledge Base	Delete all imported documents (cannot be undone)
Import to Knowledge Base	Open the Agent's dedicated knowledge base to upload documents

Use the search box at the top to search by keyword. You can also filter by status—**All Enabled**, **Not Enabled**, **Knowledge Base Retained**—or by scenario category.

Select the row itself, rather than an icon, to open the Agent's **Preview** window. You can see how it works before deciding whether to enable it.

Enable or Deactivate: Select “Apply Changes” to Take Effect

Changing a switch does not take effect immediately. All changes remain pending until you select **Apply Changes** at the bottom to submit them together.

Switch the Agents You Want to Enable or Deactivate

In the list, set the switch to **Enabled** or **Deactivated**. The row is then labeled **Pending Enablement** or **Pending Deactivation**.

Review the Change Summary at the Bottom

A fixed action bar appears at the bottom of the screen. When changes are pending, it shows the number **Pending Enablement**, the number **Pending Deactivation**, and the resulting number of **Enabled** Agents versus the limit. It also reminds you that **an Agent cannot be deactivated within 7 days of enablement**. The **Apply Changes** button becomes available at this point.

Select “Apply Changes”

Select **Apply Changes** to submit everything at once. To discard the changes, select **Cancel** to restore all switches.



After a switch is turned off, the row is labeled “Pending Deactivation,” and the counts and cooldown reminder update in the action bar below

If an individual item fails after the changes are applied, the screen identifies the reason for each failure, such as reaching the limit, an active cooldown period, or insufficient permissions. Successful changes are unaffected.

Enablement Limit and Cooldown Period

Concurrent Enablement Limit

The **Enabled Count** at the top of the page shows how many slots are currently used and the maximum allowed. The limit depends on your plan. See [Trial and Subscription Plans • Feature Differences](#).

When the limit is reached, switches for Agents that are not enabled are locked and display **Enablement limit reached. Deactivate an Agent or upgrade your plan.**

7-day Cooldown Period

Each time an Agent is enabled, it cannot be deactivated for 7 days. During this period, the list displays **Cooldown: N Days Remaining**, and the switch is locked with the message **In cooldown. Can be deactivated in N days.**

The cooldown period restarts each time the Agent is enabled. If you deactivate and then re-enable it, you must wait another 7 days before deactivating it again. We recommend using [Preview](#) to confirm that an Agent is suitable before enabling it, because you cannot swap it out during the cooldown period.

When You Want to Downgrade Your Plan

If the number of Agents currently enabled exceeds the new plan's limit, the system blocks the plan change. You must first deactivate enough Agents to meet the new plan's limit. Because deactivation is also subject to the 7-day cooldown period, allow time for the cooldown when planning a downgrade.

What Happens When You Deactivate an Agent

Deactivation does more than hide the Agent. The system removes the assistant along with its conversation channels: The assistant disappears from the [AI Assistants](#) list, and members can no longer chat with it. Its conversation channels, such as Web Chat, and group permissions are removed. Its Agent schedules stop running. Your imported knowledge base is retained, with the status [Deactivated, Knowledge Base Retained](#).

If you re-enable it later, the system creates **a new assistant** based on the latest official blueprint and **automatically reconnects the original knowledge base**, so you do not need to upload the documents again. However, conversations from before deactivation do not appear under the new assistant.

Configure Authorized Users

Enabling an Agent does not make it available to the entire organization. You must also specify which roles/groups can use it.

Select [Configure Authorized Users](#) on the Agent's row. In the drawer, select the roles/groups that can use it, then select [Save](#).

Configured Agents display [Authorized: {Role Name}](#) in the list.

Agents without authorized users display [Not Yet Available](#).



Select the roles/groups allowed to use this Agent in the drawer

If a member sees **No Permission to Use** on an enabled Agent in the **Agent Marketplace**, they have not been added to its authorized users. Return here and add them.

Import to the Dedicated Knowledge Base

You cannot change the role instructions or tools of an official built-in Agent, but **you can give it your organization's documents**. Every enabled Agent has a dedicated knowledge base.

Select **Import to Knowledge Base** on the Agent's row to open its knowledge base settings. From there, use it like any other knowledge base: upload documents, create FAQs, and configure tags. See [Knowledge Base Overview](#).

For example, you could import your company's contract templates and review guidelines into a “Contract Risk Review” Agent, or internal administrative policies into an “Administrative Process Q&A” Agent, so its answers reflect your organization's actual rules.

Clear the Knowledge Base

Select **Clear Knowledge Base** on the Agent's row to delete its imported documents.

This permanently deletes all imported documents and cannot be undone. Before proceeding, make sure the documents are backed up elsewhere.

The outcome depends on the Agent's current status:

Agent Status	After Clearing
Enabled	Only the knowledge base content is deleted. The knowledge base itself remains connected to the Agent, and you can continue importing new documents
Deactivated, Knowledge Base Retained	The entire knowledge base is deleted and disconnected. Re-enabling the Agent later creates a new, empty knowledge base

Connectors: Which Ones You Need to Configure First

Some official built-in Agents must connect to third-party services, such as Google Workspace, Microsoft 365, or Notion, to retrieve data. Cards indicate one of two authorization types:

Label	Meaning	Who Takes Action
Administrator Must Configure Connection First	The connector uses organization-level credentials	An organization administrator completes the connection setup before members can use it
No label	Personal user authorization	Each member authorizes their own account in the conversation the first time the connector is used

Administrator Must Configure Connection First does not mean that you must authorize every user. It is a one-time, organization-level setup. For connectors without this label, each member authorizes their own account, and the administrator does not need to be involved.

For more information, see [Connectors](#).

Agent Builder Overview

What is Agent Builder?

Agent Builder is the core building area of the MaiAgent platform. Here, you can combine different modules to create AI Agents that can answer questions, query data, and execute tasks.

Big Picture: How Agent Relates to Each Module



Agent Builder Module Overview

The Agent is the core brain, responsible for understanding user questions and generating responses. Other modules extend the Agent's capabilities — the Knowledge Base gives it domain expertise to reference, the Database lets it query structured data, Tools let it perform actions, Skills let it follow SOPs to handle complex tasks, and Scheduled Tasks let it run automatically on a schedule. The Crawler is one of the data sources for the Knowledge Base. When tasks become complex, Agent Teams allow multiple specialized Agents to collaborate and divide work; Hooks automatically enforce safety and quality at the entry and exit points of AI messages.

Twelve Modules at a Glance

Module	What It Does	Typical Use Cases
<u>Agent (AI Assistant)</u>	The core of conversations — understands intent and generates responses	Customer service bots, internal Q&A, business assistants
<u>Knowledge Base</u>	Provides enterprise-specific knowledge for Agent to reference	Product manuals, internal policies, FAQ
<u>Database</u>	Lets Agent query structured data using natural language via Text2SQL	Order lookup, inventory queries, report analysis
<u>Scheduled Tasks</u>	Lets Agent automatically execute tasks on a schedule	Daily summary reports, periodic data checks
<u>Tools</u>	Lets Agent call external services to perform actions	Send emails, check weather, integrate with CRM
<u>Code Interpreter</u>	Executes Python in a sandbox for file processing and computation	Generate Word/presentations, data analysis, chart creation

Module	What It Does	Typical Use Cases
<u>Browser Tool</u>	Controls a browser to navigate web pages, interact with elements, and capture screenshots	Operate websites without APIs, scrape data after login, automate web workflows
<u>Voice Agent</u>	Lets Agent serve customers through listening and speaking (Realtime STT + LLM + TTS)	Phone customer service, IVR, voice queries, hands-busy scenarios
<u>Skills</u>	Defines multi-step reasoning and decision flows for Agent	Quoting processes, troubleshooting SOPs
<u>Crawler</u>	Automatically scrapes content from web pages and imports it into the Knowledge Base	Crawl official websites, scrape public information
<u>Agent Teams</u>	Organizes multiple specialized Agents into collaborative teams with handoff and delegation	Customer service routing, order/returns/recommendations multi-agent collaboration
<u>Hook (Message Interception)</u>	Automatically intervenes before messages enter AI and before responses are sent	PII masking, malicious prompt interception, response quality control

Not Sure What to Use?

Go to [I Want to Do X — What Should I Use?](#) to see common scenario comparisons.

Recommended Build Order

If this is your first time using MaiAgent, follow this order:

```

Step 1 → Create an Agent, set up role instructions (what role it plays, how it responds)
Step 2 → Build a Knowledge Base, upload documents or FAQ (give it something to reference)
Step 3 → Expand as needed: add Tools, Database, Skills (based on what capabilities your scenario requires)
Step 4 → Connect a conversation platform to go live (Web Chat, LINE, Messenger ... )

```

You don't need to configure all modules at once. Start with Agent + Knowledge Base, then gradually expand capabilities.

Agent (AI Assistant)



Agent Core Concepts

Agent Core Workflow: LLM-Driven × System Prompt × Memory

What Is This?

An Agent is an AI conversational persona you build on MaiAgent. Each Agent is an independent chatbot with its own identity, personality, knowledge, and capabilities.

Think of an Agent as a new employee — you need to tell it:

- **Who you are** (System Prompt)
- **What you know** (Knowledge Base, Database)
- **What you can do** (Tools, Skills)

Once configured, the Agent can serve users across various conversation platforms.

Core Concepts

Large Language Model (LLM)

The Agent's comprehension and response capabilities come from the underlying Large Language Model (LLM). Think of the LLM as the Agent's "brainpower" — different models vary in comprehension, response quality, speed, and cost.

When creating an Agent, you can choose the model that best fits your scenario. MaiAgent supports multiple models, including Claude, GPT, and more — you don't need to manage the models yourself; the platform handles it for you.

Quick guide: For precise responses and complex reasoning, choose a more powerful model. For high-frequency simple Q&A or cost control, choose a faster and lighter model.

System Prompt

The system prompt defines how the Agent behaves. For example:

"You are a friendly customer service representative. Respond in Traditional Chinese with a professional yet approachable tone."

"You are a financial analysis assistant. Only answer finance-related questions. If data is uncertain, clearly state so."

A well-crafted system prompt significantly improves the Agent's response quality.

Response Mode

MaiAgent offers two response modes:

Chat Mode: Suitable for general Q&A, responds directly using Knowledge Base content

Agent Mode: Suitable for scenarios requiring multi-step reasoning, where the Agent decides what data to look up and which tools to use

Memory

Agents have memory capabilities:

Conversation Memory: Remembers what was discussed earlier within the same conversation

Long-Term Memory: Retains important information across conversations, so you don't need to reintroduce yourself every time

Advanced note: The memory feature allows the Agent to maintain contextual coherence across multiple conversation turns and recall past conversations through vector retrieval. You can adjust related parameters in the Agent's "Memory Settings."

What Do I Need to Do?

Create an Agent — Give it a name and choose a language model

Write a System Prompt — Define the Agent's role and response style

Attach Capabilities — Connect a Knowledge Base, Tools, Skills (as needed)

Test Conversations — Test the results in the built-in chat window

Connect a Platform and Go Live — Deploy to Web Chat, LINE, or other conversation platforms

Further Reading

[AI Assistant Features](#)

[Create an AI Assistant](#)

[System Prompt Design Guide](#)

[Add a System Prompt to Your AI Assistant](#)

Knowledge Base



Knowledge Base Concept

Knowledge Base RAG Flow: From Document Upload to Agent Retrieval and Response

What Is This?

A Knowledge Base is the enterprise knowledge encyclopedia for your Agent. You add company documents, FAQs, and web content to the Knowledge Base, and the Agent can answer user questions based on this data.

Without a Knowledge Base, the Agent can only respond using the language model's general knowledge. With a Knowledge Base, it can answer enterprise-specific questions like "What is our company's return policy?"

What Can Go in a Knowledge Base?

Data Type	Description	Suitable Scenarios
Document Upload	PDF, Word, Excel, TXT, etc.	Product manuals, internal guidelines, contract templates
FAQ	Question-answer pairs	Frequently asked questions, standard responses
Crawler Import	Automatically scrape content from web pages	Official website information, public data

How Is It Different from a Database?

This is the most frequently asked question:

	Knowledge Base	Database
Data Type	Unstructured (documents, articles, FAQs)	Structured (tables, fields, numbers)
Query Method	Semantic search ("What is the return policy?")	SQL queries ("What was last month's revenue?")

	Knowledge Base	Database
Best For	Knowledge-based questions, policy guidelines, operating instructions	Data queries, statistical analysis, real-time status

Simple rule of thumb: If the answer is in a "document," use a Knowledge Base; if the answer is in a "report," use a Database.

Advanced Note: The Knowledge Base uses RAG (Retrieval-Augmented Generation) technology under the hood. Uploaded documents are split into paragraphs and converted into vectors. When a user asks a question, the system finds the most relevant paragraphs and provides them to the Agent as the basis for its response.

What Do I Need to Do?

Create a Knowledge Base — Name it and select an Embedding model

Import Data — Upload documents, create FAQs, or use a crawler to scrape content

Search Test — Test the Knowledge Base with sample questions to verify it finds the correct answers

Attach to an Agent — Link this Knowledge Base in the Agent settings

Further Reading

[Knowledge Base Overview](#)

[How to Create a Knowledge Base: Basic Settings](#)

[How to Create FAQs](#)

[Document Management: Tags and Metadata](#)

[Search Test](#)

Database



Database concept

Text to SQL: Natural Language → SQL Query → Plain Language Response

What Is This?

Database enables Agents to query structured data using natural language. When a user asks "How many laptops did we sell last month?", the Agent automatically converts the question into a SQL query, retrieves the answer from the database, and responds in plain language.

The technical name for this feature is **Text to SQL**.

When Do You Need It?

When your use case involves:

Data queries: Order status, inventory quantities, sales figures

Statistical analysis: "How much did this quarter's revenue grow compared to last quarter?"

Real-time status: "How many pending tickets are there right now?"

The answers to these questions are stored in structured tables, not in documents, so you need a Database rather than a Knowledge Base.

What Databases Are Supported?

Type	Description
MaiAgent Built-in Database	A database provided by the platform where you can directly upload CSV files or manually create tables
External Database	Connect to your PostgreSQL, MySQL, SQL Server, or Oracle

Is It Secure?

The Agent only executes **read-only queries** (SELECT) and will not modify or delete your data. You can also restrict which tables and fields the Agent can access.

Advanced Note: The accuracy of Text to SQL depends on how clearly the table structure is defined. It is recommended to add Chinese descriptions to tables and fields to help the Agent understand what each field represents.

What Do I Need to Do?

Create a database connection — Select the database type and configure connection details

Set accessible tables — Choose which tables the Agent can query

Add field descriptions — Add descriptions to fields to improve query accuracy

Test queries — Test whether the Agent can query correctly using natural language

Attach to Agent — Link to the Agent as a tool

Further Reading

[Text to SQL Feature](#)

[Using MaiAgent Knowledge Base for Text to SQL](#)

[Using Supabase for Text to SQL](#)

Agent Schedule



Agent scheduling concept

Agent Scheduling: Timed Trigger → Agent Auto-Executes → Results Delivered

What Is This?

Agent Scheduling lets an Agent execute tasks automatically on a set schedule without waiting for someone to ask a question. You provide a prompt, set a schedule, and the Agent runs on time and delivers results to the specified destination.

Think of it as setting an alarm clock for the Agent — when the time comes, it gets to work automatically.

When Do You Need This?

Periodic Reports: Compile yesterday's customer service summary every morning at 9 AM

Monitoring Checks: Check system status every hour and send notifications if anomalies are detected

Data Processing: Compile last week's sales data every Monday

One-Time Tasks: Execute a specific action at a designated time

Schedule Types

Type	Description	Example
Cron	Set precise times using Cron expressions	<code>0 9 * * 1-5</code> (Monday through Friday at 9 AM)
Interval	Execute at fixed intervals	Every 30 minutes, every 2 hours
One-Time	Execute once at a specified time	2026-04-15 14:00

Execution Modes

Mode	Description	Best For
Context Mode	Retains memory across executions; each run can continue from the previous one	Tasks that require accumulated information (e.g., continuous tracking)

Mode	Description	Best For
Independent Mode	Starts fresh each time with no memory from previous runs	Standalone repetitive tasks (e.g., daily reports)

Result Delivery

After execution, results can be sent to:

Conversation: Sent to a specified conversation on a messaging platform

Webhook: POST to your external URL (integrate with Slack, LINE, etc.)

Advanced Note: Agent Scheduling only supports AI assistants in Agent mode. When the schedule runs, the Agent uses your configured prompt as instructions and leverages all its attached tools and skills to complete the task.

What Do I Need to Do?

Select an Agent — Choose an existing Agent (must be in Agent mode)

Write a Prompt — Describe what the Agent should do each time it runs

Set the Schedule — Choose the schedule type and time

Configure Delivery — Decide where results should be sent

Enable — Turn on the schedule to start automatic execution

Next Steps

For complete step-by-step instructions, see [How to Set Up Agent Scheduling](#).

How to Set Up Agent Schedule

This page demonstrates how to create an agent schedule. If you're not yet familiar with what agent scheduling does, read [Agent Scheduling](#) first.

Prerequisites

Before creating a schedule, confirm the following:


You have created an AI assistant in Agent mode — Agent scheduling only supports AI assistants in Agent mode. Non-Agent mode assistants will not appear in the dropdown.

The AI assistant is connected to at least one conversation platform — The schedule uses a conversation platform as its container to run the Agent.

You have Agent Scheduling access — If you don't see the "Agent Scheduling" menu, contact your organization admin to check permission settings.

Navigate to Agent Scheduling

From the left menu, click [AI Features](#) → [Agent Scheduling](#). You'll see the schedule list:



Agent scheduling list page showing the status, type, delivery method, and last execution time for all schedules

Column descriptions:


Column	Description
Name	The schedule's identifier
AI Assistant	The AI assistant used by this schedule
Execution Mode	Context mode or independent mode
Schedule	Schedule type (Cron, interval, one_shot) and timing
Delivery	💬 = send to conversation, 🔗 = Webhook, auto = automatically created dedicated conversation
Last Execution	Time of the most recent execution
Status	Active Paused

Create a Schedule

1. Click New Schedule

Click the **New Schedule** button in the upper-right corner to open the form. The form has four tabs: **Basic Info**, **Prompt**, **Schedule Settings**, and **Delivery Settings**.

2. Fill in Basic Info



Basic Info tab: select AI assistant, name, execution mode, and conversation platform

Field	Required	Description
AI Assistant	✓	Select an existing AI assistant (only Agent mode assistants are shown)
Name	✓	The schedule's identifier, e.g., "Daily Customer Service Summary"
Execution Mode	✓	Choose context mode or independent mode
Conversation Platform	✓	The conversation platform container used during execution; context mode uses it to accumulate Agent memory (used internally — it won't appear in the customer service conversation list)
Save Results to Context	—	Only shown in context mode. Only affects memory for the next execution (doesn't affect whether you can view results — all past responses are retained on the execution log page)


How to choose the execution mode?

Context mode: Retains memory across executions. Suitable for tasks that need to accumulate information, such as continuously tracking complaint resolution progress or monitoring data trends.

Independent mode: Starts fresh each time with no previous memory. Suitable for standalone recurring tasks, such as daily reports or periodic statistics.

The AI Assistant, Execution Mode, and Conversation Platform fields cannot be changed after creation. To change them, you must delete the schedule and create a new one. Other fields (name, prompt, schedule time, delivery method) can be edited at any time.

3. Write the Prompt



Prompt tab: describe the task the Agent should complete each time it runs

In the prompt field, describe what the Agent should do **each time it runs**. The prompt is sent to the Agent as its instruction, and the Agent will use its attached tools and skills to complete the task.

Prompt example:

```
Query all unresolved customer complaints from the past 24 hours,  
sort them by priority, and compile a bullet-point summary.  
If there are any "High" priority cases, mark them with "⚠ Urgent".
```


Tips for writing prompts:

Clearly describe the task objective and output format

If you need the Agent to use a specific tool, guide it in the prompt (e.g., "Query the orders table in the database")

Avoid relying on external state beyond the current time; each execution should be an independent task (unless using context mode)

4. Configure Schedule Timing



Cron mode: use a standard 5-field cron expression to set precise timing

There are three schedule types:

Cron


Use a standard 5-field cron expression to set a precise schedule: `minute hour day month weekday` .

Expression	Meaning
0 9 * * 1-5	Monday through Friday at 9:00 AM
0 */2 * * *	Every 2 hours on the hour
30 8 * * *	Every day at 8:30 AM
0 0 1 * *	First day of every month at midnight

Must be paired with the "Timezone" field, which defaults to Asia/Taipei (Taipei time).

Interval

Execute once every fixed interval, configured by combining four fields — days/hours/minutes/seconds:



Interval mode: combine days, hours, minutes, and seconds fields


The interval must be **at least 1 second**

For example, to run "every 30 minutes," enter 30 in the "Minutes" field and 0 in all others

Interval mode does not use the timezone field (since it's a relative time, not an absolute time)

One-time

Execute once at a specified time, after which the schedule automatically ends:



One-time execution: select a specific date and time — the schedule ends after execution

Select a date and time precise to the second

Must be paired with the "Timezone" field


After execution, the schedule remains in the list for viewing records but will not run again

Other fields:

Field	Description
Enabled	Whether to activate the schedule immediately. When off, the schedule is "Paused" and won't trigger
Maximum Executions	Limit the maximum number of times the schedule runs; it will automatically pause when the limit is reached. Leave empty for unlimited

5. Configure Delivery

Where should the results go after execution? Configure this in this tab.



Delivery settings: configure multiple conversation targets and multiple Webhook URLs

This tab configures where results are actively pushed to (specified conversations, Webhook endpoints). Even if you don't set any delivery targets, the complete results of each execution are still retained on the [Execution Log](#) page — they just won't be actively pushed elsewhere.

Conversation

Click [Add Conversation](#) and select a conversation from the dropdown. Results will be sent as a bot message. You can add multiple conversations.

Webhook

Click [Add Webhook URL](#). Results will be sent as a POST request to your external URL (must start with `http://` or `https://`). Common use cases:

Integrate with Slack

Discord

LINE

Write to your own system for further processing

Trigger other automation workflows

It's fine not to set any delivery targets — execution results are always retained on the [Execution Log](#) page. But if you want your team/colleagues to passively receive results (e.g., push to Slack), you need to set at least one conversation or Webhook, otherwise no one will be notified.

6. Save and Enable

Click [Confirm](#) to save. If the "Enabled" toggle is on, the schedule will immediately start triggering according to the configured timing.

What Happens After Saving?

After the schedule is saved successfully, you'll immediately see the following:

The modal closes and the schedule appears in the list

Status shows "Active" (green tag), or "Paused" (if the enabled toggle was off)

The "Last Execution" column is empty until the first trigger, after which it updates with the execution time

The system does two things in the background (no UI changes visible, but happening behind the scenes)

Registers the schedule with the backend scheduler to automatically trigger at the configured time

Context mode also creates an internal system conversation container to accumulate Agent memory (this container won't appear in the [Customer Service Conversations](#) list)

Wait for the schedule to trigger

Cron

Interval: Executes automatically at the configured time, e.g., cron `0 9 * * 1-5` triggers at the next Monday–Friday 9:00 AM

One-time: Executes once at the configured time

Don't want to wait? Click the [▶ Run Now](#) button in the list to trigger an immediate execution for testing

During execution

The list status remains "Active," and the "Last Execution" column updates with each run

Each execution creates a record on the [Execution Log](#) page (with the complete Agent response)

If you added conversations or Webhooks in "Delivery Settings," results are also sent to those targets

For details on where to view results, see [Where Do Results Go?](#) below

When the task ends

Reaches the "Maximum Executions" limit → automatically switches to "Paused" and stops triggering
One-time execution completes → the schedule remains in the list but won't trigger again

Where Do Results Go?

After a schedule runs, the Agent's response appears in the following locations. **The Execution Log page is the primary entry point for viewing history** — regardless of execution mode, every run leaves a complete record here. Other locations depend on your delivery settings.

1. Execution Log Page (Primary Entry Point)

Each execution generates a record. This is the most reliable way to view historical results.

How to access:

Agent Scheduling list → Click  **View Execution Log** on the schedule's row → Enter the execution history page → Click  **View Details** on a specific record

You'll see:

Execution status (Success

Failed

Partial Success)

Start time, completion time

Complete Agent response (even if subsequent delivery failed, the content is still preserved here)

Error messages on failure

2. Your Specified Conversations (Conversations Added in Delivery Settings)

Execution results are sent as Agent messages into these conversations, just like messages seen by customers/colleagues.

How to find:

Customer Service Conversations → Enter the conversation you added in delivery settings → The message will be at the latest position.

3. Webhook URLs (URLs Added in Delivery Settings)

Execution results are sent as HTTP POST to your external URL. **Body is JSON:**


```
{
  "content": "Complete response content generated by the Agent"
}
```

Technical details:

Content-Type: `application/json`

Timeout: 30 seconds (the Webhook endpoint must respond with 2xx within 30 seconds, otherwise it's treated as failed)

Failures are recorded in the execution log's "Error Message" field (execution log status becomes "Partial Success"), but no retry is attempted

What Does the "Save Results to Context" Toggle Do?

This toggle **only affects the memory for the next execution in context mode** — it doesn't affect whether you can view results.

Setting	Effect on Next Execution
On (default)	The previous prompt and Agent response are carried forward as conversation history for the next run, allowing the Agent to continue the context
Off	No messages are retained per execution, effectively behaving like independent mode

Regardless of the toggle state, **the Execution Log page always retains the complete response from each run** — this toggle doesn't affect your ability to view history.

Summary: Results Visibility Reference Table

Execution Mode Setting	Execution Log	Specified Conversations	Webhook	Agent Cross-run Memory
Context mode + Save results to context (default)	✓	✓ If configured	✓ If configured	✓ With memory
Context mode + Don't save results	✓	✓ If configured	✓ If configured	✗ No memory
Independent mode	✓	✓ If configured	✓ If configured	✗ No memory

Why is there no "Dedicated Conversation" entry point? Context mode creates a dedicated backend conversation to accumulate memory, but this conversation currently does not appear in the **Customer Service Conversations** list — it's an internal system container, not a conversation for human agents to handle. To view historical results, go to the **Execution Log**.

Full Example: Daily Customer Service Summary

Below is an end-to-end example showing how to create a schedule that automatically compiles the previous day's customer service summary every morning at 9:00 AM and sends the results to Slack.

Scenario

Every morning before work, you want the AI assistant to compile all unresolved complaints from yesterday, create a summary, and send it to the team's Slack #customer-support channel so the customer service manager can see it as soon as they open Slack.

Preparation

- An AI assistant in Agent mode, with tools attached that can query complaint data (e.g., database tool or customer service system API tool)
- A connected conversation platform (as the schedule's execution container)
- A Slack incoming webhook URL (to receive results)

Configuration

Basic Info:

Field	Value
AI Assistant	Customer Service AI Manager
Name	Daily Complaint Summary - 9:00
Execution Mode	Independent mode (compile independently each day, no need to remember the previous day)
Conversation Platform	Internal Customer Service

Prompt:

```
Query all customer complaints with status "Unresolved" or "In Progress"
from the past 24 hours (ending at today's midnight),
sort by priority (High → Medium → Low) and compile a summary:
```

```
## Format
```

```
### 🟠 High Priority (X cases)
```

```
- [#Case ID] Customer Name – one-line description
```

```
### 🟡 Medium Priority (X cases)
```

```
...
```

```
### 🟢 Low Priority (X cases)
```

```
...
```

```
End with a summary line: "Total X cases pending, of which X high-priority cases
need to be handled today."
```

Schedule Settings:

Schedule type: **Cron**

Cron expression: **0 9 * * 1-5** (Monday through Friday at 9:00 AM)

Timezone: **Asia/Taipei**

Enabled: On

Maximum executions: Leave empty (run continuously)

Delivery Settings:

Conversation: Don't add extra conversations (the Execution Log page already retains the Agent's complete output — no need to send it to a customer service conversation and take up space)

Webhook URL: <https://hooks.slack.com/services/T000000000/B000000000/XXXXXXXXX>

After Saving

Monday at 9:00 AM (first trigger): The schedule triggers automatically and the Agent starts querying and compiling

Within ~30 seconds (depending on Agent processing speed):

The team's Slack #customer-support channel receives the summary via Webhook push (POST body is `{"content": "..."}`)

The "Last Execution" in the schedule list updates to **2026-04-20 09:00:12**

Want to review what the Agent produced? Two places:

Execution Log page (primary entry): **Agent Scheduling** → Click the schedule's ⌚ → Click 👁 → View the complete bot response and token usage

Slack: Go directly to the #customer-support channel (this is the Webhook endpoint you configured)

Tuesday, Wednesday... ongoing execution: Triggers automatically every weekday at 9:00 AM with no further human intervention needed

Troubleshooting

Slack didn't receive a message → First check the Execution Log page to confirm the execution status. If "Failed," check the error message. If "Success" but Slack didn't receive it, verify the Webhook URL is correct and check for Slack rate limits

Summary content doesn't match expectations → Check the Agent's actual output in the execution log details, then adjust the prompt

Schedule didn't trigger → See the "FAQ" section below

Manage Existing Schedules

Back in the schedule list, the action area on the right side of each row provides these functions:

Icon	Function	Description
	Edit	Modify schedule content (except AI Assistant, Execution Mode, and Conversation Platform)
	View Execution Log	Enter the schedule's execution history page
	Run Now	Trigger an immediate execution without waiting for the scheduled time
 	Pause Enable	Toggle the schedule's active status
	Delete	Delete the schedule. Execution logs are also removed

Run Now

"Run Now" is the fastest way to test if a schedule is configured correctly. After clicking:

The system triggers an immediate execution without affecting the original schedule timing

A "Schedule triggered" notification appears


Check "View Execution Log" to observe the execution results

Pause and Enable

After pausing, the schedule won't trigger, but settings and execution logs are preserved. You can re-enable it at any time.

View Execution Logs

Click the  icon in the list to enter the execution log page:



Execution Log page: filter by status and date range, and view results of each execution

Filters

Status: Pending

Running

Success

Failed

Date Range: Specify start and end dates

Column Descriptions

Column	Description
Status	Pending (queued), Running, Success, Failed
Start Time	When the Agent started execution

Column	Description
Completion Time	When execution ended
Error Message	If execution failed, displays the error reason
Bot Response	Response content generated by the Agent (full content available in details)

View Execution Details

Click the 👁 icon on the right side of each row to open the details modal, where you can view the complete bot response and error messages:



Execution details: status, start/completion time, complete bot response, and error messages

FAQ

- ▶ Why didn't the schedule trigger?
- ▶ Why can't I find my AI assistant in the dropdown?
- ▶ The schedule ran — where do I see what the Agent produced?
- ▶ After turning on "Save Results to Context," where do I see that conversation?
- ▶ What's the data format received by Webhook?
- ▶ Independent mode vs. context mode — what's the actual difference?

Further Reading

Agent Scheduling — Conceptual introduction to agent scheduling

Agent (AI Assistant) — Create an AI assistant in Agent mode

Platform Integration: Website — Set up a conversation platform (using website as an example)

Tools



Tools concept

Tools: Agents call external services through Function Calling

What is this?

Tools enable Agents to "take action" — calling external services to perform actions or retrieve real-time information.

An Agent without tools can only answer questions using its existing knowledge. With tools, it can check live weather, create a customer record in a CRM, send emails, or fetch the latest data from third-party APIs.

How are tools different from skills?

	Tools	Skills
Nature	A callable external service	A multi-step reasoning workflow
Granularity	Single action (e.g., "check weather", "send email")	Multi-step SOP (e.g., "complete quotation process")
Who decides when to use it	The Agent determines when to call it	The Agent triggers the workflow within a skill based on context
Does it use tools	A tool is the smallest unit itself	A skill can chain multiple tools together

Quick rule of thumb: If it's a single capability (query data, perform an action), make it a tool. If it's an entire workflow (first query A, then decide whether to do B), make it a skill.

Tool Types

API Tool

Directly calls a REST API. You configure the URL, HTTP method, and parameters, and the Agent can use it during conversations.

Best for: Internal company APIs, third-party service APIs, simple HTTP requests

MCP Tool

Connects to external services through MCP (Model Context Protocol). MCP is a standardized protocol that allows Agents to integrate with various services more easily.

Best for: Services that already have an MCP Server, external integrations requiring more complex interactions

Built-in Tool

Tools provided by MaiAgent by default, such as AI image generation. No additional configuration required.

Advanced note: Tools are powered by Function Calling technology under the hood. The Agent determines whether to call a tool based on the conversation context, automatically assembles parameters, sends the request, interprets the returned results, and replies to the user in natural language.

What do I need to do?

Create a tool — Choose the type (API MCP) and configure the connection details

Define parameters — Describe the tool's purpose and parameters so the Agent knows when to use it

Test — Verify the tool can be called and returns results correctly

Attach to an Agent — Configure this tool in the Agent settings

Further Reading

[Tool Feature Overview](#)

[Create an MCP Tool](#)

[Create an API Tool](#)

[Built-in AI Image Generation Tool](#)

[Configure Tools for AI Assistants](#)

Code Interpreter (Code Execution)

What Is This?

Code Interpreter lets an Agent execute Python code in a **secure sandbox environment**, completing tasks programmatically.

With Code Interpreter enabled, the Agent can do more than just reply with text — it can write code to perform calculations, process files you upload, generate Word/PowerPoint/Excel documents, and return the results for you to download.

How Is It Different from Tools?

	Tools	Code Interpreter
Nature	Calls pre-defined external APIs	Writes and executes Python on the fly in a sandbox
Flexibility	Limited to what the API provides	Anything Python can do
Typical use cases	Send emails, query CRM, integrate third-party services	Computation, data processing, file generation, charting

Quick rule of thumb: Need to call an external API → Tools; need to compute, process files, or generate documents → Code Interpreter.

What Can It Do?

Data analysis & computation: CSV/Excel calculations, statistics, pivot tables, data cleaning

Chart creation: Turn data into line charts, bar charts, pie charts, and more

Document generation: Produce Word, PowerPoint, Excel, and PDF files for download

File processing: Parse PDFs, convert formats, batch-process uploaded files

Sandbox Environment

The execution environment is a **secure sandbox**, isolated from production systems — it won't affect any other data

Supports standard Python along with common data and document processing packages

Files generated by the Agent appear in the **Output** panel within the chat interface and can be downloaded directly

When to Use

Finance/sales teams ask the Agent to organize Excel files and produce weekly reports

HR/admin teams ask the Agent to turn survey results into charts and generate PowerPoint presentations

Legal/contract teams ask the Agent to extract key fields from PDFs and output Word summaries

Any conversation scenario that requires ad-hoc computation or on-the-fly file generation

What Do I Need to Do?

Go to Agent settings — Open the settings page of the AI Assistant you want to enable it for

Enable Code Interpreter — Turn it on in the tools or advanced capabilities section

Test — Ask the Agent to perform a task that requires computation or file generation in a conversation, and verify that the Output panel allows normal downloads

Browser Tool (Browser Operations)

What Is This?

Browser Tool lets an Agent **control a browser**: navigate to web pages on its own, click buttons, fill out forms, scroll the page, and take screenshots to "see" what's currently displayed before deciding what to do next.

This capability is also commonly referred to as **Browser Use** or **Computer Use** in the industry. Common related terms include “browser tool,” “browser operations,” “browser automation,” and “web automation.” MaiAgent's Browser Tool is a built-in version integrated directly into the Agent — once enabled, it can be used right within a conversation without connecting any external services.

What Can It Do?

Navigate & explore: Open a specified URL, follow links across pages

Interact with page elements: Click, double-click, type text, press keys, scroll

Visual understanding: Captures a screenshot after each action; the Agent reads the screen to decide the next step

Multi-step tasks: Sequential workflows like login → search → filter → retrieve data

How Is It Different from Tools and Code Interpreter?

	Tools (API MCP)	Code Interpreter	Browser Tool
Acts on	Backend services APIs	Sandbox programming environment	Real browser & web pages
Can see the screen?	No (JSON responses only)	Can see sandbox output	Yes (screenshots)
Best for	Structured data queries, well-defined API calls	Computation, file processing, document generation	Services with a web UI but no API

Quick rule of thumb: The target has an API → Tools; you need to compute or generate files → Code Interpreter; you can only interact via a web page → Browser Tool.

When to Use

Internal systems with no public API but a Web UI you can operate
Scraping information from third-party websites that require login or interaction
Automating multi-page, multi-step web workflows (search, filter, download)
Anything "a person would do in a browser" that you want the Agent to handle automatically

Caveats

Optimized for Claude models: Browser Tool relies on model capabilities for interpreting visual screenshots. It works best with the Claude series; other models may produce inconsistent results.

Same-origin pages work best: Cross-domain iframes and pages with strict CSP may have limited interactivity.

Execution takes time: Each step requires loading the page, taking a screenshot, and reasoning. Multi-step tasks will be slower than a single API call — design your scenarios accordingly.

What Do I Need to Do?

Go to Agent settings — Open the settings page of the AI Assistant you want to enable it for

Enable Browser Tool — Turn it on in the tools or advanced capabilities section

Use a Claude model (recommended) — For the best visual interpretation results

Test — Ask the Agent to perform a task that requires web interaction in a conversation, and verify that the browser panel displays and interacts correctly

Voice Agent

What Is This?

Voice Agent enables AI Assistants to interact through **voice**: users speak, and the AI responds by speaking back. It integrates real-time voice models, STT (Speech-to-Text), TTS (Text-to-Speech), and interruption control under the hood, making it suitable for phone customer service, IVR, voice assistants, and similar scenarios.

Once enabled, the AI Assistant gains a "Voice Call" interface where users can tap the microphone to start a conversation.

Three Interaction Modes

Different modes suit different needs, depending on latency, voice customization flexibility, and tool support.

Mode	How It Works	Best For
Realtime	Uses a voice model for direct real-time voice conversation	Scenarios requiring the lowest latency where the default voice is acceptable
Realtime + TTS	Real-time voice model paired with custom TTS output	Scenarios needing a custom brand voice while maintaining a real-time feel
STT + LLM + TTS	Traditional pipeline: Speech-to-Text → LLM → Text-to-Speech	Scenarios requiring full tool support, where slightly higher latency is acceptable and maximum flexibility is desired

Quick decision: For the fastest and most natural experience → Realtime; for a branded voice → Realtime + TTS; for maximum flexibility and tool support → STT + LLM + TTS.

Interruption Control (Turn Handling)

A critical aspect of the voice conversation experience — can the user interrupt the AI while it's speaking? How does the system determine whether the user is actually speaking versus background noise or filler words?

Adjustable parameters:

Parameter	Description
Minimum speech duration (seconds)	The user's speech must last at least this many seconds to be recognized as "I want to interrupt"
Minimum word count	The user must say at least this many words to trigger an interruption

Note: `min_duration` and `min_words` are only effective in **STT + LLM + TTS** mode. In Realtime mode, interruption detection is handled internally by the voice model.

Conversation States

During a voice conversation, the interface displays the current state:

- Listening:** The AI is receiving the user's voice input
- Thinking:** The AI is processing (searching the knowledge base, using tools, generating a response)
- Responding:** The AI is replying with voice output
- Initializing:** The connection has just been established and is being prepared

How Is This Different from Text Chat?

	Text Assistant	Voice Agent
Input	Keyboard typing, file attachments	Microphone voice input, possibly DTMF key presses
Output	Text, Markdown, images, files	Voice, with a text transcript appended afterward
Latency requirements	Seconds-level is acceptable	Must be milliseconds-level to feel natural
Best for	Detailed information needs, reviewing conversation history	Immediate responses, hands-busy situations, phone channels

Best Use Cases

- Phone customer service:** Replace traditional IVR — AI directly answers calls, listens to questions, and provides answers
- Voice query systems:** Customers call to check orders, balances, or insurance policies
- Voice FAQ:** Common questions answered by simply speaking

Driving

hands-busy scenarios: Users can't type but need AI assistance

Accessibility needs: More friendly for users who have difficulty typing

Limitations

Tool support: Currently, **Realtime** and **Realtime + TTS** modes only support **MCP Tools**; to use API Tools, switch to STT + LLM + TTS mode

Knowledge base search: Voice Agent searches all linked knowledge bases and **cannot be restricted to specific documents only**

Microphone permissions: Users must grant browser microphone access before the conversation can begin

What Do I Need to Do?

Go to AI Assistant settings — Navigate to the settings page of the Agent you want to enable voice for

Select voice agent mode — Choose one of the three: Realtime

Realtime + TTS

STT + LLM + TTS

Configure provider and settings — Select the corresponding STT

TTS

Realtime provider and JSON configuration for your chosen mode

Adjust interruption control (optional) — In STT + LLM + TTS mode, adjust minimum duration and minimum word count

Test — Open a voice call in the interface and confirm that the Listening, Thinking, and Responding states all work smoothly

Further Reading

[Voice Customer Service Scenarios Overview](#)

[IVR Customer Service Intent Recognition](#)

[Voice Call Summary](#)

[Voice Call Quality Inspection](#)

Skills



Skills concept

Skills: Define multi-step SOPs for Agents to follow during conversations

What Is This?

A skill is a multi-step reasoning workflow for an Agent. You define an SOP in natural language, and the Agent follows this workflow step by step when it encounters a relevant scenario in conversation.

If tools are an Agent's "individual capabilities," skills are "chaining multiple capabilities into a complete workflow."

An Example

Suppose you want your Agent to handle "customer quotes":

Without skills: The Agent might give scattered responses and miss important information.

With skills: You define a quoting SOP —

- 1. First confirm the product model the customer needs
- 2. Look up the product price list (using a tool)
- 3. Confirm the customer's quantity requirements
- 4. Apply discount rules based on quantity
- 5. Generate a quote summary
- 6. Ask the customer if adjustments are needed

The Agent follows this workflow to guide the conversation, ensuring no step is missed.

What Can Be Used Inside a Skill?

Skills can reference tools and knowledge bases already mounted on the Agent, so data lookups or action execution during the workflow happen automatically.

Combination	Description
Tools	Used when the skill workflow needs to call external services

Combination	Description
Knowledge Base	Used when the skill workflow needs to look up reference materials
Resource Files	Skills can include reference documents that the Agent can consult during execution

When Should You Use Skills?

Workflows with fixed SOPs: Quoting, complaint handling, problem diagnosis

Multi-turn information gathering: Booking workflows (need to ask for date, time, number of people)

Conditional branching: Taking different processing paths based on different situations

Ensuring no steps are missed: Form-filling guidance, review workflows

When Don't You Need Skills?

Simple Q&A (a knowledge base is sufficient)

Only need to look up one piece of data or perform one action (just use a tool)

Advanced Note: Skills are defined as plain-text prompts, not flowcharts or code. This means you can describe an entire workflow in natural language, and the Agent will understand and execute it. Skills also support uploading resource files as supplementary references.

What Do I Need to Do?

Create a skill — Write manually or upload a skill definition file

Describe the workflow — Write out the complete processing steps in natural language

Attach resources — Add reference documents or tools if needed

Mount to an Agent — Enable this skill in the Agent settings

Test — Simulate user conversations to verify the workflow runs smoothly

Crawler



Crawler concept

Crawler: Automatically scrape web content and import it into the Knowledge Base

What Is This?

The crawler automatically scrapes content from web pages and imports it into the Knowledge Base. Simply provide a URL, and it will crawl the text content of that page (or the entire website) and convert it into knowledge that the Agent can reference.

When Do You Need It?


Website content: Automatically import product descriptions and service information from your company website

Public information: Scrape publicly available web pages such as regulations, announcements, and technical documents

Continuous updates: When website content changes, periodically re-crawl to keep the Knowledge Base up to date

Relationship Between Crawler and Knowledge Base

The crawler is **one of the data sources** for the Knowledge Base. Its role is:



The content scraped by the crawler ultimately enters the Knowledge Base. The Agent does not interact with the crawler directly.

What Do I Need to Do?

Enter a URL — Tell the crawler which web page or website to scrape

Set the scope — Scrape a single page only, or recursively crawl subpages

Run the crawl — Start the crawler and wait for it to complete

Import into Knowledge Base — Import the crawl results into a specified Knowledge Base

Verify the content — Check whether the imported content is correct and complete

Further Reading

[How to Use the Crawler \(Scrape Data\) Feature](#)

I Want to Do X, What Should I Use?



Decision Guide

Scenario → Module Mapping: Quickly Find the Right Starting Point

Not sure which module to use? Start from your need and find the right starting point.

Scenario Comparison Table

Let the Agent Answer Questions

I want to...	Use	Why
Let the Agent answer product-related questions	Knowledge Base	Upload product manuals so the Agent can find answers
Let the Agent answer frequently asked questions	Knowledge Base (FAQ)	Create FAQ pairs for precise matching of common questions
Let the Agent check order status	Database	Orders are in structured tables, requiring Text to SQL
Let the Agent query real-time information (weather, exchange rates)	Tools	Needs to call external APIs to get real-time data

Let the Agent Take Action

I want to...	Use	Why
Let the Agent send notification emails	Tools	Integrate with an Email API; the Agent triggers the send
Let the Agent create customer records in CRM	Tools	Integrate with a CRM API; the Agent executes the creation
Let the Agent handle complaints following an SOP	Skills	Define a complete workflow to ensure no steps are missed
Let the Agent guide customers through booking	Skills	Multi-turn conversation to collect information, following a defined process

Let the Agent Automate

I want to...	Use	Why
Automatically generate a report every morning	Agent Scheduling	Set a scheduled execution to automatically deliver results
Check system status every hour	Agent Scheduling	Set an interval execution, combined with Tools to check status
Automatically import website content into the Knowledge Base	Crawler	Scrape web content and import it into the Knowledge Base

Common Module Combinations

Most scenarios don't use just a single module. Here are common combinations:

Smart Customer Service

Agent + Knowledge Base + FAQ

The most basic and common combination. The Agent uses the Knowledge Base to answer product questions, and FAQ handles high-frequency questions.

Business Assistant

Agent + Knowledge Base + Database + Tools

Can look up product information (Knowledge Base), check inventory and pricing (Database), and help customers place orders (Tools).

Automated Reporting

Agent + Tools + Agent Scheduling

The Agent uses Tools to retrieve data, Agent Scheduling triggers it on a timer, and results are automatically pushed to Slack or LINE.

Process-Guided Bot

Agent + Skills + Tools + Knowledge Base

Skills define the complete SOP. During the process, the Agent queries the Knowledge Base for reference materials and uses Tools to execute actions.

Still Not Sure?

Start with the simplest setup: **Agent + Knowledge Base**. Get the Agent answering questions first, then decide what capabilities to add based on actual user feedback after going live.


AI Assistant

AI Assistant Features

Purpose and Use Cases of AI Assistants

AI assistants can be widely applied across diverse industries and use cases, helping enterprises automate repetitive tasks and improve operational efficiency, allowing teams to focus on high-value work and core strengths.

Below are common examples of AI assistant applications in enterprise workflows:



AI assistants can be applied across various industries, leveraging intelligent automation and real-time interaction to help enterprises optimize processes, improve service quality, and enhance user experience. The following table lists common AI assistant use cases by industry for reference.

Industry	Use Cases
E-commerce & Retail	Intelligent customer service, product recommendations, return process handling
Finance & Insurance	Policy inquiries, financial planning recommendations, risk Q&A
Education	Course teaching assistant, learning diagnostics, language practice

Industry	Use Cases
Healthcare	Health consultations, appointment scheduling, patient education support
Government	Public service inquiries, policy Q&A, public feedback
Manufacturing	Internal knowledge management, maintenance operation guides
Travel & Tourism	Itinerary recommendations, guided tours, real-time Q&A support

Four Stages of MaiAgent AI Assistant Setup

The MaiAgent AI assistant setup process consists of four stages. Below is a brief description of the goals and features for each stage. Complete step-by-step instructions are provided in the following sections.

1. Create a New AI Assistant

Customize the AI assistant to your needs — select an appropriate language model, configure RAG (Retrieval-Augmented Generation) sources, and define the AI assistant's role and mission to build an intelligent assistant tailored to your business scenario.

2. Provide Reference Materials for the AI Assistant

Build a solid knowledge foundation for the AI assistant by configuring a knowledge base, creating FAQ entries, or importing web crawler data. This further improves the accuracy and usefulness of responses.

3. Launch the AI Assistant

Decide whether to make the AI assistant publicly available or restrict it to internal use. For public deployment, integrate it through your company website, LINE, Messenger, or other channels for flexible integration with existing platforms.

4. Track AI Assistant Performance

Use conversation logs, response quality metrics, webhooks, and usage analytics for performance tracking, providing a basis for continuous optimization and experience improvement.



MaiAgent

AI 助理設定流程圖

1 建立新的 AI 助理

2 提供 AI 助理參考資料

3 讓 AI 助理正式上線

4 追蹤 AI 助理運作成效

基本設定

設定知識庫

選擇串接對話平台

所有對話

RAG 設定

FAQ 常見問題

對外

Web

LINE

Message

回覆品質控管

回答樣式設定

爬蟲

對內

內部問答

Webhook

權限設定

設定前台呈現 (限網站)

使用分析

嵌入網站・成功上線

Before you start building an AI assistant, consider reading these helpful articles!

 [How to Choose an LLM \(Large Language Model\)?](#)

 [What is RAG \(Retrieval-Augmented Generation\)?](#)

 [What is a System Prompt?](#)

Create an AI Assistant

1. Create an AI Assistant

Go to **AI Features** in the left sidebar, then select **AI Assistant**, and click **+ Create AI Assistant** in the upper right corner.



2. Name Your AI Assistant

Select the "Basic Settings" tab, and enter the name of your AI Assistant in the **AI Assistant Name** field. You can name it based on its primary task, such as XX AI Customer Service, XX Regulation Inquiry Helper, or XX Project Smart Assistant.

A single account can create multiple AI Assistants (subject to quantity limits based on your subscription plan).



3. Select RAG to Make Your AI Assistant Smarter and More Accurate

What Is RAG

Think of RAG as a combination of "a conversational assistant + a skilled librarian who excels at finding information."

A standard AI Assistant is like a person with excellent memory who is great at storytelling, but can only share knowledge they have previously learned. However, when an AI Assistant is equipped with RAG technology, it is like this assistant **goes to the library to find the latest information** before answering a question, then organizes the findings into its own words and delivers a clear response.

In the MaiAgent platform, this "library" is our **Knowledge Base**. The AI Assistant uses RAG technology to retrieve relevant information from the Knowledge Base, making responses more accurate, timely, and tailored to your needs.

The Knowledge Base setup will be explained in detail in the next section.

MaiAgent RAG not only incorporates the RAG technology mentioned at the OpenAI Developer Conference but also combines various classic NLP algorithms with proprietary retrieval techniques. Based on internal datasets compared with OpenAI RAG response accuracy, both achieve 95% response precision.

RAG Settings

Select the **RAG Settings** tab, and choose a different RAG (Retrieval-Augmented Generation) from the **RAG** dropdown menu. If you have no special requirements, **the default is MaiAgent RAG**.



Compared to OpenAI RAG, MaiAgent RAG offers more additional features, can be flexibly applied to different deployment needs, handles more diverse data processing formats, and provides a more powerful retrieval and generation experience. For a detailed comparison, see [What Is RAG? Comparison Between MaiAgent RAG and OpenAI RAG](#)

Configure FAQ Priority Response

To ensure the accuracy and consistency of your AI Assistant's responses, you can enable the "FAQ Priority Response" feature. When this feature is enabled, if both FAQ content and other document content exist in the Knowledge Base, the LLM will prioritize using FAQ content to answer questions, ensuring responses align with your established standard answers.

Feature Benefits:

Improved response accuracy: When both FAQ and other documents exist in the Knowledge Base, verified FAQ content is used first, reducing potentially inaccurate AI-generated answers

Ensured response consistency: All users receive the same standard FAQ answers, maintaining service quality

Reduced hallucination risk: FAQ standard answers are referenced first, preventing the AI from generating fabricated or incorrect information

Setup Steps:

Select the **Response Mode Settings** tab

In the **System Prompt** field, enter the following instruction:

Please prioritize using FAQ content to answer user questions

基本設定

回答模式設定

知識庫設定

工具

權限設定

進階設定

回答模式

☐ RAG 問答

反應快速，通常 2 秒內輸出答案，適合簡易客服與資訊查詢，無法調用外部工具

☒ Agent 模式

具備自主規劃能力，可調用外部工具與資料庫，適合用來完成複雜任務，速度比「RAG 問答」慢

> 其他模式

啟用內嵌引用功能

☒

啟用內嵌引用功能，會在回應中插入 [1][2] 格式的引用標記

啟用工具搜尋功能

☒

啟用工具搜尋功能，當 AI 助手尋找到可用的工具來輔助回答

角色指令

① 使用提示

- 明確定義 AI 助理的身份、回答範圍和語氣風格，可大幅提升回覆品質
- AI 助理無法自動取得當前時間，需至「對話平台 → 功能設定 → 時區設定」手動啟用

快速插入提示詞

請優先使用faq的內容回答使用者問題

資料庫 URL

請選擇

請輸入

After enabling this feature, when both FAQ content and other document content exist in the Knowledge Base, the AI Assistant will prioritize using FAQ content to answer questions. If no relevant FAQ content is found, the AI Assistant will still respond based on other Knowledge Base documents or general knowledge.

4. Select a Model to Give Your AI Assistant a Smart Brain!

Purpose of Language Model Selection


The performance of each AI Assistant largely depends on its brain -- the language model (LLM) it uses. In this step, you can choose different types of models based on your needs, which will affect response speed, comprehension ability, and the depth and naturalness of answers.

Choosing the right model is like enabling high-performance mode for your AI Assistant, providing the best experience tailored to your use case!

Key Factors When Choosing a Large Language Model

Language Model Settings

Select the **Basic Settings** tab, and choose a different large language model from the **LLM Model** dropdown menu. If you have no special requirements, **the default is Claude 4.5 Haiku**.



5. Create a System Prompt for Your Use Case

To make your AI Assistant better suited to different application needs, you can set up a "System Prompt" to make the AI's response style and content more aligned with the scenario.

[What Is a System Prompt? System Prompt Example Templates? AI Tool for Generating Role Assistants](#)

Hallucination-Free Generative AI Response Mechanism

MaiAgent's "Hallucination-Free Generative AI Response Mechanism" ensures high accuracy when the AI answers questions. When facing uncertain or out-of-scope questions, it honestly expresses its limitations rather than generating fabricated answers, providing users with a more reliable and trustworthy AI interaction experience. The importance for various industries and public sector applications is described below:

Industry Applications

Financial Industry:

When handling investment advice and risk assessments, the AI must provide analysis based on verified data, avoiding false information that could lead to incorrect investment decisions. When information is insufficient or uncertain, the system clearly states so, ensuring the reliability of investment decisions.

Healthcare Industry:

When assisting with medical diagnosis and drug consultation, the AI system must strictly follow established medical knowledge and must not generate suggestions that could mislead patients. For novel or unverified medical information, the system clearly indicates that further professional consultation is needed.

Manufacturing Industry:

In applications such as production process optimization and quality control, the AI must provide recommendations based on actual production data and verified methods, avoiding production losses due to inaccurate predictions.

Education Industry:

When assisting with teaching and answering student questions, the AI needs to provide accurate knowledge rather than incorrect information that could mislead learning. For complex or ambiguous concepts, the system honestly acknowledges its understanding limitations.

Legal Industry:

When providing legal information and advice, the AI must be based on existing regulations and case law, rather than offering speculative advice that may carry legal risks. The system clearly identifies matters that require further confirmation by a professional lawyer.

Customer Service:

When handling customer inquiries, the AI must provide accurate product information and service descriptions. For questions that cannot be confirmed, it immediately refers them to relevant professionals, avoiding customer misdirection.

Public Sector Applications

Government Policy Consultation:

When providing citizens with policy information and service guidelines, the AI must answer based on the latest and correct regulations and administrative procedures, avoiding outdated or incorrect information. When encountering complex questions requiring professional judgment, the system clearly advises citizens to seek assistance from the relevant department.

Public Service Decision-Making:

When assisting government in evaluating public infrastructure and social welfare decisions, the AI must analyze based on real data and research. Uncertain predictions must be clearly stated to ensure the reliability of policy-making.

Emergency Response Management:

When handling natural disasters, public health, and other emergencies, the AI system must provide accurate information and guidance. It must not generate false information that could mislead the public and affect disaster response effectiveness.

Select the Appropriate Response Mode to Build Prompts

Select the **"Response Mode"** tab. From free-form conversation to highly structured responses, it meets various business needs. Each mode has its own features and applicable scope, and you can flexibly choose based on your actual use case.



Response Mode: General

Applicable Scenarios

The AI answers freely based on the question, generating the most suitable response according to context and Knowledge Base content. Suitable for most Q&A scenarios.

Operation Steps

Select the **Response Mode Settings** tab, choose **"General (Default)"** as the response mode, and enter the System Prompt you have defined for this AI Assistant in the **System Prompt** field. For the output format, you can choose to output plain text or JSON format.



Use Case: Website Customer Service Assistant

To create a website customer service assistant for "MaiAgent - AI Assistant Development Platform," you can enter the AI's response settings in the "System Prompt" field, clearly defining its response style and scope of responsibilities.



Response Mode: Template

When you select **Response Mode: Template**, the probability of hallucinated responses from your AI Assistant is reduced to zero, because the AI Assistant will strictly reply based on the Knowledge Base and

FAQ content you have created, using a classification-based approach to answer questions rather than having the LLM generate content. It uses a template system to generate answers, guaranteeing 100% hallucination-free responses.

Applicable Scenarios

Scenarios requiring unified response formats, such as standard consultation processes and reports.

Case Study

Suppose you want to create a "Tainan City Government 1999" Customer Service AI Assistant to respond to citizens' municipal-related questions in real time.

1. Select the **Response Mode Settings** tab, and choose **"Template"** as the response mode



2. Fill in the Response Template (using the template system to generate answers)

If you have not yet created a Knowledge Base and FAQ, the default response template instructions are as follows.

It is recommended to first create a Knowledge Base or FAQ to use the response template, and the format **must be tabular files such as Excel, CSV, json, or jsonl**.

回應範本
請將範本置於以下格式中，以便系統辨識。

範本範例如下，請填入FAQ 內容：

```
[begin]
[question] [question]
[answer]
[agency]
[category]
[publish_date]
[end]
```

請開始輸入範本：

In this example, I have an AI Assistant called "Tainan 1999," and I have added "Tainan 1999 FAQ" to the Knowledge Base.

The FAQ fields include Question, Answer, Agency, Category, and Publish Date.

It is recommended to upload tabular files such as **Excel, CSV, json, or jsonl**, and the fields must have **headers and corresponding content**.

	A	B	C	D	E
1	問題	答案	機關	分類	發布時間
2					
3					
4					
5					
6					
7					
8					
9					
10					
11					
12					
13					
14					
15					
16					
17					
18					

Now return to editing the response template, and click **Initialize Response Template** in the upper right corner. You will see that the system has generated response examples based on the documents you just uploaded to the Knowledge Base.

回應範本
請將範本置於以下格式中，以便系統辨識。

範本範例如下，請填入FAQ 內容：

```
[begin]
[question] [question]
[answer]
[agency]
[category]
[publish_date]
[end]
```

請開始輸入範本：

系統已根據您上傳的知識庫文件，自動生成回應範本，請參考範本格式，填入FAQ 內容。

範本範例如下，請填入FAQ 內容：

```
[begin]
[question] [question]
[answer]
[agency]
[category]
[publish_date]
[end]
```

請開始輸入範本：

Now you can edit and format the response text.

The

[] {} portions are system instructions, so you only need to handle the text portions.

📌 **Loop: Enter the corresponding document filename (specify which document to use for response generation)**

I modified the following parts:

[x] "Opening statement" content

[x] "Question" content

[x] "Category" content

[x] Removed "Agency"

[x] Closing statement content

[x] Line break formatting

回覆模板

建議新增知識庫或 FAQ 以使用回覆模板

感謝市民提出的問題，根據我的查詢，您的問題與下方資料相關

[loop: 台南 1999 FAQ.xlsx]

{@loopIndex}.

您詢問的問題 {問題}

答案: {答案}

所屬分類: {分類}

發布時間: {發布時間}

[endloop]

謝謝您的詢問，希望可以解決您的問題。

When you return to the AI Assistant Q&A interface and ask a related question, the AI Assistant's response will completely follow the **template format** and **FAQ document content** you just edited.

台南市政 1999

2025/04/23 12:24:36

什麼是農地重劃？

2025/04/23 12:31:22

台南市政 1999

已使用知識庫

感謝市民提出的問題，根據我的查詢，您的問題與下方資料相關

1. 您詢問的問題: 什麼是農地重劃？
答案: 農地重劃係指將一定區域內不合經濟利用的農地加以重新整理，予以交換分合，劃分整理成一定標準區塊，並同時配合興修水利，改良灌溉排水，配置農水路，使每一區塊農地均能直接臨路、直接灌溉及直接排水，以改善生產環境，達到增進農地利用的一種土地改良方式。
所屬分類: 地政
發布時間: 2024-11-13T00:00:00
謝謝您的詢問，希望可以解決您的問題。

引用

台灣 1999 FAQ.xlsx 引用節點

2025/04/23 12:31:28

社宅的租期限制

2025/04/23 12:32:19

台南市政 1999

已使用知識庫

感謝市民提出的問題，根據我的查詢，您的問題與下方資料相關

1. 您詢問的問題: 社宅的租期限制？
答案: 租賃契約之期限最長為三年，承租人於租賃期限屆滿時仍符合承租資格者，得申請續租，租賃及續租期限合計最長不得超過六年。但符合住宅法第四條規定者，租賃期限得延長為十二年。
所屬分類: 社會住宅
發布時間: 2024-11-05T00:00:00
謝謝您的詢問，希望可以解決您的問題。

引用

請輸入您想問的問題

Powered by MAIA GENT

3. Fill in the Unable to Reply Template

If the AI Assistant determines there is no relevant data, it will respond based on the **Unable to Reply Template**.

Finally, click the **Save** button to complete.



Response Mode: Hybrid

Applicable Scenarios

Combines general responses with templates, using template format for some content while freely answering the rest. Suitable for scenarios requiring partially structured responses.

Operation Suggestions

The System Prompt content may conflict with "Response Mode: Template" at this point, so when selecting "Response Mode: Hybrid," it is recommended to write the System Prompt content with general principles and guidelines.

For example, the task of the System Prompt, interaction principles, communication attitude, etc.



Response Mode: Workflow

Applicable Scenarios

Suitable for specific task-oriented scenarios, such as knowledge management, data summarization, and brainstorming.

Operation Steps

Go to the "Response Mode Settings" tab and select "Workflow" as the response mode.\ This mode is suitable for task-oriented application scenarios, such as:

Knowledge Management: Assist in compiling, retrieving, or maintaining internal knowledge data

Information Summarization: Quickly organize document highlights or meeting notes

Planning and Writing: Assist in brainstorming ideas, drafting proposals, etc.

Select the corresponding workflow module based on your actual business needs to maximize the effectiveness of your AI Assistant for specific tasks.



Use Case: Copywriting Assistant

Suppose you are a marketing professional at a food company, writing promotional copy for a newly launched healthy snack.

First, select "Writing Assistant" under "Workflow."



Next, go to the AI Assistant Q&A interface and submit your writing request. The AI Assistant will guide you through filling in the key information needed for the copy, such as:

Copy topic

Writing style

Target audience

Word count

The AI will generate copy options that match your scenario and communication needs based on your settings, helping you brainstorm quickly and improve writing efficiency.

三

寫作助理測試

⋮

擬定一份宣傳健康零食的廣告文案

2025/04/23 12:57:36

✎ 寫作助理測試

謝謝你的回覆：我們可以開始進一步了解你的偏好，以便更好地擬定這份宣傳健康零食的廣告文案。
首先，請問你希望這篇廣告文案的主題是什麼？例如，是否想強調健康、口感、或是其他特點？

2025/04/23 12:57:42

主題強調健康兼具美味
風格幽默有趣
字數不超過兩百字

2025/04/23 12:58:40

✎ 寫作助理測試

謝謝你的回覆：我們已經了解了你對於廣告文案的主題、風格和字數的偏好。接下來，我們想進一步了解你的其他偏好。
請問你希望這篇廣告文案使用哪種語言？例如，中文、英文或其他語言？此外，你對於目標受眾有什麼特定的想法嗎？例如，年齡層或性別？

2025/04/23 12:58:51

中文，受眾為青壯年人

2025/04/23 12:59:29

✎ 寫作助理測試

謝謝你的回覆：我們已經了解了你對於語言和目標受眾的偏好。接下來，我們想進一步了解你的其他偏好。
請問你對於這篇廣告文案的數據準確性有什麼要求嗎？例如，是否需要引用具體的數據或研究結果來支持文案中的主張？此外，你對於文章的長度有沒有特別的要求？例如，是否希望有特定的字數範圍？
隨時可以結束偏好調查。

2025/04/23 12:59:38

結束偏好調查 開始寫作

2025/04/23 13:58:48

🗣️ 請輸入你想問的問題

🔍

Powered by

ATAKORNY

獨特賣點：我們的產品如何改變你的飲食

想要改變飲食，卻又不想成為那個在派對上默默啃生菜的“健康怪人”？別擔心，我們的產品就像是你飲食中的超級英雄，讓你在享受美味的同時，還能輕鬆保持健康！想像一下，當你咬下那口香脆的健康零食，瞬間感受到的不是罪惡感，而是滿滿的幸福感，彷彿在嘴裡開了一場派對！

我們的產品不僅僅是食物，它們是你生活中的小助手。幫助你在忙碌的日常中，輕鬆攝取所需的營養。無論是工作間隙的小點心，還是健身後的能量補充，我們的美味選擇都能讓你在享受的同時，悄悄地為身體加分。想要擺脫那些無聊的飲食選擇？我們的產品就像是你的飲食魔法師，讓每一口都充滿驚喜！

所以，別再猶豫了，讓我們的產品成為你飲食的秘密武器，讓健康與美味攜手共舞，讓你在每一餐中都能感受到生活的樂趣！

產品介紹：健康又美味的選擇

在這個快節奏的生活中，健康與美味似乎總是難以兼得，但現在，這一切都將改變！想像一下，當你咬下那一口香脆的零食，卻不必擔心卡路里的負擔，這就是我們最新推出的健康美味產品的魅力所在！

這款產品不僅選用天然食材，還經過精心調配，讓每一口都充滿驚喜。無論是香濃的堅果、清新的水果，還是富含纖維的穀物，這些成分不僅能滿足你的味蕾，還能為你的身體提供所需的營養。想要在工作間隙來一點小點心？這款產品絕對是你的最佳選擇，讓你在享受美味的同時，還能保持活力滿滿！

而且，這款產品的包裝設計時尚又便攜，無論是在辦公室、健身房，還是戶外野餐，都能輕鬆攜帶，隨時隨地享受健康美味。別再猶豫了，讓我們一起告別那些無趣的零食，迎接這款健康又美味的選擇，讓生活更加多姿多彩！

用戶評價：真實的改變

在這個快節奏的生活中，健康與美味似乎總是難以兼得，但現在，這一切都將改變！想像一下，當你品嚐到一口香濃的健康零食，嘴裡的美味讓你忍不住想要再來一口，而心裡卻知道這是對身體的最佳選擇。這就是我們的秘密武器——全新推出的健康美味零食，讓你在享受美味的同時，還能輕鬆保持健康！

這款零食不僅低糖、低卡，還富含纖維和天然成分，讓你在享受美味的同時，無需擔心卡路里的負擔。想要在工作時提神？想要在運動後補充能量？這款零食都能滿足你的需求！而且，獨特的口味組合，讓每一口都像是在舌尖上跳舞，讓你忍不住想要分享給朋友。

許多用戶已經開始體驗這款零食帶來的真實改變，從體重管理到提升能量，無不讚不絕口。現在就加入這場健康美味的革命，讓你的生活從此不再平凡！

結尾：行動起來，享受健康美味

想要在健身房裡成為超級英雄，卻又不想放棄美食的快樂？別擔心，健康飲食不再是你心中的「禁忌之食」！想像一下，當你咬下那一口香脆健康的健康零食，瞬間感受到的不是罪惡感，而是滿滿的幸福感，彷彿在嘴裡開了一場派對！



請輸入您想問的問題



Powered by



Response Mode: Agent

Applicable Scenarios

In daily business operations, employees frequently need to compile data, perform data analysis, and create reports and analytical summaries in response to routine business needs or tasks assigned by managers. Common query questions include:

"Which product had the highest sales last month?"

"List all customers with sales exceeding 100,000"

"What is the revenue trend over the past three months?"

If non-technical personnel handle these questions, they often need to rely on data teams to write SQL queries, making the process time-consuming and inefficient.

Now, through MaiAgent's Agent mode, the system can use the Text to SQL tool to automatically convert natural language questions into corresponding SQL syntax and query the database in real time, quickly delivering analysis results.

For Text to SQL feature details, see: [Text to SQL Feature](#)

This feature is particularly suitable for scenarios requiring **real-time queries and data insights**, such as report analysis, operational metric tracking, and data queries. It enables non-technical users to easily access data and achieve a more intuitive and efficient data-driven decision-making process.

Text to SQL Operation Steps

Go to the "Response Mode Settings" tab and select **"Agent"** as the response mode



Upload database content or select a database URL

```
{% hint style="info" %  
}
```

MaiAgent supports:

MySQL

PostgreSQL

Oracle DB

Microsoft SQL Server (MSSQL)

The maiagent option applies Excel files you have already uploaded to the MaiAgent Knowledge Base
{% endhint %
}

Use Case: E-commerce Product Sales Data Query

Suppose you are a marketing professional at an e-commerce platform and want to quickly query product sales data.

First, select **"Agent"** as the response mode in the "Response Mode Settings" tab.



Next, you can upload an Excel file to the Knowledge Base, and the system will automatically convert it into a queryable database format.

For detailed instructions on using MaiAgent Knowledge Base for Text to SQL, see: [Using MaiAgent Knowledge Base for Text to SQL](#)

1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31	32	33	34	35	36	37	38	39	40	41	42	43	44	45	46	47	48	49	50	51	52	53	54	55	56	57	58	59	60	61	62	63	64	65	66	67	68	69	70	71	72	73	74	75	76	77	78	79	80	81	82	83	84	85	86	87	88	89	90	91	92	93	94	95	96	97	98	99	100	101	102	103	104	105	106	107	108	109	110	111	112	113	114	115	116	117	118	119	120	121	122	123	124	125	126	127	128	129	130	131	132	133	134	135	136	137	138	139	140	141	142	143	144	145	146	147	148	149	150	151	152	153	154	155	156	157	158	159	160	161	162	163	164	165	166	167	168	169	170	171	172	173	174	175	176	177	178	179	180	181	182	183	184	185	186	187	188	189	190	191	192	193	194	195	196	197	198	199	200	201	202	203	204	205	206	207	208	209	210	211	212	213	214	215	216	217	218	219	220	221	222	223	224	225	226	227	228	229	230	231	232	233	234	235	236	237	238	239	240	241	242	243	244	245	246	247	248	249	250	251	252	253	254	255	256	257	258	259	260	261	262	263	264	265	266	267	268	269	270	271	272	273	274	275	276	277	278	279	280	281	282	283	284	285	286	287	288	289	290	291	292	293	294	295	296	297	298	299	300	301	302	303	304	305	306	307	308	309	310	311	312	313	314	315	316	317	318	319	320	321	322	323	324	325	326	327	328	329	330	331	332	333	334	335	336	337	338	339	340	341	342	343	344	345	346	347	348	349	350	351	352	353	354	355	356	357	358	359	360	361	362	363	364	365	366	367	368	369	370	371	372	373	374	375	376	377	378	379	380	381	382	383	384	385	386	387	388	389	390	391	392	393	394	395	396	397	398	399	400	401	402	403	404	405	406	407	408	409	410	411	412	413	414	415	416	417	418	419	420	421	422	423	424	425	426	427	428	429	430	431	432	433	434	435	436	437	438	439	440	441	442	443	444	445	446	447	448	449	450	451	452	453	454	455	456	457	458	459	460	461	462	463	464	465	466	467	468	469	470	471	472	473	474	475	476	477	478	479	480	481	482	483	484	485	486	487	488	489	490	491	492	493	494	495	496	497	498	499	500	501	502	503	504	505	506	507	508	509	510	511	512	513	514	515	516	517	518	519	520	521	522	523	524	525	526	527	528	529	530	531	532	533	534	535	536	537	538	539	540	541	542	543	544	545	546	547	548	549	550	551	552	553	554	555	556	557	558	559	560	561	562	563	564	565	566	567	568	569	570	571	572	573	574	575	576	577	578	579	580	581	582	583	584	585	586	587	588	589	590	591	592	593	594	595	596	597	598	599	600	601	602	603	604	605	606	607	608	609	610	611	612	613	614	615	616	617	618	619	620	621	622	623	624	625	626	627	628	629	630	631	632	633	634	635	636	637	638	639	640	641	642	643	644	645	646	647	648	649	650	651	652	653	654	655	656	657	658	659	660	661	662	663	664	665	666	667	668	669	670	671	672	673	674	675	676	677	678	679	680	681	682	683	684	685	686	687	688	689	690	691	692	693	694	695	696	697	698	699	700	701	702	703	704	705	706	707	708	709	710	711	712	713	714	715	716	717	718	719	720	721	722	723	724	725	726	727	728	729	730	731	732	733	734	735	736	737	738	739	740	741	742	743	744	745	746	747	748	749	750	751	752	753	754	755	756	757	758	759	760	761	762	763	764	765	766	767	768	769	770	771	772	773	774	775	776	777	778	779	780	781	782	783	784	785	786	787	788	789	790	791	792	793	794	795	796	797	798	799	800	801	802	803	804	805	806	807	808	809	810	811	812	813	814	815	816	817	818	819	820	821	822	823	824	825	826	827	828	829	830	831	832	833	834	835	836	837	838	839	840	841	842	843	844	845	846	847	848	849	850	851	852	853	854	855	856	857	858	859	860	861	862	863	864	865	866	867	868	869	870	871	872	873	874	875	876	877	878	879	880	881	882	883	884	885	886	887	888	889	890	891	892	893	894	895	896	897	898	899	900	901	902	903	904	905	906	907	908	909	910	911	912	913	914	915	916	917	918	919	920	921	922	923	924	925	926	927	928	929	930	931	932	933	934	935	936	937	938	939	940	941	942	943	944	945	946	947	948	949	950	951	952	953	954	955	956	957	958	959	960	961	962	963	964	965	966	967	968	969	970	971	972	973	974	975	976	977	978	979	980	981	982	983	984	985	986	987	988	989	990	991	992	993	994	995	996	997	998	999	1000	1001	1002	1003	1004	1005	1006	1007	1008	1009	1010	1011	1012	1013	1014	1015	1016	1017	1018	1019	1020	1021	1022	1023	1024	1025	1026	1027	1028	1029	1030	1031	1032	1033	1034	1035	1036	1037	1038	1039	1040	1041	1042	1043	1044	1045	1046	1047	1048	1049	1050	1051	1052	1053	1054	1055	1056	1057	1058	1059	1060	1061	1062	1063	1064	1065	1066	1067	1068	1069	1070	1071	1072	1073	1074	1075	1076	1077	1078	1079	1080	1081	1082	1083	1084	1085	1086	1087	1088	1089	1090	1091	1092	1093	1094	1095	1096	1097	1098	1099	1100	1101	1102	1103	1104	1105	1106	1107	1108	1109	1110	1111	1112	1113	1114	1115	1116	1117	1118	1119	1120	1121	1122	1123	1124	1125	1126	1127	1128	1129	1130	1131	1132	1133	1134	1135	1136	1137	1138	1139	1140	1141	1142	1143	1144	1145	1146	1147	1148	1149	1150	1151	1152	1153	1154	1155	1156	1157	1158	1159	1160	1161	1162	1163	1164	1165	1166	1167	1168	1169	1170	1171	1172	1173	1174	1175	1176	1177	1178	1179	1180	1181	1182	1183	1184	1185	1186	1187	1188	1189	1190	1191	1192	1193	1194	1195	1196	1197	1198	1199	1200	1201	1202	1203	1204	1205	1206	1207	1208	1209	1210	1211	1212	1213	1214	1215	1216	1217	1218	1219	1220	1221	1222	1223	1224	1225	1226	1227	1228	1229	1230	1231	1232	1233	1234	1235	1236	1237	1238	1239	1240	1241	1242	1243	1244	1245	1246	1247	1248	1249	1250	1251	1252	1253	1254	1255	1256	1257	1258	1259	1260	1261	1262	1263	1264	1265	1266	1267	1268	1269	1270	1271	1272	1273	1274	1275	1276	1277	1278	1279	1280	1281	1282	1283	1284	1285	1286	1287	1288	1289	1290	1291	1292	1293	1294	1295	1296	1297	1298	1299	1300	1301	1302	1303	1304	1305	1306	1307	1308	1309	1310	1311	1312	1313	1314	1315	1316	1317	1318	1319	1320	1321	1322	1323	1324	1325	1326	1327	1328	1329	1330	1331	1332	1333	1334	1335	1336	1337	1338	1339	1340	1341	1342	1343	1344	1345	1346	1347	1348	1349	1350	1351	1352	1353	1354	1355	1356	1357	1358	1359	1360	1361	1362	1363	1364	1365	1366	1367	1368	1369	1370	1371	1372	1373	1374	1375	1376	1377	1378	1379	1380	1381	1382	1383	1384	1385	1386	1387	1388	1389	1390	1391	1392	1393	1394	1395	1396	1397	1398	1399	1400	1401	1402	1403	1404	1405	1406	1407	1408	1409	1410	1411	1412	1413	1414	1415	1416	1417	1418	1419	1420	1421	1422	1423	1424	1425	1426	1427	1428	1429	1430	1431	1432	1433	1434	1435	1436	1437	1438	1439	1440	1441	1442	1443	1444	1445	1446	1447	1448	1449	1450	1451	1452	1453	1454	1455	1456	1457	1458	1459	1460	1461	1462	1463	1464	1465	1466	1467	1468	1469	1470	1471	1472	1473	1474	1475	1476	1477	1478	1479	1480	1481	1482	1483	1484	1485	1486	1487	1488	1489	1490	1491	1492	1493	1494	1495	149
---	---	---	---	---	---	---	---	---	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	-----	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	-----

官網銷售額最高的三個品項是什麼，排除運費

2025/04/23 14:58:43

 電商營業額分析助手

根據查詢結果，官網銷售額最高的三個品項（排除運費）依序為：

1. 6201可可吻9：總銷售額為15,746元
2. 5107經典黑9：總銷售額為9,834元
3. 5208巧克力5.5：總銷售額為9,376元

這三項產品在官網的銷售表現最為突出，其中6201可可吻9的銷售額明顯領先其他品項。

2025/04/23 14:58:54   



請輸入你想問的問題



Powered by  MAIADENT

6. Pre-assign Assistant Permissions

RBAC-Based Permission Management Architecture and Description

Select the **Permission Settings** tab to configure which members the assistant should be pre-assigned to for access. By default, all roles are selected. You can modify the settings based on your use case.



7. Complete Your AI Assistant Setup


After completing the above steps, click the blue **Confirm** button at the bottom right of the dialog to finish the setup.

Role Prompt Design Guide

You can use the [System Prompt Generator AI tool](#) created by MaiAgent to help you refine your system prompt

What Is a "System Prompt"?

Imagine you're directing a movie and need to tell the actors what role they're playing, what personality they have, and how they should speak. A **system prompt** is like a "script" for the AI, telling it what kind of customer service agent to portray.



Glossary of Technical Terms

Before we begin, let's get familiar with a few important concepts:

AI Agent

Analogy: Like a virtual employee you've hired

Explanation: An AI system capable of autonomously executing tasks; in customer service scenarios, it's your digital customer service representative

Prompt

Analogy: An "instruction manual" for the AI

Explanation: The instructional text that tells the AI what to do and how to do it

Context

Analogy: The "memory" and "background information" of a conversation

Explanation: The relevant information the AI needs to understand the current conversation scenario

Token

Analogy: A "unit" through which AI processes text, like breaking sentences into small chunks

Explanation: The basic unit of text processing for AI, affecting processing efficiency and cost

Why Is the System Prompt So Important?

AI Customer Service Without a System Prompt

User: "Why hasn't my order arrived yet?"

AI: "According to system data, your order status is processing."

Feels like talking to a robot — cold and impersonal

AI Customer Service With a Good System Prompt

User: "Why hasn't my order arrived yet?"

AI: "I completely understand your concern! Let me check your order status right away.


Based on the tracking information, your package is currently at the distribution center,

and it's expected to arrive tomorrow afternoon. Would you like me to provide the tracking number?"

Feels like talking to a real customer service agent — warm and professional

The Four Pillars of a System Prompt

AI Role Instruction Framework



1. Identity Definition (Who)

Tell the AI "who you are"

Bad — Vague identity

```
You are a customer service agent
```

Good — Clear identity

```
You are "Mei," a senior customer service specialist at MaiAgent Company,  
with 5 years of customer service experience, specializing in product consultation  
and problem resolution.
```

2. Tone and Style (How)

Tell the AI "how to speak"

Tone Selection Guide:

Professional and formal: Suitable for finance, legal, and healthcare industries

Friendly and approachable: Suitable for retail, dining, and lifestyle services

Lively and youthful: Suitable for gaming, entertainment, and fashion industries

Warm and caring: Suitable for education, health, and social services

Example: Friendly and Approachable Style

When responding, use a warm and friendly tone, use polite forms of address to show respect, include phrases like "happy to help" for a personal touch, but avoid being overly enthusiastic or using too many exclamation marks.

3. Professional Capabilities (What)

Tell the AI "what you can do"

Your professional capabilities include:

- Product feature introductions and usage guidance
- Order inquiries and logistics tracking
- Return and exchange policy explanations
- Initial technical problem diagnosis
- Account and payment issue resolution

When encountering issues beyond your capabilities, proactively transfer the customer to a human agent.

4. Behavioral Guidelines (Rules)

Tell the AI "what you can and cannot do"

Mandatory principles:

- Protect user privacy and never disclose personal information
- Admit when you don't know something and never fabricate answers
- When handling complaints, stay calm – apologize first, then resolve
- Keep responses under 100 words for conciseness

Strictly prohibited:

- Making promises that cannot be fulfilled
- Arguing with or contradicting users
- Disclosing internal company information
- Processing refund requests beyond your authority

Hands-On: Design Your First System Prompt

Step 1: Analyze Your Business Needs

Consider these questions:

Who are your customers? (age, occupation, usage habits)

What are your product/service highlights?

What questions do customers ask most frequently?

What kind of service experience do you want to deliver?

Step 2: Choose the Right Role Configuration

E-commerce Customer Service Example

```
You are "Helper," a professional online shopping consultant.  
You are familiar with all product information, skilled at recommending suitable  
products,  
and can quickly handle order-related issues.  
Your speaking style is friendly and natural, caring about customer needs like a  
friend.
```

Healthcare Customer Service Example

```
You are "Health Helper," a professional healthcare service consultant.  
You have basic medical knowledge and can provide information about appointments,  
examination procedures, etc.,  
but you never provide medical diagnosis advice.  
Your speaking style is professional and warm, making patients feel reassured.
```

Step 3: Write the Complete System Prompt

System Prompt Template

```
# Role Identity
You are [Role Name], [Job Title] at [Company Name].
[Role background and professional experience]

# Service Objective
Your primary task is [core task],
helping customers [achieve goal] through [service approach].

# Language Style
- Tone: [Choose appropriate tone]
- Vocabulary: [Specific vocabulary requirements]
- Length: [Response length limit]

# Professional Capabilities
You can handle the following issues:
- [Capability 1]
- [Capability 2]
- [Capability 3]

# Behavioral Guidelines
Must follow:
- [Guideline 1]
- [Guideline 2]

Strictly prohibited:
- [Prohibited action 1]
- [Prohibited action 2]

# Special Situation Handling
When encountering [special situation], please [handling approach].
```

Advanced Tips: Make Your System Prompt More Vivid

1. Add Personality Traits

Basic version:

```
You are a customer service agent responsible for answering questions.
```

Vivid version:

You are "Zhi," a curious tech enthusiast.
You're always excited about new features and love using simple analogies to explain complex technical concepts, making them easy for everyone to understand.

2. Context-Aware Instructions

Adjust responses based on different scenarios:

```
# Context-Aware Instructions
- When a customer expresses dissatisfaction: respond with empathy first, then provide a solution
- When a customer asks a technical question: explain using everyday analogies
- When a customer hesitates to purchase: provide objective advice without pressuring a sale
- When a customer praises the product: express gratitude and ask if they need further assistance
```

3. Dynamic Adjustment Mechanism

```
# Conversational Adaptability
Adjust your style based on the customer's responses:
- If the customer uses formal language, be correspondingly formal
- If the customer seems impatient, be more concise and efficient
- If the customer is chatty, you can slightly extend the topic
```

4. Version Control Friendly

Add version numbers and changelogs

Maintain backward compatibility, ensuring existing workflows still function correctly after updating the system prompt

Example:

```
Version: v1.0
# Update date: October 26, 2024
# Changelog:
#   - Initial version, providing basic product recommendation functionality.
#   - Recommends products based on user preferences.
#   - Response tone set to friendly and professional.
```


Test and Optimize Your System Prompt

Phase 1: Basic Testing

Test whether the AI's responses meet expectations using common questions

Test Question Examples:

"What is your return policy?"

"When will my order arrive?"

"Is this product right for me?"

Phase 2: Boundary Testing

Test the AI's performance in extreme situations

Boundary Scenario Examples:

When a customer is emotionally agitated

When encountering a question with no known answer

When asked to do something beyond authority

Phase 3: Continuous Optimization

Collect Data

Customer satisfaction scores

Issue resolution rate

Rate of transfers to human agents

Adjustment Strategies

Supplement the knowledge base based on common questions

Adjust language style based on customer feedback

Refine behavioral guidelines based on failure cases

System Prompt Length Control Tips****

Best Practice Guidelines

The ideal system prompt length should be kept within 500-2000 words (approximately 200-800 tokens), with the optimal range typically being 800-1500 words.

Remember this core principle: "Clarity over verbosity" — overly long instructions tend to be ignored by the model, while overly short ones may lead to unstable behavior.

Why Does Length Matter?

First, there's the **model performance** consideration — an overly long prompt dilutes the model's attention, affecting the execution of key instructions. Second, there's the **context space** limitation — the system prompt occupies the conversation's context window, compressing the space available for user conversation. Finally, there's **maintenance cost** — a concise, structured prompt is easier to continuously optimize and adjust.

Three Writing Tips

Structured design: Clearly segment content and place core instructions at the beginning.

Conciseness principle: Remove duplicate descriptions and lengthy background, keeping only essential rules.

Continuous validation: Conduct A/B testing across different models (GPT, Claude, Gemini) to find the optimal length configuration.

Practical Considerations

Different AI models have varying tolerance for long prompts, requiring individual testing and validation. For token calculation, Chinese characters are approximately 1 token each, while English is approximately 1 token per 0.75 words. Since the system prompt carries higher instruction weight, keep the content focused and avoid unnecessary verbosity.

Frequently Asked Questions and Solutions

Q1: How can I make the AI's responses more natural?

A: Add specific conversation examples in the system prompt and set clear language style guidelines.

Q2: What if the AI keeps answering irrelevant questions?

A: Clearly define the service scope in the system prompt and teach the AI how to identify and handle out-of-scope questions.

Q3: How do I balance professionalism and approachability?

A: Adjust based on your target audience. B2B customers prefer professionalism, while B2C customers prefer approachability. Design A/B tests to find the optimal balance.

Conclusion: Build Your Custom AI Customer Service

Designing a system prompt is like training a new employee — it requires patience and continuous adjustment. Remember these key principles:

Clear positioning: Clearly tell the AI who it is

Consistent style: Maintain a unified language style

Clear boundaries: Clearly define what can and cannot be done

Continuous optimization: Continuously improve based on actual usage

Start designing your first system prompt now! Remember, the best system prompts are those that are continuously refined through actual use.

Add a Role Prompt to an AI Assistant

Once you have designed your system prompt, you can add it in the AI Assistant settings on MaiAgent:

Add a System Prompt

1. Go to the AI Assistant Settings Page

Select the AI Assistant you want to configure and click Settings.



2. Switch to the Response Mode Settings Page



Set the mode to Agent, General (Default), or Hybrid mode:

The "Template" and "Workflow" modes do not support system prompt configuration.




3. Enter Text in the System Prompt Section



Once configured, the AI Assistant will act according to the system prompt you defined, becoming your personalized AI Assistant.

Multi-Language Support

MaiAgent provides multi-tiered language support: **Web Chat supports 13 interface languages, the knowledge base supports content processing in all languages, and the admin dashboard offers Chinese and English interfaces.** This enables your AI assistant to engage in natural, fluent conversations with users worldwide, breaking language barriers and expanding your business reach.



Supported Languages

MaiAgent Web Chat currently supports the following 13 languages:

Chinese

Traditional Chinese (Traditional Chinese) - zh-TW

Simplified Chinese (Simplified Chinese) - zh-CN

International Common Languages

English (English) - en

Southeast Asian Languages

Thai (Thai) - th

Malay (Malay) - ms

Vietnamese (Vietnamese) - vi

Indonesian (Indonesian) - id

Filipino (Filipino) - fil

Burmese (Burmese) - my

Khmer (Khmer) - km

Lao (Lao) - lo

East Asian Languages

Japanese (Japanese) - ja

Korean (Korean) - ko

Continuous Updates: We continue to add more language support based on user demand. If you need support for a specific language, please contact our technical team.

 Language Configuration Methods

1. Configure via JavaScript

Refer to [**Technical Manual — Web Chat**](#)

2. Web Chat Settings

Configuration Steps:

Click  in the top right corner of the window



2. Click Traditional Chinese



3. Select your preferred language



🌐 Multi-Language Feature Highlights

1. Intelligent Language Detection & Response

Auto-Detection (Chinese & English): Supports automatic detection of Chinese and English; the system responds in the corresponding language.

Multi-Language Configuration: Supports multiple language systems; you can switch and select the display language in settings. Once selected, conversations respond in the chosen language.

High Accuracy: Based on advanced NLP technology, language detection accuracy reaches 98%.

2. Natural Language Understanding

Semantic Understanding: Not just translation, but genuine understanding of semantics across different languages

Cultural Adaptation: Considers expression habits across different cultural backgrounds

Slang Recognition: Understands common slang and colloquial expressions

3. Localized Responses

Word Choice: Uses wording and expressions that match local conventions

Format Adjustment: Adjusts date, number, and currency formats based on different languages

📖 Knowledge Base & Interface Language Guide

I. Interface Language Support Levels

MaiAgent Admin Dashboard

Supports 3 interface languages: Traditional Chinese, Simplified Chinese, English

Applicable for system management and configuration operations

Web Chat Interface

Supports all 13 languages for interface display

Includes voice recognition and voice response features

Covers: Chinese (Traditional & Simplified), English, Thai, Japanese, Korean, Malay, Vietnamese, Indonesian, Filipino, Burmese, Khmer, Lao

II. Knowledge Base Content Processing Capabilities

Full-Language Document Support

The knowledge base can parse and process document content in any language

No document language restrictions — supports all global language formats

Intelligent retrieval can accurately search across multi-language knowledge bases

Upload Steps:

1. Go to Knowledge Base Management

2. Select Add Document

3. Upload documents in any language (no language restrictions)

4. The system automatically detects the document language and builds multi-language indexes

III. Cross-Language Q&A Capabilities

Core Features

Multi-Language Queries: Support free-form questions in 13 languages

Full-Language Retrieval: Search for relevant content across the complete multi-language knowledge base

Intelligent Translation: Automatically translate retrieval results into the user's preferred language

Accurate Responses: Ensure completeness and accuracy of answers

Workflow

Receive Query - User asks a question in any supported language (Thai, Japanese, Korean, etc.)

Intelligent Retrieval - System searches across languages for the most relevant content

Content Integration - LLM analyzes and integrates multi-language data

Language Conversion - Translates into the user's selected display language

Output Answer - Provides a complete response

🚀 Use Cases

1. Southeast Asian Market Expansion

Regional Coverage: Supports 8 major Southeast Asian languages, covering over 600 million people

Localized Service: Understands cultural differences and business customs across countries

Cross-Border E-Commerce: Provides native-language shopping experiences for customers in different countries

2. International Customer Service Center

Global Support: Covers customer service needs across major Asian markets

24/7 Service: Multi-language customer support across time zones

Cultural Sensitivity: Understands communication styles and politeness norms across different cultures

3. Multinational Enterprise Internal

Employee Support: Provides internal Q&A services for multinational staff

Knowledge Sharing: Breaks language barriers to promote knowledge flow

Training Assistance: Multi-language training content and Q&A

4. Travel & Tourism Industry

Tourist Services: Provides real-time assistance for tourists from different countries

Attraction Information: Multi-language attraction information and guided tour services

Booking Services: Supports multi-language reservation and inquiry services

AI Customer Service Quality Management

Target Audience: Customer service managers, quality management personnel, customer service trainers

1. Quick Start: Three Quality Metrics for AI Customer Service

How to View Evaluation Report Scores

Path: AgentOps (sidebar) → AI Assistant Monitoring

In the table, you can directly view the three major scoring metrics for each conversation. Click "View" to see full details.

Why Is Evaluation Needed?

Just like reviewing call recordings of customer service agents, we also need to check the quality of AI responses.\n\nThe system automatically scores each conversation, helping you quickly identify issues.

Three Core Metrics

Metric	Plain Language Explanation	Scoring Criteria
Faithfulness Score	Is the information provided by the AI correct? Does it make things up or hallucinate?	85+  60-84  Below 60 
Answer Relevancy Score	Did the AI answer the customer's actual question?	85+  60-84  Below 60 
Context Precision Score	Did the AI find the correct reference materials and respond precisely to the context?	85+  60-84  Below 60 

Quick Assessment Method

All three metrics > 80 →  This response is good
Any metric < 60 →  Immediate improvement needed
Two or more < 70 →  Systemic issue, comprehensive review needed

2. How to Read Evaluation Reports

Report Example

```
Conversation ID: #20240120-001
Customer Question: "Do you still have the black trench coat in XL?"
AI Response: "The black trench coat is currently in stock. XL is available for
order."

Evaluation Results:
├─ Faithfulness Score (faithfulness_score): 45  (Said it was in stock, but it
actually wasn't)
├─ Answer Relevancy Score (answer_relevancy_score): 90  (Did answer the stock
question)
└─ Context Precision Score (context_precision_score): 70  (Found trench coat
data, but size information was not precise enough)

Diagnosis: AI provided incorrect inventory information
```

Three Common Issue Types

Issue A: Low Faithfulness Score (< 60)

Symptoms: The information from the AI is incorrect or fabricated

Common Causes:

Reference data is outdated (prices, inventory, policies have been updated)

Contradictory data (different documents say different things)

AI "guesses" the answer instead of relying on database content

Impact: Customers may receive incorrect information, leading to complaints

Issue B: Low Answer Relevancy Score (< 60)

Symptoms: The AI did not answer what the customer actually asked

Common Causes:

AI provided a lot of information but missed the key point
Off-topic response with irrelevant content
Only provided background without giving an actual answer

Impact: Customers need to ask again, reducing satisfaction

****Issue C: Low Context Precision Score (< 60)****

Symptoms: The AI referenced the wrong materials or was not precise enough

Common Causes:

Inaccurate keyword search
Confused data from different products/categories
Unclear titles or categories in the reference materials

Impact: Even if the AI intends to answer well, using the wrong data leads to wrong answers

3. Real-World Cases: Common Issues in the Apparel Industry

Case 1: Wrong Price (Low Faithfulness Score)

Problem Identified

Customer Question: "How much is this down jacket?"

AI Response: "This down jacket is priced at NT\$ 3,990."

Evaluation Results:

Faithfulness Score (faithfulness_score): 38 ❌

Answer Relevancy Score (answer_relevancy_score): 95 ✅

Context Precision Score (context_precision_score): 85 ✅

Actual Situation:

The product had been repriced to NT\$ 2,990 (after discount), but the AI still quoted the original price.

****Improvement Plan****

Step 1: Check the Data

Open the "Product Price List" in the Knowledge Base
Found that the price had not been updated

Step 2: Update the Data

❌ Problematic Data

Down Jacket Series

- Classic Down Jacket: NT\$ 3,990
- Long Down Jacket: NT\$ 4,990

✅ Corrected Data

Down Jacket Series (2024 Winter Prices)

Product Name: Classic Down Jacket

- Original Price: NT\$ 3,990
- Discounted Price: NT\$ 2,990 (effective 2024/1/1)
- Promotion Period: 2024/1/1 - 2024/2/28

Product Name: Long Down Jacket

- Original Price: NT\$ 4,990
- Discounted Price: NT\$ 3,990 (effective 2024/1/1)
- Promotion Period: 2024/1/1 - 2024/2/28

Last Updated: 2024/1/15

Step 3: Instruct the AI

Add the following to "AI Response Guidelines":

"When answering price questions, always confirm that the price is the latest one and state the applicable period. If there is a promotional discount, provide both the original price and the discounted price."

Case 2: Off-Topic Response (Low Answer Relevancy Score)

Problem Identified

Customer Question: "I bought something that doesn't fit. Can I return it?"

AI Response: "We offer a variety of sizes, including S, M, L, and XL. Each garment has a size chart on the product page. We recommend checking the size chart before purchasing."

Evaluation Results:

Faithfulness Score (faithfulness_score): 90 ✅

Answer Relevancy Score (answer_relevancy_score): 35 ❌

Context Precision Score (context_precision_score): 60 ⚠️

Diagnosis:\

The customer asked "Can I return it?" but the AI talked about "How to choose a size," completely missing the core question.

Improvement Plan

Step 1: Adjust Response Guidelines

Add the following to "AI Response Guidelines":

Core Principle: Answer the customer's core question first

Bad example:

Customer asks "Can I return it?"

AI answers "We have comprehensive after-sales service..." ❌ (Didn't say yes or no)

Good example:

Customer asks "Can I return it?"

AI answers "Yes. Within 7 days of delivery, if the item is unused, you can apply for a return..." ✅ (Direct answer)

Step 2: Restructure the Data

❌ Problematic Data (too unfocused)

After-Sales Service Description

Our company values customer rights and provides comprehensive after-sales service.

Please review the size chart carefully when purchasing to choose the right size.
If you have any questions, feel free to contact customer service...

✅ Corrected Data

Return and Exchange Policy

Q: Can I return an item?

A: Yes. Within 7 days of delivery, if the item is unused and tags are intact, you can apply for a return.

Return Conditions:

- ✓ Item is unused, unwashed, and not tried on for more than 10 minutes
- ✓ Tags and packaging are intact
- ✓ Return request submitted within 7 days of delivery
- ✗ Sale items and undergarments are non-returnable

Return Process:

1. Call customer service at 0800-XXX-XXX or fill out the online form
2. Customer service will provide the return address and tracking number
3. Refund will be issued to the original payment method within 7-14 days

Q: Can I exchange an item?

A: Yes. Within 7 days of delivery, you can exchange for a different size or color (same product only) at no charge.

Exchange Process:

1. Contact customer service to confirm stock availability
2. Ship the item back (same conditions as returns)
3. New item will be shipped within 3 days of receiving the returned item

Key Improvements:

- ✓ Starts with a direct "Yes" or "No" answer
- ✓ Uses Q&A format for clarity at a glance
- ✓ Lists clear conditions to avoid disputes

Case 3: Wrong Product Referenced (Low Context Precision Score)

Problem Identified

Customer Question: "What material is the black knit top made of?"

AI Response: "This knit top is made of 100% pure cotton, soft and comfortable, suitable for all seasons."

Evaluation Results:

Faithfulness Score (faithfulness_score): 88 ✓

Answer Relevancy Score (answer_relevancy_score): 90 ✓

Context Precision Score (context_precision_score): 48 ✗

Actual Situation:

The black knit top is 70% wool + 30% polyester. The AI referenced the "white knit top" data (100% pure cotton) instead.

****Improvement Plan******Step 1: Check Data Labels**

Problematic file name:

```
knit-top.pdf
```

Issue: All knit tops are in the same document, making it difficult for the AI to distinguish between them.

Step 2: Improve Data Structure**✓ Solution A: Separate Files**

```
Product Data/  
├─ knit-top_black_model-A001.pdf  
├─ knit-top_white_model-A002.pdf  
└─ knit-top_gray_model-A003.pdf
```

✓ Solution B: Clear Headings

Knit Top Product Information

Black Knit Top (Model: A001)

- Color: Black
- Material: 70% Wool + 30% Polyester
- Suitable Season: Fall/Winter
- Care Instructions: Hand wash, do not tumble dry

White Knit Top (Model: A002)

- Color: White
- Material: 100% Pure Cotton
- Suitable Season: All seasons
- Care Instructions: Machine washable, tumble dry on low

Gray Knit Top (Model: A003)

- Color: Gray
- Material: 50% Wool + 50% Acrylic
- Suitable Season: Fall/Winter
- Care Instructions: Dry clean only

Step 3: Instruct the AI

Add the following to "AI Response Guidelines":

"When a customer mentions a product's color or model number, always verify that the reference data corresponds to the correct color and model. Different colors of the same product may have different materials and specifications."

4. Three-Step Improvement Plan

When you identify issues, follow this process:

```
Low score detected
↓
Step 1: Update data content (most important)
↓
Step 2: Adjust AI response guidelines
↓
Step 3: Escalate to the technical team (if necessary)
```

Step 1: Update Data Content

When to use:

- ✓ Low Faithfulness Score (data is incorrect or outdated)
- ✓ Low Context Precision Score (data is disorganized or poorly labeled)

Checklist:

- ☐ Is the data up to date?
- ☐ Are prices, inventory, and policies correct?
- ☐ Is data for different products clearly separated?
- ☐ Are headings clear? (So the AI can find them easily)
- ☐ Is the content presented in bullet points or tables? (Rather than large blocks of text)

Data Quality Examples:

✗ Poor Data

Return Policy

Some products can be returned, but certain conditions must be met.
Some special products cannot be returned. Please check before purchasing.
Contact customer service if you need to return an item.

✓ Good Data

Return Policy

Returnable Products:

- ✓ General apparel (tops, pants, outerwear)
- ✓ Accessories (bags, hats, scarves)

Non-Returnable Products:

- ✗ Undergarments, swimwear
- ✗ Sale items (50% off or more)
- ✗ Customized items

Return Conditions (all must be met):

1. Within 7 days of delivery
2. Item is unused (tags intact, no signs of wear)
3. Packaging is intact

Return Process:

1. Call customer service at 0800-XXX-XXX
2. Provide your order number
3. Customer service will provide the return address
4. Ship the item back (registered mail recommended)
5. Refund issued within 7-14 days of receiving the item

Contact Information:

- Customer Service Hotline: 0800-XXX-XXX (09:00-21:00)
- Online Chat: Bottom right corner of the website
- Email: service@example.com

Step 2: Adjust AI Response Guidelines

When to use:

- ✓ Low Answer Relevancy Score (off-topic responses)
- ✓ Low Faithfulness Score (AI guessing or hallucinating)

AI Response Guidelines Template:

AI Customer Service Response Guidelines

Core Rules

1. Answer the core question first
 - Customer asks "Can I/Is it possible" → Answer "Yes" or "No" first
 - Customer asks "How much" → State the price first
 - Customer asks "How to" → Provide the steps first
2. Only state what you are certain about
 - All information must come from reference materials
 - If unsure, say "This requires assistance from a human agent"
 - Never guess or assume
3. Pay attention to details
 - Do not confuse colors, sizes, or model numbers
 - Verify prices are current
 - Clearly state promotion periods

Response Format

Policy Questions (returns, membership, promotions)

Paragraph 1: Directly answer "Yes" or "No"

Paragraph 2: State the conditions (in bullet points)

Paragraph 3: Tell the customer what to do (process or contact info)

Product Questions (price, material, stock)

Paragraph 1: Directly answer the question (price/material/availability)

Paragraph 2: Supplementary product info (specs, sizes, colors)

Paragraph 3: Purchase link or next steps

Process Questions (how to buy, return, exchange)

Paragraph 1: Summarize the process (3-5 steps)

Paragraph 2: Explain each step in detail

Paragraph 3: Notes or contact information

Prohibited Behaviors

- ✗ Do not say "usually," "generally," or "approximately" (be specific)
- ✗ Do not confuse information from different products
- ✗ Do not omit important details (price, size, deadline)
- ✗ Do not speculate about information the customer hasn't provided

Examples

✓ Good Response:

Customer: "Can I return this jacket?"

AI: "Yes. Within 7 days of delivery, you can apply for a return if the following conditions are met:

- Item is unused, tags are intact
- Packaging is undamaged
- Not a sale item

Return process: Please call customer service at 0800-XXX-XXX, and we will provide the return address and instructions. Refunds are processed within 7-14 business days to the original payment method."

✗ Bad Response:

Customer: "Can I return this jacket?"

AI: "Our company values customer rights and provides comprehensive after-sales service. We recommend reviewing the product description carefully and choosing the right size before purchasing. If you have any questions, feel free to contact customer service..." (Did not directly answer whether a return is possible)

Step 3: Escalate to the Technical Team

When to use:

Context Precision Score is consistently low

The same issue occurs repeatedly

No improvement after adjusting data and guidelines

Escalation Content:

Issue Type: Low Context Precision Score

Issue Description:

When customers ask about "black" products, the AI frequently retrieves data for "white" or other colors.

Scope of Impact:

Approximately 15% of product-related queries are affected

Attempted Improvements:

- ✓ Separated product data by color into individual files
- ✓ Explicitly labeled colors in headings
- ⚠ Issue is not fully resolved

Suggested Technical Adjustment:

Request the system to more accurately recognize "color" keywords

Attachments:

- test_cases_color_queries.csv (100 test questions)
- current_results.csv (current system retrieval results)
- expected_results.csv (expected correct results)

5. Daily Management Checklist

Daily Checks

When issues are found:

If the same type of issue occurs ≥ 3 times
→ Handle immediately (update data or adjust guidelines)

If it involves price or policy errors
→ Urgent fix, complete within the same day

If it's an occasional issue
→ Log and observe, add to discussion agenda

Response Quality Tracking

1. Data Review

This Week's Statistics:

- Total conversations: __ sessions
- Average Faithfulness Score (faithfulness_score): __ points
- Average Answer Relevancy Score (answer_relevancy_score): __ points
- Average Context Precision Score (context_precision_score): __ points
- Anomalous conversations: __ sessions (____%)

2. Issue Analysis

Top 3 High-Frequency Issues:

1. _____ (__ occurrences) - Which metric is low?
2. _____ (__ occurrences) - Which metric is low?
3. _____ (__ occurrences) - Which metric is low?

3. Improvement Actions

This week's tasks:

- ☐ Update __ documents (Owner: __)
- ☐ Adjust __ response guidelines (Owner: __)
- ☐ Escalate __ technical issues (Owner: __)

Next week's goals:

- Reduce anomalous conversation rate to < ____%
- All metrics averaging > __ points

Appendix A: Issue Diagnosis Quick Reference

Score Status	Possible Cause	Improvement Method
Low Faithfulness Score	Outdated or incorrect data, AI hallucination	Step 1: Update data content
Low Answer Relevancy Score	AI gives off-topic responses	Step 2: Adjust response guidelines
Low Context Precision Score	AI references wrong data or lacks precision	Step 1: Improve data labeling
Multiple metrics are low	Systemic issue	Steps 1+2, Step 3 if necessary

Improvement Priority Order

First Priority: Faithfulness Score < 60

→ May give customers incorrect information or hallucinated content, leading to complaints

Second Priority: Answer Relevancy Score < 60

→ Poor customer experience, requires repeated inquiries

Third Priority: Context Precision Score < 60

→ Although the issue is less obvious, it will affect quality over time

Appendix B: System Evaluation Metrics Reference

Primary Metrics (No Ground Truth Required)

These three metrics are the core of this guide and can be directly applied to daily customer service conversation evaluations:

Name	English Full Name	Description
Faithfulness Score	Faithfulness	Evaluates whether the AI response aligns with database content, or if it hallucinates or fabricates information
Answer Relevancy Score	Answer Relevancy	Evaluates whether the AI response is relevant to the customer's question, or if it's off-topic
Context Precision Score	Context Precision	Evaluates whether the AI response precisely addresses the context and references the correct materials

Advanced Metrics (Ground Truth Required)

The following metrics require pre-prepared "ground truth" answers and are suitable for test case evaluations:

Name	English Full Name	Description
Answer Correctness	Answer Correctness	Compares AI response against the ground truth to evaluate correctness
Answer Similarity	Answer Similarity	Evaluates semantic similarity between the AI response and the ground truth
Context Recall	Context Recall	Evaluates whether the system retrieved all necessary reference materials

Other Available Metrics (DeepEval)

The system also supports the following additional evaluation metrics for more comprehensive quality checks:

Name	English Name	Description
Bias Detection	Bias	Detects whether responses contain biased or discriminatory content
Toxicity Detection	Toxicity	Detects whether responses contain inappropriate or offensive content
Hallucination Detection	Hallucination	Detects whether the AI generates content that contradicts facts
Contextual Relevancy	Contextual Relevancy	Evaluates whether retrieved reference materials are relevant to the question

Usage Recommendations

- Daily Monitoring:** Use the three primary metrics (Faithfulness Score, Answer Relevancy Score, Context Precision Score)
- Test Evaluations:** Use advanced metrics with prepared ground truth answers for systematic evaluations
- Quality Assurance:** Enable bias and toxicity detection to ensure responses comply with corporate standards

FAQ

Q1: I'm not technical — can I still manage AI customer service?

A: Absolutely! Just like managing customer service staff, you only need to:

- Review evaluation reports daily to identify problematic conversations
- Check whether data is correct and complete
- Adjust the AI's "response guidelines" (just like training customer service scripts)

Q2: How are the scores calculated? Does the AI grade itself?

A: No. The scoring is performed automatically by a dedicated "evaluation system," like having another AI serve as "quality control" to check the first AI's responses.

Q3: Are all three metrics equally important? Can I just look at one?

A: We recommend looking at all three, as they reflect different issues:

Faithfulness Score (`faithfulness_score`): Whether the AI aligns with database content, whether it hallucinated

Answer Relevancy Score (`answer_relevancy_score`): Whether the AI understood the question and gave a relevant response

Context Precision Score (`context_precision_score`): Whether the AI precisely addressed the context and found the right reference materials

Looking at only one may cause you to miss important issues.

****Q4: How long until I see improvement after making changes?****

A:

Data updates: Take effect immediately (improvements visible the same day)

Response guideline adjustments: Take effect immediately

Technical adjustments: Require 2-4 weeks (depending on complexity)

Conclusion

Managing AI customer service is just like managing a real customer service team:

- ✓ **Regularly check quality** (review evaluation reports)
- ✓ **Continuously update knowledge** (update data content)
- ✓ **Optimize response scripts** (adjust response guidelines)
- ✓ **Track improvement results** (monitor score changes)

By following this guide, you can continuously improve your AI customer service — even without technical expertise!

Cross-Conversation Search

Cross-conversation search lets an AI assistant search recent conversations with the same contact in the same inbox when responding. It is useful when service history needs to carry over, but users frequently start new conversations.

This feature is disabled by default. Enabling it only expands the AI assistant's search scope. It does not mix content from different contacts or inboxes.

Enable Cross-Conversation Search


In the left menu, click “[AI Assistant](#).”

Select the AI assistant you want to configure and open “[Settings](#).”

Open “[Advanced Settings](#).”

Under the conversation memory settings, enable “[Cross-Conversation Search](#).”

Under “[Lookback Period \(Days\)](#),” enter a value from 1 to 365 days. The default is 90 days.



Enable cross-conversation search under Advanced Settings and configure the search range for recent conversations

Click “**Save.**”

A longer lookback period expands the searchable history. Configure it according to your service requirements and comply with your organization's data use and privacy policies.

Use Case

MaiAgent's customer service assistant tracks after-sales cases. A customer provided their order status in a conversation last week and starts a new conversation today to ask for an update. The administrator enables cross-conversation search for the customer service assistant and sets the lookback period to 30 days. After the settings are saved, the assistant can use the conversation history tool to find relevant content from the same customer's recent conversations and respond based on those records.

Scope

Cross-conversation search includes:

The same contact

The same inbox

Regular chat conversations that have not been deleted

Conversations whose latest message falls within the configured lookback period

For conversations within the lookback period, the entire message history is included, not just messages sent during that period. The system selects up to 200 of the most recently active conversations. The current conversation is always included in the search scope, regardless of the configured lookback period.

When this feature is disabled, the assistant searches only the current conversation. Cross-conversation search and “Vector Memory” are independent settings and can be enabled or disabled separately.

Custom Maximum Output Tokens

Feature Introduction

The custom max output tokens feature allows you to precisely control the maximum length of each AI assistant response. By adjusting this setting, you can have the AI assistant generate brief summaries or detailed explanations based on different use cases, while effectively managing token consumption costs.

About Tokens Tokens are the basic units used by AI models to process text. Generally speaking: - Chinese: 1 character is approximately 1.5-2 Tokens - English: 1 word is approximately 1-1.5 Tokens - Punctuation marks also count toward the token count

Configure Max Output Tokens

Navigate to Settings Page

Click "**AI Assistant**" in the left menu

Select the AI assistant you want to configure

Click the "**Settings**" button

Switch to the "**Advanced Settings**" tab

Adjust Token Limit

Select Token Limit Mode

"**Use Model Maximum**": Directly apply the maximum output tokens supported by the selected model (default)

"**Custom Amount**": Specify the limit yourself

Enter Custom Value (when "**Custom Amount**" is selected)

Enter the target value in the number input field

The available range depends on the selected model, and the screen will display the allowed range in real time (e.g., "Range: 512 - 64000 tokens" as shown below)

Save Settings

Click the "**Save**" button

Settings take effect immediately without restarting the AI assistant



AI assistant advanced settings page: Choose "Use Model Maximum" or "Custom Amount" to control the max output token limit

Important Notes - Setting the value too low may cause responses to be truncated, affecting completeness - Setting the value too high will increase token consumption costs - It is recommended to adjust based on actual needs and find the optimal balance through testing

Recommended Settings

Below are recommended token limit settings for different use cases:

Use Case	Recommended Tokens	Description
Quick Q&A	256 - 512	Suitable for quick answers, such as FAQ bots Approximately 100-200 Chinese characters
General Customer Service	512 - 1024	Balances detail and cost Approximately 200-400 Chinese characters
Professional Consultation	1024 - 2048	Provides detailed explanations and solutions Approximately 400-800 Chinese characters
Content Generation	2048 - 4096	Generates long-form articles or reports Approximately 800-1600 Chinese characters

Use Case	Recommended Tokens	Description
Code Generation	1024 - 2048	Produces complete code snippets Including comments and explanations

Use Case Examples

Scenario 1: FAQ Quick Response Bot

Requirement: Common customer questions need short and clear answers, avoiding lengthy explanations

Configuration:

- Max output tokens: 256-512
- System prompt: "Please answer questions in a brief and clear manner, no more than 3 sentences"
- Expected result: Quick responses, reduced token consumption

Example conversation:

```
User: What are your business hours?  
AI: Our business hours are Monday to Friday from 9:00 AM to 6:00 PM,  
Saturday from 10:00 AM to 5:00 PM, and we are closed on Sundays.
```

Scenario 2: Technical Support Detailed Explanation

Requirement: Users encounter technical issues and need detailed step-by-step instructions

Configuration:

- Max output tokens: 1024-2048
- System prompt: "Please provide detailed step-by-step instructions to ensure users can follow along"
- Expected result: Complete troubleshooting guide

Example conversation:

User: What should I do if the software won't start?

AI: Please follow these steps to troubleshoot:

1. Check System Requirements
 - Operating System: Windows 10 or above
 - Memory: At least 4GB
 - Disk Space: At least 2GB
2. Reinstall the Software
 - Remove the existing version (Control Panel > Programs)
 - Download the latest version [link]
 - Run the installer as administrator
3. Check Firewall Settings
 - Add the software to the exception list
 - Allow the software to access the network

If the above steps don't resolve the issue, please provide a screenshot of the error message and we will assist you further.

Scenario 3: Content Summary Generation

Requirement: Extract key summaries from long articles

Configuration:

- Max output tokens: 512-768
- System prompt: "Please organize key points into 3-5 bullet points, each no more than 2 sentences"
- Expected result: Concise summaries that are easy to quickly understand

Relationship with Other Settings

System Prompt

Even with a high token limit, the system prompt can still constrain response length:

Good example:

"Please answer concisely unless the user requests a detailed explanation"

Bad example:

"Please answer every question as thoroughly as possible" (may lead to overly long responses)

Model Selection

Different models support different maximum output token counts. When you switch the model used by the AI assistant, the customizable range will change accordingly, and the range hint displayed on the screen reflects the upper and lower limits of the currently selected model. When "Use Model Maximum" is selected, the system will automatically apply the maximum value for that model.

Cost Optimization Recommendations

Calculate Token Consumption

Token consumption per conversation includes:

- **Input Tokens:** User question + conversation history + system prompt
- **Output Tokens:** AI response (limited by this setting)

Example calculation:

```
Input: User question 50 Tokens + Conversation history 200 Tokens + System prompt
100 Tokens = 350 Tokens
Output: AI response (limit set to 512 Tokens)
Total: 350 + 512 = 862 Tokens (maximum)
```

Optimization Strategies

Adjust by scenario

Use a low token limit for simple questions

Only increase the limit for complex questions

Avoid unnecessary long responses

Explicitly require conciseness in the system prompt

Use multi-turn conversations instead of single long responses

Regularly review usage

Use the "Usage Analytics" tab to view conversation word count and other usage trends

Adjust settings to optimize costs

For detailed cost calculations and optimization recommendations, please refer to [Usage Calculation](#).

FAQ


Q: After setting the token limit, will the AI response be forcibly truncated?

A: Yes. When the AI-generated response reaches the token limit, it will stop at that point, which may result in incomplete sentences. Recommendations:

- Set an appropriate limit (not too low)
- Include in the system prompt: "End naturally before reaching the character limit"

Q: How do I check the average token usage of my AI assistant?

A: Go to the AI assistant "Settings" page and select the "Usage Analytics" tab. You can view trend charts by date range for "Conversation Word Count", "Conversation Count", "Message Count", "Average Conversation Interactions", and "User Satisfaction". The "Conversation Word Count" represents the cumulative word count generated by the assistant. Word count is positively correlated with token count and can serve as a reference for token consumption.



AI assistant "Usage Analytics" tab: View trend charts by date range for conversation word count, conversation count, message count, average conversation interactions, and user satisfaction

Q: Does setting a higher token limit slow down response speed?

A: There is a slight impact. Longer responses require more generation time, but the difference is usually minimal (within a few seconds). The main factors affecting speed are:

- Model selection
- Network connection speed
- Server load

Q: Do different languages require different token limit settings?

A: It is recommended to adjust accordingly:

- **Chinese:** The same amount of text consumes more Tokens, so consider slightly increasing the limit
- **English:** Token efficiency is higher, so a lower limit can be used
- **Multi-language environments:** It is recommended to set a higher limit to ensure flexibility

Q: Does the token limit affect knowledge base retrieval results?

A: It does not directly affect retrieval, but it affects how retrieval results are presented:

- **High token limit:** Can reference more knowledge base segments
- **Low token limit:** Only references the most relevant segments

Testing and Adjustment

Testing Process

Set Initial Value

Start with the recommended value (e.g., 1024)

Conduct Test Conversations

Enter typical user questions

Observe response length and completeness

Evaluate Results

Is the response complete?

Is there unnecessary verbosity?

Is token consumption reasonable?

Adjust Gradually

Too long → Lower the token limit

Too short/truncated → Increase the token limit

Continuous Monitoring

Regularly review usage

Fine-tune based on actual needs

Testing Examples

Test 1: Token limit 256

User: Please introduce your company

AI: We are a company focused on AI technology, providing enterprise-level chatbot solutions. [Response truncated]

Evaluation: Too short, need to increase

Test 2: Token limit 1024

User: Please introduce your company

AI: We are a company focused on AI technology... [Complete detailed introduction]... Feel free to contact us to learn more.

Evaluation: Appropriate, adopt this setting

Related Features

Extended Features - [Create an AI Assistant](#) - [System Prompt Design Guide](#) - [Usage Calculation](#) - [Usage Analytics](#)


Knowledge Base

Knowledge Base Overview

What Is the Knowledge Base?

After creating an AI assistant, the next step is to set up the **knowledge base**, which serves as the data source for the AI assistant when answering questions.

How the knowledge base works:



Knowledge Base Operation Diagram

Think of the knowledge base as an open-book exam. The AI assistant finds the most relevant content from the knowledge base to use as the basis for its answers.

Although the knowledge base can hold a large amount of data, it is still recommended to upload content that is highly relevant to the assistant's tasks. If the content is irrelevant, current retrieval technology may still find unrelated information to use as answer context, leading to inaccurate responses.

What Can the Knowledge Base Do?

With the following knowledge base features, you can:

Tag-Based Document Management

You can add multiple tags to each document to create a systematic classification structure

Example: Add `#tent` `#beginner` `#setup-guide` tags to the "Tent Manual"

Selective Document Permissions

You can set different document access permissions based on user identity to provide differentiated services

Example: Product price lists are only accessible in VIP member conversations, while cost data is only accessible in internal conversations

? Create FAQ

You can organize frequently asked questions into a standard Q&A set to ensure response consistency and completeness

Example: "How many years is the tent warranty?" → Standard answer: "We provide a 3-year warranty service"



Document Metadata Input

You can add detailed information such as version, update date, and applicable scope to documents for easier future maintenance

Example: Document version v2.1, update date 2024/11/28, applicable product "Mountain Series Tents"



Search Testing

You can test retrieval results for different questions to verify whether the AI finds the correct data segments

Example: Enter "tent setup steps" → Check whether the AI finds the correct setup guide segment

How to Upload to the Knowledge Base?


For example, suppose a knowledge base administrator at a publishing company wants an AI assistant to search the contents of product e-books. The administrator uploads EPUB e-books to the product knowledge base. Once processing is complete, the AI assistant can search the books for information to answer questions, just as it can with PDF or Word documents.

Go to the left sidebar "[AI Features](#)" → "[Knowledge Base](#)"

Find the knowledge base you want to upload to and click "[Edit](#)" in the Actions column (if you haven't created one yet, please refer to [How to Create a Knowledge Base: Basic Settings](#))

On the settings page, click the "[Knowledge Base](#)" tab

Click the blue "[Upload File](#)" button in the upper-right corner of the page to upload the data you want the AI assistant to know



The upload dialog lists EPUB and other supported file formats

Different plans provide different knowledge base storage capacities. If your files are large or you have a lot of data, the system may need some time to process after uploading.



You can check the processing progress in the "Processing Status" column.

ID	设备名称	连接名称	设备类型	连接类型	IP地址	连接策略	连接	操作
907faceb-7021-4835-8263-1417671c1422 (pc)	907faceb-7021-4835-8263-1417671c1422 (pc)	pc1	13.146	连接中	192.168.1.100 (192.168.1.100)	2023-05-25 14:15:12		
907faceb-7021-4835-8263-1417671c1422 (pc)	907faceb-7021-4835-8263-1417671c1422 (pc)	pc1	13.146	连接中	192.168.1.100 (192.168.1.100)	2023-05-22 10:25:16		

Currently supported file types for the knowledge base: Spreadsheet files: .xls, .xlsx, .csv, .ods Word processing files: .doc, .docx, .odt, .pdf, .md, .txt Presentation files: .ppt, .pptx, .odp Web files: .html, .htm Data formats: .json, .jsonl E-books: .epub Audio files: .wav, .mp3, .m4a, .aac Video files: .mp4

View Data

Data Chunk Description

When data is uploaded to the knowledge base, it is divided into chunks. This provides the following benefits:

Improved retrieval accuracy: Smaller chunks match queries more precisely, avoiding interference from irrelevant content

Optimized response speed: Faster processing with reduced computational resource consumption

Enhanced content relevance: Each chunk focuses on a single topic, improving specificity

Easier maintenance and management: Easier to update, track, and control quality

Compliance with AI model limitations: Avoids exceeding input length limits, ensuring information completeness

Edit and View Chunk Content

You can view the segmented content of your documents:

Go to the document page in the knowledge base, and click "Edit"

View the segmented content by clicking on the chunk you want to view

The screen will display the complete content of the chunk in Markdown format.



Edit Source Document

On the edit page, click "View Document"



Enter the edit page, where the text box will display the document in Markdown format

You can edit your source document in the text box. All file formats can be edited using this method. After finishing, click "**Save**" to save your edits.






Data Management Reminder It is recommended to back up important content before editing After editing the source document, the system will re-segment the data into chunks

How to Create a Knowledge Base: Basic Settings

Basic Settings

General Information

You can define the name of the knowledge base and add a description on the following page:



Retrieved Chunks

The number of retrieved chunks represents the maximum number of data chunks the AI assistant will reference when answering. The system default is "12", meaning the AI assistant will retrieve the 12 most relevant chunks for each response.

You can increase or decrease the number of retrieved chunks to adjust the amount of information the AI assistant references when answering.

Similarity Threshold

The similarity threshold excludes chunks that are not sufficiently relevant to the question. Chunks with a similarity score below the threshold will not be retrieved. Enter a value from to in increments of . Leave it blank or set it to to disable filtering.

A higher threshold narrows the retrieval results, but it may also exclude useful content. Make small adjustments each time and compare the search results using the same set of questions.

Use case: Customer service manager Ching notices that search tests for "return conditions" occasionally include unrelated delivery information. She first records the original retrieval results, then raises the


similarity threshold and saves the setting. Finally, she tests again using the same question. If the return policy chunk still appears while the delivery information is excluded, she can keep this setting.

Go to **Knowledge Bases** and click "Edit" on the right of the knowledge base you want to adjust.

Switch to **Retrieval Settings**.

Enter a value in "Similarity Threshold (0-1)."

Click "Save," then use Search Test to retest the original question.



Adjust the similarity threshold in Retrieval Settings

What Is a Parser?

A parser enables the system to "understand" the content in uploaded documents, making it searchable, editable, or convertible to other formats.

Document Parsers

When uploading PDF, Word, or other documents, you can choose from the following four parsers:

MaiAgent Parser (Default): Low cost, fast speed, suitable for plain text documents, supports 22 formats

MaiAgent Parser (Online): Uses LLM, can OCR-parse text in images, supports 20 formats

MaiAgent Parser (Offline): OCR + AI semantic understanding of images, best structure preservation, deployable on-premises, supports 20 formats

Vision Parser: AI visual understanding, best image parsing results, supports 7 formats

Speech-to-Text Parsers

When uploading audio files, you can choose from the following four speech-to-text parsers:


Azure Speech: Real-time transcription with high accuracy

Whisper (Groq): Fastest speed, low cost


Whisper (OpenAI): Stable and reliable cloud solution

Whisper (Offline): Fully local deployment, free with data privacy protection

After audio parsing is complete, you can view the transcript via "View Document" and download the transcript file in TXT or SRT format:




Transcript view



Download transcript (supports TXT and SRT formats)

For a detailed comparison of each parser, refer to: [Technical Manual - Parser Tools](#)

If issues occur during data parsing, you can click the [Re-parse] icon to have the parser reprocess the data.



How to Adjust Documents with Complex Layouts

For documents such as product catalogs and specification sheets that contain two-column layouts, tables spanning multiple pages, and images, correct parsing does not guarantee that complete content will be retrieved for every question. We recommend checking each stage in order: parsing results → retrieval results → response results.

Choose a parser suited to the content first: Use the default parser for plain-text documents. When fields and table structures must be preserved, prioritize comparing the results from MaiAgent Parser (Offline). If the questions focus on labels in images, the appearance of parts, or relationships between images and text, compare the results from Vision Parser.

Inspect the parsed Markdown: Open "**View Document**" on the document page and confirm that the reading order of the left and right columns, table titles, column names, row data, and image descriptions remain in the same section. If content is interleaved or missing at this stage, switch parsers and parse the document again. If necessary, convert the source document to a single-column layout, or save important tables separately as `.xlsx`, `.csv`, or Markdown files and upload them.

Run search tests with representative questions: Test the product name, exact fields in tables, and questions that require comparisons across rows. Confirm that the required data appears in the retrieved chunks. If the correct chunk falls outside the retrieval range, gradually increase the number of "Retrieved Chunks" and compare the results. Increasing the number of chunks also introduces more irrelevant content, so we do not recommend setting it to the maximum immediately.

Make each chunk independently understandable: If a table is split, add the product name, specification category, field names, and units to each section. For data that requires highly complete responses, you can also create an FAQ that organizes frequently asked comparison questions into complete question-and-answer pairs.

Confirm the purpose of images: A Parser can recognize text or visual semantics in images, but this does not guarantee that responses will automatically display the original images. If users must see an image,

confirm that the parsed Markdown retains it and that the chunk containing the image is retrieved. Also test how it is displayed in the actual conversation channel being used. If you only need specification details from an image, include the key text, figure number, and image description in the document.

Adjust only one setting at a time and rerun the [search test](#) with the same set of questions. This lets you determine whether the improvement comes from the parser, document content, or retrieval settings.

Retrieval Model Settings

In the knowledge base settings, you can select the Embedding model and Reranker model you want to use.




Embedding Model

Embedding is like translating human language into a "numerical language" that AI can understand, enabling computers to comprehend the true meaning of text. This process is called "vectorization." Different Embedding models have different characteristics, such as the languages they handle best and the deployment environments they support. Different model settings in the knowledge base can be used to adjust the vectorization results when documents are uploaded. You can choose the most suitable Embedding model for different scenarios.

You can freely choose from multiple Embedding models:

For details on Embedding model differences, refer to: [Technical Manual - Embedding Models](#)



Reranker Model

A Reranker acts like a professional judge that re-evaluates which data best answers the customer's question from the initial search results. What is the difference between using a Reranker and not using one?

When a customer asks: "What tent is suitable for beginners? Budget under 8,000 NTD"

Without Reranker:

AI may respond:

"We have tents at various price points. Products at 8,000 NTD include ... "
(May mention advanced models, not targeted enough for beginner needs)

With Reranker:

AI responds:

"Here are some specially recommended tents under 8,000 NTD for beginners ... "
(Precisely targeted for beginners + budget + product recommendations)

When search result reranking is enabled, the AI assistant will re-sort the retrieved knowledge base chunks and respond based on the most relevant documents.

For details on the Reranker model, refer to: [Technical Manual - Reranker Models](#)



In summary, using Embedding combined with Reranker enables the AI assistant to understand the knowledge you provide, review content importance after retrieving chunks, and respond using the knowledge most relevant to the question.

Associated AI Assistants

Multiple AI Assistants Sharing a Knowledge Base

Associating AI assistants means authorizing specific AI assistants to use this knowledge base. If you have two AI assistants:

- Product Customer Service AI
- Order Customer Service AI

When both need to answer return-related questions, you can associate both AI assistants under the "Return Policy" knowledge base settings:





Diagram of associating multiple AI assistants

Select the AI assistants to associate



Click to add the AI assistant

After adding, it will appear in the selected AI assistants area. Click "Save" in the bottom-right corner to complete the association.



Once associated, both AI assistants can share the "[Return Policy](#)" knowledge base and respond based on the same content. For ongoing maintenance, you only need to **update one knowledge base** to ensure the AI assistants use the latest data.

One AI Assistant Using Multiple Knowledge Bases

In addition to sharing knowledge bases, a single AI assistant can also use multiple knowledge bases.


Go to the AI assistant page, select the AI assistant you want to configure, and click Settings



Go to Model Settings and click "Select Knowledge Base"



Select the knowledge bases to use and click Confirm. The selected knowledge bases will appear in the list





Finally, click "Save" and the AI assistant will be able to use multiple knowledge bases

How to Use the Crawler (Web Scraping) Feature

Feature Purpose and Value

As a business professional, you may often receive instructions from management to reference or compile regulatory data from certain public websites.

If you have a technical background or engineering support, you might be able to write a scraping script to automatically extract data. However, for non-technical users, the only option is usually to manually organize pages one by one — a time-consuming, labor-intensive process that is also prone to missing key information.

In such cases, you can leverage MaiAgent's web scraping feature to quickly extract website content through a No-Code approach, automatically creating structured data. This significantly improves information organization efficiency, allowing you to invest your time in higher-value core business tasks.

How to Scrape?

To create a scraping request, you can:

Create a Page Scraping Request

Go to the left sidebar "AI Features" → "Web Scraping", and click the "+ Create Page Scraping Request" button in the upper right corner.

Enter the URL

Enter the URL of the page you want to scrape, and click the [Confirm] button.

Please note that the URL cannot exceed 200 characters. If the status has not changed, you can click the refresh button in the upper right corner to update the status.

建立頁面爬取請求

※ URL

le.mori-mori.com.tw/%E6%96%B0%E6%89%8B%E9%9C%B2%E7%87%9F/

取消

確認

基本

表 爬蟲 (Beta)

品質管理

Webhook

使用分析

即時對話平台

介紹 自定義爬蟲模式

介紹 全螢幕登入示範

介紹 全螢幕登入示範

介紹 自定義爬蟲模式

介紹 簡便快捷模式

介紹 自定義爬蟲模式

建立頁面爬取請求

URL	爬取狀態	建立時間	操作
https://article.mori-mori.com.tw/%E6%96%B0%E6%89%8B%E9%9C%B2%E7%87%9F/	完成	2025-08-19 18:23:1	刪除
https://article.mori-mori.com.tw/%E6%89%8B%E9%9C%B2%E7%87%9F/	完成	2025-08-19 17:57:29	刪除

共 2 條資料 1 10 條/頁

Review Scraped Data

Once the status shows as completed, click "Import" on the right side to review the scraped data entries.

Select Data

Check the boxes on the left to select the data you want to import into the knowledge base. After selecting, click the "Import" button, and the data will be automatically imported into the AI assistant's knowledge base.

To browse more data entries on the same page, you can click "10 entries/page" in the lower right corner to expand the viewing range.



In the knowledge base, you can see the data presented as .md files, and just like regular data, you can configure tags and metadata for them.



Web Scraping Considerations

Ensure you have permission to scrape the target website's content

It is recommended to test with a small range of data before performing large-scale scraping

After scraping is complete, you can verify data quality through the [Search Testing feature](#)

Regularly update scraped data to maintain information timeliness

How to Create FAQ


How to Upload FAQ

Select Knowledge Base

Go to the left sidebar "AI Features" → "Knowledge Base", and click on the knowledge base where you want to create FAQ (if you haven't created one yet, please refer to [How to Create a Knowledge Base: Basic Settings](#)).

Add Q&A Content

Enter the knowledge base page, select the "FAQ" tab, and click the "+Create FAQ" blue button in the upper right corner to start adding Q&A content.



Create FAQ

The FAQ editor supports multiple answer formats:

Basic text formatting (bold, underline, italic, etc.)

Equations

Source code (HTML, XML)

Embedded links (videos, websites)

[illegible]

[illegible]

After creation, you can click preview to confirm the answer content:

建立 FAQ



問題

線上支付方式有哪些

答案

14px

只支援 LINE PAY、APPLE PAY，其餘一概不接受

取消

確認

預覽

僅能使用 LINE PAY、APPLE PAY，其他線上支付方式不接受

100%

After adding, click "**Confirm**" to complete the FAQ creation. The AI assistant will then be able to reference your provided answers when responding:



When customers ask questions related to online payment methods, the LLM will respond based on the FAQ content:



This way, you can restrict the AI assistant's response scope and prevent it from providing incorrect or non-compliant answers.

Other Use Cases

Product Operation Guidance Q&A

Using **MaiAgent platform usage** as an example, suppose you need to help new users understand how to get started with the MaiAgent system. You can add the following Q&A in the FAQ section:

Question: What should I do if I don't know how to use MaiAgent?

When the answer is too long, you can click preview to confirm the full content





Answer preview screen

The AI assistant will explain MaiAgent's available resources and getting started steps based on your provided answer, and provide the links in your answer for users to reference:




Knowledge Introduction


Using **LLM technology introduction** as an example, suppose you run a technology education platform and need to explain what Large Language Models (LLM) are to the general public. You can add the following Q&A in the FAQ section:

Question: What is the working principle of LLM?

You can enter video links within the answer



Click Insert/Edit Media



Link input illustration



Media insertion preview

The AI assistant will respond based on the answer you set, and provide the video reference link for users to access:



FAQ Management and Maintenance Tips

Content Timeliness Management

Event information updates: Remove or update related content immediately after promotions or limited-time offers end

Policy synchronization: When company policies (such as return/exchange conditions, terms of service) change, ensure FAQ content is updated accordingly

Product information maintenance: New product launches, specification changes, discontinued products, and other information should be reflected in related Q&A promptly

Data-Driven Optimization

Answer quality improvement: Optimize answer content for high-frequency questions by providing more detailed and practical information

Gap analysis: Discover common questions not covered by FAQ through customer feedback and customer service records

Quality Control Mechanisms

Content accuracy checks: Regularly verify that links, contact information, and technical data in FAQ are correct

Tone consistency: Ensure all FAQ content aligns with brand tone and communication style

Multi-language version synchronization: If multi-language support is needed, ensure all language versions are kept in sync

You can check the hit count of FAQ entries to identify the most frequently asked question categories and update them regularly. For hit count details, please see: [Search Testing](#)

問題	答案	標題	元資訊	命中次數	操作
如何設置Wi-Fi 密碼？	如何設置 Wi-Fi 密碼？	如何設置 Wi-Fi 密碼？		5	詳情 編輯 刪除
如何設置Wi-Fi 密碼？	如何設置 Wi-Fi 密碼？	如何設置 Wi-Fi 密碼？		5	詳情 編輯 刪除
如何設置Wi-Fi 密碼？	如何設置 Wi-Fi 密碼？	如何設置 Wi-Fi 密碼？		5	詳情 編輯 刪除
如何設置Wi-Fi 密碼？	如何設置 Wi-Fi 密碼？	如何設置 Wi-Fi 密碼？		5	詳情 編輯 刪除
如何設置Wi-Fi 密碼？	如何設置 Wi-Fi 密碼？	如何設置 Wi-Fi 密碼？		5	詳情 編輯 刪除

This management mechanism ensures the AI assistant always provides accurate, timely, and user-relevant responses, enhancing the overall customer service experience.

Document Management: Tags & Metadata

Since a Knowledge Base may contain a large number of documents or FAQs to manage, MaiAgent provides tag management and metadata management features, allowing you to categorize and organize large volumes of documents and FAQs.

Tag Management

When your Knowledge Base contains a large amount of data, the tagging system helps you quickly organize documents, FAQ attributes, or control reference material permissions. For example:

Product Data Knowledge Base (categorized by product information)

- └─ #Tent
- └─ #4-Person Tent
- └─ #3-Season Tent
- └─ #SnowPeak

Tutorial Content Knowledge Base (categorized by permission level)

- └─ #Non-Member
- └─ #Regular Member
- └─ #VIP

Add Tag

Click the Add button



Enter the tag name



After clicking Add, the tag you just entered will be displayed as follows, including ID, name, etc.

The ID can be used for granting document access permissions. For details, see [Technical Manual — Query Metadata Control Items](#)



Add Tags to Documents

After adding tags, go to the document page.

Click Edit File Information



Select tags

Add corresponding tags based on the document's requirements (you can select multiple tags).



The added tags will then be displayed next to the corresponding document:



Add Tags to FAQs

If you want to add hierarchical categorization or tag management for FAQs, you can:

Go to the FAQ management interface within the Knowledge Base, select the FAQ you want to edit, then click the Edit button



Select the tags to add, then click "**Confirm**"





The added tags will appear in the FAQ field.



Tag Access Control

You can control what different users see based on different tags. For example, VIP customers can view more advanced product tutorials, while non-members can only see general tutorials.

Filtering can be applied to both FAQs and documents.

AND Filter

The AND filter opens content that matches all tag conditions simultaneously.

For example: If a conversation is a tutorial-related conversation with a non-member, you can use tag filtering to only open documents that simultaneously match the **"Non-Member"** and **"Tutorial"** tags for the AI Assistant to use.

As shown: After filtering, the AI Assistant can only use the beginner camping document, avoiding disclosure of irrelevant information in non-member conversations.




OR Filter

For conversations with members, you can use the OR filter to allow documents matching either Regular Member or Non-Member to be available for the AI Assistant to reference:



Members can see more document content than non-members, achieving tiered access control.

You can switch between filtering mechanisms by clicking "AND" or "OR" in the tag filter.



For filtering results, refer to: [Internal Q&A Features](#)

Document Metadata Management

What Is Metadata?

Metadata is like a document's "ID card," recording detailed information about the document and helping the system better manage and use these documents.

Example: Beginner Camping.pdf

```
Metadata settings:
├─ Document version: v2.1
├─ Created date: 2024-03-15
├─ Last updated: 2024-11-20
├─ Next review: 2025-05-20
├─ Updated by department: Marketing
└─ Version status: Reviewed and approved for release
```

Why Use Metadata?

When data is complex and voluminous, using metadata helps you maintain data quality more conveniently and quickly by:

Improving Management Efficiency

Quickly identify document status and version information

Stay more organized when managing large volumes of documents in bulk
Reduce time spent on manual searching and verification

Ensuring Content Quality

- Use the latest document versions
- Track document update history and version status
- Ensure the AI Assistant uses the latest and reviewed content
- Set periodic review reminders to maintain data quality

Add Metadata to Documents

Go to the document page and click Edit File Information:



- Switch to the Metadata tab and enter the metadata mappings
- MaiAgent uses key-value pairs to store and retrieve metadata

請輸入文件名稱或 ID

全部狀態

初級露營.pdf



標籤

元資料

目前元資料

新增元資料

文件版本

v2.1

+ 新增元資料

元資料可以儲存文件的額外資訊，如作者、版本等。

保留字包含「url」，會覆蓋當預覽被引用文件時開啟的位置

取消

確定



After entering, click the "[Add Metadata](#)" button to add the key-value pair you just entered (you can add multiple metadata entries at once).



After adding, you can see the metadata you just added appearing in the field.

Add Metadata to FAQs

Go to the FAQ management interface within the Knowledge Base, select the FAQ you want to edit, then click the Edit button



Switch to the Metadata page, enter the values, then click Add Metadata

After entering all values, click "Confirm"

Custom URL Mapping

`url` is a reserved word in the metadata system. If you set `url: https:// ... your-specified-URL`, the URL you set will override the original file location when viewing citation snippets in conversations, opening your specified URL instead.

Example Scenario: Brand Spring Promotional Campaign

Original Settings:

```
title: "2025 Spring Sale Event Details"
content: "This event includes 20% off storewide, bonus gifts for minimum purchase amounts, and many more offers ..."
url: https://www.yourstore.com/spring-sale-2025
```

Usage Effect: When a user asks "What offers are available in the spring sale?" in a conversation:

The AI will cite this knowledge snippet to answer the question

When the user clicks the citation source, it will not open the internal file location

Instead, it will navigate directly to <https://www.yourstore.com/spring-sale-2025>, the official campaign page

The user can immediately see the latest event details, purchase buttons, countdown timers, and other dynamic content

Practical Scenarios:

Website Announcements: `url: https://www.company.com/announcements/system-maintenance`

Product Spec Pages: `url: https://www.product.com/specs/model-x1`

Customer Service Policies: `url: https://support.company.com/return-policy`

Event Registration: `url: https://events.company.com/register/webinar-2025`

The benefit of this setup is making AI responses more actionable -- not only providing information but also guiding users directly to relevant pages to take action, improving user experience and conversion rates.


Before Setting url

As shown below, before setting the url value, clicking a citation node and clicking any node (e.g., Beginner Camping.pdf) allows you to view the original document (you can configure whether to allow users to download original document files in the conversation platform):

Conversation Platform - Allow Citation Document Download Settings




View citation nodes



View original file location

After Setting url (Using Google as an Example)



Click the cited node file (Camping Precautions.txt) again. Instead of showing a document preview, it will open the Google homepage:



After setting the `url` field in metadata, when users see the AI Assistant citing that document in a conversation, they can click directly to navigate to the link you set for more complete information.

Move Knowledge Base Files

When you need to reorganize a knowledge base—for example, by department or product line, or to archive outdated content—use the **Move Files** feature to move files in batches to another knowledge base in the same organization. Moving files **does not reprocess or vectorize them**, and their file IDs remain unchanged. This makes it fast to move even large numbers of files without incurring additional processing fees.

You can move up to 1,000 files at a time. Process larger sets in multiple batches. Moves run in the background, so you can continue with other tasks while they are in progress.

1. Open the Move Window

There are three entry points. All of them open the same **Move Knowledge Base Files** window. The only difference is whether the system preselects the source knowledge base or files based on the entry point.

Entry Point 1: Upper Right of the Knowledge Base List

Click **Knowledge Bases** in the left menu to open the knowledge base list, then click **Move Files** in the upper-right corner. This entry point does not preselect a source and is suitable for planning moves among several knowledge bases at once.



Entry Point 1: The **Move Files** button in the upper-right corner of the knowledge base list

Entry Point 2: Move Icon in the Knowledge Base List

In the knowledge base list, click the move icon in the **Actions** column for a knowledge base. This entry point **preselects that knowledge base as the source**, making it suitable when you already know which knowledge base you want to organize.



Entry Point 2: The move icon in the **Actions** column of the knowledge base list

Entry Point 3: Select Files in a Knowledge Base's Document List

Open the **Documents** list on a knowledge base settings page. Select the files to move, then click **Move Files**. This entry point **locks the current knowledge base as the source** and preselects the files you just selected. After choosing the destination, you must still click **Move In** to add them to the pending move list. This is useful when selecting files to move while browsing documents.



Entry Point 3: On the knowledge base settings page, select files in the document list and click **Move Files**

2. Understand the Move Window

The **Move Knowledge Base Files** window uses a two-column layout. The left column contains files in the source knowledge base, and the right column contains files in the destination knowledge base. Use the buttons between them to move files from one column to the other:



Overview of the Move Knowledge Base Files window

Component	Location	Description
Source Knowledge Base drop-down menu	Upper-left area of the window	Select the knowledge base where the files are currently located
Destination Knowledge Base drop-down menu	Upper-right area of the window	Select the knowledge base to which the files will be moved. The destination must use the same vector settings as the source. When you submit the move, the system also checks that you have delete permission for both the source and destination
Swap button	Between the two drop-down menus	Quickly swap the source and destination without selecting them again
Search box	One in each column	Search by file name to quickly find files in long lists
Select all checkbox	Top of each column's file list	Select all files currently shown in the list at once (after a search, Select All selects only the search results)
Move In button (→)	Between the two columns	Add files selected in the left (source) column to the pending move list

Component	Location	Description
Move Back button (←)	Between the two columns	Add files selected in the right (destination) column to the reverse pending move list. Both directions can be processed in the same batch
Pending move indicator	File row	Files added to the pending move list are marked Pending Move In or Pending Move Back ; they have not actually been moved yet
Persistent notice	Bottom of the window	File IDs remain unchanged • No reprocessing or vectorization • 1,000-file limit per batch , reminding you of the feature's three key characteristics

The Edit permission shown in the window does not mean that move authorization has passed. When you submit, you must still have delete permission for both knowledge bases. Until you click **Confirm Move**, all Move In/Move Back actions only edit the pending move list. No files have actually moved, so you can safely adjust the list as often as needed.


3. Procedure

Select the Source and Destination Knowledge Bases

At the top of the window, select the **Source Knowledge Base** and **Destination Knowledge Base**. The destination drop-down excludes the source itself and filters knowledge bases by vector settings. Use the swap button between the menus to quickly exchange the source and destination. Seeing a knowledge base in the menu does not guarantee that the move can be completed; delete permissions for both the source and destination are checked when you submit.

Select Files and Add Them to the Pending Move List

Select files in the left (source) column, then click **Move In** (→) to add them to the pending move list. You can also select files in the right (destination) column and click **Move Back** (←). Moves in both directions can be submitted together in the same batch. For a long file list, filter it with the search box before using Select All.




Selected files have been added to the pending move list. Each file row is marked Pending Move In or Pending Move Back

A file's existing tags and Metadata are cleared after it is moved because tag definitions belong to the original knowledge base. To retain them, first create tags with the same names in the destination knowledge base, then apply the tags again after the move. When you add files to the pending move list, the system displays a warning that Tags for N files will be cleared.

Confirm the Move

The bottom of the window shows the number of pending moves. After confirming the count, click **Confirm Move**.



The pending move count appears at the bottom. Click **Confirm Move** to submit

After submission, the move runs in the background and the screen displays **Move in Progress**. You can work on other tasks in the meantime.



The move runs in the background without interrupting other tasks

View Move Results

When the move finishes, the **Move Results** page presents three summary cards: **Completed**, **Skipped** (already in the destination knowledge base), and **Failed** (the file remains in its original knowledge base). It also lists the reason for each file.



Move Results page: 1 completed, 0 skipped, and 0 failed. Click Move Another Batch to continue

If any files fail, click **Retry This Batch**. Retrying is safe: files that have already moved are skipped and are not processed twice.

To continue organizing other files, click **Move Another Batch** to return to the selection screen.

Move results are retained for 7 days. During that time, you can return to review them or retry a batch.

4. Limitations and Important Notes

Item	Description
Permissions	You need delete permission for both the source and destination knowledge bases. Edit permission alone is insufficient. Organization owners automatically have full permissions and require no additional configuration
Scope	Files can only be moved between knowledge bases in the same organization. Knowledge bases in other organizations do not appear in the drop-down menu
Vector settings	The destination knowledge base must use the same vector settings as the source. Knowledge bases with different settings do not appear in the drop-down menu
Files being processed	Files that are still being processed cannot be moved. Wait for processing to finish before moving them
Per-batch limit	1,000 files; process larger quantities in multiple batches
Files with the same name	The system does not deduplicate or rename files. If the destination knowledge base already contains a file with the same name, both files will exist after the move. Check before moving them
Tags and Metadata	Always cleared after a move and must be reapplied in the destination knowledge base (see the instructions below for how to preserve them)

Why Must the Destination Knowledge Base Use the Same Vector Settings?

Moving files can avoid reprocessing and vectorization because the system directly reuses each file's existing processed results and vector data, changing only its ownership. Vector data is generated from a knowledge base's vector settings. If the destination uses different settings, vectors from the source cannot be retrieved correctly in the destination. To prevent this, the **Destination Knowledge Base** drop-down only lists knowledge bases with the same vector settings as the source.

If you need to move files to a knowledge base with different vector settings, **upload the original files again in the destination knowledge base** so the system can process and vectorize them using its settings.

Why Are Tags Cleared, and How Can I Preserve Them?

Tag definitions—their names and purposes—belong to individual knowledge bases and are not shared across the organization. Tags in the source knowledge base do not exist in the destination, so a file's original tags and Metadata are always cleared after it is moved.

If tags are important to your retrieval or management workflows, follow these steps:

Before moving files, record or take a screenshot of their current tags.

Create tags with the same names in the destination knowledge base.

Move the files.

After the move, apply the tags to the files again in the destination knowledge base's document list.

5. Error Scenarios and Troubleshooting

Scenario	System behavior	Recommended action
Some files do not meet the requirements at submission time (for example, they are still being processed)	The entire batch is rejected. No files are moved, and the reason for each file is listed	Resolve or remove the problematic files based on the reasons shown—for example, remove files still being processed from the pending move list—then submit again
Files are being processed by a previous move batch	The system reports that they cannot be added and provides a Remove These N Files button	Click Remove These N Files to move the remaining files first, then process these files after the previous batch finishes
Some files fail during the move	Failed files remain in their original knowledge base , and the results page lists the reason for each file	Click Retry This Batch on the results page. Files already moved are automatically skipped, and only failed files are retried

The reject entire batch design is intentional. It prevents some files in a batch from moving while others fail, which would leave the knowledge base in a difficult-to-track state. Remove files that do not meet the requirements before submitting so the batch can complete successfully.

6. Use Cases

Financial Services: Split a Shared Knowledge Base by Department

A bank initially stores credit card, mortgage, and wealth management documents in a single **Customer Service Knowledge Base**. As the number of documents grows, its customer service AI assistant occasionally cites credit card promotions when answering mortgage questions. An administrator creates three knowledge bases—**Credit Cards**, **Mortgages**, and **Wealth Management**—using the same vector settings, then moves hundreds of documents to the corresponding knowledge bases by department. The files do not need to be reprocessed or vectorized, and no additional processing fees are incurred. After the split, each department's AI assistant only mounts its own knowledge base, producing more focused answers.

Manufacturing: Organize Catalogs and Specifications by Product Line

A sales inquiry assistant at an electronic components manufacturer originally uses one **Product Documents** knowledge base containing catalogs and specifications for connectors, cables, and power modules. During a reorganization, an administrator filters for connector-related files using the search box in the document list, selects them, and clicks **Move Files** (Entry Point 3) to move them into a new **Connector Product Line** knowledge base. The administrator then repeats the process for the other product lines. Because file IDs remain unchanged, internal records that reference those IDs are unaffected.

Education: Archive Content at the End of a Semester

A university's admissions and academic affairs AI assistant knowledge base contains many course announcements and calendars for the current semester. At the end of the semester, an administrator moves outdated documents to a **Past Semesters Archive** knowledge base. This keeps the active knowledge base streamlined and prevents the AI assistant from citing outdated information, while the archived content remains available to internal search assistants. Repeating the process each semester creates a consistent knowledge base lifecycle management workflow.

Healthcare: Reorganize Patient Education Materials by Department

A hospital initially stores its patient education materials in a single knowledge base, then decides to split them by department, such as family medicine, orthopedics, and obstetrics and gynecology, so that each department's LINE patient education assistant mounts the appropriate knowledge base. Because most patient education files have tags such as **Patient Education** and **Forms**, the administrator first creates tags with the same names in each department's knowledge base, then reapplies them after moving the files so tag filtering continues to work.

Retail and Hospitality: Rotate Seasonal Campaign Content

After a peak-season campaign ends, a hotel chain moves room type and campaign documents for its **Summer Promotion** from the online customer service knowledge base to a **Campaign Archive** knowledge base, while moving documents for the new campaign in. Because files can be moved in both directions in the same batch, the old and new campaign content can be swapped in one operation.

7. FAQ

Q: Are Files Reprocessed or Vectorized When Moved? Are There Any Fees?

No. A move directly reuses the existing processed results and vector data, changing only the file's ownership. It is therefore fast and does not incur additional processing fees. This is what the persistent

notice at the bottom of the window means: **File IDs remain unchanged** • **No reprocessing or vectorization** • **1,000-file limit per batch**.

Q: What If I Move a File to the Wrong Place?

Open the move window again, swap the source and destination using the button between them, and move the file back. Note that if the file originally had tags, they were cleared during the first move and are not automatically restored when you move it back. You must apply them again.

Q: Why Can't I Find the Destination Knowledge Base in the Drop-Down Menu?

The **Destination Knowledge Base** menu excludes the source itself and filters by vector settings. If a knowledge base is missing, check whether it belongs to **another organization** (cross-organization moves are not supported) or uses **different vector settings**. Even if you can select the destination, the system checks that you have **delete permission** for both the source and destination when you confirm the move. If you lack permission, contact the organization owner to adjust your role permissions.

Q: Will an Existing File with the Same Name in the Destination Be Overwritten?

No. The system does not overwrite, deduplicate, or automatically rename files. Both files with the same name will exist after the move. To avoid duplicates, check the destination's file list before moving or manually delete the unnecessary copy afterward.

Q: What If I Need to Move More Than 1,000 Files?

Process them in multiple batches. Use the search box and Select All in the left column to split files into batches by keyword. After each batch completes, click **Move Another Batch** to continue.

Q: How Long Are Move Results Available?

Move results are retained for 7 days. During that time, you can return to review the result for each file and click **Retry This Batch** for failed batches.

PDF Document Translation

Feature Overview

MaiAgent provides AI-powered PDF document translation, allowing you to translate PDF files in your knowledge base into a specified language. After translation, two versions are produced:

Translated version: Contains only the translated content in the target language

Bilingual side-by-side version: Displays the original text and translation side by side for easy comparison


Supported languages include: Traditional Chinese, Simplified Chinese, English, Japanese, Korean, Spanish, French, German, Portuguese, Russian, Arabic, Italian, Dutch, Polish, and Turkish.

Currently, only PDF format is supported for document translation, and the PDF must contain extractable text content (scanned/image-only PDFs are not supported).

How to Translate a PDF Document

1. Open the Document Details Page


In the **Documents** tab of your knowledge base, click the PDF file you want to translate to open the document details page.



PDF document details page with the "Translate PDF" button in the upper right corner

2. Open Translation Settings

Click the **Translate PDF** button in the upper right corner. If the file already has a translation record, the translation results dialog will appear first — click **Translate Again** to open the settings. For first-time translations, the translation settings dialog opens directly.



Translation settings dialog

3. Configure Translation Parameters

Field	Description
Source Language	The language of the original text. Defaults to "Auto Detect" — the system identifies it automatically
Target Language	The target language for translation (required)
Large Language Model	Select the LLM model to use for translation
Custom Terminology/Prompt	(Optional) Guide the translation style or specify terminology mappings, e.g., "Translate 'Hermione' as 'Miao-Li'"

The custom prompt supports up to 5,000 characters and is ideal for specifying proper nouns, brand names, or particular translation styles.

4. Start Translation


Click **Confirm** to start the translation task in the background. The status updates in real time during translation, and results are displayed automatically upon completion.

View Translation Results

After translation is complete, click the **Translate PDF** button to open the translation results dialog, which contains three tabs:

Result Preview

Preview the bilingual side-by-side PDF with the original text and translation displayed side by side.



Translation results — Bilingual side-by-side PDF preview


Download Files

Three file download options are available:

Original file: Download the untranslated original PDF

Translated version: A PDF containing only the target language translation

Bilingual side-by-side version: A PDF with the original text and translation displayed side by side



Download files — Three formats available

Processing Details


View complete statistics for the translation task:

Translation configuration: Source/target languages, LLM model used

Processing statistics: Token consumption, number of LLM calls, processing time breakdown

File information: File size and character count for both the original and translated files

Timestamps: Creation, start, and completion times along with the final status



Processing details — Token consumption and time statistics

Translation Version Management

A single PDF file can be translated multiple times (e.g., into different languages or using different models), and each translation is saved as a separate version.

Use the **Select Translation Version** dropdown at the top of the translation results dialog to switch between versions. The version format is: **Date | Source Language → Target Language | Model Used**.

To translate again, click the **Translate Again** button to open a new translation settings dialog.

Search Testing

Why Do You Need Search Testing?

Ensure Response Accuracy

The quality of AI assistant responses directly affects customer satisfaction

Incorrect or irrelevant answers damage brand trust

Avoid Information Gaps

Prevent important product details or safety instructions from being overlooked

Optimize the efficiency of knowledge base usage

Improve Q&A Consistency


Ensure the same question receives a consistent answer at different times

Avoid response instability caused by knowledge base updates

Through search testing and segment preview features, you can instantly review the quality of data segments used by the AI assistant to ensure your Q&A standards.


How to Get Started?

Go to the Knowledge Base page, select the knowledge base you want to view, and enter its settings



Click on Search Test to enter the page

Type the content you want to search in the dialog box: "What should I pay attention to when choosing a campsite?", then click the search button. The AI assistant will search for the most relevant documents based on your input.



View search results

The right side will display the segments that the AI assistant retrieves when answering. As shown in the figure, there are 12 segments (to configure the number of retrieved segments, refer to: [How to Create a Knowledge Base: Basic Settings](#))



The right side shows the retrieved content used to answer this question. Each segment displays:

Word count: The number of words covered by the segment

Hit count: The number of times the AI assistant referenced this segment when answering (only official Q&A sessions are counted; test Q&A sessions are not included in hit count calculations)

A higher hit count indicates that users ask more questions related to this segment.

Search test records past test Q&A sessions. You can quickly click on history records to retest, ensuring that the same topics maintain or improve data reference quality.

知識庫搜尋測試

輸入搜尋內容，測試知識庫的回傳效果

我要去露營，可以給我一些參考資料嗎

17 / 200


搜尋

搜尋歷史 清除歷史

我要去露營，可以給我一些參考資料嗎 露營要注意什麼呢？ 露營要注意啥

Check Knowledge Segment Quality

Through search testing, you can determine which segments the AI assistant uses during reasoning, and then check whether the content of these segments needs optimization, such as:



If you find many blank segments in the search results, and these blank segments all belong to the file "Beginner Camping.pdf", you can then check whether the content of "Beginner Camping.pdf" is correct and update it immediately **before going live**.

Search Test by Metadata

What Is Metadata Search?

The Metadata search feature allows you to perform precise queries based on document attribute information, rather than being limited to the document content itself. This feature is particularly suitable for:

- Document classification queries:** Search by document type, tags, categories, and other information
- Attribute filtering:** Filter by document creation time, author, source, and other attributes
- Precise targeting:** Quickly find document collections that meet specific criteria

How to Use Metadata Search?

On the search test page, you can:

- Select search mode:** Switch to "Metadata Search" mode
- Specify search criteria:** Enter the metadata field and value you want to search
- View search results:** The system will display document segments that match the metadata criteria

Search examples:

Search for documents in a specific category: `category:Product Description`

Search for documents within a specific time range: `date:2025-01`

Search for documents by a specific author: `author:Customer Service Department`

Advantages of Metadata Search

1. More Precise Results

Query directly against document attributes, reducing irrelevant results

Suitable for scenarios that require precise filtering

2. Faster Targeting

No need for full-text search; quickly locate target documents

Especially suitable for knowledge bases with large volumes of documents

3. Combined with Content Search

Can be used in combination with regular content search

First narrow down the scope with metadata, then pinpoint with content search

Practical Use Cases

Scenario 1: Product Document Management

Search for all documents in a specific product line

Example: `product_line:Camping Equipment`

Scenario 2: Time Range Query

Find recently updated documents

Example: `updated_date:2025-11`

Scenario 3: Document Type Filtering

View only specific types of documents (e.g., operation manuals, FAQs, etc.)

Example: `doc_type:FAQ`

Search Testing Helps You

Test Important Questions Regularly

Test core business-related questions frequently

Retest after knowledge base updates

Record and compare improvements in test results

Identify Content Gaps

Discover knowledge gaps: Test common questions and find queries with no matching results

Example: Testing "tent waterproof rating explanation" reveals no relevant segments → waterproof specification documents need to be added

Evaluate Content Completeness

Check answer depth: Confirm whether the AI answer covers all aspects of the question

Example: Query "camping safety precautions" → Check whether it covers weather, wildlife, fire safety, and other aspects

Question Testing Tips

You can:

Ask the same question in different ways

Test complex scenarios and multi-condition queries

Verify the retrieval effectiveness of technical terms and product names

Use these testing tips to comprehensively evaluate the accuracy and completeness of responses.

Database

Text to SQL Feature

What is Text to SQL?

Text to SQL (also known as Text2SQL) is an intelligent tool that automatically converts **natural language questions** (everyday human language) into **SQL database query statements**. Simply put, it enables the AI Assistant to "**understand human language**" and directly operate on databases.

To learn more about Text2SQL, see: [Technical Manual — Text to SQL](#)

Imagine you are a convenience store owner:

Traditional approach:

You: "Help me check how many bottles of cola were sold yesterday"

Employee: "Boss, you need to teach me how to use the POS system to query..."

You: "Click here, select there, enter conditions..."

Requires hands-on instruction, which is time-consuming.

With Text2SQL:

You: "Help me check how many bottles of cola were sold yesterday"

AI Employee: "Sure!" (*automatically generates the query method and calls the database*) → "87 bottles of cola were sold yesterday"

Core Features of Text to SQL

Natural language question → AI understanding & analysis → SQL query statement →
Execute query → Return results

When you ask the AI Assistant "Find the best-selling beverage from yesterday," the AI Assistant first analyzes your question, understands that you want to find a "beverage" that sold the "best" "yesterday," then generates SQL syntax to query the database, and finally tells you the result: "cola."

What Text to SQL Can Help AI Accomplish

Use Cases

Medical Clinic

```
Doctor asks: "Find today's appointments for diabetic patients"
AI responds: "There are 5 diabetic patients with appointments today:
- Mr. Wang 10:00 follow-up visit
- Mrs. Li 14:30 follow-up checkup
- Ms. Chen 16:00 blood sugar monitoring"
```

School Management

```
Teacher asks: "How many students failed this exam?"
AI responds: "This math exam:
- Total students: 45
- Failed: 8 (17.8%)
- The list of students requiring a makeup exam has been compiled"
```

Retail Chain Store

```
Store manager asks: "Compare the monthly revenue of three branches"
AI responds: "November revenue comparison by branch:
- Taipei store: $1,200,000 (↑15%)
- Taichung store: $950,000 (↑8%)
- Kaohsiung store: $800,000 (↓3%)"
```

Advantages of Text to SQL

Efficiency Improvement

Traditional approach: Requires an SQL expert to write queries → 30 minutes

Text2SQL: Ask in natural language → 3 seconds

Lower Barrier

No need to learn complex SQL syntax

No need to memorize database structure

Anyone can query data

Real-Time Interaction

Ask and get answers immediately

Supports follow-up questions
Dynamically adjust query conditions

How to Use Text to SQL in MaiAgent

Before starting, you need to prepare: an enterprise relational database or Excel files (uploaded to the Knowledge Base) If using a database, it is recommended that you: Ensure the data table structure is clear Create appropriate indexes Maintain data quality

1. Access AI Assistant Settings

Select the AI Assistant you want to configure
Switch to the Response Mode settings page





2. Switch the AI Assistant's Response Mode to Agent Mode

Make sure to switch to Agent Mode, otherwise the AI Assistant cannot use the Text to SQL feature. For detailed information about each Response Mode, refer to: [Create AI Assistant](#)



3. Enter the Enterprise Database URL

Use the dropdown menu to select the database service used by your enterprise.

Enter the database service URL used by your enterprise to allow the MaiAgent system to connect to the database.

- MaiAgent supports: MySQL PostgreSQL Oracle DB Microsoft SQL Server (MSSQL) - The maiagent option applies to Excel files you have already uploaded to the MaiAgent Knowledge Base



Make sure the URL format is correct and includes the necessary connection information, such as: hostname, port, database name, username, and password.

MaiAgent

Refer to: [Using MaiAgent Knowledge Base for Text to SQL](#)

Microsoft SQL Server(MSSQL)

Connect to an existing MSSQL database by pasting the MSSQL database connection string

Note: Ensure the database URL is accessible by the MaiAgent service



MySQL

Connect to an existing MySQL database by pasting the MySQL database connection string

Note: Ensure the database URL is accessible by the MaiAgent service



Oracle

Connect to an existing Oracle database by pasting the Oracle database connection string

Note: Ensure the database URL is accessible by the MaiAgent service



PostgreSQL

Connect to an existing PostgreSQL database by pasting the PostgreSQL database connection string

Note: Ensure the database URL is accessible by the MaiAgent service



4. Click Save to Store the Settings



With this setup, the AI Assistant can help you quickly query your inventory, employee information, and more, and compile well-organized reports and trends for you.

Troubleshooting

Connection failed: Check the database URL format and network connectivity

Query error: Verify that table names and field names are correct

Insufficient permissions: Check database user permission settings

Slow response: Check query complexity and consider adding indexes

Text2SQL turns the AI Assistant into a database expert, allowing anyone to quickly obtain business insights using natural language, greatly improving the efficiency of data-driven decision making!

Create & Configure a Database

MaiAgent's **Database Management** feature lets you centrally manage the database connections used by your AI Assistants on the platform. Combined with the Text to SQL feature, AI Assistants can query database data directly using natural language.

For a conceptual introduction to Text to SQL, see: [Text to SQL Feature](#)

MaiAgent supports two database types:

MaiAgent Database: Upload Excel

CSV

JSON files and the system automatically creates data tables. Ideal for users without their own database


External Database: Connect to an existing PostgreSQL, MySQL, MSSQL, or Oracle database

Create a Database

1. Go to the Database Management Page

Navigate to the "**AI Features**" section from the left sidebar and click "**Database**." The page has two tabs: **MaiAgent Database** and **External Database**.

Click the "**Add Database**" button in the upper right corner to begin.



Database management page showing the MaiAgent database list

If the button is grayed out and unclickable, your role has not been granted permission to create databases. Contact your organization administrator to set up permissions.

2. Select the Database Type

Choose the type of database you want to create by clicking the corresponding card:


MaiAgent: Suitable for uploading spreadsheet files; no connection string required

PostgreSQL

MySQL

MSSQL

Oracle: Connect to an existing external enterprise database




Step 1: Select the database type

3. Fill In Basic Information

MaiAgent Database

Field	Required	Description
Database Name	Yes	e.g., "2024 Annual Sales Data"
Description	No	Additional notes about the database's purpose or content

MaiAgent databases do not require a connection string — the system handles data storage automatically. After creation, upload your spreadsheet files.



Step 2: MaiAgent database configuration screen

External Database


Field	Required	Description
Database Name	Yes	e.g., "Enterprise ERP Database"
Description	No	Additional notes about the database's purpose or content
Connection String	Yes	The database connection URL
Specified Tables	No	Restrict which tables the AI Assistant can query, separated by commas

Connection string formats are as follows:

```

PostgreSQL: postgresql://username:password@host:port/database_name
MySQL:      mysql://username:password@host:port/database_name
MSSQL:      mssql+pymssql://username:password@host:port/database_name
Oracle:     oracle+cx_oracle://username:password@host:port/database_name
  
```


Make sure the connection string format is correct and includes all required connection information. The database URL must be accessible by the MaiAgent service.



Step 2: External database (PostgreSQL) configuration screen

After entering the connection string, click the "**Test Connection**" button to verify the connection:

Connection successful: MaiAgent can access your database normally

Connection failed: Check the connection string format, network connectivity, and whether the database allows external connections

If your database contains a large number of tables, enter the table names you want to expose in the "**Specified Tables**" field (e.g., `customers, orders, products`) to improve query efficiency and security.

4. Set Group Permissions

The final step is to set which groups can access this database. Use the transfer widget to move the groups that need authorization from the left panel to the right panel.

If you don't need to set group permissions right now, you can skip this step and adjust it later from the database settings page.

5. Complete Creation

After confirming all settings, click the "Create" button. Upon successful creation, you will be taken to the database settings page for further configuration.

After creation, don't forget to associate the database with an AI Assistant so the AI Assistant can use this database for queries. See: [Associate AI Assistants and Group Permissions](#)

Upload Files to Create Data Tables (MaiAgent Database)

After creating a MaiAgent database, go to the settings page, switch to the "Table Management" tab, and click the "Upload Table File" button.



Table Management tab for uploading files and creating data tables

Supported File Formats

Format	Extensions	Description
Excel	<code>.xlsx</code> , <code>.xls</code>	Each worksheet (Sheet) creates a separate data table
CSV	<code>.csv</code>	Comma-separated values file
JSON	<code>.json</code> , <code>.jsonl</code>	JSON format data file

After uploading, the system automatically parses the file structure and creates data tables in the background. Each data table displays a processing status:

Status	Description
Pending	File has been uploaded, waiting for system processing
Processing	System is parsing the file and creating the data table
Completed	Data table created successfully; the AI Assistant can query it
Failed	An error occurred during creation; hover over the error icon to see the reason

File Format Guidelines

- [x] The first row must contain **column names**
- [x] **Data types** within each column should be consistent (e.g., the price column should contain only numbers)
- [x] Avoid **merged cells**
- [x] Avoid **duplicate column names**
- [x] Each worksheet (Sheet) in an Excel file should contain only **one table**
- [x] It is recommended to annotate data formats after column names, e.g., `price(int)`, `date(datetime)`

If a data table fails to create, common causes include: incorrect or corrupted file format, special characters in column names, duplicate column names, or inconsistent data types within the same column.

Edit Column Descriptions

Column descriptions are a **critical setting** for improving Text to SQL query accuracy. The AI Assistant references column descriptions to understand the meaning of the data, enabling it to generate more precise SQL queries.

For example, if your data table has a column called `status`, the AI Assistant may not know what it represents. But if you add a description like "Order status, possible values: pending, completed, cancelled," the AI Assistant can correctly understand the column's meaning.

How to Edit

MaiAgent Database: In the "Table Management" tab, click "Edit" on a data table to edit the table description and column descriptions. The dialog shows completion progress (e.g., "3/10 columns have descriptions").

External Database: Under the "Basic Information" tab, in the data table list below, click to expand a table and enter descriptions directly in the column description field.

Writing Tips

Good column descriptions can significantly improve AI query accuracy:

Column Name	Poor Description	Better Description
<code>amt</code>	Amount	Total order amount (NTD), including tax and shipping
<code>cust_lvl</code>	Level	Customer level. A = VIP customers with annual spending over 1 million, B = annual spending 100K-1M, C = others
<code>reg</code>	Region	Customer geographic region. Possible values: North, Central, South, East, Offshore Islands

Manage Existing Databases

Edit a Database

In the database list, click "Edit" in the action menu to enter the settings page.



Database settings page: Basic Information tab

The settings page contains the following tabs:

Tab	Description	Applicable Type
Basic Information	Edit name, description, enabled status, connection string	All
Table Management	Upload files, manage data tables and column descriptions	MaiAgent only
Associated AI Assistants	Specify which AI Assistants can use this database	All
Role Permissions	Set group access permissions	All

Enable Disable

In the database list, use the "Enable" toggle to quickly switch the status. When disabled, AI Assistants temporarily cannot query this database, but settings and data are preserved.

Delete a Database

Deleting a database is an irreversible operation. MaiAgent Database: After deletion, all data tables created by uploading files will also be removed External Database: Only the connection configuration in MaiAgent is removed; the original database and its data are not affected

Text to SQL with MaiAgent Database

To have the AI Assistant help you analyze spreadsheet data such as Excel and CSV files, simply upload the files to the MaiAgent Database. The AI Assistant can then query the data using natural language and produce statistical analysis results.

For an introduction to the Text to SQL concept, see: [Text to SQL Feature](#)


Step 1: Create a MaiAgent Database

Navigate to "AI Features" > "Database" from the left sidebar

Click "Add Database"

Select "MaiAgent Database"

Enter a database name (e.g., "2024 Annual Sales Data") and click Create



Select MaiAgent Database and enter a name

For detailed creation steps, see: [Create and Configure Database](#)

Step 2: Upload Spreadsheet Files

After creation, go to the database settings page:

Switch to the "Table Management" tab

Click "Upload Table File"

Drag and drop or select your spreadsheet file



Upload files in the Table Management tab

You can download the sample file (a food sales file with 100 items in .xlsx format) to test and experience the AI Assistant's data analysis capabilities.

After the upload is complete, wait for the processing status to show "Success", which indicates the data table has been created successfully.

Supported formats: .csv, .xlsx, .xls, .json, .jsonl. Each worksheet in an Excel file will be created as a separate data table.

Table Format Notes

- [x] Each worksheet (Sheet) must contain only one table
- [x] The first row must be field names
- [x] It is recommended to annotate data formats after field names, such as: `Price(int)`, `Date(datetime)`
- [x] Data types within each field should be consistent
- [x] If there are special formats (e.g., currency, percentages), standardize the formatting before uploading

To improve AI query accuracy, it is recommended to add descriptions to data tables and fields. See: [Create and Configure Database — Edit Field Descriptions](#)

Step 3: Link AI Assistant

Link the database to your AI Assistant so it can query the data:

Switch to the "[Linked AI Assistants](#)" tab

Select the AI Assistant to link from the left side and move it to the right

Click "[Save](#)"



Link the database to the AI Assistant

Ensure the AI Assistant's Response Mode is set to Agent Mode, otherwise the Text to SQL feature cannot be used. For detailed Response Mode settings, see: [Create AI Assistant](#)

Step 4: Chat with the AI Assistant

Once the above steps are complete, you can go to the AI Assistant's conversation page and describe your needs directly using natural language.

For the AI Assistant conversation interface, see: [Choose an Integration Platform](#) to select the most suitable conversation platform for your needs.

Data Analysis

The Text to SQL tool can help you analyze data in spreadsheets. For example: "Find the best-selling dairy product from last year"



When you make a request, the AI Assistant will automatically invoke the Text-to-SQL tool.

This tool acts like a translator, converting your natural language into SQL syntax that the computer can understand, allowing the AI Assistant to extract data from the database.

The AI Assistant automatically adds conditional logic, sorting criteria, and other query constraints to convert your request into a more complete query command.

AI Assistant's Query Reasoning: Find the best-selling dairy product from last year, sorted by sales quantity in descending order. It automatically adds descending order conditions to better match your needs.

AI Assistant's Response (click the tool response to view):





Result

The AI Assistant found all sales data for "Dairy Products" along with other field data for those items. Each row represents a data record, including: Product ID, Product Name, Product Category, Sales Amount, Sales Quantity, etc.

Based on this data, the AI Assistant determined that the best-selling dairy product is "Cheese", with a sales quantity of 110 and sales amount of 40,150 TWD.

Report Creation (Canvas Mode)


The Text to SQL tool can also help you create polished reports. This feature requires enabling Canvas Mode:

Go to "AI Assistant Settings Page > Response Mode Settings" and find **Agent Mode**





Change Agent Mode from Standard Mode to **Canvas Mode**



After saving the settings, you can chat with the AI Assistant. For example: "Please generate a visualized revenue report with statistics"



AI Assistant's Query Reasoning: Calculate revenue statistics for each product in the food company, including total revenue, average revenue, revenue standard deviation, revenue distribution ranges, revenue share by category, and other statistical metrics. It automatically generates the calculation results needed for the report to better match your needs.

AI Assistant's Response (click the tool response to view):

```
[('Category Revenue Share - Staples', '10.72%'), ('Category Revenue Share - Dairy', '11.97%'), ('Category Revenue Share - Frozen', '9.06%'), ('Category Revenue Share - Canned', '10.25%'), ('Category Revenue Share - Meat', '13.23%'), ('Category Revenue Share - Fruits & Vegetables', '13.34%'), ('Category Revenue Share - Condiments', '4.51%'), ('Category Revenue Share - Snacks', '14.53%'), ('Category Revenue Share - Beverages', '12.40%'), ('Total Products', '100'), ('Average Revenue', '19546.06'), ('Minimum Revenue', '1400'), ('Maximum Revenue', '63364'), ('Revenue Std Dev', '15736.83'), ('Total Revenue', '1954606')]
```

Result

The AI Assistant received the automatically computed results for all required statistical items from the tool.

The AI Assistant used this data as the basis for visualization and created an interactive report interface using Canvas.


Click the "Use Canvas" box to view the content



The report display is shown below. You can:

Toggle between viewing the source code or the rendered view

Switch report content by statistical category (this varies depending on each AI Assistant's output)



When switching to other categories, you can view the category revenue share pie chart and distribution analysis bar chart:

營收統計視覺化報表

詳細的營收數據統計分析與視覺化展示

總覽統計

分類分析

分布分析

總營收

NT\$ 486,262

所有商品營收總和

📈 +12.5%



平均營收

NT\$ 40522

中位數: NT\$ 38,990



營收範圍

NT\$ 37,605

最高: 62,805 | 最低: 25,200



標準差

NT\$ 12679

營收變異程度



🕒 分類營收占比



分類統計詳情



🔗 Open Sandbox

營收統計視覺化報表

詳細的營收數據統計分析與視覺化展示

總覽統計

分類分析

分布分析

總營收

NT\$ 486,262

所有商品營收總和

📈 +12.5%



平均營收

NT\$ 40522

中位數: NT\$ 38,990



營收範圍

NT\$ 37,605

最高: 62,805 | 最低: 25,200




標準差

NT\$ 12679

營收變異程度



營收分布區間



營收散點分析



Open Sandbox

Additional Tips

Data quality matters: Ensure your spreadsheet data is clean and accurate, as this directly affects the accuracy of the AI Assistant's analysis results. Data you do not want the AI to access must be deleted: Hidden data in spreadsheets can still be parsed and provided to the AI Assistant. Therefore, data you do not want exposed cannot simply be hidden — it must be deleted.

Try different questions: By asking questions in different ways, you can uncover more valuable insights from your data.

Make good use of Canvas Mode: Canvas Mode allows you to easily create various visualized reports for a more intuitive presentation of data analysis results.

Make good use of field descriptions: Adding descriptions to data table fields can significantly improve AI query accuracy.

Link AI Assistants & Group Permissions

After creating a database and configuring the data tables, you need to **link the database to an AI Assistant** so the AI Assistant can use the Text to SQL feature to query data.

Link AI Assistants

Steps

Go to the database settings page and switch to the "**Link AI Assistants**" tab

The left side of the page shows available AI Assistants, and the right side shows linked AI Assistants

Select the AI Assistants you want to link from the left side and click the arrow button to move them to the right side

Click "**Save**" when finished



Link AI Assistants tab: available assistants on the left, linked assistants on the right

You can use the search field to quickly find a specific AI Assistant. When there are many AI Assistants, the list is displayed with pagination.

Make sure the AI Assistant's response mode is switched to Agent Mode, otherwise the Text to SQL feature cannot be used. For details on response modes, see: [Create AI Assistant](#)

Multi-Database Query

An AI Assistant can be linked to multiple databases. When a user asks a question, the AI Assistant automatically determines which database to query based on the question content:

```
User: "What were the sales last month?"  
→ AI Assistant automatically queries the "Sales Database"  
  
User: "How many employees are currently in the Marketing department?"  
→ AI Assistant automatically queries the "HR Database"
```

To help the AI Assistant accurately distinguish between different databases, give each database a name that clearly describes its purpose, and provide detailed descriptions and field descriptions.

Group Permission Management

Group permissions let you control database access scope for different teams in your organization, ensuring each team can only see databases relevant to them.

Configure Group Permissions

Go to the database settings page and switch to the "[Group Permissions](#)" tab

Use the transfer component to move the groups you want to authorize from the left side to the right side

Click "[Save](#)" when finished

Group permissions and "Link AI Assistants" are independent settings. Even if a database is linked to a specific AI Assistant, if the user's group does not have access to that database, the user still cannot query it.

Permission Levels

Database access is determined by the following conditions together:

Can a user query the database through an AI Assistant?

- ✓ Database is enabled
 - └─ ✓ Database is linked to the AI Assistant
 - └─ ✓ User's group has access to the database
 - └─ ✓ Query allowed

Database owner (creator) can access all databases they created, not restricted by group permissions

Non-owners need group permissions to access databases

Tools

Tools Feature Overview

What Are Tools?

Tools are like **plugins** or **extensions** for an AI assistant, enabling it to do more than just chat. For example, if an AI assistant has a **"Check Weather"** tool, it can tell you today's temperature; if it has a **"Play Music"** tool, it can play music directly for you.

By letting users define a set of "tools" that the AI assistant can use, the AI assistant is able to:

Understand complex user requests.

Automatically determine when a specific tool needs to be used.

Automatically generate the parameters required to call that tool.

This enables the AI assistant to go beyond generating text responses and actually perform a variety of tasks, such as:

Querying real-time information: Retrieve the latest stock prices, weather forecasts, flight statuses, etc. from databases or APIs.

Executing external operations: Call booking system APIs, control smart home devices, send emails or messages.

Processing files: Read, write, or analyze local or cloud documents.

Integrating with other software: Operate CRM systems, project management tools, or other enterprise applications.

How Tools Work

A basic tool invocation workflow includes the following steps:

Define tools:

Users must first define the relevant parameters for the tools.

Configure which tools the AI assistant can use.

Each tool must include:

A clear **Name**.

An easy-to-understand **Description** that explains the tool's purpose.

Detailed **Parameter specifications**, including each parameter's name, data type, whether it's required, etc.

User asks a question:

The user makes a request to the AI assistant in natural language.

Example: "Check tomorrow's weather in Taipei for me."

Model reasoning and tool selection:

The LLM inside the AI assistant analyzes the intent of the user's request.

The model determines whether a tool is needed and which tool from the available list to use.

Example: The model determines that weather information is needed and selects the tool named `get_weather`.

Generate tool call parameters:

The model generates a structured output (typically in JSON format) containing the tool name and its required parameters.

Example:

```
json { "name": "get_weather", "arguments": { "city": "Taipei", "date": "tomorrow" } }
```

5. Application executes the tool:

- * The AI assistant's backend application receives and parses the JSON instruction generated by the model.

- * The application executes the corresponding function or calls the external API based on the tool name and parameters in the instruction.

- * *Example:* The backend calls the weather query API, passing "Taipei" and "tomorrow" as parameters.

6. Return the result to the model:

The application returns the result of the tool execution (typically also in JSON format) back to the AI assistant's model.

Example:

```
json { "temperature": "25°C", "condition": "Sunny" }
```

7. Model generates the final response:

- * The model receives the tool execution result and incorporates it into the final natural language response.

- * *Example:* "Tomorrow's weather in Taipei is expected to be sunny with a temperature of about 25°C."

Key Advantages of Tools

Extend AI assistant capabilities: Break free from the limitation of only generating text, enabling the AI assistant to access real-time information and perform real-world tasks.

Improve reliability and accuracy: Through structured calls and responses, ensure task instructions are clear and precise, reducing the risk of model "hallucinations" or operational errors.

Enable complex automated workflows: Design AI assistants that can autonomously complete multi-step, cross-system tasks, dramatically improving efficiency (e.g., automatically planning a travel itinerary and booking flights and hotels).

More natural interaction experience: Users only need to describe their needs in natural language, and the AI assistant can understand and translate them into precise system operations.

Tool Types Supported by MaiAgent

The following major types are currently supported:

API Tools

The most commonly used type. Used to connect to and call external HTTP/HTTPS API services.

Common applications: Retrieving weather information, querying external databases, triggering webhooks, integrating with third-party services, etc.

Required configuration: API endpoint URL, HTTP method, request headers, parameter schema.

MCP Tools

Model Context Protocol (MCP) enables collaboration between servers, clients, and hosts through a standardized protocol.

Applicable scenarios: Enabling the AI assistant to call external tools for more complex and practical tasks.

Required configuration: MCP server URL, parameters, environment variables, etc.

Create an MCP Tool

What Is MCP?

MCP, or Model Context Protocol, is used to integrate multi-cloud platform services or execute local client applications.

Think of tool usage like electrical plugs — the plug must be connected to power before the service can be used. Traditionally, each LLM develops its own tool usage methods, like plugs of different standards that require multiple socket types. If you want both OpenAI and Claude to use Google Calendar services, developers need to build multiple scripts for different AIs.

When Google Calendar updates, developers may need to:

- Update 3 different integration versions simultaneously

- Maintain 3 separate sets of technical documentation

This leads to extended development timelines and excessive costs.

MCP Solution

MCP is a standardized protocol designed to solve this problem through:

- Standardized tool definition formats

- Unified communication protocols

- Consistent error handling mechanisms

These approaches simplify the development process.

With MCP, developing a tool requires only:

- Building an MCP Server (one-time development)

- Defining standardized tool specifications

- Implementing unified business logic

All AI platforms that support MCP can use it



Before and after MCP integration

MCP is fundamentally a protocol designed for **standardizing AI tool integration**, enabling development teams to "build once, use everywhere," significantly reducing the complexity and maintenance costs of the AI assistant tool ecosystem.

Quick MCP Tool Setup

1. Access the Tool Management Interface

From the left navigation panel, go to the "AI Features" section, then click "Tools." On the tool list page, click the "+ Add Tool" button in the upper right corner.



Click "+ Add Tool" to get started

2. Select the Tool Type

Select **MCP** as the tool type.

新增工具

* 工具類型

MCP

API

MCP

MCP 配置

* 指令或伺服器 URL

請輸入

參數 (以逗號分隔)

例：param1, param2

* 環境變數 (JSON) ⓘ

{ }

取消 確認

3. Set the Display Name

Set a clear display name for the tool. In this example, set it to **Composio mcp for google calendar**.

編輯工具

工具類型

MCP

顯示名稱 ①

Composio mcp for google calendar

MCP 配置

指令或伺服器 URL

https://mcp.composio.dev/composio/server/23b54607-75be-46d3-93de-67bc2e2a82d7/mcp?include_composio_help

參數 (以逗號分隔)

key: param1, param2

環境變數 (JSON) ①

取消 確認

Purpose: This name is displayed on the platform interface for all users to see.

Recommendation: Choose a name that clearly describes the tool's main function for easy user understanding. There are no strict formatting requirements for this name.

4. Fill In the MCP Configuration

a. MCP Server URL

Purpose: MaiAgent currently accepts external MCP servers. By providing the MCP server's service URL, the AI assistant can call the MCP service to connect to external applications.

Format:

Enter the full URL (e.g., https://mcp.dev/maiagent/mcp_service).

Note: This field is required.

For instructions on obtaining an MCP URL, refer to [Technical Manual — Remote MCP Service Overview](#)

Paste your MCP server URL here, and the system will automatically retrieve the list of connected tools from that server:

編輯工具

工具類型

MCP

顯示名稱 ⓘ

Composio mcp for google calendar

4 MCP 配置

指令或伺服器 URL

https://mcp.composio.dev/composio/server/23b54607-75be-46d3-93de-67bc2e2a82d7/mcp?include_composio_help

參數 (以逗號分隔)

例：param1, param2

環境變數 (JSON) ⓘ

{}

取消 確認

b. 🧩 MCP Command Arguments (mcp_args)

If you don't need to set specific environment variables, this field can be left empty.

Purpose: Define parameter names that need to be passed when executing MCP commands or calling MCP services. The content is automatically generated by the AI assistant.

Format: Use JSON array format, where each element is a string representing a parameter.

Example (JSON array):

```
[  
    "--user",  
    "admin",  
    "--config",  
    "/path/to/config.yaml"  
]
```

* **At runtime:** The AI assistant passes these parameters to the MCP tool in order

* If you enter a comma-separated string (e.g., arg1,arg2,arg3), the system will attempt to parse it as a parameter list. To avoid ambiguity, use JSON arrays.

編輯工具

工具類型

MCP

顯示名稱 ①

Composio mcp for google calendar

MCP 配置

指令或伺服器 URL

5 https://mcp.composio.dev/composio/server/23b54607-75be-46d3-93de-67bc2e2a82d7/mcp?include_composio_help

參數 (以逗號分隔)

例: param1,param2

環境變數 (JSON) ①

{}

取消 確認

Configure MCP command arguments

c. 🌿 MCP Environment Variables (mcp_env)


If you don't need to set specific environment variables, this field can be left empty.

Purpose: Set required environment variables for the MCP command execution environment.

Format: Must be a valid JSON object where keys are environment variable names and values are environment variable contents (strings).

Example:

```
json { "API_KEY": "{{SECRET_MCP_API_KEY}}", "REGION": "us-west-1", "DEBUG_MODE": "true" }
```




Configure MCP environment variables

5. Find "Allowed Tools (JSON Array)" and Click Refresh

Purpose: Specify which sub-tools the AI assistant is authorized to use under this MCP client. A single MCP client may provide multiple different features or sub-tools.

Auto-detect/Leave empty: If this field is left empty or not provided, the system will attempt to auto-detect all available sub-tools when first connecting to the MCP client, and allow all detected sub-tools by default. If you want to restrict the AI assistant to only specific sub-tools, explicitly list them here.

After clicking, the system automatically retrieves the tool content connected to that server and displays it in the list:



編輯工具



允許的工具 (JSON 數組) ⓘ

重新取得

GMAIL_CREATE_EMAIL_DRAFT



GMAIL_DELETE_DRAFT



GMAIL_DELETE_MESSAGE



GMAIL_FETCH_EMAILS



GMAIL_FETCH_MESSAGE_BY_MESSAGE_ID




GMAIL_GET_ATTACHMENT



取消


確認



MCP-defined tool names

6. 📁 Save the Tool

After confirming all settings are correct, scroll to the bottom of the page and click the "**Confirm**" button. Your new tool is now created!



⚠ Important Reminders

Connection Testing

After creating the tool, test the MCP connection to ensure it works properly
Verify tool functionality in a test environment

Permission Management

Carefully select allowed tools to avoid authorizing unnecessary features
Regularly check tool usage status

Troubleshooting

If the connection fails, check whether the MCP server URL is correct
Confirm that environment variables and parameter formats meet the requirements

Customize Tool Message Display

When an AI assistant calls a tool, you can control how much execution information appears in the conversation and format the tool results for easier reading. These settings apply to API, MCP, and Function tools.

Members with permission to edit the tool can modify the display copy, display mode, and output rendering template. Members without edit permission can only view these settings.

Use Case

A customer service manager creates an API tool for looking up trial requests. She wants customer service agents to see the lookup progress without expanding all the technical details. On the tool's **Display Copy** tab, she selects **Compact** and configures the text displayed when execution finishes. When the assistant later calls the tool, customer service agents see a brief execution status instead of the tool's complete input and output.


Configure the Display

Open the Tool Editor

From the left menu, go to **AI Features** → **Tools**. Find the tool you want to configure in the tool list, then click the edit icon on the right.

Open the Display Copy Tab

In the **Edit Tool** window, select **Display Copy**.



Display copy and display mode settings for a tool

Select a Display Mode

Select a mode based on how much information conversation users need to see:

Mode	Effect in the conversation
Follow conversation platform settings	Uses the conversation platform's default display behavior. The current default is Full .
Full	Displays complete tool execution information.
Compact	Displays only concise execution information.
Hidden	Does not display the tool execution block.

Hidden only hides the tool execution block in the conversation; it does not disable the tool. To prevent the AI assistant from calling the tool, go to the AI assistant's Tools settings and remove the tool from its configuration.

Customize the Display Copy

In the **Display Copy** section, click **Add** and configure the text displayed while the tool is running and after it finishes. The copy can use `{{variable}}` to insert a top-level parameter from the tool input, such as `{{query}}`. If the variable cannot be found, an empty string is displayed.

Configure the Output Rendering Template

To format the content returned by the tool, click **Add** in the **Output Rendering Template** section. The template can use `{{output}}` to insert the complete tool result or `{{variable}}` to insert a top-level parameter from the tool input. If the variable cannot be found, an empty string is displayed. If the template is left blank, the raw JSON is displayed.

When finished, click **Confirm** to save.

Recommendations for Selecting a Mode

Select **Full** when users need to verify the tool results.

Select **Compact** and provide clear display copy when users only need to know that the system is processing the request.

Select **Hidden** when execution details should not appear in the conversation.

After completing the setup, assign the tool to a test AI assistant and send a test question. Confirm that the display copy and output template match the fields actually returned by the tool.

Contacts MCP Credential Setup

I. Feature Background



What Is the MCP Credential Feature?

The MCP credential feature is like "setting up a dedicated pass for each customer." When multiple customers need to use the same tool to query their own data, each person needs their own "key" (credential) to ensure they can only see their own data and not anyone else's.

Simple Analogy

Think of an apartment building — all residents use the same elevator (MCP tool), but each person must swipe their own access card (credential) to enter their own floor.

Why Is This Feature Needed?

Protect Customer Privacy — Each customer uses their own credentials, ensuring data is not mixed together

Individualized Service — Different customers can have different usage permissions

Easy Management — When a customer's credentials expire or need replacement, you only need to update that one customer's settings

Use Case Examples

Scenario 1: Personal Data Queries

Customer A and Customer B both need to query their own purchase records, but each has different account credentials

Scenario 2: Department Permission Management

Different departments use the same query tool but have different viewing permissions

Scenario 3: Tiered Service


VIP customers can query more detailed data, while regular customers can only query basic information

II. Feature Location & Preparation

Where to Find This Feature

Navigation Path

Contact Management → Click Edit Contact → Switch to the "MCP Credentials" tab



Important Reminder

You must first create and save the contact before you can configure MCP credentials (to ensure credentials are correctly bound to the contact)

Preparation Before Setup

Before starting the configuration, confirm the following:

MCP tools have already been created in the system (AI Features → Tools → MCP)



- 2. The contact to receive credentials has been created and saved
- 3. You have obtained the contact's credential information (usually provided by the customer or authorized personnel)

Permission Note

System administrators can flexibly assign "Tools" and "Contacts" feature permissions based on organizational needs.

III. Operating Instructions

3.1 View Available Tools

Steps

- Open the contact you want to edit
- Click the "MCP Credentials" tab at the top
- You will see a list of all configurable tools

Tool List Description

Each tool displays:

Tool Name (e.g., "Customer Database Query")

Tool Description (if available, shown in gray text below the name)

Edit Button (pencil icon on the right)

How to Identify Configuration Status

Already configured → A green "Configured" label appears to the right of the tool name

Not yet configured → No label appears to the right of the tool name

3.2 Add or Modify Credentials

Step 1: Open the Settings Area

Click the "Edit" button (pencil icon) to the right of the tool to expand the settings area

Step 2: Fill In Headers (Credential Information)

Paste the credential information in the text box using JSON format:

```
{"Authorization": "Bearer your-key-code-here"}
```

If you're unsure about the format, click the "**Format JSON**" button to auto-format

What Are Headers?

Headers are like "identification documents" that contain the contact's dedicated key. Just like you need to bring your ID when going to the bank, the system needs these Headers to confirm "who this person is."

What Is JSON?

JSON is a data format that looks like: {"field_name": "content"}. Don't worry — usually someone will provide this to you directly, and you just need to copy and paste it.

Step 3: Important Notes

Headers are **required** (marked with a red *)

The entered Headers will **replace** the tool's default settings

Important: Different MCP tools require different Headers formats. Make sure the credential format you received matches the tool's requirements

Step 4: Save or Cancel

Save Settings

After confirming everything is correct, click the **"Save"** button in the bottom right corner

The system will validate the format

After saving successfully:

The editing area will automatically collapse

A **green "Configured" label** will appear to the right of the tool name (only appears for the first time)

Cancel Editing

If you don't want to save, click "Cancel"

3.3 Reset Credentials (Clear Configured Credentials)

When Would You Use This?

When credentials are configured incorrectly, or when the credentials are no longer needed, you can reset them.

How to Operate

Click the tool's "Edit" button to expand the editing area

If the tool already has credentials configured (has a green "Configured" label), a red **"Reset Credentials"** button will appear in the bottom left corner

Click to confirm, and the credentials will be cleared

After clearing, the green "Configured" label to the right of the tool name will disappear

Important Reminder

Reset **cannot be undone**

It is recommended to copy and back up the Headers content before resetting

If you only need to modify the content, edit it directly — there's no need to reset

IV. Operation Examples

Example 1: Set Up Database Query Credentials for a Customer

Scenario

Customer A needs to query their own purchase records, and you have received a set of credential codes.

Headers Content

```
{
  "apikey": "eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9 ... ",
  "Authorization": "Bearer eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9 ... "
}
```

Steps

Find contact "Customer A" and click Edit

Switch to the "MCP Credentials" tab

Find the "Customer Database Query" tool and click the Edit button

Paste the credential codes

Click "Format JSON" to ensure proper formatting

Click Save

Example 2: Set Up Dedicated Query Credentials for a Department

Scenario

The sales department needs to access customer data but can only view basic information, not financial data.

Headers Content

```
{
  "X-API-Key": "dept_sales_abc123",
  "X-Department": "sales"
}
```

Field Descriptions

X-API-Key — The sales department's dedicated key

X-Department — Identifies this is being used by the sales department

Example 3: Set Up Credentials with Multiple Fields

Headers Content

```
{
  "Authorization": "Bearer token_xyz",
  "X-Client-ID": "client_12345",
  "Content-Type": "application/json"
}
```

Field Descriptions

- Authorization — The main authentication key
- X-Client-ID — Client ID number
- Content-Type — Data format specification (usually fixed to this value)

V. Contact MCP vs. Tool MCP Differences

Feature Comparison Table

Comparison Item	Tool → MCP	Contact → MCP Credentials
What Is It	Creates the tool itself	Sets up dedicated credentials for individual contacts
Scope	Available system-wide	Only for a specific contact
Configuration	• Tool URL • Default credentials • Available functions	• Only this person's credentials • Overrides default credentials
When to Use	When connecting a new tool for the first time	When each customer needs different credentials
Who Can Operate	Anyone with permissions	Anyone with permissions
Where to Configure	Tool management page	Contact editing page
Common Uses	• Add a new query tool • Set up tool basics	• Each customer uses their own credentials • VIP customers have special permissions

Simple Analogy: Computer Lab

Imagine a school computer lab:

Tool MCP (Set up by authorized personnel)

Install software in the computer lab
Configure all computers to use the software
Determine which features the software has

Contact MCP Credentials

Create individual accounts for each student
Each student logs in with their own account
Different students can have different usage permissions

Practical Workflow

```
Step 1: Create MCP Tool (by personnel with tool creation permissions)
  ↓
Create "Customer Database Query" tool
- Set tool URL
- Set default credentials (if applicable)
  ↓
Step 2: Set up dedicated credentials for each contact (by personnel with
credential configuration permissions)
  ↓
Customer A: Configure A's dedicated credentials
Customer B: Configure B's dedicated credentials
Customer C: Configure C's dedicated credentials
  ↓
Result: Everyone uses the same tool but with their own credentials
```

Permission Flexibility

System administrators can determine who can create tools and who can configure credentials based on organizational needs. These can be the same people or managed separately.

Key Takeaways

Tool MCP — You need the "tool" first before it can be used

Contact MCP Credentials — Then give each person their "key"

VI. FAQ

Q1: Why can't I see the "MCP Credentials" tab?

Possible Reasons

You are on the "Add Contact" page (you need to save the contact first, then re-enter edit mode to see it)

No MCP tools have been created in the system yet

The contact data hasn't been saved yet

Solutions

Fill in and save the contact's basic information first

Confirm that MCP tools exist in the system (check with authorized personnel)

Re-enter edit mode

Q2: What should I do if the Headers format is incorrect?

Common Error Examples

✗ Wrong: `{'key': 'value'}` (used single quotes)

✗ Wrong: `{key: value}` (no quotes)

✗ Wrong: `{"key": "value",}` (trailing comma)

✓ Correct: `{"key": "value"}`

Solutions

Click the "Format JSON" button — the system will auto-check

If there's still an issue, copy the credentials again

Or ask the credential provider to re-confirm the format

Q3: How do I know what to fill in?

How to Obtain Credentials

Ask the customer to provide their credential information

Ask the responsible personnel to provide or help create them

Ask the colleague responsible for the tool

Refer to other successfully configured contacts (but don't copy directly!)

⚠ Don't fill in random values

Credentials have specific formats and content. Entering incorrect values will prevent the tool from working and may cause data errors.

Q4: Why don't I see the "Configured" label?

Display Conditions

Only tools that have already been configured with credentials show the green label

When configuring for the first time, the label won't appear until after saving successfully

The label appears only after a successful save

No labels on any tools in the list?

This means the contact hasn't had any MCP credentials configured yet — this is normal.

Q5: Can't find the "Reset Credentials" button?

Display Conditions (all must be met)

The tool already has credentials configured (green "Configured" label appears to the right of the tool name)

You have clicked the "Edit" button to expand the editing area

Button Location

After expanding the editing area, the red "Reset Credentials" button is in the bottom left corner; the "Cancel" and "Save" buttons are in the bottom right corner.

Tool doesn't have credentials configured yet?

The "Reset Credentials" button won't appear because there are no credentials to reset. The bottom left corner will be empty.

Q6: What happens when I click "Reset Credentials"?

What Happens

The contact's dedicated credentials will be cleared

The green "Configured" label to the right of the tool name will disappear

If the tool has default credentials, those will be used instead (but usually there are none)

 **Deletion cannot be undone**

Recommended Approach

Copy and back up the Headers content first

Make sure you really want to delete

If you only need to modify the content, use "Edit" instead

Q7: Can multiple contacts use the same set of credentials?

Technically, yes

The system won't prevent you from copying the same credentials to different people

But it's not recommended

You won't know who is using it

Data may get mixed together

When something goes wrong, you won't know whose issue it is

Correct Approach

Each contact uses their own credentials

Keep data clearly separated

Easier to track and manage

Q8: The tool list is empty, showing "No available tools"?

Possible Reasons

No MCP tools have been created in the system yet

All tools are "global" type (global tools don't require individual credential configuration)

Solutions

Ask personnel with tool creation permissions to confirm whether MCP tools have been created

Confirm that the tool type settings are correct

If everything looks normal but you still can't see any tools, contact technical support

Appendix: Operation Checklist

Pre-Setup Confirmation

- ☐ The contact has been created and saved
- ☐ MCP tools are available in the system
- ☐ Credential information has been obtained (from the customer or responsible personnel)
- ☐ Credential format has been verified as correct

During Setup Confirmation

- ☐ Headers format is correct (use the Format button to check)
- ☐ All required fields are filled in
- ☐ No extraneous or incorrect content
- ☐ No error messages appear after formatting

Post-Setup Confirmation

- ☐ Green "Configured" label appears to the right of the tool name
- ☐ Tool has been tested and works correctly
- ☐ Setup information has been recorded (date, contact, tool name)
- ☐ Relevant personnel have been notified of the completed setup, if needed

Create an API Tool

You can use MaiAgent's [Tool Builder AI Assistant](#) to help you create API tools

API tools are used to integrate external services and automate operational workflows.


What Is an API?

API (Application Programming Interface) is a bridge for communication between different software systems. Simply put, it's like a **"waiter" in the software world**, helping different programs exchange information and execute functions.

Imagine you're dining at a restaurant:

- You:** The customer who needs food (the application)
- Kitchen:** The place that prepares food (the service provider system)
- Waiter:** The person who relays messages between you and the kitchen (**API**)

You don't need to go into the kitchen yourself — just tell the waiter what you want. The waiter conveys your request to the kitchen, then brings the finished meal to you.



API workflow diagram

API tools can help you automate standardized workflows, set specific response formats, and retrieve information from your systems such as:

E-Commerce Customer Service Automation

Customer asks about order → API queries order status → Automatically replies with delivery progress

Marketing Campaign Management

```
New product listed → Automatically updates website → Sends EDM → Promotes on social media
```

Online Course Platform:

```
Student asks about course progress → API queries learning records →  
Automatically replies with completion percentage and next class time
```

Through API tools, the AI assistant evolves from a simple chatbot into an intelligent assistant capable of executing actual business workflows, significantly improving work efficiency and automation.

Quick API Tool Setup

1. Access the Tool Management Interface


From the left navigation panel, go to the "AI Features" section, then click " Tools." On the tool list page, click the " Add Tool" button in the upper right corner.



Click "+ Add Tool" to get started

2. Select the Tool Type

Select API as the tool type.



The screenshot shows a 'Tool Configuration' dialog box with the following fields and options:

- Tool Type (工具類型):** A dropdown menu with 'API' selected. A blue circle with the number '2' highlights this field.
- Tool Name (工具名稱):** A text input field containing 'google_calendar_retriever'.
- Description (描述):** A large text area with the placeholder text '請輸入'.
- API Configuration (API 配置):**
 - API URL (API URL):** A text input field containing 'https://api.example.com/endpoint'.
 - HTTP Method (HTTP 方法):** A dropdown menu with 'GET' selected.

At the bottom right, there are two buttons: '取消' (Cancel) and '確認' (Confirm).

3. Set the Display Name

Set a clear display name for the tool. In this example, set it to **google calendar**.

編輯工具

+ 工具類型

API

+ 顯示名稱 ⓘ

google calendar

+ 工具名稱 ⓘ

google_calendar_retriever

+ 描述 ⓘ

請輸入

API 配置

+ API URI ⓘ

https://api.example.com/endpoint

HTTP 方法

GET

取消 確認

Purpose: This name is displayed on the platform interface for all users to see.

Recommendation: Choose a name that clearly describes the tool's main function for easy user understanding. There are no strict formatting requirements for this name.

4. Set the Tool Name

Next is the "**Tool Name**" field.

Purpose: This name is the unique identifier used by the AI assistant internally to call and identify this tool.

Naming Rules (Important):

Must be in English.

Can only contain:

Lowercase letters (a-z)

Uppercase letters (A-Z)

Numbers (0-9)

Underscores (`_`)

Hyphens (`-`)

Examples: `get_weather_forecast` , `database-query-tool`

The example below is set to `google_calendar_retriever`

編輯工具

+ 工具類型

API

+ 顯示名稱 ⓘ

google calendar

+ 工具名稱 ⓘ

google_calendar_retriever

+ 描述 ⓘ

請輸入

API 配置

+ API URI ⓘ

https://api.example.com/endpoint

HTTP 方法

GET

取消 確認

API tool name definition

5. Write the Tool Description

In the "**Tool Description**" field, provide a clear and detailed tool description.

Importance: A good description helps the AI assistant more accurately understand:

The tool's function and purpose.

When to use this tool.

How to interpret the tool's output results.

Suggested Content: Describe what the tool does, what it takes as input, what it outputs, and any usage notes.

編輯工具

工具類型

API

顯示名稱 ⓘ

google_calendar

工具名稱 ⓘ

google_calendar_retriever

描述 ⓘ

當用戶詢問具體的行程安排、時間可用性或需要管理日曆活動時，應調用 Google Calendar 工具。包括查詢特定日期的會議安排（「我明天有什麼會議？」）、檢查時間衝突（「週三下午 2 點有空嗎？」）、創建或修改行程（「幫我安排下週一的會議」）等。如果用戶只是討論時間管理概念或詢問一般性時間問題，則不需要調用此工具。主要是區分用戶是否需要存取真實的個人日曆數據。

API 配置

API URL ⓘ

https://api.example.com/endpoint

HTTP 方法

請選擇

取消

確認

Tool description

6. Detailed API Configuration

a. API URL

Enter the full URL of the target API endpoint (including `http://` or `https://`).

Example: `https://api.opencalendar.org/data/2.5`

編輯工具

6

API URL ①

https://api.opencalendar.org/data/2.5

HTTP 方法

GET

標題 (JSON) ①

{ \"Content-type\": \"application/json\" }

參數結構 (JSON Schema) ①

{ \"type\": \"object\", \"properties\": { } }

取消 確認

b. HTTP Method


Select the HTTP verb required by the API service from the dropdown:

GET : Typically used for retrieving resources.

POST : Typically used for creating new resources or submitting data.

PUT : Typically used for fully replacing or updating resources.

DELETE : Typically used for deleting resources.



c. 📄 Headers

Headers are like the "**envelope**" of a letter. Before seeing the actual data content, they tell the receiver some important information. **Without correct headers, API requests may fail authentication, or the receiver may not be able to parse the data correctly.**

Common Uses:

Authentication (**Authorization** , **X-API-Key**)

Specify content type (**Content-Type**)

Specify accepted response format (**Accept**)

To add headers:

Click "**+ Add Header**" to define HTTP headers sent with the request.

Format: Must be a valid JSON object where keys are header names and values are header content (strings).

Example:

```
{
  "Content-Type": "application/json; charset=utf-8",
  "Authorization": "Bearer {{SECRET_API_TOKEN}}",
  "Accept": "application/vnd.github.v3+json"
}
```

編輯工具

+ API URL ① HTTP 方法

8

+ 標頭 (JSON) ①

```
{
  "Content-Type": "application/json; charset=utf-8",
  "Authorization": "Bearer {{SECRET_API_TOKEN}}",
  "Accept": "application/vnd.github.v3+json"
}
```

+ 參數結構 (JSON Schema) ①

```
{
  "type": "object",
  "properties": {}
}
```

取消 確認

Configure required HTTP request headers

d. 🔍 Query Parameters

Query parameters are query strings appended after the question mark in a URL (for example, `?format=json&limit=10`) and are sent with every request. They are suitable for fixed query conditions or API keys that can only be passed through the query string.

Click "[Available Context Variables](#)" to expand the list of currently supported variables (contact ID, conversation ID, inbox ID, and others).

Format: A valid JSON object where each key is a parameter name and each value is the parameter content (a string).

Values can be fixed text, or they can use `{{parameter_name}}` to reference a parameter declared in the Parameters Schema and supplied by the AI assistant when the tool is called. This lets you configure query conditions where the first part is fixed and the remaining value is provided by the user. See [Reference AI assistant parameters in Query Parameters](#) below.

Example:

```
{
  "$format": "json",
  "$filter": "ProductName eq '{{product_name}}'"
}
```



The Query Parameters field and the Available Context Variables list

e. Parameters Schema

The parameters schema is like an "order form" — it tells the AI assistant what data it can request from the API and how to make the request.

Core Configuration: Define which parameters the AI assistant can or must provide when calling this tool, as well as the format of those parameters.

Format: Use standard **JSON Schema** format.

Key Elements:

`type: "object"` : Indicates the parameters are an object.

`properties` : Defines the object for each parameter.

Parameter name (e.g., `"search"`): The corresponding object contains that parameter's details.

`type` : The parameter's data type (`string`, `integer`, `number`, `boolean`, `array`, `object`).

`description` : A description for the AI assistant explaining the parameter's meaning.

`default` (optional): The parameter's default value.

`enum` (optional): If the parameter value can only be one of specific options, list them here.

`required` : An array containing the names of all **required** parameters.

Example (video search tool):

```
{
  "type": "object",
  "properties": {
    "limit": {
      "type": "integer",
      "minimum": 1,
      "description": "Maximum number of results to return"
    },
    "fields": {
      "type": "string",
      "description": "Comma-separated list of fields"
    },
    "search": {
      "type": "string",
      "description": "Search keyword"
    }
  },
  "required": ["search"]
}
```

編輯工具

×

× API URL ①

https://api.opencalendar.org/data/2.5

HTTP 方法

POST

× 標頭 (JSON) ①

```
{
  "Content-Type": "application/json; charset=utf-8",
  "Authorization": "Bearer {[SECRET_API_TOKEN]}",
  "Accept": "application/vnd.github.v3+json"
}
```

9

× 參數結構 (JSON Schema) ①

```
{
  "description": "回傳指定日期工作表",
  "fields": {
    "type": "string",
    "description": "以逗號分隔的欄位清單"
  },
  "search": {

```

取消 確認

Use JSON Schema to precisely define API parameters

7. 📁 Save the Tool

After confirming all settings are correct, scroll to the bottom of the page and click the "Confirm" button. Your new tool is now created!

編輯工具

✖

✖

✖ API URL ①

https://api.opencalendar.org/data/2.5

HTTP 方法

POST

✖ 標頭 (JSON) ①

```
{
  "Content-Type": "application/json; charset=utf-8",
  "Authorization": "Bearer ([SECRET_API_TOKEN])",
  "Accept": "application/vnd.github.v3+json"
}
```

✖ 參數結構 (JSON Schema) ①

description: 回傳尚未啟用的表

```
{
  "fields": {
    "type": "string",
    "description": "以逗號分隔的欄位清單"
  },
  "search": {
```

10

取消

確認

Reference AI assistant parameters in Query Parameters

Many third-party APIs use query conditions made up of a single string containing fixed syntax and a user-provided value. Examples include the OData-style `$filter=ProductName eq 'Chai'` and the `q=` parameter used by search APIs. These APIs do not expect a separate parameter; instead, they require the value to be embedded in a fixed string.

Enter `{{parameter_name}}` in a Query Parameters value. When the AI assistant calls the tool, MaiAgent substitutes the value of the parameter with the same name from the Parameters Schema before sending the request. The person configuring the tool defines the fixed structure, while the AI assistant only supplies the value, so syntax correctness no longer depends on the model.

Use case

A procurement department wants its internal AI assistant to directly answer questions such as, "What are the current unit price and inventory for a product?" The data comes from an OData query service provided by the company's product catalog system. Previously, each API query required custom code.

Now, a procurement specialist can create an API tool, set the query condition to `ProductName eq '{{product_name}}'`, and assign the tool to the AI assistant. When a colleague asks, "What are the unit price and inventory for Chai?", the AI assistant automatically inserts `Chai` into the query condition, calls the API, and replies with the product name, unit price, and inventory.

Setup steps

The following example uses the public Northwind sample data service (OData), so you can follow the steps directly:

Create the tool and enter the basic information

Follow the steps in "Quick API Tool Setup" above to add an API tool. Enter a display name (for example, `Product Price Lookup`), the tool name `product_price_lookup`, a description, and a prompt. Set the `API URL` to:

```
https://services.odata.org/V4/Northwind/Northwind.svc/Products
```

Keep the `HTTP Method` set to `GET`.

Declare the parameter that the AI assistant will provide in the Parameters Schema

Declare `product_name` in the `Parameters Schema (JSON Schema)`, and use `description` to tell the AI assistant what to provide:

```
{
  "type": "object",
  "properties": {
    "product_name": {
      "type": "string",
      "description": "The complete English product name, such as Chai"
    }
  },
  "required": ["product_name"]
}
```

Reference the parameter in a Query Parameters value

Enter the fixed query structure in `Query Parameters (JSON)`, and reference the parameter you just declared by using `{{product_name}}` in the value:


```
{
  "$format": "json",
  "$select": "ProductID,ProductName,UnitPrice,UnitsInStock",
  "$filter": "ProductName eq '{{product_name}}'"
}
```

Click "**Confirm**" to save the tool.

Assign the tool to an AI assistant and test it

Follow [Assign Tools to an AI Assistant](#) to add this tool to an AI assistant in Agent mode and save it. Then, use the left navigation bar to open "**Agent Marketplace**", switch to "**Organization-created**", and select the AI assistant to start a conversation. Enter "What are the unit price and inventory for Chai?" The AI assistant displays "Tool used" and replies with Chai's unit price and inventory.



The AI assistant inserts the user-provided product name into the query condition before calling the API

Substitution rules

Rule	Description
Variable names must match the Parameters Schema	<code>{{product_name}}</code> corresponds to the parameter name under <code>properties</code> in the Parameters Schema. The name is case-sensitive.
Only values are substituted	Templates apply only to Query Parameters values. Keys remain unchanged and are not substituted.
Values are inserted verbatim	Parameter values are always inserted as literal text with no numeric conversion. For example, the leading zero in the business registration number <code>04595257</code> is preserved.
Missing values stop the request	If the AI assistant does not provide the parameter for a call, or the value is empty, the tool stops with an error that identifies the missing parameter. It does not send <code>{{product_name}}</code> as-is.
Context variables take priority	If a parameter has the same name as a context variable, such as <code>contact_id</code> or <code>conversation_id</code> , the context variable's value is always used. A value provided by the AI assistant cannot override it.
Referenced parameters are still sent separately	A parameter referenced by a template is still sent as an independent query parameter according to the existing behavior (for example, <code>product_name=Chai</code> is also sent). Most APIs ignore extra parameters.
Parameter names are preserved	Query parameter keys containing underscores or uppercase letters, such as <code>Business_Accounting_NO</code> , are preserved and are not rewritten by the system.

Values in "**Headers (JSON)**" also support `{{parameter_name}}` references to parameters in the Parameters Schema. The same rules as Query Parameters apply.

Existing tool settings that do not contain `{{ }}` are completely unaffected. If the system previously rewrote a Query Parameter name in an existing tool (for example, changing `Business_Accounting_NO` to `business__accounting_no`), reopen the tool, restore the correct name, and save it. When you reopen the editing window, Query Parameters values are masked as ; saving directly does not remove the original values.

Important Reminders

Connection Testing

After creating the tool, test the API to ensure it works properly

Use testing tools to verify functionality, such as:

POSTMAN

Your company's own API testing platform

Permission Management

Regularly check tool usage status and permission access settings

Built-in AI Image Generation Tool

Overview

The MaiAgent system integrates four top-tier AI image generation engines, providing a comprehensive solution ranging from everyday creative work to professional design. Simply enable the settings and start generating images immediately.

See: [Configure Tools for Your AI Assistant](#)

Complete Tool Comparison

Tool	Language Support	Core Strengths	Best For
Gemini 2.0	Traditional Chinese + English	Context understanding, conversational editing	Everyday creation, Chinese-language needs
GPT Image	English	Professional quality, multi-round refinement	Brand design, professional use
DALL-E 3	English	Fast generation, concept visualization	Quick prototyping, supplementary illustrations
Imagen 4.0	English	Photorealistic rendering, product visualization	Commercial photography, product showcase

Real Generation Examples

Gemini Native - Everyday Creation Example

Prompt: "A cute orange cat sitting on a windowsill, with sunlight streaming through the window onto it"



Highlights: Perfectly understands Chinese descriptions, warm and natural colors, ideal for everyday creative needs

GPT Image - Professional Design Example

Prompt: "Professional logo design: minimalist coffee cup with transparent background"



Highlights: Transparent background, refined lines, suitable for brand applications

DALL-E 3 - Quick Concept Example

Prompt: "Quick concept sketch: futuristic city skyline with flying cars"



Highlights: Fast generation, clear concepts, ideal for creative brainstorming

Google Imagen - Product Photography Example

Prompt: "Professional product photography: sleek smartphone with studio lighting"



Highlights: Photorealistic quality, professional lighting, commercial grade

Usage Guide and Best Practices

Basic Usage

Use Case	Example Prompt	Recommended Engine
Everyday Creation	Draw a cute puppy	Gemini Native
Professional Design	Design a modern minimalist logo with transparent background	GPT Image
Quick Prototyping	Quickly generate a concept for a website homepage	DALL-E 3
Product Showcase	Create a professional product photography image	Google Imagen

Multi-Round Iterative Refinement (GPT Image)

```
Round 1: "Design a coffee shop logo"
Round 2: "Change the color to dark brown"
Round 3: "Add some steam effects"
Round 4: "Make it more minimalist overall"
```

Image Reference Editing (Gemini Native)

"Based on this image, change the background to a beach scene"
"Keep the person unchanged, only modify the clothing color"
"Add some flowers to this scene"

Practical Application Guide

Scenario 1: Social Media Content Creation

Need: Create images for Instagram posts

Recommended: Gemini Native

Example Prompt: Create a cozy cafe scene suitable for an IG post

Scenario 2: Corporate Brand Design

Need: Design a company logo and brand assets

Recommended: GPT Image

Example Prompt: Design a tech company logo, minimalist modern style, transparent background

Scenario 3: Product Showcase Images

Need: E-commerce product hero images

Recommended: Google Imagen

Example Prompt: Professional product shot of wireless headphones on white background

Scenario 4: Creative Ideation and Prototyping

Need: Quickly visualize creative concepts

Recommended: DALL-E 3

Example Prompt: Concept art for a mobile app interface design

FAQ and Solutions

Quality-Related Questions

Q: How do I get the highest quality images?

A: Use GPT Image or Google Imagen and provide detailed descriptions:

Specific style requirements (e.g., "professional photography style")

Detailed scene descriptions (lighting, angle, atmosphere)

Clear quality requirements (e.g., "high resolution," "commercial quality")

Q: Why doesn't the generated image match my expectations?

A: Tips for optimizing your prompts:

Use specific rather than abstract descriptions

Specify a clear art style

Describe composition and viewpoint requirements

Describe color and lighting effects

Feature Usage Questions

Q: How do I generate images with transparent backgrounds?

A: Explicitly mention "transparent background" in your description:

```
Design a logo with a transparent background  
Create an icon with a transparent background
```

Q: Can I modify an already generated image?

A: Yes! Use Gemini Native's image reference feature:

```
Based on the image above, change the sky to sunset colors  
Keep the composition the same, only modify the person's clothing
```

Gemini 3 Image Generation Tool

Feature Overview

The Gemini 3 Image Generation Tool is an AI image generation feature powered by MaiAgent's integration with the Google Gemini 3 Pro Image model. This tool enables AI Assistants to generate high-quality images based on users' text descriptions, supports both Traditional Chinese and English prompts, and allows style transfer and image editing through reference images.

Gemini 3 Pro Image is Google's latest multimodal AI model with powerful image generation capabilities. It can generate high-quality, contextually appropriate images based on contextual understanding and world knowledge reasoning.

Key Features

Feature Type	Description	Use Cases
Text-to-Image	Generate images based on text descriptions	Marketing materials, social media posts, product visualization
Reference Image Editing	Perform style transfer or content modification based on reference images	Brand visual consistency, image style adjustment
Conversational Image Editing	Continuously adjust and refine images across multi-turn conversations	Iterative design, detail fine-tuning
Chinese Prompt Support	Directly describe desired images in Traditional Chinese	Chinese-speaking users can use it without translation

Enable the Gemini 3 Image Generation Tool

Prerequisites


- Confirm your organization plan supports image generation
- Confirm AI feature permissions are enabled
- Prepare the AI Assistant that will use this tool

Enable Steps

Go to the AI Assistant Settings Page

- Click "AI Assistant" in the left menu
- Select the AI Assistant to configure

Click "Settings"



Navigate to "AI Assistant" from the left menu, select the assistant to configure, and click Settings


Configure Tool Settings

Switch to the "Tools" tab

Click the "Select Tools" button

Find "Nano Banana Pro Image Generation" in the available tools list and check it

Click "Confirm"



Switch to the "Tools" tab to view the enabled image generation tool



After clicking "Select Tools", find "Nano Banana Pro Image Generation" in the built-in tools and check it to enable

Save Settings

Click the "**Save**" button in the bottom right corner

Wait for the system to confirm the settings are complete

Notes - Image generation consumes token quota. Please be aware of your plan limits - Generated images default to 4K high-quality output

How to Use

Use via Chat Interface



Use the image generation tool in the MaiGPT chat interface. The AI Assistant generates images based on your description

Enter an Image Description

Type your desired image description in the chat box

Examples:

"Generate a landscape image of a sunset beach for me"

"Draw a cat wearing a spacesuit"

"Design a minimalist-style coffee shop logo"

"Create a watercolor painting of a mountain landscape"

Use Reference Images

Upload an image in the conversation as a reference

Combine with text instructions to describe the desired modifications

Examples:

Upload a photo and say: "Convert this photo to a watercolor painting style"

Upload a product image and say: "Design a promotional poster for this product"

Conversational Editing

After generating an image, request modifications in subsequent messages

Examples:

"Change the background to blue"

"Add some stars"

"Make the overall atmosphere warmer"



Conversational editing: The user requests changing the shopping bag to pink, and the AI Assistant regenerates the modified image based on context memory

The more detailed the description, the better the generated image quality. It is recommended to include details such as style, color, composition, and atmosphere.

Use Case Examples

Scenario 1: Marketing Material Creation

Requirement: Quickly generate visual materials for social media posts

Conversation Example:

User: Generate an Instagram-suitable image for me. The theme is spring new product launch, with a fresh and bright style, featuring cherry blossom and shopping bag elements.

AI: [Generated image] Created a spring new product launch promotional image for you, featuring cherry blossom petals and stylish shopping bag elements with an overall fresh and bright style.

User: Great, but can you change the shopping bag color to pink?

AI: [Generated updated image] The shopping bag has been adjusted to pink, making it more consistent with the cherry blossom theme.



Actual operation screen: The user inputs marketing material requirements, and the AI Assistant instantly generates a spring promotional image with cherry blossom and shopping bag elements

Scenario 2: Product Visualization

Requirement: Generate visual representations for product concepts not yet in production

Conversation Example:



User: I want to design a smartwatch with a round watch face, silver metal band, and deep blue dial. Please generate a concept image for me.

AI: [Generated image] Created a smartwatch concept image with a round watch face paired with a silver metal band and deep blue dial.



Actual operation screen: The AI Assistant generates a smartwatch concept image based on the description, featuring a round watch face, silver metal band, and deep blue dial

Scenario 3: Style Transfer

Requirement: Convert existing images to different artistic styles

Conversation Example:

User: [Upload photo] Convert this photo to comic style

AI: [Generated image] Your photo has been converted to comic style, preserving the original composition and main elements.



Actual operation screen: Converting Taipei 101 tower to Japanese anime style, preserving the main architectural outline and applying comic-style rendering

Prompt Writing Tips

Effective Prompt Writing

Include specific descriptions:

- Subject: What content you want to generate
- Style: Realistic, cartoon, watercolor, oil painting, graphic design, etc.
- Color: Main color tones and color preferences
- Composition: Close-up, panoramic, bird's-eye view, etc.
- Atmosphere: Warm, cool, lively, calm, etc.

Good examples:

- "An orange Shiba Inu sitting on a park bench in autumn, with golden maple leaves in the background, warm afternoon light, realistic photography style"
- "Minimalist-style tech company logo, using blue and white, geometric design"

Avoid being too vague:

- "Draw a picture" (lacks specific description)
- "A nice image" (no clear direction)

FAQ

Q: What languages does the Gemini 3 Image Generation Tool support for prompts?

A: It supports Traditional Chinese and English prompts. You can directly describe the desired image content in Chinese.

Q: How many tokens does it cost to generate one image?

A: Image generation consumes more tokens than text-only conversations. The actual consumption depends on the prompt complexity and image quality settings. It is recommended to monitor your organization's token usage to avoid exceeding plan limits.

Q: Can I use reference images?

A: Yes. Upload an image in the conversation as a reference, and the AI Assistant will incorporate it into the generation process. This is especially useful for style transfer, image editing, or creating based on existing materials.

Q: What is the resolution of generated images?

A: Images are generated in 4K high quality by default, suitable for most use cases.

Q: Can I continuously modify images across multi-turn conversations?

A: Yes. The AI Assistant remembers the conversation context. You can request detail adjustments, style modifications, or element additions in subsequent messages to progressively refine the image.

Q: What should I do if image generation fails?

A: Possible causes and solutions:

- **Prompt violates content policy:** Adjust the description to avoid inappropriate content
- **System is busy:** Try again later
- **Insufficient token quota:** Check the remaining quota of your organization plan

Related Documentation

Further Reading - [Tool Feature Overview](#) - [Configure Tools for AI Assistants](#) - [Built-in AI Image Generation Tool](#) - [Create an AI Assistant](#)

Text to SQL with Supabase

What Is Supabase?

Supabase is an open-source platform service designed to simplify modern application development. Its core features include:

Database: Store various types of data.

Real-time updates: Your application is notified immediately when data changes.

Account management: Manage user accounts and passwords.

Built-in authentication mechanism: A built-in authentication mechanism that simplifies user identity management and provides multiple authentication methods.

Auto-generated API (data access endpoints): Access your database data through simple methods.

Benefits of Integrating Supabase

Cross-table queries: Supabase supports multi-table relationships, allowing you to query relationships between Table A and Table B.


Indexing for improved query performance: You can create indexes on frequently queried content to help the AI assistant find what you need more precisely.

Set Up Your Supabase

1. Create a Supabase Account

First, go to the [Supabase website](#) and click "Sign in Start your project."

If you haven't registered yet, please create an account first to proceed with the following steps




After logging in, you can create a new organization or use an existing one. Within the organization, create a project — each project has its own independent database.





2. Navigate to the Database Page

After entering the project, click Database > Tables in the left navigation bar to add data







Create a New Table

Click "New Table" to create a table and name it:





Supabase offers several ways to create a new table:




Manually add columns: Suitable for designing the table structure from scratch. You can add columns one by one and set each column's data type, default value, etc.

Import .csv/.tsv or plain text: Suitable for quickly creating tables, especially when you already have existing data.

Notes:

The first row of a plain text file must contain column names, with columns separated by commas (CSV) or tabs (TSV).

Here, choose to import data: click "**Import data from CSV**" and paste a tab-separated text file. Scroll down to see the table result. After import is complete, click "Save."



Select import method




Import plain text



Preview result


Primary Key

After the import is complete, you are redirected to the settings page where you must specify a primary key. A primary key is like an ID number — it uniquely identifies each record. Here, we select Customer ID as the primary key.



Foreign Key

Scroll down to see the Foreign Key settings. A Foreign Key is like an address that points to the source of a record or more detailed data associated with it.



Suppose we have two tables: "Customers" and "Orders."

Customers Table:

CustomerID - Primary Key

CustomerName

Phone

Address

Orders Table:

OrderID - Primary Key

CustomerID - Foreign Key (references CustomerID in the Customers table)



OrderDate

TotalAmount

In this example, the "CustomerID" in the Orders table is a foreign key that references the primary key "CustomerID" in the Customers table. Through this foreign key, we can identify which customer placed each order.

In this example, the foreign key should be placed in the Order Table: Relationship direction: One customer → can have multiple orders (one-to-many relationship) One order → belongs to only one customer Foreign key principle: > The foreign key should be placed on the "many" side Therefore: The Order Table has a customer_id (foreign key), because it is the "many" side in the one-to-many relationship The Customer Table does not need to store order information

Therefore, we set up the Order Table to link the Customer ID in the Customer Table to the Customer ID in the Order Table.



After the relationship is configured, click "Save" to establish the database relationship.

3. Setup Complete

Once the tables are created, you have a complete database and can query data within it using SQL syntax.



How to Create a Supabase Tool

To use the Supabase tool on MaiAgent, you need to create it as an MCP tool so the AI assistant can access Supabase functionality:

For tool introduction, see: [Tool Overview](#)

Create a Server on an MCP Service Platform and Connect to Supabase

For how to connect MCP tools, see: [Remote MCP Service Overview](#). Currently, only the [Composio Platform](#) supports Supabase integration.

Enable Available Features in Composio

When connecting to Composio, enable basic database operations such as create, delete, query, etc.

Composio has basic options pre-selected under Important by default. You do not need additional configuration — just make sure the Important options remain checked.

Configure Actions for Supabase

Select the actions you want to enable for this MCP server using your connected Supabase account

Selected Actions:

7/30

Search actions...

☒ Important 7

☒ Execute project database query

☒ Create an organization

☒ Create new project

☒ Generate TypeScript types

☒ Get Table Schemas

☒ List all organizations

☒ List all projects

☐ Secrets 7

☐ DestructiveHint 22

☐ OpenWorldHint 54

Back

Create Server

Add the Configured Supabase Tool to the Available Tools List

For how to create MCP tools, see: [Create MCP Tools](#)

Supabase Tool Connection URL

After creating the server service on the MCP service platform, process the URL as follows:

Original URL (obtained from the MCP Server):\



https://backend.composio.dev/v3/mcp/12345678/mcp?include_composio_helper_actions=true

Remove "[?include_composio_helper_actions=true](#)" entirely

New URL (to paste on the MaiAgent tools page):\nhttps://backend.composio.dev/v3/mcp/12345678/mcp

You must remove the above content; otherwise, the AI assistant will not be able to use the Supabase tool correctly.

After the tool is created, go to the AI assistant settings and add the Supabase tool to the assistant's available tools list:



Make sure to click Save on the tools page; otherwise, the AI assistant still cannot use the Supabase tool.

Using the Supabase Tool

With MaiAgent AI assistant and the Supabase tool, simply describe the data you want to query in everyday language, and Supabase automatically generates the corresponding SQL syntax to extract the information you need from the relational database.

Sample Database

1. Customers Table

Column Name	Primary Key	Required	Description
CustomerID	Yes	Yes	Unique customer identifier
CustomerName	No	No	Customer company or individual name
CustomerType	No	No	Customer category (e.g., Retailer, Restaurant, Distributor)
ContactName	No	No	Primary contact name
Phone	No	No	Contact phone number
Email	No	No	Email address
Address	No	No	Customer address
Region	No	No	Geographic region (e.g., North, South)
CustomerLevel	No	No	Customer importance level (A, B, C)

Primary Key: CustomerID

Foreign Key Relationships: None

2. Orders Table

Column Name	Primary Key	Required	Description
OrderID	Yes	Yes	Unique order identifier
CustomerID	No	No	References the Customers table
OrderDate	No	No	Order creation date
DeliveryDate	No	No	Expected or actual delivery date
PaymentMethod	No	No	Payment method (Cash, Credit Card, Wire Transfer)
OrderStatus	No	No	Order processing status
TotalAmount	No	No	Order total amount (numeric)
ShippingFee	No	No	Shipping cost

Primary Key: OrderID

Foreign Key Relationships:

CustomerID → Customers.CustomerID (The CustomerID in Orders references the CustomerID in the Customers table)

3. Products Table

Column Name	Primary Key	Required	Description
ProductID	Yes	Yes	Unique product identifier
ProductName	No	No	Product name
Description	No	No	Detailed product description
Category	No	No	Product category
Brand	No	No	Product brand
Size	No	No	Product specification or size
Cost	No	No	Product cost (numeric)
Price	No	No	Product selling price (numeric)

Column Name	Primary Key	Required	Description
StockQuantity	No	No	Current stock quantity (numeric)

Primary Key: ProductID


Foreign Key Relationships: None

Scenario 1: Track Incomplete Orders

Database state: The Orders table has two incomplete orders: OR003 and OR004.

Filter	Sort	Insert	Add RLS policy	Enable Realtime	Role postgres		
訂單編號 (OrderID)	ntMethod)	訂單狀態 (OrderStatus)	總金額 (TotalAm				
OR001		已完成	12500				
OR002		已完成	8000				
OR003		處理中	25000				
OR004		已確認	15000				
OR005		已完成	5000				

Natural language input: In the AI assistant's chat, simply type: "Tell me which orders are not yet completed." The AI assistant automatically invokes the tool and generates a structured SQL query.



Supabase automatic query: The AI assistant automatically invokes the Supabase tool, converting your natural language into an SQL query, such as:



Combining the query results with the AI assistant's analysis, the AI assistant responds with the following order information, sorted by priority:



Through the collaboration of the Supabase tool and the AI assistant, you can easily track incomplete orders and receive analysis and prioritization suggestions from the AI assistant, enabling more efficient order processing and improved customer satisfaction.

Scenario 2: Query Contact Information for Incomplete Orders

Order OR003 is still being processed, with customer CU003 and the contact person being **Huang Procurement**

Supabase automatic relational query: Although the conversation with the AI assistant is about the Orders table, through the configured foreign key relationships, Supabase knows that the **Customer ID** here references the **ID** in the **Customers** table, so the returned content comes from the **Customers** table.



AI assistant response: The AI assistant then provides the correct contact information and other contact methods based on its analysis.



Through the Supabase tool, you can fully leverage table relationships in the database to easily extract information from multiple related tables and receive a customer name list from the AI assistant. The clearly defined relationships between tables ensure data consistency and reliability, making query results more trustworthy.

Additional Notes:

You can adjust the natural language input based on your actual needs, such as: "Tell me about today's incomplete orders" or "Tell me about VIP customers' incomplete orders." The Supabase tool can accurately parse and execute these queries.

The AI assistant can further integrate other information, such as inventory status and logistics information, to provide more comprehensive order analysis.

The tool can only access data stored in your database. If you have analysis needs, make sure to upload the data to the database before starting the analysis.

Configure Tools for an AI Assistant

The tools you have created can now be selected in AI assistant settings, giving your AI assistants more powerful capabilities. This guide explains how to configure the tools you have created for a specific AI assistant.

Configure Tools

1. Enter AI Assistant Settings

First, navigate to the "AI Features" section from the left navigation bar, then click "AI Assistants". This will display a list of all AI assistants created in your account.





Enter the settings page of a specific AI assistant

3. Switch the AI Assistant's Response Mode

Important prerequisite: Only AI assistants set to Agent Mode can use tools. If your AI assistant is in another mode (e.g., General (Default) response mode), you cannot configure or use tools for it.

On the AI assistant's settings page, switch to the "Response Mode Settings" page tab.

Check: Confirm whether the currently selected mode is "Agent Mode"

Change: If it is not currently in Agent Mode, please **switch** it to Agent Mode.



Make sure the AI assistant is set to Agent Mode to use tools

4. Navigate to the Tool Configuration Page

On the AI assistant's settings page, find the "**Tools**" page tab.



Find the tool configuration page

5. Select Tools to Enable

On this tool configuration page, click the "**+ Select Tools**" button in the upper right corner. This will open a list of pre-built tools, where you can perform the following actions.

Actions:

Browse the list and find the tools you want this AI assistant to be able to use.

Check the checkbox next to the tools you want to enable.

You can select any number of tools based on the AI assistant's tasks and needs.

You can select MaiAgent built-in tools or tools you have defined yourself



Directly check the tools you want this AI assistant to use

Tip: Make sure to only configure tools that the AI assistant actually needs and is authorized to use. Configuring too many irrelevant tools may affect the AI's judgment accuracy and response efficiency.

6. Review Selected Tools

After completing your selections, you can review the tools configured for the AI assistant once more to ensure the correct tools have been configured for this AI assistant.



Confirm the correct tools have been configured for the AI assistant

Control System Functions

In addition to the tools you select, AI assistants have system functions enabled by default, such as searching the current conversation history, reading attachments, and generating download links. You can disable these functions individually based on each AI assistant's needs to reduce unnecessary token usage and make the assistant's behavior more predictable.

For example, a customer service supervisor creates an AI assistant that only answers frequently asked questions and does not process attachments. She opens the assistant's advanced settings, disables attachment reading and parsing, and saves the changes. The assistant can still answer knowledge base questions, but it will no longer attempt to parse files uploaded by customers.

1. Open Advanced Settings

Open the AI assistant you want to adjust and select **Advanced Settings** from the left side of the settings page.

2. Adjust System Functions

Scroll to **System Functions**, then enable or disable the following functions as needed:


Search conversation history

Get surrounding message context

Read plain-text attachments

- Parse attachment content
- List available parsers
- Generate download links

After a function is disabled, the assistant can no longer use it, which may affect related responses.



Enable or disable system functions based on the AI assistant's purpose

3. Save the Settings

Click **Save** in the lower-right corner to apply the settings.

The last four functions in the list only work in **Agent Mode**. If you cannot change these toggles, go to **Response Mode Settings** and switch to Agent Mode first.

7. Save AI Assistant Settings

After completing your tool selections, **make sure to** click the "**Save**" button to save the changes you have made to the AI assistant settings.




Click "Save" to apply the tool configuration changes

Configure MaiAgent Built-in Tools

MaiAgent has four built-in AI image generation tools that you can apply to your AI assistants based on your needs, with no additional setup required:

For a comparison of the four tools, please refer to: [MaiAgent Built-in AI Image Generation Tools](#)



Now, your selected AI assistant has the ability to use these enabled tools. When interacting with this assistant, if a relevant scenario is triggered, the AI will attempt to call these tools to complete tasks or retrieve information.

Agent UI

1. What Is Agent UI?

Traditional AI assistants can only reply with plain text. **Agent UI** lets you design visual "information cards" for your AI assistant — when users ask questions, the AI assistant doesn't just talk, it can also send beautifully structured cards with images and buttons.

The currently supported card format is **FlexMessage**, which can render rich card UIs in Web Chat and LINE channels.


Agent UI Highlights: The AI assistant automatically determines when to send a card based on the conversation context, fills in the corresponding information, and users can tap card buttons to continue interacting.

Comparison: Plain Text vs. Agent UI Cards

Scenario	Plain Text Reply	Agent UI Card Reply
Industrial Equipment Spec Inquiry	"IBX-7800 Industrial Motherboard, Intel i7, DDR5, NT\$28,500"	Spec sheet (CPU/I/O/Operating Temp) + Stock tag + Request Quote Download Spec Sheet buttons
Credit Card Statement Inquiry	"Current spending \$6,508, payment due 06/15"	Transaction details + Current period/Rewards/Amount due + Pay Now button
Medical Appointment Confirmation	"Cardiology 06/03 14:00, Registration No. R2026053-088"	Doctor/Time/Location + Check-in QR Code + View Waiting Status button
Electronic Boarding Pass	"JX 802, TPE 09:25 → NRT 13:35, Seat 12A"	Flight segment/Class/Seat/Baggage + Boarding QR Code + Online Seat Selection button

Live Conversation Demo

Whether it's B2B manufacturers inquiring about product specs, customers checking appointment information, or users viewing boarding passes directly in a conversation — the AI assistant automatically sends the corresponding card based on the conversation context, presenting all information clearly at once.



Manufacturing example: Customer asks "Recommend an industrial computer motherboard," AI assistant returns a product card with CPU, I/O, operating temperature specs and a "Request Quote" button



Medical clinic example: User checks appointment information, AI assistant returns an appointment confirmation card with department, time, location, and QR Code check-in, plus a "View Waiting Status" button



Airline example: User checks flight information, AI assistant presents an electronic boarding pass card with flight segment, class, seat, baggage allowance, and boarding QR Code

For more industry scenarios, see "[5. Use Case Examples: Six Industry Scenarios](#)" below.

2. Core Concepts

Fixed Fields vs. AI-Generated Content

Each field in a card can be set to either a "fixed default value" or "let AI fill in dynamically from the conversation context":

Field	Configuration	Use Case
Fixed Value (Default)	Enter fixed text directly	Brand name, card title, fixed labels (e.g., <code>TRANSACTION</code> , <code>BOOKING PASS</code>)
AI-Generated (AI Prompt)	Describe in natural language what AI should	Last four digits of customer account, spending amount, amount due — any information that needs to be retrieved from the

Field	Configuration	Use Case
	fill in	knowledge base or conversation

Example (Financial Statement Card):

Field	Default Value (Fixed)	AI Prompt (Dynamic)
Title Label	TRANSACTION (fixed)	—
Card Title	Credit Card Statement (fixed)	—
Statement Period last four card	—	Retrieve statement period and digits from the billing database
Description Transaction List for this period	—	List major merchants and amounts
Total Spending period from	—	Retrieve total spending for this the billing database
Rewards from the	—	Retrieve current period rewards membership system
Amount Due deducting rewards	—	Calculate amount due after from total spending
Payment Deadline the billing	—	Retrieve payment deadline from database
Button Text	Pay Now (fixed)	—
Payment Link link using	—	Generate personalized payment statement number

What Can Card Buttons Do?

FlexMessage cards support two types of interactive buttons:

URL Button

Opens a specified URL when clicked.

Use Cases:

- Manufacturing: Open product technical spec PDF, schedule a Demo form
- Finance: Jump to statement payment page, credit card application form
- Education: Open admissions brochure, registration link

- Medical: View waiting status, cancel appointment page
- Airlines: Online seat selection, add-on baggage or meals

Configuration: In the "Link" field's AI prompt, enter "Retrieve the corresponding page link from the knowledge base," and AI will automatically populate the correct URL.

Prompt Button

Automatically sends a message to the AI assistant when clicked, and AI continues the conversation.

Workflow (Education enrollment scenario example):

```
User: "Tell me about courses or admissions"
    ↓
AI returns a course progress card (In-progress courses
  Pending assignments
  Open enrollment)
    ↓
User clicks "Enroll Now" button
    ↓ (Button triggers prompt: "I want to enroll in the summer internship")
AI continues: "What is your current department and year? What is your preferred
  internship field?"
```

Use Cases:

- Trigger follow-up processes (enrollment, appointments, applications, reservations)
- Ask "View full details" or "View other room types" to expand more information
- Collect additional information AI needs to continue serving the user

3. Creating an Agent UI (FlexMessage)

Navigate to the Agent UI Page


In the left menu, click **AI Features** → **Agent UI**, then click **Add Agent UI** in the top right corner.



Agent UI management list showing created card tools and their associated AI assistants


Select Card Type

Currently available: **FlexMessage Card Information Display**. Click to enter the configuration page.



Select card type — currently supports FlexMessage

Fill In Basic Information



After filling in basic information, you can select a template and configure AI fields

Field	Description
Display Name	Name shown in the admin dashboard for easy identification (e.g., <code>Finance/Banking/Securities Product Card</code>)
Tool Name	English name for AI model identification. Naming rules: English letters, underscores, no spaces (e.g., <code>finance_product_card</code> , <code>medical_clinic_card</code> , <code>airline_flight_card</code>)
Description	Describe the purpose of this card tool (optional)
Prompt	Tell the AI assistant when to use this card (e.g., "When users ask about credit card statements, payment information, or current spending details, return the statement card")

The prompt is the trigger key. The more specific you write it, the better AI knows when to proactively send a card instead of just replying with text.

Select a Template

Click a preset template on the left to preview the card style on the right in real time. After confirming, click **Apply This Template**.


12 Built-in Templates and Applicable Industries:

Template	Use Case	Suitable Industries
Restaurant	Store recommendations, reservations	Food & Beverage, Tourism
Hotel	Room type display, booking	Hotels, Accommodations
Local Search	Location search, nearby recommendations	Tourism, Retail
Real Estate	Property introduction, viewing appointments	Real Estate, Rentals
Apparel	Product display, spec introduction	Retail, E-commerce, Manufacturing
Transit	Routes/Flights/Schedules	Transportation, Airlines, Logistics
Shopping	Product comparison, shopping cart	Retail, E-commerce
Menu	Item display, menu	Food & Beverage, Retail
Receipt	Statements, order receipts, transaction details	Finance, Retail, E-commerce
Ticket	Tickets, boarding passes, appointment confirmations	Airlines, Events, Medical Registration
Social	Profile cards, social profiles	Education, HR, Social
TODO app	Task lists, course progress, work order status	Education, Customer Service, Manufacturing Work Orders
Custom	Paste custom FlexMessage JSON	Any scenario requiring custom layouts

Selecting a template is just the starting point. After applying, you can freely adjust the content and logic of each field in the "AI Field Settings."

Configure AI Fields

After applying a template, the system automatically parses all editable fields in the card:



Each field can have a fixed default value, or an AI prompt for AI to dynamically generate content

Each field has three properties:

Type: The field's data type (text, image, button, link)

Default Value: Fixed content to display (if an AI prompt is also set, AI-generated content takes priority)

AI Prompt: Tells AI what to fill in this field, e.g., "Based on the user's query, retrieve the corresponding product name

account information

department from the knowledge base"

The **live preview** on the right updates with your settings, so you can confirm the card's visual appearance.

Save

After confirming all settings, click **Save** in the top right corner. The Agent UI tool is now created and will appear in the list.

4. Binding to an AI Assistant


After creating an Agent UI, you need to attach it to an AI assistant so the assistant knows it can use this card.

Go to AI Assistant Settings

Navigate to **AI Features** → **AI Assistants**, select the AI assistant you want to configure, and click **Settings**.

Open the Agent UI Tab

In the settings page tab bar, click **Agent UI**.



In the AI assistant settings page, you can manage all bound Agent UI tools

Add Agent UI

Click **Select Agent UI** and choose the card tool to add from the list.

A single AI assistant can bind **multiple** Agent UIs

Each Agent UI's **prompt** determines when AI uses which card

Save

Click **Save** in the bottom right corner. The settings take effect immediately.

5. Use Case Examples: Six Industry Scenarios

From manufacturing, finance, and education to medical, hospitality, and airlines — every industry can design custom cards tailored to their needs, enabling AI assistants to respond to users in the most intuitive way. Below are actual conversation demos from six typical enterprise scenarios.

All examples in this section have been built and tested on the MaiAgent platform. A single AI assistant can bind multiple Agent UIs simultaneously, and AI will automatically select the appropriate card based on the user's question.

Multi-Card Management for One Assistant

Configure multiple industry scenarios in the admin dashboard and attach them to the same AI assistant:



One AI assistant bound with six industry cards — AI automatically selects the appropriate card based on the user's question

Example Configuration Tutorial: Financial Statement Card

Below is a complete walkthrough using the "Finance — Statement Inquiry" as an example, demonstrating the entire process from zero configuration to live conversation. Other industries can follow the same approach.

Scenario Description

A bank wants members to receive a statement card (with transaction details, amount due, payment deadline, and a "Pay Now" button) when they ask about their current statement on LINE or Web Chat, instead of a plain text reply.

Configuration Details

Agent UI Settings:

Field	Setting
Display Name	Finance/Banking/Securities Product Card
Tool Name	finance_product_card
Prompt	When users ask about current spending, payment information, credit card statements, or transaction details, return the statement card
Template	Receipt

AI Field Settings:


Field	Default Value	AI Prompt
Title Label	TRANSACTION	—
Card Title	Credit Card Statement	—
Statement Period Description	—	Retrieve statement period and last four card digits from the billing database
Transaction List	—	List major merchants and amounts for this period
Total Spending	—	Retrieve total spending for this period from the billing database
Rewards	—	Retrieve current period rewards from the membership system
Amount Due	—	Calculate amount due after deducting rewards from total spending
Payment Deadline	—	Retrieve payment deadline from the billing database
Statement Number	—	Retrieve statement number from the billing database
Button Text	Pay Now	—

Field	Default Value	AI Prompt
Payment Link	—	Generate personalized payment link using statement number

Logic for choosing "Fixed Value" vs. "AI Prompt": Use fixed values for content that doesn't change with conversations, such as brand labels and button text. Use AI prompts for content that needs to be retrieved from systems, such as customer accounts, spending amounts, and payment links.

Live Conversation Result

After the user asks a question, the AI assistant automatically triggers the card tool, retrieves data from the billing database and membership system, and fills in the fields:




User enters "Recommend a credit card or financial product" (or "Check my statement for this period"), AI assistant returns a statement card with transaction details, amount due, and payment deadline — user can click "Pay Now" to complete payment directly

Six Industry Scenario Overview

Once you understand the configuration logic above, you can apply the same approach to different industries. Here are six typical scenarios:

1. Tech/Electronics Manufacturing: Product Catalog and Spec Inquiry


When customers or sales teams inquire about industrial computers, embedded systems, AOI inspection equipment, and other products, the AI assistant automatically returns product cards with complete specs, pricing, stock status, and "Request Quote" "Schedule Demo" buttons, accelerating the B2B sales process.



User: "Recommend an industrial computer motherboard" — AI assistant returns a product card with complete specs including CPU, I/O, and operating temperature

2. Finance/Banking/Securities: Statements and Product Recommendations


For financial scenarios such as statements, credit card recommendations, financial plans, and insurance inquiries, the AI assistant can integrate internal system data to return structured financial information cards with "Pay Now" "Apply" CTAs.



User: "Recommend a credit card or financial product" — AI assistant returns a statement card with current transaction details, rewards, and payment deadline, with a "Pay Now" button

3. Educational Institutions: Course Progress and Enrollment

Students can instantly check course progress, pending assignments, completed grades, and admissions information. Schools can also use the AI assistant to push campus announcements and registration links.



User: "Tell me about courses or admissions" — AI assistant returns a course progress card with in-progress courses, pending assignments, and open enrollment statuses

4. Medical Clinics: Appointment Registration and Clinic Information

For scenarios including department introductions, clinic schedules, appointment registration, QR Code check-in, and health education information — patients can complete appointments and check waiting status directly on LINE or Web Chat, reducing call center workload.



User: "Show me departments and clinic schedules" — AI assistant returns an appointment confirmation card with doctor, time, location, and QR Code, providing "View Waiting Status" and "Cancel Appointment" buttons

5. Hospitality/Hotel Industry: Room Booking and Marketing Recommendations

For hospitality scenarios including room type recommendations, restaurant reservations, event registration, and membership benefits — visual cards display pricing, ratings, promotions, and booking CTAs, improving conversion rates.



User: "Recommend hotels or reservation info" — AI assistant returns room type and restaurant recommendations simultaneously, with ratings, limited-time offers, and booking/reservation buttons

6. Airlines: Flight Information and Electronic Boarding Passes

For scenarios including flight inquiries, class comparisons, baggage regulations, and online seat selection, the AI assistant can directly present electronic boarding pass QR Codes, seats, baggage allowances, and other key information — no need to open a separate app.



User: "Check flight and seat information" — AI assistant returns an electronic boarding pass card with TPE→NRT flight segment, seat, baggage allowance, and boarding QR Code

6. FAQ

- What are the naming rules for tool names?
- How should I write an effective prompt?
- What happens if the AI prompt for a field is not filled in?
- Can a single AI assistant have multiple Agent UIs attached?
- Can Agent UI be used on LINE channels?

Skills

Skills Feature Overview

What Is a Skill?

A skill is like a **professional certification** for an AI assistant — each skill learned gives the AI assistant an additional professional capability. If tools are the "screwdrivers in the AI assistant's toolbox," then skills are "a complete set of Standard Operating Procedures (SOPs)" that not only tell the AI assistant what tools to use, but also **when to use them, how to use them, and how to respond after using them**.

Imagine you're running a customer service center:

Without skills:

Customer: "I want to return an item"

AI agent: "Sure, what's your order number?" (*can only handle basic Q&A*)


With skills:

Customer: "I want to return an item"

AI agent: "Sure!" (*automatically activates the return process skill*) → Query order → Verify return eligibility → Generate return form → "Your return request has been created. The return number is RT-20260324, and processing is expected to complete within 3-5 business days."

Core Components of a Skill

Each skill consists of three parts:



Instructions: Detailed instructions in Markdown format, like an SOP that guides the AI assistant through task completion step by step

Attached Tools: Tools required by the skill (MCP tools, API tools, etc.) that enable the AI assistant to perform actual operations

Resource Files: Attachment data that the skill can reference during execution (only available for skills created via upload)

How Skills Work

A complete skill workflow is as follows:

Create a skill:

Define the skill's name and description (the description determines when the AI triggers this skill)

Write detailed instructions (SOP in Markdown format)

Bind the required tools

Bind to an AI assistant:

Select the skills to use in the AI assistant's settings

An AI assistant can have multiple skills bound to it

User asks a question:

The user makes a request to the AI assistant in natural language

Example: "Check how many vacation days I have left"

AI assistant evaluates and activates the skill:

The AI assistant matches the user's question against the descriptions of bound skills

Determines which skill should handle the question

Automatically expands the complete instructions for that skill

Execute instructions and call tools:

The AI assistant follows the steps in the skill instructions one by one

Automatically calls the skill's attached tools when needed (querying databases, calling APIs, etc.)

Generate the final response:

The AI assistant consolidates the results returned by tools

Generates a structured response according to the reply format defined in the skill instructions

Key Advantages of Skills

Standardize Complex Tasks

Without skills: The AI assistant can only answer based on general knowledge, resulting in inconsistent quality

With skills: Every execution follows the SOP, ensuring consistent response quality

Precise Triggering and Task Division

Through trigger conditions in descriptions, the AI assistant can automatically determine when to use which skill

Different skills handle different tasks, achieving modular capability division

Reusable and Shareable

A single skill can be bound to multiple AI assistants
Skills can be exported as `.skill` files and shared with other organizations

Encapsulated Instructions and Tools

Bundle "what to do" (instructions) and "what to use" (tools) together
No need to repeatedly write extensive process details in the AI assistant's role instructions

Differences Between Skills and Tools

Comparison	Tools	Skills
Nature	A single external capability (API, MCP)	A complete execution workflow
Contents	API endpoint MCP server	Instructions + Tools + Resources
Analogy	Screwdriver	Assembly manual + Screwdriver
Trigger method	AI decides on its own whether to call it	Matched based on trigger conditions in the description
Use cases	Single actions (check weather, send email)	Multi-step workflows (return processing, vacation calculation)

Skills and tools are complementary: skills invoke tool capabilities through "attached tools," while tools provide the actual execution capabilities that skills need.

Practical Application Scenarios

Enterprise Customer Service

Skill: Return Processing Workflow

Trigger: When a customer requests a return or exchange

Steps: Query order → Verify return eligibility → Create return form → Reply with processing result

Attached tools: Order Query API, Return System API

Government Agencies

Skill: Vacation Day Calculation

Trigger: When a user asks about vacation days or annual leave days

Steps: Confirm personnel category → Search regulatory examples → Compare and calculate → Produce result

Attached tools: Knowledge base search

E-commerce Platforms

Skill: Customer Data Collection

Trigger: When a customer expresses purchase intent or requests a quote

Steps: Guide them to provide name/company/contact information → Write to CRM

Attached tools: Google Sheet API, CRM tool

Marketing Teams

Skill: Email Sending Service

Trigger: When a user requests to send an email or contact a customer

Steps: Confirm recipient → Compose email content → Send

Attached tools: Gmail MCP tool

Skills Management

If you're not yet familiar with what skills are, please read: [Skill Overview](#)

Creating Skills

Click [Skills](#) in the left sidebar to enter the skills management page, then click the [Add Skill](#) button in the upper right corner.

The system provides two creation methods:

Method 1: Manual Creation

Suitable for writing skill instructions from scratch.

In the add skill dialog, click the [Switch to manual creation](#) link at the bottom to switch to manual mode

Fill in the [Basic Settings](#) tab:


Skill Name (required): Give the skill an easily identifiable name

Description (required): Briefly describe the skill's purpose



Manual skill creation — Basic Settings tab

Switch to the **Instructions** tab and write instructions in Markdown format. You can also click the **Apply Template** button to use a system-provided default template as a starting point



Instructions tab — Supports Markdown format and template application

(Optional) Switch to the **Attached Tools** tab to select the tools required by the skill

Click **Save** to complete creation

Method 2: Upload Skill Package

Suitable for importing pre-packaged **.skill** or **.zip** files.

In the add skill dialog (upload mode by default), drag and drop the **.skill** or **.zip** file into the upload area, or click the area to select a file



Upload skill package — Supports .skill or .zip format

After the upload is complete, the system automatically parses the `SKILL.md` file to extract the skill name, description, and instructions

After successful creation, you are automatically redirected to the edit page where you can further adjust the settings

The skill package must contain a SKILL.md file in the following format: markdown --- name: Skill Name description: Skill Description --- This is where the AI assistant's execution instructions go...

Editing and Configuring Skills

In the skill list, click the `Edit` icon for the skill to enter the skill settings page.

Basic Settings

Edit the skill's name and description.



Edit skill — Basic Settings

Instructions

Edit the skill instructions in Markdown format. These instructions serve as the execution guide when the AI assistant activates the skill.


Click [Apply Template](#) to load a default instruction template.

Applying a template will overwrite the current instructions. Please confirm before proceeding.

Skill Package (Upload-created skills only)

If the skill was created by uploading a `.skill` `.zip` file, this tab displays:

- Skill package info:** File name, size, upload time
- Preview:** Expand to view the file structure within the skill package
- Download:** Export the skill package file
- Re-upload skill package:** Upload a new version to replace the existing skill package



Edit skill — Instructions (with practical examples)

Attached Tools


Bind the required tools to the skill so the AI assistant can call them when executing the skill.

Click the **Select Tools** button to open the tool selection dialog

Check the tools to attach (multiple selection supported), then click Confirm

Selected tools appear in the list, showing the tool name, type (API MCP), and description

To remove a tool, click the delete icon on the tool card



Attached Tools tab

After completing the configuration, click the **Save** button in the lower right corner of the page.

Exporting Skills

In the skill edit page's **Skill Package** tab, click the **Download** button to export the skill as a **.skill** file.

Manually created skills can also be exported. The system automatically packages the instructions and resources into a .skill file. Exported skill packages can be shared with other organizations for import.

Binding Skills to an AI Assistant

Skills must be bound to an AI assistant to take effect:

Go to the [AI Assistant](#) page, select the target AI assistant, and enter its settings

Find the **Skills** setting in the model configuration area

Select the skills to bind and click Confirm

Save the AI assistant settings

Once bound, the AI assistant automatically determines whether to activate a skill based on the user's question during conversations and executes the corresponding task according to the skill instructions.

Controlling Skill Visibility and Permissions by Role

If you want members with specific roles to see only the skills relevant to their work, you can use role permission management to assign skills to each role individually.

For example, suppose a company's system administrator wants sales representatives to be able to read the “Customer Summary” skill but not modify or delete it. The administrator can edit the sales role, assign this skill, and grant read permission only. Once configured, sales representatives can see “Customer Summary” in the skill list without receiving permission to modify or delete it.

How It Works

Unassigned means hidden: Skills created by the organization must be assigned to a role before members with that role can see them in the skill list. Unassigned skills are completely hidden from that role

Granular permissions: Each “role × skill” assignment can have its **Read**, **Edit**, and **Delete** permissions configured independently, allowing read-only scenarios such as “can read but cannot modify”

Create permission: Each role has its own “Allow skill creation” toggle, which controls whether members with that role can create skills

Owners are unrestricted: Organization owners automatically have full access to all skills and do not require assignments

Global skill exception: Built-in global skills are visible to all members and cannot be restricted through role assignments

Configuration Steps

Open Role Settings

From the left sidebar, select “**Organization Settings**” → “**Organization Overview**,” then click “**Edit Role Permissions**.” Select the role you want to configure, then switch to the “**Assign Skills**” tab.

Set Skill Creation Permission


Select or clear “**Allow skill creation**” as needed to determine whether members with this role can create skills. Changes are saved immediately, and a success message appears after the update is complete.

Assign Skills

Click the “Assign Skills” button, select the skills you want to make available to this role, and confirm your selection.

Set Skill Permissions

In the assigned skills list, configure the read, edit, and delete permissions for each skill individually, or remove the assignment. Permission changes are saved immediately when selected or cleared. A success message appears after the update, and there is no separate “Save” button on the page.



Manage skill creation permission and individual skill permissions on the “Assign Skills” tab

分配技能

可選技能 共 1 項

搜尋可選技能

☐ 全選 已選 0 項

客戶摘要 (文件示範)

新增技能 >

< 移除技能

已選技能 共 1 項

搜尋已選技能

☐ 全選 已選 0 項

客戶摘要 (文件示範)

取消 確認

1

Select the skills to assign to the role from the available skills

For information about creating roles and assigning members, see [Role Permission Management](#).

Deleting Skills

In the skill list, click the **Delete** icon for the skill, then confirm to delete it.

After deleting a skill, AI assistants that had this skill bound will no longer be able to use it. This action cannot be undone.

Agent Teams

What is Agent Teams?

1. What is Agent Teams?

In the past, a single AI assistant had to handle everything: customer service routing, order inquiries, returns and exchanges, product recommendations — all crammed into one system prompt, one knowledge base, and one set of tools. As tasks piled up, this "jack-of-all-trades" assistant started dropping the ball — the system prompt grew too long, knowledge bases interfered with each other, and costs stayed high.

Agent Teams lets you replace "one generalist" with "a group of specialists." You can compose multiple existing AI assistants into a team, each with its own role (for example: a routing agent, an order agent, a returns agent, and a recommendation agent), then use a visual canvas to define "when should a question be handed to whom."

Users see **a single conversation entry point**; behind the scenes, the team automatically routes questions to the most suitable agent, as if the user were interacting with an entire professional team.

Value to you: Each agent retains its own independent knowledge base, tools, and model settings without interference. Only the agent actually handling the question loads the relevant content, saving tokens and making each agent more focused and accurate. When you add a new business scenario, you simply add an agent node and draw a connection — no need to rewrite existing agents.

2. How is this different from "single agent + skills"?

Aspect	Single Agent (with Skills)	Agent Teams
Division of work	One agent handles all scenarios	Multiple agents each responsible for their area of expertise
Knowledge & tools	All attached to the same agent, prone to interference	Each agent independently owns its own knowledge base, tools, and model
Model cost	Only one model for all tasks	Use lightweight models for simple tasks and powerful models for complex reasoning, configured separately
Collaboration	Cannot hand off or delegate between agents	Supports handoff and delegate between agents

Aspect	Single Agent (with Skills)	Agent Teams
Scalability	New scenarios require modifying the existing agent's settings	Just add a new node and draw a connection

3. Three interaction types: Handoff, Delegate, Auto

Agents in a team collaborate through **connections (directed arrows)**. Each connection can be configured with an interaction type that determines "how control flows once the question reaches the other agent." The three types are explained below, with accompanying diagrams.

Handoff

Handoff **transfers control of the entire conversation**: the source agent determines that someone else should handle this matter and passes the conversation to the target agent, who **takes over** and continues talking with the user — like a customer service rep transferring a call to another department, after which the new person takes over.

Best for: Parallel division of work where each agent handles its own domain. For example, a routing agent determines the user's intent and hands the conversation off to an order inquiry agent or a returns agent.



Handoff: The source agent transfers control, and the target agent takes over the conversation with the user

Delegate

Delegate **outsources a subtask**: the source agent asks the target agent to handle something specific, the target agent completes it and **returns the result**, while control remains with the source agent — like a colleague looking up some data for you and reporting back, with you still being the one who compiles the final response.

Best for: Hierarchical collaboration. For example, a customer service agent receives an inventory check request, delegates the inventory query to a warehouse agent, and then compiles the response for the customer.



Delegate: The source agent sends a subtask to the target agent and compiles the response after getting the result back

Auto (Default)

Auto lets the AI **decide on the spot** whether to use handoff or delegate based on the current situation. Use this when you are not sure which type to choose or want to keep things flexible.



Auto: The AI decides on the spot whether to hand off or delegate

Communication volume and cost differences across the three types

When multiple agents collaborate, the most common concern is: **will token costs spiral out of control?**
The key lies in how much content is passed between agents across the three types:

Aspect	Handoff	Delegate	Auto
Control	Transferred to the target agent	Stays with the source agent	Decided by the AI on the spot
Content passed between agents	The target agent reads the accumulated conversation context when taking over	The source agent only sends "one subtask" to the target agent — concise and well-scoped	Depends on the type chosen
Token/cost characteristics	Grows with conversation length and number of handoffs; only one agent runs at a time	Adds a sub-agent run on top of the current turn (one extra LLM call), but the sub-agent's input is concise and scope is controllable	Falls on whichever type is selected, harder to predict in advance
Predictability	Medium	High (task scope is well-defined)	Low (varies by situation)

Costs are controllable and observable. The team has multiple built-in safeguards: Max iterations (default 25): Limits how many times agents can go back and forth in a single conversation, preventing infinite loops. Overall token limit (default 200000): The token budget cap for the entire team processing a single message. Timeout, and per-node configurable per-call token limits. Context strategy: Turn off "visible original user input" for a node so that agent only receives the summarized output from upstream agents rather than the full conversation, further saving tokens. Execution traces: Every execution records total tokens and cost breakdown per agent, making it clear where tokens are spent (see [Team execution traces](#)).

The conversation "remembers" who is currently serving. After a handoff, the receiving agent continues serving the subsequent conversation until it determines another handoff is needed; the next message does not go back to the entry point for re-routing.

4. Typical use cases

E-commerce customer service team: A routing agent determines intent and **hands off** to order returns recommendation agents, each with their own knowledge base and tools.

Enterprise internal helpdesk: HR agent + IT agent + Finance agent, automatically routing employee questions to the appropriate desk.

Technical support team: L1 agent performs initial troubleshooting, **delegates** to L2 agent for deeper analysis, and the result is returned to L1 for a compiled response.

5. How it works (concepts)

Agent Teams uses a **directed graph** to describe the team:

Nodes (agents): Each node on the graph is an existing AI assistant, carrying its own system prompt, knowledge base, tools, and model.

Entry node: When the team receives a message, processing starts from this node. A team has exactly one entry node.

Connections (directed arrows): Define "from which agent, under what conditions, using which interaction type, to which agent."

Guardrails: The team has upper limits for max iterations, overall token limit, and timeout to prevent agents from looping endlessly.

When you create a team, the system automatically creates an **entry agent** and a corresponding **Web Chat channel** for the team, allowing you to chat with the entire team directly through the channel.

6. Next steps

Ready to create a team? See [Create and orchestrate teams](#) — how to create a new team, drag in agent nodes, draw connections to set interaction types, designate an entry point, and start a conversation.

Want to see how the team operates? See [Team execution traces](#) — which agents were involved in each execution, how many tokens each used, and how long it took.

Prerequisite: Agent Teams must be enabled by MaiAgent for your organization. If you do not see the "Teams" menu under "AI Features" in the admin panel, contact the MaiAgent team to activate it.

Create and orchestrate teams

This page explains how to build an Agent Team from scratch: create a team, add agents on the canvas, configure connections and interaction types between them, designate the entry node, and start chatting with the entire team.

Prepare the AI assistants you want to add to the team first. Each node in a team is an existing AI assistant, so create your specialized assistants (e.g., routing, order, returns agents) under "AI Assistants" first, then come back here to compose them into a team.

1. Create a team

1. Go to the team list

From the left menu, navigate to **AI Features** > **Teams**. The list shows all teams in your organization, including **team name**, **number of agents**, and **status**. Use the search bar in the upper right to find a team.



Team list: view each team's name, number of agents, and activation status

2. Create a team

Click **Create Team** in the upper right corner, fill in the **Team name** (required) and **Description**, then create.



新建團隊

① 團隊名稱為必填，說明選填

* 團隊名稱

電商客服團隊

說明

分流、訂單查詢、退換貨、商品推薦，各由專屬助理協作處理


取消 確認

Create team: enter the team name and description

After successful creation, the system automatically creates an entry agent and a corresponding Web Chat channel for this team, so you can chat with the entire team right away.

2. Orchestrate the team on the canvas

Click into a team to enter the default **Canvas Orchestration** view. The screen is divided into three areas: the **Agent List** on the left, the **Canvas** in the center, and the **Settings Panel** on the right (appears when you select a node or connection).



Zoomed-in team graph: the entry node "Customer Service Routing Agent" uses "handoff" to route to Order Inquiry Returns

Product Recommendation agents, and the Order Inquiry agent "delegates" to the Inventory Query agent. Each node displays the agent name, model, and Agent/RAG mode; the entry node is labeled "Entry" (the actual screen also has the Agent List on the left and the Settings Panel on the right)

Canvas orchestration is best done on a desktop browser. On small screens like mobile devices, the canvas area will show a prompt asking you to switch to desktop for orchestration.

1. Add agent nodes

In the **Agent List** on the left, find the agent you want to add and **drag it onto the canvas**, or **click** it to add. Agents already in the team are labeled "Already in team" — clicking them will navigate to the corresponding node on the canvas. Use the search box at the top to filter when you have many agents.

Each node displays the agent name, the model it uses, and whether it is in **Agent** or **RAG** mode for easy identification.

If node positions are messy, click **Auto Layout** on the canvas to have the system rearrange them. Node positions are saved automatically.

2. Configure nodes

Click a node to open the **Node Settings** panel on the right. You can adjust:

Setting	Description
Entry node	When enabled, the team starts processing conversations from this node. A team can have only one entry node.
Per-call token limit	Limits the token count for a single response from this node (optional).
Max iteration count	Limits the number of iterations for a single processing cycle on this node (optional).
Context strategy	Controls what this agent can see: " Visible original user input " determines whether it can read the user's original message (when off, it only sees the processed output from upstream agents); " Visible output sources " lets you specify which nodes' outputs it can see.

目前設定的節點： 客服分流助理

```
graph LR; A[客服分流助理 Claude 4.5 Flanex Agent] -- "轉交 - 使用者想要訂單..." --> B[訂單查詢助理 Claude 4.0 Sonnet Agent]; A -- "轉交 - 使用者要退貨、退..." --> C[退貨貨助理 Claude 4.0 Sonnet Agent]; A -- "轉交 - 使用者想採購..." --> D[商品推薦助理 Claude 4.0 Flanex Agent];
```

節點設定

Chatbot

客服分流助理

連線

→ 訂單查詢助理

移交

→ 退貨貨助理

移交

→ 商品推薦助理

移交

入口節點

團隊收到訊息時，從此節點開始處理對話

☒

每次呼叫 Token 上限

最大迭代次數

可見原始使用者輸入

此節點可讀取使用者的原始輸入，而非僅上游節點的輸出

☒

儲存

刪除

Node settings: designate the entry node, set token/iteration limits, and configure the context strategy

3. Draw connections and set interaction types

Drag from one node's connection point to another node to create a **directed connection** (source -> target). Click the connection to open the **Connection Settings** panel on the right:

Setting	Description
Trigger condition description	Describe in plain language "when should this connection be used." This description is provided to the AI as a basis for routing decisions — writing it clearly helps improve routing accuracy.
Interaction type	Choose Auto , Handoff , or Delegate (see the table below for differences).
Inherit source agent's tools	When enabled, the target agent can use all tools and resources from the source agent.

Interaction types (descriptions as shown in the settings panel):

Interaction type	Description
Auto (default)	The LLM automatically determines the most suitable interaction type.
Handoff	Hand off control — the source agent goes idle and the target agent takes over the conversation.
Delegate	Delegate a task — the target agent executes and returns the result to the source.

目前設定的連線：委派（訂單查詢 → 倉儲查詢）

✖

訂單查詢助理

Claude 4.5 Sonnet - Agent

↓

✖

倉儲查詢助理

Claude 4.5 Sonnet - Agent

互動方式

☐ 自動（預設）
由 LLM 自動判斷最適合的互動方式

☐ 移交
移交控制權 — 來源代理休眠，目標代理接管對話

☒ 委派
委派任務 — 目標代理執行後回傳結果給來源

觸發條件描述

需要查詢商品庫存數量時，委派倉儲查詢助理查詢後回傳

繼承來源 Agent 的工具

開啟後，目標 Agent 可使用來源 Agent 的所有工具和資源

☒

Connection settings: fill in the trigger condition description, select the interaction type (auto/handoff/delegate), and optionally inherit the source agent's tools

Watch for semantic warnings at the top of the canvas. When the team graph has issues, a yellow notice appears at the top of the canvas, such as: "No entry node set," "Multiple entry nodes — keep only one," "Duplicate connections," or "Connections form a cycle." It is recommended to resolve these warnings before starting a conversation.

3. Configure basic information and guardrails

In the team page, switch to the **Basic Information** view to adjust team-level settings:


Basic information: Team name, description, and whether it is enabled.

Guardrails:

Max iterations (default **25**) — The maximum number of times agents can go back and forth in a single conversation, preventing infinite loops.

Overall token limit (default **200000**) — The token budget cap for the entire team processing a single user message.

Timeout (default **600** seconds) — The hard timeout for the entire team's processing flow.



Basic information and guardrails: adjust team name, activation status, and max iterations, overall token limit, and timeout

4. Chat with the team

Once orchestration is complete and the team is "Enabled," you can chat with the entire team through the team's **Web Chat channel**. Users face a single entry point, and the system automatically routes questions to the appropriate agent.

When the team performs handoffs or delegates between agents, the chat interface displays transition messages such as "**Handing off to,**" "**Delegating to,**" or "**Switching to,**" letting users know which agent is currently serving them.

5. Delete a team

In the danger zone at the bottom of the [Basic Information](#) view, you can [Delete Team](#). Deleting a team also removes its entry agent, and **this action cannot be undone** — please confirm before proceeding.

6. Next steps

Once the team is running, you can check the [Team execution traces](#) to see which agents were involved in each execution, how many tokens each used, and how long it took — confirming the team is dividing work as expected.

Team execution traces

Teams automatically hand off and delegate between agents, and the process is seamless to users. **Execution traces** let you go back and review "which agents were involved and how many resources each consumed" for every conversation, helping you confirm the team is dividing work as expected and making it easy to troubleshoot and optimize.

1. View execution traces

In the team page, switch to the **Execution Traces** view to see a list of every execution for this team. Each record includes:

Field	Description
Trace ID	The unique identifier for this execution.
Status	"Running," "Completed," "Failed," or "Timed out."
Iterations	The total number of back-and-forth exchanges between agents during this execution.
Token count	The total token usage for this execution.
Duration	The total time from receiving the message to completing the response.

ID	Status	Iteration Count	Token Count	Cost	Creation Time
tr-427088-1-1-4	已完成	1	2,407,123 out	4.05	2023-10-10
tr-427088-1-2-4	已完成	2	1,000,000 in, 2,000,000 out	10.25	2023-10-10
tr-427088-1-3-4	已完成	3	1,000,000 in, 2,000,000 out	10.25	2023-10-10
tr-427088-1-4-4	已完成	4	1,000,000 in, 2,000,000 out	10.25	2023-10-10
tr-427088-1-5-4	已完成	5	1,000,000 in, 2,000,000 out	10.25	2023-10-10
tr-427088-1-6-4	已完成	6	1,000,000 in, 2,000,000 out	10.25	2023-10-10
tr-427088-1-7-4	已完成	7	1,000,000 in, 2,000,000 out	10.25	2023-10-10
tr-427088-1-8-4	已完成	8	1,000,000 in, 2,000,000 out	10.25	2023-10-10
tr-427088-1-9-4	已完成	9	1,000,000 in, 2,000,000 out	10.25	2023-10-10
tr-427088-1-10-4	已完成	10	1,000,000 in, 2,000,000 out	10.25	2023-10-10

Execution traces: status, iteration count, token count, and duration for each team execution

2. Expand step details

Expanding a single record reveals the **Step Details** — which agent nodes this execution **passed through in order**, along with the details of each step:

Item	Description
Step number	Which step this is in the overall flow and which agent executed it.
Input	A summary of the content this agent received at this step.
Output	A summary of the content this agent produced at this step.
Cost breakdown	The token usage for this step (input tokens, output tokens, embedding tokens), letting you see exactly which agent the cost is attributed to.

The screenshot displays the AgentOps dashboard. On the left is a sidebar with navigation options like 'All Agents', 'Agent Groups', 'Agent Traces', 'Agent Settings', 'Agent Tools', 'Agent Logs', 'Agent Metrics', 'Agent Alerts', 'Agent Audit', 'Agent Help', 'Agent Settings', 'Agent Tools', 'Agent Logs', 'Agent Metrics', 'Agent Alerts', 'Agent Audit', 'Agent Help'. The main area shows a table of execution traces with columns: Trace ID, Agent, Status, Iterations, Tokens, Cost, and Time. Below the table, a detailed view of a specific trace is shown, including a 'Steps' section with a list of steps and their details.

Trace ID	Agent	Status	Iterations	Tokens	Cost	Time
tr-4f7088-b010-4...	Agent A	Success	1	2,407,122	\$0.00	2023-06-29 14:00
tr-802207f4e4e4-0f...	Agent B	Success	6	11,084,111	\$0.00	2023-06-29 14:01

Steps:

- Step 1: [Details]
- Step 2: [Details]
- Step 3: [Details]
- Step 4: [Details]

At the bottom, there is a pagination bar showing '1' of 12 items.

Step details: see which agents were involved step by step, with input/output summaries and token cost breakdown for each

Not seeing any traces? Execution traces are only generated after the team is actually triggered by a conversation. If the list shows "No execution traces yet," chat with the team through the channel first, then come back to check.

3. How to use these traces to optimize your team

Routing not as expected: If questions are being sent to the wrong agent, go back to the canvas orchestration and adjust the connection's "trigger condition description" to be more explicit about "when to use this connection."

Iterations or duration too high: Agents may be handing off back and forth too many times. Check for unnecessary connections, or adjust the team's "max iterations" and per-node settings.

Cost concentrated on one agent: Use the "cost breakdown" to find the agent consuming the most tokens, and consider switching it to a lighter model or narrowing its knowledge base and tool scope.

AgentOps

AgentOps Overview

What is AgentOps?

AgentOps (Agent Operations) is the AI Assistant operation management module provided by the MaiAgent platform, focusing on testing, evaluating, and monitoring the actual performance of AI Assistants. With AgentOps, you can systematically manage AI Assistant quality and ensure consistently excellent service experiences.

Why Do You Need AgentOps?

Building an AI Assistant is just the first step — continuous monitoring and optimization are essential for long-term success. AgentOps helps you:

Ensure Consistent Quality

- * Establish standardized testing mechanisms
- * Track AI Assistant performance trends
- * Identify and fix quality issues promptly

Improve Operational Efficiency

- * Automate testing processes to save manpower
- * Quickly identify performance bottlenecks
- * Optimize system resources and costs

Data-Driven Decision Making

- * Quantify AI Assistant performance metrics
- * Compare the effectiveness of different settings or models
- * Develop optimization strategies based on actual data

AgentOps Core Features

1. Test Dataset Management

Create and manage test case collections to verify AI Assistant response quality.

Key Features:

- * Create multiple test datasets
- * Manage test cases and expected responses

- * Support search and categorization
- * Team collaboration on test case maintenance

Use Cases:

- * Verify response accuracy after Knowledge Base updates
- * Establish standard testing workflows
- * Collect frequently asked questions as test benchmarks

For details: [Test Dataset Management](#)

2. Automated Evaluation

Use test datasets to run automated evaluations and generate detailed quality reports.

Key Features:

- * One-click batch testing
- * Calculate success rates and response times
- * Generate detailed pass/fail reports
- * Compare performance across different periods or AI Assistants

Use Cases:

- * Regular quality checks
- * Regression testing before and after updates
- * Compare effectiveness of different AI Assistants or models

For details: [Automated Evaluation](#)

3. AI Assistant Monitoring

Monitor AI Assistant operations in real time and analyze every conversation in detail.

Key Features:

- * View input/output message content
- * Track reply time and processing time
- * Monitor token usage and costs
- * Analyze quality scores and user feedback

Use Cases:

- * Real-time system operation monitoring
- * Identify performance issues and anomalies
- * Track costs and usage
- * In-depth analysis of individual conversations

For details: [AI Assistant Monitoring](#)

4. Evaluation Insight Report

After batch testing is complete, the system automatically analyzes the results and generates an improvement recommendations report — no need to manually interpret each scoring data point.

Key Features:

- * Automatically identify problem patterns in test results
- * Suggest handling priorities based on severity and scope of impact
- * Provide concrete improvement directions, not just scores
- * Support multilingual reports

Use Cases:

- * Success rate falls below expectations and you need to know where the problems are
- * A single test run produces many failed cases and you need to decide which to fix first
- * You need to explain the current quality status and improvement plan to your team or management

For details: [Evaluation Insight Report](#)

5. Tool Execution Records

Track the complete history of AI Assistant external tool invocations and quickly pinpoint the cause of tool execution failures.

Key Features:

- * Record input parameters, execution results, duration, and success/failure status for each tool invocation
- * Retain detailed error messages to help pinpoint failure causes
- * Track usage count and success rate for each tool
- * Cover API tools, Text-to-SQL tools, crawler tools, and custom tools

Use Cases:

- * AI Assistant cannot answer questions that require querying external systems
- * Tools fail intermittently and you need to determine whether it's a timeout, authentication, or parameter issue
- * Evaluating whether a tool needs configuration adjustments or replacement

For details: [Tool Execution Records](#)

AgentOps Workflow

Sub-feature Relationship Map

AgentOps has five sub-features divided into two tracks: **Pre-launch** uses test datasets and automated evaluation to establish quality benchmarks, while **Post-launch** uses monitoring and tool records to observe real operations. The two tracks form a cycle by "feeding issues discovered in production back into the test datasets."

1. Test Dataset Management	Create reusable test case collections — the starting point of the entire testing chain.	<u>Test Dataset Management</u>
2. Automated Evaluation	Batch test AI Assistants using test datasets, producing success rates and average response times.	<u>Automated Evaluation & AI Assistant Monitoring</u>
3. Evaluation Insight Report	Automatically analyze evaluation results into problem patterns and improvement suggestions, ranked by priority.	<u>Evaluation Insight Report</u>
4. AI Assistant Monitoring	Inspect real conversations one by one after launch — inputs, outputs, duration, tokens, and scores.	<u>Automated Evaluation & AI Assistant Monitoring</u>
5. Tool Execution Records	Track AI Assistant external tool invocation parameters, results, and failure causes.	<u>Tool Execution Records</u>

Sequence and dependencies of each sub-feature:

Sub-feature	Stage	Prerequisites	Output	Next Step
Test Dataset Management	Pre-launch	Common questions and expected responses	Reusable test case collections	Hand off to Automated Evaluation
Automated Evaluation	Pre-launch	Test datasets, AI Assistant to test	Success rate, average response time, per-case pass/fail	Hand off to Evaluation Insight Report for analysis
Evaluation Insight Report	Pre-launch	A completed automated evaluation	Problem patterns, priorities, concrete improvement directions	Adjust Knowledge Base or role instructions based on suggestions, then test again
AI Assistant Monitoring	Post-launch	Real conversations from a live AI Assistant	Input/output, duration, token usage, scores for each conversation	Feed discovered issues back into test datasets
Tool Execution Records	Post-launch	AI Assistant with configured tools	Tool invocation parameters, results, duration, and error messages	Fix tool settings, or add a test case

Two commonly confused points: "Automated Evaluation" tests questions from the test datasets, while "AI Assistant Monitoring" examines conversations actual users have had — the former is a pre-launch quality gate, the latter is post-launch observation. "Evaluation Insight Report" is not a standalone feature — it is an extension of Automated Evaluation: you must first complete a batch test before an insight report is available.

Standard Quality Management Flow

1. Create Test Datasets
↓
Collect common questions and expected responses
Build structured test cases
2. Run Automated Evaluations
↓
Run tests regularly or after updates
Generate success rate and quality reports
3. Analyze Evaluation Results
↓
Identify failed cases
Find issues that need improvement
4. Optimize AI Assistant
↓
Update Knowledge Base content
Adjust AI settings and prompts
5. Monitor Actual Operations
↓
Use AI Assistant Monitoring to track real conversations
Continuously ensure quality stability
6. Continuous Improvement
↓
Discover new issues from monitoring data
Update test datasets
Repeat the evaluation and optimization cycle

Quick Start Steps

Week 1: Build Foundational Tests

1. Collect 20-30 frequently asked questions
2. Create your first test dataset

3. Run the initial automated evaluation
4. Record the current success rate as a baseline

Week 2: Continuous Optimization

1. Analyze failed test cases
2. Update the Knowledge Base or adjust AI settings
3. Re-run evaluations to verify improvements
4. Start using AI Assistant Monitoring to observe real conversations

Week 3: Build Habits

1. Run automated evaluations once a week
2. Check monitoring data daily to identify anomalies
3. Continuously add test cases to cover more scenarios
4. Create quality reports to track long-term trends

AgentOps Integration with Other Features

AgentOps + Knowledge Base Management

Integration Use Cases:

- * Run automated evaluations immediately after Knowledge Base updates
- * Supplement Knowledge Base content when monitoring data reveals knowledge gaps
- * Create test cases based on real conversations to optimize the Knowledge Base

Related documentation: [Build an Enterprise Knowledge Base](#)

AgentOps + Usage Analytics

Complementary Analysis:

- * **Usage Analytics:** Provides overall trends and statistics (conversation volume, satisfaction, etc.)
- * **AgentOps:** Provides in-depth quality analysis and individual conversation details

Combined use enables you to:

- * Detect anomalies from overall trends
- * Identify root causes from monitoring details
- * Comprehensively evaluate AI Assistant performance

Related documentation: [Usage Analytics](#)

AgentOps + Response Quality Control

Integration Use Cases:

- * AgentOps provides technical quality monitoring

- * Response Quality Control provides manual review and annotation
- * Together they form a complete quality assurance system

Related documentation: [Response Quality Control](#)

Best Practices

1. Establish a Regular Evaluation Schedule

Daily Check (5 minutes):

- * Review AI Assistant Monitoring to identify anomalies

Weekly Evaluation (30 minutes):

- * Run core test datasets
- * Check whether success rates remain stable
- * Update or add test cases

Monthly Analysis (2 hours):

- * Run a full evaluation
- * Export monitoring data for in-depth analysis
- * Produce quality reports
- * Plan optimization goals for the next month

2. Set Quality Baselines

Set reasonable quality targets for your AI Assistant:

Metric	Baseline	Target	Description
Automated Evaluation Success Rate	≥ 85%	≥ 95%	Standard test case pass rate
Average Response Time	< 5 seconds	< 3 seconds	User-perceived response speed
User Satisfaction	≥ 80%	≥ 90%	Like total feedback ratio
Token Usage	-	-	Set limits based on budget

3. Establish an Issue Tracking Workflow

When quality issues are discovered:

```
Issue Discovered
↓
Record in test dataset (prevent recurrence)
↓
Analyze root cause (Knowledge Base
  AI settings
  model capability)
↓
Implement improvements
↓
Run evaluations to verify effectiveness
↓
Continue monitoring to confirm stability
```

4. Team Collaboration

Establish AgentOps responsibility assignments:

AI Manager:

- * Overall quality monitoring
- * Develop optimization strategies
- * Regular evaluations and reporting

Knowledge Base Administrator:

- * Update the Knowledge Base based on evaluation results
- * Supplement missing knowledge content

Technical Staff:

- * Analyze performance issues
- * Optimize system settings
- * Handle technical anomalies

5. Continuous Improvement Cycle

The core of AgentOps is continuous improvement:

Measure: Obtain data through evaluations and monitoring

Analyze: Identify issues and improvement opportunities

Act: Implement optimization measures

Verify: Confirm improvement effectiveness

Standardize: Incorporate successful practices into workflows

Q1: What is the difference between AgentOps and "Usage Analytics"?

Usage Analytics:

- * Overall trends and statistics (conversation volume, word count, satisfaction)
- * Suitable for understanding AI Assistant usage overview
- * Focuses on business metrics

AgentOps:

- * In-depth quality analysis and individual conversation details
- * Suitable for technical quality management
- * Focuses on technical metrics (success rate, response time, cost)

The two are complementary — we recommend using them together.

Q2: How often should evaluations be run?

Recommended frequency:

- * **Core features:** Once a week
- * **Full tests:** Once a month
- * **After major updates:** Run immediately
- * **When issues are found:** Test and verify at any time

Adjust based on your AI Assistant's importance and update frequency.

Q3: How many test cases should a test dataset contain?

Recommended amounts:

- * **Minimum:** 20 test cases (covering core features)
- * **Recommended:** 50-100 test cases (balancing coverage and efficiency)
- * **Comprehensive:** 200+ test cases (for large or critical systems)

Start with core features and expand gradually.

Q4: What success rate is considered acceptable?

It depends on the AI Assistant's purpose:

- * **High-risk applications** (finance, healthcare): $\geq 95\%$
- * **Customer support:** $\geq 90\%$
- * **General conversation:** $\geq 85\%$
- * **Experimental features:** $\geq 80\%$

The key is to establish a baseline and continuously improve.

Q5: Does AgentOps affect real users?

No. AgentOps testing and monitoring run in isolated environments or in the background, without interfering with the actual user experience.

Q6: Can I create multiple test sets for different business areas?

Yes. We recommend creating multiple test datasets for different business areas or features:

- * Product inquiry test set
- * Order processing test set
- * Technical support test set
- * ...and more

This allows you to track quality performance separately for each area.

Get Started with AgentOps

Ready to start using AgentOps? We recommend following this order:

Create Test Datasets: Collect common questions and create your first test set

Run Automated Evaluations: Test your AI Assistant and review success rates and response times

Read the Evaluation Insight Report: Understand where the problems are and what to prioritize

Use AI Assistant Monitoring: Gain in-depth understanding of actual operations after launch

View Tool Execution Records: If your AI Assistant has configured tools, track whether tool invocations are working properly

With AgentOps, you can establish a systematic quality management workflow to ensure your AI Assistant consistently delivers excellent service.

Test Dataset Management

What Is a Test Dataset?

A test dataset is a collection of pre-prepared test cases used to validate the response quality and accuracy of AI assistants. By creating standardized test datasets, you can:

Ensure Response Quality

- * Verify that the AI assistant's responses to common questions meet expectations
- * Ensure responses don't degrade after knowledge base updates
- * Continuously track response quality trends

Improve Efficiency

- * Batch test multiple questions without manual verification one by one
- * Create reusable test cases to save testing time
- * Quickly identify issues and make timely adjustments

Team Collaboration

- * Centrally manage all test cases
- * Different team members can jointly maintain test data
- * Track who created and last updated test cases

How to Access Test Dataset Management

Go to "**AgentOps**" in the left sidebar, then click "**Test Datasets**" to see the test dataset management page.

Test dataset management interface

The page displays the following information:

- * **Test Dataset ID:** A unique identifier automatically generated by the system
- * **Test Dataset Name:** The name you assigned to the dataset
- * **Test Case Count:** The total number of test cases in the dataset
- * **Created By:** The user who created this dataset
- * **Created At:** The date and time the dataset was created
- * **Last Updated:** The time the dataset was last modified
- * **Action Buttons:** Edit, delete, and other functions

Creating a Test Dataset

1. Start Creating

Click the "**Create Dataset**" button in the upper right corner to start creating a new test dataset.

2. Set Basic Information

In the dialog that appears, enter the following information:

Test Dataset Name (required)

- * Use a meaningful name, such as: "Product Query Test Set" or "Customer Service FAQ Test"
- * This makes it easier to manage and identify purposes later

Test Dataset Description (optional)

- * Describe the purpose and coverage of this test set
- * Example: "For testing product specification query functionality, containing 50 common product questions"

Select AI Assistant (required)

- * Choose the AI assistant to test
- * A test dataset can correspond to one or more AI assistants

3. Add Test Cases

After creating the test dataset, you need to add test cases. Each test case contains:

Test Question

- * Questions that users might ask
- * Recommended to cover various phrasings and scenarios

Expected Response

- * How the AI assistant should answer
- * Can be keywords, sentences, or complete paragraphs
- * The system compares whether the actual response contains the expected content

Test Case Examples:

Question: What is your return policy?

Expected Response: Unconditional returns within 7 days of purchase, items must remain new and unused

Question: How do I apply for a refund?

Expected Response: Apply for a refund through the order management page in the member center, refund will be processed within 3-5 business days after approval

Question: Who pays for return shipping?

Expected Response: We cover shipping for defective products, buyers cover shipping for personal preference returns

Managing Test Datasets

Search Test Datasets

Use the search bar at the top of the page to quickly find target datasets by:

- * Test dataset ID
- * Test dataset name
- * Description keywords

After entering keywords, the system filters and displays matching datasets in real time.

Edit Test Datasets

Click the "**Edit**" button in the action column to:

- * Modify the test dataset name and description
- * Add, edit, or delete test cases
- * Adjust the order of test cases

Regularly review and update test cases to ensure test content stays aligned with actual business needs

Delete Test Datasets

Click the "**Delete**" button in the action column to delete test datasets that are no longer needed.

Deleted test datasets cannot be recovered. Proceed with caution. If the dataset has been used in automated evaluations, the related evaluation records will be retained, but the same test cannot be run again.

Test Dataset Best Practices

1. Create Structured Test Categories

Group test cases by business needs:

Product-related knowledge base tests:

- Product Query Test Set (50 cases)
- Price Inquiry Test Set (30 cases)
- Inventory Verification Test Set (20 cases)

Customer service workflow tests:

- Order Processing Test Set (40 cases)
- Return & Exchange Workflow Test Set (35 cases)
- Membership Service Test Set (25 cases)

2. Cover Multiple Phrasings

The same concept can be phrased in different ways. Include various phrasings in your test cases:

Formal phrasing: Could you please tell me your business hours?

Colloquial phrasing: What time do you open?

Short phrasing: Business hours

Full sentence: I'd like to know what your business hours are, from when to when

3. Include Edge Cases

Test extreme situations and ambiguous questions:

- * Incomplete questions: "price," "what do I do"
- * Questions with typos: "prodcut specs"
- * Compound questions: "How much does it cost? Can I return it? What colors are available?"
- * Out-of-scope questions: Test how the AI assistant responds

4. Regularly Update Test Data

When new features launch: Add corresponding test cases

After knowledge base updates: Check whether related test cases need adjustment

When issues are discovered: Immediately add test cases for that scenario to prevent recurrence

Regular reviews: Review test cases monthly or quarterly to ensure they still reflect current conditions

5. Set Appropriate Test Quantities

Recommended test cases per dataset:

- * **Minimum 20 test cases:** Ensures coverage of basic scenarios
- * **Recommended 30-50 test cases:** Balances test coverage and execution time
- * **100+ cases for special situations:** Suitable for complex business logic

Frequently Asked Questions

Q1: What is the relationship between test datasets and automated evaluations?

Test datasets are the foundation of automated evaluations. You need to create a test dataset first, then select the test dataset and AI assistant in the "Automated Evaluation" feature. The system automatically runs the tests and generates an evaluation report.

For more information, see: [Automated Evaluations](#)

Q2: Can one test dataset be used to test multiple AI assistants?

Yes. You can select the same test dataset with different AI assistants during automated evaluation to compare how different AI assistants perform on the same test cases.

Q3: Does the expected response need to match exactly?

No. The system checks whether the AI assistant's actual response "contains" the key content of the expected response. You can set keywords or important sentences, and the test passes as long as the actual response includes that content.

Q4: How do I import a large number of test cases?

Currently, you can batch-add test cases by copying and pasting. Future versions will support Excel or CSV file import — stay tuned.

Q5: Can test datasets be shared with other organization members?

Yes. All members within the same organization can view and use test datasets. You can identify who created a dataset through the "Created By" column to facilitate team collaboration.

Q6: If I delete an AI assistant, will the related test datasets disappear?

No. Test datasets are managed independently. Even if you delete an AI assistant, the test datasets are retained. You can use the same test dataset to test other AI assistants.

Next Steps

After creating test datasets, you can:

- * Go to [Automated Evaluations](#) to run tests and view results
- * Use [AI Assistant Monitoring](#) to track the AI assistant's actual operational performance
- * Adjust knowledge base content or AI assistant settings based on evaluation results


Automated Evaluation & AI Assistant Monitoring

Automated Evaluation

The automated evaluation feature allows you to use pre-built test datasets to automatically test AI Assistant response quality. The system sends test questions to the AI Assistant, compares actual responses with expected responses, and generates detailed evaluation reports.

Access Automated Evaluation

Go to "[AgentOps](#)" in the left sidebar, then click "[Automated Testing](#)".



测试ID	测试名称	操作	测试集	AI助手	成功率	操作
60d1449b-b9e0-4d11-8f31-6b0000000000	测试集: 文件操作 - 文件管理测试 - 基础: 测试 Agent 处理文件操作...	AI 文件管理测试	测试集: 文件操作	MaxAgent 智能助理	100%	操作
1a259121-5109-47f5-b101-6b0000000000	测试集: 文件操作 - 基础: 测试集...	基础: 测试 Agent 处理文件操作...	测试集: 文件操作	MaxAgent 智能助理	100%	操作
90d1449b-b9e0-4d11-8f31-6b0000000000	AI 文件管理测试 - 测试集...	测试集: 文件操作	测试集: 文件操作	MaxAgent 智能助理	100%	操作
71d1449b-b9e0-4d11-8f31-6b0000000000	测试集: 文件操作 - 基础: 测试集...	基础: 测试 Agent 处理文件操作...	测试集: 文件操作	MaxAgent 智能助理	100%	操作
20d1449b-b9e0-4d11-8f31-6b0000000000	测试集: 文件操作 - 基础: 测试集...	基础: 测试 Agent 处理文件操作...	测试集: 文件操作	MaxAgent 智能助理	100%	操作
30d1449b-b9e0-4d11-8f31-6b0000000000	测试集: 文件操作 - 基础: 测试集...	基础: 测试 Agent 处理文件操作...	测试集: 文件操作	MaxAgent 智能助理	100%	操作
40d1449b-b9e0-4d11-8f31-6b0000000000	测试集: 文件操作 - 基础: 测试集...	基础: 测试 Agent 处理文件操作...	测试集: 文件操作	MaxAgent 智能助理	100%	操作
50d1449b-b9e0-4d11-8f31-6b0000000000	测试集: 文件操作 - 基础: 测试集...	基础: 测试 Agent 处理文件操作...	测试集: 文件操作	MaxAgent 智能助理	100%	操作
60d1449b-b9e0-4d11-8f31-6b0000000000	测试集: 文件操作 - 基础: 测试集...	基础: 测试 Agent 处理文件操作...	测试集: 文件操作	MaxAgent 智能助理	100%	操作
70d1449b-b9e0-4d11-8f31-6b0000000000	测试集: 文件操作 - 基础: 测试集...	基础: 测试 Agent 处理文件操作...	测试集: 文件操作	MaxAgent 智能助理	100%	操作
80d1449b-b9e0-4d11-8f31-6b0000000000	测试集: 文件操作 - 基础: 测试集...	基础: 测试 Agent 处理文件操作...	测试集: 文件操作	MaxAgent 智能助理	100%	操作
90d1449b-b9e0-4d11-8f31-6b0000000000	测试集: 文件操作 - 基础: 测试集...	基础: 测试 Agent 处理文件操作...	测试集: 文件操作	MaxAgent 智能助理	100%	操作

Automated testing list showing success rate and average response time for each test run

The page displays all evaluation records, including evaluation name, test dataset, AI Assistant, success rate, average response time, and creation time.

Create and Run a Test

Ensure you have created a test dataset (refer to [Test Dataset Management](#))

Click the "[Create Test](#)" button

Fill in the evaluation name and description, then select the test dataset and AI Assistant

Click "[Start Evaluation](#)" and the system will automatically execute all test cases

Evaluation execution time depends on the number of test cases. Typically, 50 test cases take approximately 2-3 minutes.

Use Custom Evaluation Metrics

In addition to built-in metrics, you can define your team's own scoring criteria using natural language. For example, Mai, a customer service manager, wants to verify whether the customer service assistant remains polite and empathetic. She creates a "Customer Service Tone" metric, enters the evaluation criteria, and sets a passing threshold. After the evaluation is complete, she can view each case's "Customer Service Tone" score in the results and identify responses that need adjustment.

When creating an evaluation, complete the basic settings, including the name, test dataset, AI Assistant, and evaluation model.

In the "[Custom Metrics](#)" section, click "[Add Custom Metric](#)", then enter an easily identifiable metric name and evaluation criteria.

Adjust the passing threshold according to your team's quality standards, then click "[Confirm](#)" to start the evaluation.

單輪 Q&A

多輪對話

* 名稱

請輸入名稱

描述

請輸入描述

* 選擇測試集

請選擇測試集

* 選擇AI助理

請選擇AI助理

* 選擇評測模型 ⓘ

請選擇評測模型

最大輪數 ⓘ



6

評估指標

逐輪忠實度 (Turn Faithfulness)

每輪回覆是否忠於檢索到的內容



逐輪相關性 (Turn Relevancy)

每輪回覆是否切題



對話完成度 (Conversation Completeness)

整段對話是否達成使用者的目標



模擬語言

繁體中文



取消

確認

After switching to multi-turn conversations, you can set the maximum number of turns, three evaluation metrics, and the simulation language

建立評估

✕

★ 選擇AI助理

請選擇AI助理

★ 選擇評測模型 ⓘ

請選擇評測模型

自訂指標 ⓘ

+ 新增自訂指標

客服語氣

評估準則 評估步驟

回覆是否有禮貌、清楚且具同理心

15 / 2000

通過門檻 0.50

取消

確認

Set evaluation criteria for customer service tone using natural language

The system compares names without regard to letter case. A custom metric cannot duplicate any of the 11 built-in metrics. If a name conflicts, use a name that expresses your team's evaluation objective.

Evaluate Multi-turn Conversations

Multi-turn evaluation is ideal for testing tasks that require back-and-forth confirmation. For example, Mei, a quality manager, wants the assistant to ask for the order number first, confirm the reason for the return, and then explain the next steps. She creates a multi-turn test case, enters the scenario and expected result, and runs the evaluation. When it is complete, she can expand the conversation transcript to review every response and the Knowledge Base content retrieved during each turn, instead of looking only at the final answer.

First, go to "[Test Datasets](#)" and open the target test dataset. Add a scenario, expected result, and simulated user persona to the multi-turn test case.

Return to "[Automated Testing](#)", click "[Create Evaluation](#)", and switch to "[Multi-turn Conversation](#)".

Select the test dataset, AI Assistant, and evaluation model. Set the maximum number of turns, evaluation metrics, and simulation language, then start the evaluation. At least one evaluation metric must be enabled.

After the evaluation is complete, open the results and select an individual case to expand its details. Review the metric scores, conversation transcript, and "[Retrieved Content](#)" for each turn. If an individual case fails, its error is displayed separately without hiding the results of other cases.

[DIGIEN-DEWU] 使用選擇後填填驗證多輪結果檢查內容。

成功率 0%

總得分數 0.67

賽以第 1/1

最大輪數 2

指標分數

Turn Relevancy 1.00

Turn Faithfulness 1.00

Conversation Completeness 0.00

已通過 未通過

顧客要退還所訂單車頭燈袋，助理需先取得訂單編號與原因。
預期結果：助理應先查詢單車頭燈袋詳情並提供協助。

Turn Relevancy 1.00
得分原因: 0.00
The score is 1.0 because there were no listed irrelevant; on average, two tags contain information that is relevant to the user's requests.

Turn Faithfulness 1.00
得分原因: 0.00
The score is 1.0 because all evaluated claims were consistent with the retrieved truth, and no contradictions, unsupported statements, false claims, or image-related errors were found.

Conversation Completeness 0.00
得分原因: 0.00
The score is 0.0 because the AI response only said 「您好，請提供頭燈袋詳情，並協助您查詢單車頭燈袋詳情。」 and directed the user to contact merchant support, while ending 「請提供單車頭燈袋詳情，請提供單車頭燈袋詳情。」 and 「請提供單車頭燈袋詳情，請提供單車頭燈袋詳情。」. It therefore rather overplayed the requested return using the provided order number and reason rather than fully follow the requested information.

對話語句 (4)

助理問問

您好，請提供單車頭燈袋詳情，並協助您查詢單車頭燈袋詳情。

助理問問

您好，請提供單車頭燈袋詳情，並協助您查詢單車頭燈袋詳情。

助理問問

您好，請提供單車頭燈袋詳情，並協助您查詢單車頭燈袋詳情。

助理問問

您好，請提供單車頭燈袋詳情，並協助您查詢單車頭燈袋詳情。

Expand a case to review the assistant's response and the Knowledge Base content used in each turn

If you no longer need an ongoing multi-turn evaluation, click "**Cancel Evaluation**" in the evaluation details and confirm. The system stops processing subsequent cases after the current case is complete. Scores and transcripts for completed cases are retained, the status is shown as "**Partially Completed**", and unprocessed cases are shown as "**Pending Evaluation**".



Completed cases are retained after cancellation, while unprocessed cases remain pending evaluation

View Evaluation Results

Click on an evaluation record to view the detailed report:



Test details page: success rate, AI insights summary, metric statistics, and improvement suggestions

The detailed report includes:

- Success Rate:** Percentage of test cases that passed
- AI Insights:** Automatically generated evaluation summary and improvement suggestions
- Metric Statistics:** Quality scores, response relevance, and other metrics
- Test Case Details:** Expected response, actual response, score, and status for each question

Success Rate Reference Benchmarks:

AI Assistant Type	Recommended Success Rate
Product Inquiry Assistant	≥ 95%
Customer Support Assistant	≥ 90%
General Conversation Assistant	≥ 80%

For more information about insight reports, refer to: [Evaluation Insight Reports](#)

Manage Evaluation Records

Search and Filter: Filter records by test dataset, AI Assistant, or keywords

Re-run: After modifying the Knowledge Base or AI settings, re-run tests to verify improvements

Export: Export evaluation results in Excel format

Edit the Evaluation Name and Description

After a test is complete, you can still edit its evaluation name and description without deleting or rerunning the test. The test dataset, AI Assistant, evaluation results, and execution status remain unchanged.

For example, after Mai created a weekly customer service quality test, she noticed that the week number in the evaluation name was incorrect. She opened the edit dialog from the automated testing list, corrected the name, and added the test objective. After she saved the changes, the list immediately displayed the updated information while retaining the original test results.

Open the Automated Testing List


From the left menu, go to “AgentOps” → “Automated Testing” .

Open the Edit Dialog

Find the evaluation you want to update, then click the “Edit” icon in the “Actions” column for that row.

Edit and Save

Edit the “Evaluation Name” or “Description” , then click “Confirm” . The evaluation name is required and can contain up to 200 characters; the description may be left blank.



The edit dialog is prefilled with the current evaluation name and description


You can edit the name and description of evaluations that are running, completed, or failed. This action only updates the identifying information and does not rerun the test.

AI Assistant Monitoring

AI Assistant Monitoring provides real-time conversation operational data, allowing you to gain deep insight into the processing details, performance metrics, and quality scores of each conversation.

Access AI Assistant Monitoring

Go to "[AgentOps](#)" in the left sidebar, then click "[AI Assistant Monitoring](#)".



AI Assistant Monitoring interface showing detailed technical metrics for each conversation

Monitoring Field Descriptions

Field	Description
User Input Message	The question sent by the user to the AI
Output Message	The AI Assistant's response content
AI Assistant	The name of the AI Assistant that handled the conversation
User Feedback	Thumbs up 👍 or thumbs down 👎
Faithfulness Score	Whether the response is faithful to the Knowledge Base content
Response Relevance Score	The degree of relevance between the response and the question
Response Time	The total time for the AI to generate a response
LLM Processing Reasoning Time	The time spent on LLM reasoning and response generation
Total Token Count	Total tokens consumed in the conversation (including question and answer)

Field	Description
LLM	The name of the language model used
User	The user who initiated the conversation

View Image Search Usage

When an AI Assistant uses image search in the Knowledge Base, the system separately records image embedding usage. You can review an individual conversation record to see which resources were used for that response:

Go to [AgentOps](#) → [AI Assistant Monitoring](#), then switch to [Conversation Records](#).

Find the response you want to review, then click [Details](#) on the far right.

Scroll down to view the Token and Credit breakdown. Standard text embedding is listed under [Embedding](#). If the response used image search, [Image Embedding](#) is also displayed, and the Token and Credit totals include this item.



Usage breakdown for an individual conversation; “Image Embedding” appears only when the response actually used image search

If **Image Embedding** is not displayed, it does not mean the record is missing. This item appears only when the response actually incurs image embedding usage. When image search is not used, the screen continues to show the standard embedding breakdown.

For example, Mai's customer service manager asks the product assistant to identify the appearance of a component from catalog images. After testing, she opens the conversation details and sees the image embedding usage and corresponding Credits. This allows her to distinguish the cost of image search from standard text retrieval before deciding whether to adjust the image catalog or assistant settings.

Search and Filter


- Keyword Search:** Search input/output messages or usernames
- LLM Filter:** Select a specific language model to compare performance across different models
- AI Assistant Filter:** Select a specific assistant to track its operational status
- Time Range:** Select last 7 days, 30 days, 90 days, or a custom date range
- Export:** Export monitoring data in Excel or CSV format

Dashboard Metrics and Error Rate

Switch to the **“Dashboard”** tab to view overall service metrics and their percentage changes compared with the preceding period of the same length:

Metric	Description
Conversations	The number of responses from AI Assistants during the selected period. Each AI response in a conversation counts once; responses handled by human agents and automatically scheduled runs are excluded (see the definition below for details)
Error Rate	Percentage of responses with system errors (see the calculation below)
Average Response Time	Average time from receiving a message to completing the response
Average TTFT	Average time from when a user sends a message until the first token appears

The dashboard also provides charts for conversation volume, response time, error-rate trends, and LLM model distribution. The **“Assistant Service Ranking”** compares error rates and response speeds across assistants.



Dashboard metric cards and trend charts

How Is the Conversation Count Calculated?

The dashboard's "Conversations" metric and conversation volume trend count **each response from an AI Assistant**, rather than the number of conversations opened:

Scenario	Counted?	Count
A user enters the service but does not send a message	No	0
A user sends 1 message and the AI responds once	Yes	1
In the same conversation, a user asks 5 questions and the AI responds 5 times	Yes	5
A user sends a message, but a human agent takes over and responds	No	0
A user sends a message, but an error occurs while the AI is responding	Yes	1
An automatically scheduled task runs	No	0

To see how many conversations were opened, view "Conversation Count" on the AI Assistant's "Usage Analytics" page instead. It counts 1 when a user sends the first message, regardless of how many turns the conversation contains, and includes conversations handled by human agents. These metrics have different definitions, so the dashboard's "Conversations" count is generally higher than the "Conversation Count" in "Usage Analytics." See [Usage Analytics](#) for details.

How Is the Error Rate Calculated?

Error rate = Number of responses with system errors ÷ Total number of responses × 100%

Only **system-level** errors are counted, such as model call failures and timeouts. The following are **not counted** as errors to avoid overestimating the error rate:

Responses interrupted by the user (Client Interrupt)

Normal blocking by content safeguards (Hook)

The error rate measures service stability, not response quality. If the error rate rises abnormally, review the assistant's conversation records to identify the cause, such as model issues or external tool integration failures. To assess response quality, review faithfulness and answer relevance scores together with users' thumbs-up and thumbs-down feedback.

Monitoring Best Practices

Daily Review: Check conversations from the last 24 hours to identify abnormal response times or errors

Identify Performance Bottlenecks:

- * Response time > 10 seconds: Check Knowledge Base retrieval efficiency or consider a faster LLM
- * High token usage: Evaluate whether the System Prompt or conversation history can be shortened

Quality Issue Tracking:

1. Use keyword search to find problematic conversations
2. Analyze root causes (insufficient Knowledge Base
AI misunderstanding
model limitations)
3. Add problem cases to the test dataset and run automated evaluations to verify fixes

FAQ

Q: What is the difference between Automated Evaluation and AI Assistant Monitoring?

	Automated Evaluation	AI Assistant Monitoring
Purpose	Periodic quality testing	Real-time operational monitoring
Data Source	Predefined test datasets	Actual user conversations
Key Metrics	Success rate, response time	Performance, cost, quality scores
Best For	Quality verification, regression testing	Daily monitoring, troubleshooting

Q: How often should evaluations be run?

Recommendation: Core features weekly, full tests monthly, and immediately after major updates.

Q: Do evaluations affect actual users?

No. Automated evaluations run in an isolated environment and do not interfere with actual user conversations.

Q: How long is monitoring data retained?

The default retention period is 90 days. You can periodically export important data for long-term storage.

Evaluation Insights Report

This guide explains how to use the Evaluation Insights feature to quickly understand your AI assistant's test results and receive system-generated improvement recommendations.

What Are Evaluation Insights?

Evaluation Insights is an extension of AgentOps automated testing. After completing a batch test, the system not only displays metric scores but also automatically analyzes test results and generates an "Insights Report."

This report tells you:

How the test results performed overall

Where the main issues are

How to make specific improvements

Which items to prioritize

Think of it as having a senior consultant review your test report and provide you with an improvement recommendation document.

Core Advantages of Evaluation Insights



Automated Analysis

No need to manually interpret complex scoring data — the system automatically identifies problem patterns and compiles them into an easy-to-understand report.



Priority Recommendations

The system recommends which items to address first based on issue severity and scope of impact, making your improvement efforts more efficient.



Specific Improvement Directions

It doesn't just tell you "there's a problem" — it also provides concrete improvement suggestions, such as which settings to adjust and what content to supplement.



Multi-Language Support

Insights reports support multiple languages including Traditional Chinese, Simplified Chinese, and English. You can choose the language that best suits your team's needs.


How to Use the Evaluation Insights Feature

Step 1: Complete a Batch Test

Go to "AgentOps" → "Automated Testing"

Run a batch test (refer to [AgentOps User Guide](#))

Wait for the test to complete



The screenshot displays the 'Automated Testing' section of the AgentOps interface. It features a table with columns for Test ID, Test Name, Status, Success Rate, Average Response Time, and Action. The table lists several tests, including 'Web File Integrity Test', 'Web File Integrity Test', 'Web File Integrity Test', 'Web File Integrity Test', 'Web File Integrity Test', 'Web File Integrity Test', 'Web File Integrity Test', 'Web File Integrity Test', 'Web File Integrity Test', and 'Web File Integrity Test'. Each row shows the test's progress and a link to view details.

测试ID	测试名称	状态	成功率	平均响应时间	操作
TestID14463996401	Web文件完整性测试 - 删除Kubernetes配置文件	成功	100%	100ms	查看详情
TestID14463996402	Web文件完整性测试 - 删除Kubernetes配置文件	成功	100%	100ms	查看详情
TestID14463996403	Web文件完整性测试 - 删除Kubernetes配置文件	成功	100%	100ms	查看详情
TestID14463996404	Web文件完整性测试 - 删除Kubernetes配置文件	成功	100%	100ms	查看详情
TestID14463996405	Web文件完整性测试 - 删除Kubernetes配置文件	成功	100%	100ms	查看详情
TestID14463996406	Web文件完整性测试 - 删除Kubernetes配置文件	成功	100%	100ms	查看详情
TestID14463996407	Web文件完整性测试 - 删除Kubernetes配置文件	成功	100%	100ms	查看详情
TestID14463996408	Web文件完整性测试 - 删除Kubernetes配置文件	成功	100%	100ms	查看详情
TestID14463996409	Web文件完整性测试 - 删除Kubernetes配置文件	成功	100%	100ms	查看详情
TestID14463996410	Web文件完整性测试 - 删除Kubernetes配置文件	成功	100%	100ms	查看详情

Automated testing list showing success rate and average response time for each test

Step 2: View Evaluation Results

Go to the test result detail page

Review each scoring metric

Find the "Insights Report" section



Test detail page: includes success rate, AI insights summary, metric statistics, and improvement suggestions

Step 3: Select the Report Language

In the Insights Report section, you can choose the display language for the report:

Traditional Chinese: Suitable for Taiwan-based teams

Simplified Chinese: Suitable for Mainland China-based teams

English: Suitable for international teams or situations requiring English reports

After selecting a language, the system automatically generates the insights report in the corresponding language.

Step 4: Read the Insights Report

The insights report typically includes the following content:

1. Overall Evaluation Summary

Test pass rate

Key strengths

Areas needing improvement

2. Issue Analysis

Answer relevance issues

Safety risks (bias, offensiveness, hallucinations)

Response speed issues

3. Improvement Suggestions

Priority items (labeled as "High Priority," "Medium Priority," or "Low Priority")

Specific improvement steps

Expected improvement outcomes

Step 5: Make Improvements Based on Suggestions

Address issues in order of the priority listed in the report

Adjust AI assistant settings based on the specific improvement steps

After making adjustments, run the test again to verify the results

How Evaluation Insights Are Calculated

The system considers the following factors to generate the insights report:

Scoring Metrics: The levels and distribution of each score

Failed Cases: The number and reasons for failed cases

Problem Patterns: The frequency of recurring similar issues

Best Practices: Industry standards and optimization recommendations

Generating the insights report requires some processing time, typically completed within 1-2 minutes after the test finishes.

Practical Use Cases



Medical Clinic Customer Service

Test Results:

* Pass rate: 72%

* Hallucination score: High

Insights Report Suggestion:

"The test found that some responses contained unverified medical advice. It is recommended to explicitly restrict the role instructions: 'Do not provide diagnostic or treatment advice; only provide appointment scheduling and clinic information.'"



Enterprise Internal Knowledge Management

Test Results:

* Pass rate: 88%

* Answer relevance: Low

Insights Report Suggestion:

"The test found that some responses were too brief or did not fully answer the question. It is recommended to add more detailed operating steps and examples to the knowledge base."

FAQ

Q1: Can the insights report be downloaded?

Currently, the insights report can only be viewed within the system. You can copy the report content and compile it into an internal document.

Q2: Does the insights report update automatically with new test results?

Yes, each time a new batch test is run, the system generates a new insights report reflecting the latest test results.

Q3: Can I switch the language of the insights report?

Yes, you can switch the report language at any time, and the system will regenerate the report in the selected language.

Q4: Do I have to implement all the suggestions in the insights report?

Not necessarily. You can prioritize "High Priority" items based on actual needs and available resources. The system's suggestions are reference directions — you can adjust them according to your business requirements.

Q5: Why didn't my test generate an insights report?

Possible reasons:

- * Too few test cases (at least 10 are recommended)
- * The test has not finished yet
- * The system is still processing (please wait 1-2 minutes)

Tool Execution Records

Overview

Tool execution records let you track the complete history of an AI assistant's external tool invocations, including API calls, database queries, web scraping operations, and more. When an AI assistant needs to perform tasks beyond the scope of conversation, it uses tools to do so. Tool execution records help you understand tool usage, identify causes of execution failures, and optimize tool configurations.

With comprehensive execution records, you can identify the root cause of tool failures within 5 minutes, reducing troubleshooting time by 80% and ensuring the AI assistant can reliably deliver advanced functionality.

Features

With the tool execution records feature, you can:

Complete Execution History

The system automatically records every tool invocation, including input parameters, execution results, duration, and success or failure status.

Use case: An e-commerce platform's order query tool is invoked 500 times daily and needs success rate and response time tracking

Actual benefit: Discovered 10% of queries failed due to timeouts; after adjusting timeout settings, the success rate improved to 98%

Quickly Identify Failure Causes

When a tool execution fails, the system records detailed error messages to help you quickly understand what went wrong.

Use case: A weather query API suddenly starts failing frequently and the cause needs to be identified quickly

Actual benefit: Error records revealed an expired API key; service was restored within 5 minutes of updating it

Tool Usage Statistics

View invocation counts, success rates, average execution times, and other statistics for each tool.

Use case: An enterprise has configured 10 tools and wants to know which are used most frequently and which have the most issues

Actual benefit: Discovered 3 tools were never used; removing them reduced system complexity, and other tools saw a 15% improvement in response speed

Optimize Tool Configurations

Based on execution record analysis, adjust tool parameters, optimize invocation logic, and improve error handling.

Use case: A database query tool frequently times out because the query scope is too large

Actual benefit: After adjusting query limits based on record analysis, average execution time dropped from 8 seconds to 3 seconds

Tool Types

MaiAgent supports multiple tool types, and tool execution records track the execution status of all types:

API Tools

Call external APIs to retrieve real-time data, such as weather queries, exchange rate queries, inventory queries, etc.

Common execution issues:

- * Expired or invalid API keys
- * Incorrect request parameter formats
- * API service timeout or no response
- * Exceeded API call limits

Text-to-SQL Tools

Convert natural language questions into SQL queries to retrieve data from databases.

Common execution issues:

- * SQL syntax errors
- * Query scope too large causing timeouts
- * Database connection failures
- * Insufficient permissions to access specific tables

Web Scraping Tools

Scrape real-time information from specified websites, such as latest news, product prices, etc.

Common execution issues:

- * Target website unreachable
- * Web page structure changes causing parsing failures
- * Blocked or rate-limited by the target website
- * Scraping timeout

Custom Tools

Enterprise-developed tools for specific functions, such as order processing, member verification, etc.

Common execution issues:

- * Internal service errors
- * Parameter validation failures
- * Business logic errors
- * System integration issues

Usage Steps

Step 1: Access Tool Execution Records

Enter the admin dashboard and click [AgentOps](#) → [Tool Execution Records](#) in the left sidebar
View the list of recent tool execution records

ID	Tool Name	User	Execution Content	Status	Action
6939985-1-10-418	WebAgent-小助手	Mr. Agent-MUP	why you don't get MicrosoftNews.htm	成功	🔍
4342710-1-4-1-4-10	WebAgent-小助手	Mr. Agent-MUP	why you don't get MicrosoftNews.htm	成功	🔍
6070460-1-10-1-10	WebAgent-小助手	Mr. Agent-MUP	why you don't get MicrosoftNews.htm	成功	🔍
2778291-1-10-1-10	WebAgent-小助手	Mr. Agent-MUP	why you don't get MicrosoftNews.htm	成功	🔍
4125138-8-10-1-10	WebAgent-小助手	Mr. Agent-MUP	why you don't get MicrosoftNews.htm	成功	🔍
4164510-1-10-1-4-10	WebAgent-小助手	Mr. Agent-MUP	why you don't get MicrosoftNews.htm	成功	🔍
6006840-1-1-1-1-10	WebAgent-小助手	Mr. Agent-MUP	why you don't get MicrosoftNews.htm	成功	🔍
5033305-1-10-1-10	WebAgent-小助手	Mr. Agent-MUP	why you don't get MicrosoftNews.htm	成功	🔍
5497777-10-1-4-10	WebAgent-小助手	Mr. Agent-MUP	why you don't get MicrosoftNews.htm	成功	🔍
6630240-1-10-1-4-10	WebAgent-小助手	Mr. Agent-MUP	why you don't get MicrosoftNews.htm	成功	🔍

Tool execution records list page

Step 2: Filter Records

Use filter criteria to find execution records that need attention:

Time range:

Today, Last 7 days, Last 30 days

Or custom start and end dates

Tool type:

All tools

API tools

Text-to-SQL

Web scraping tools

Custom tools

Execution status:

All

Success

Failure

In progress

Specific tool:

Select a specific tool name from the dropdown

Click "Apply Filters" to view matching records

Priority Items to Check Recommended priorities: 1. Failed execution records 2. Records with execution time exceeding 10 seconds 3. Execution status of frequently used tools

Step 3: View Execution Details

Click any record in the list

Expand the detailed information, including:

Basic Information:

- * Tool name
- * Execution time (when it was invoked)
- * Execution duration (how many seconds it took)
- * Execution status (Success/Failure)

Input Parameters:

- * Parameter content passed by the AI to the tool
- * Check whether parameters are correct

Execution Result:

- * Complete result returned by the tool
- * Shows retrieved data on success
- * Shows error message on failure

Associated Conversation:

- * The conversation record that triggered this tool invocation
- * Trace back to the user's original question

Error Message (if any):

- * Detailed error description
- * Error code or exception type
- * Stack trace information (if applicable)

Step 4: Analyze Failure Causes

When you find failed execution records, follow these steps to analyze:

Check the error message:

Read the detailed error description

Identify whether it's a configuration issue, permission issue, or service error

Check input parameters:

Confirm whether the parameters passed by the AI are correct

Verify the parameter format meets tool requirements

Review the associated conversation:

Understand what question the user asked

Determine whether a specific question type caused the failure

Compare with successful cases:

Find successful execution records for the same tool

Compare the differences between successful and failed cases

Step 5: Optimize Tool Configuration

Based on execution record analysis, take improvement actions:

Update tool settings:

Correct incorrect API keys or connection information

Adjust timeout settings

Optimize query parameters or constraints

Adjust AI prompts:

If parameters are frequently incorrect, add clear guidance in the prompt

Explain the correct usage and parameter format of the tool

Improve error handling:

Add user-friendly error messages for common error scenarios

Set up automatic retry mechanisms

Remove unused tools:

If a tool has not been used for a long time, consider removing it to simplify the configuration

Step 6: Export Execution Records

Click the "Export" button on the tool execution records page

Select the export format (CSV or Excel)

Select the time range and filter criteria

Download the report for further analysis or archiving

Exported records can be used for: In-depth data analysis and trend tracking
Reporting issues to the technical team
Regular reviews of tool usage efficiency
Compliance auditing and record keeping

Use Cases

Case 1: Diagnosing Tool Failures

Scenario: An e-commerce platform's inventory query tool suddenly starts failing in large numbers, and customer service cannot check product inventory for customers.

Steps:

1. Go to tool execution records and filter by "Inventory Query Tool" + "Failure"
2. Discover 50 failed records in the past 2 hours
3. View the error message in the failed records: "Database connection timeout"
4. Check the database service status and discover the database is under maintenance
5. Temporarily disable the tool and notify customer service to use manual queries
6. Re-enable after database recovery and monitor the success rate returning to normal

Quantified Benefits:

- * Identified root cause within 5 minutes
- * Avoided hours of troubleshooting by the technical team
- * Promptly notified customer service to use backup methods, reducing customer wait times

Case 2: Optimizing Tool Performance

Scenario: The order query tool frequently runs slowly, affecting the conversation experience.

Steps:

1. Export the order query tool's execution records from the past 30 days
2. Analyze the execution time distribution and find:
 - Average execution time: 7.5 seconds
 - 25% of queries exceed 10 seconds
 - Maximum: 18 seconds
3. Review the input parameters of slow queries and find most are querying "all orders from the past year"
4. Optimize the tool configuration:
 - Limit query scope to the past 3 months

- Add pagination with a maximum of 20 records per request
5. After optimization, average execution time drops to 3.2 seconds

Quantified Benefits:

- * Execution time reduced by 57% (7.5 seconds → 3.2 seconds)
- * Queries exceeding 10 seconds dropped from 25% to 5%
- * Significantly improved conversation experience and customer satisfaction

Case 3: Discovering Unused Tools

Scenario: An enterprise configured 12 tools for the AI assistant and wants to understand actual usage to simplify the configuration.

Steps:

1. Export all tool execution records from the past 90 days
2. Count each tool's invocation count:
 - Order query: 1,500 times
 - Member query: 800 times
 - Promotion query: 300 times
 - Product specification query: 200 times
 - Other 8 tools: 0 times
3. Review the configuration of unused tools and find:
 - Some tools have unclear definitions, so the AI doesn't know when to use them
 - Some tools have overlapping functionality
 - Some tools are outdated
4. Remove 6 unused tools and optimize the descriptions of 2 low-usage tools

Quantified Benefits:

- * Simplified configuration, improving AI tool selection accuracy by 20%
- * System response speed improved by 10% (reduced tool selection judgment time)
- * Lower maintenance costs

Case 4: Monitoring API Usage

Scenario: An enterprise uses a third-party weather API with a monthly invocation limit and needs to monitor usage to avoid exceeding the quota.

Steps:

1. Export the weather query tool's execution records every Monday
2. Track the invocation count trend:
 - Week 1: 320 times
 - Week 2: 450 times
 - Week 3: 680 times
 - Week 4: Estimated 900 times (monthly limit: 2,000 times)

3. Discover rapidly growing invocation volume and analyze the cause:
 - Increased frequency of customer weather inquiries
 - Some unnecessary duplicate queries
4. Take action:
 - Add a caching mechanism to avoid duplicate queries for the same city within 1 hour
 - Optimize the prompt to prevent the AI from over-relying on the weather tool
5. After optimization, weekly invocations drop to 400, ensuring the quota is not exceeded

Quantified Benefits:

- * Avoid exceeding the API quota, which would cause service disruption
- * 40% reduction in API invocation costs
- * Maintain service quality while controlling costs

Frequently Asked Questions

Q: How long are tool execution records retained?

A: The system retains tool execution records for 3 months by default. It is recommended to regularly export important data for long-term storage and analysis. For longer retention periods, contact customer service to upgrade your plan.

Q: Why did a tool fail but the conversation seems normal?

A: When a tool execution fails, the AI assistant may:

1. Tell the user "Unable to query at this time, please try again later"
2. Use knowledge base information to answer instead (if related content exists)
3. Ask the user for more information and retry

Although the conversation continues, functionality is limited. It is recommended to track failure causes through execution records and fix them to ensure stable tool operation.

Q: How can I reduce the tool execution failure rate?

Recommended optimization approaches:

Configuration level:

- * Ensure API keys, database connections, and other configurations are correct
- * Set reasonable timeout periods (recommended 10-15 seconds)
- * Add error retry mechanisms

Prompt level:

- * Clearly describe when to use tools and the parameter format
- * Add examples to help the AI invoke tools correctly

Monitoring level:

- * Regularly check execution records
- * Promptly identify and fix issues
- * Track success rate trends

Q: What if input parameters in execution records contain sensitive data?

A: The system automatically masks sensitive fields in request headers: common credential headers (such as `Authorization` and `X-API-Key`) and all headers supplied by contact credentials are shown as `***`. Input parameters and the request body are not automatically masked. If your tools handle other sensitive data (such as customer personal information), it is recommended to:

1. Avoid having the AI pass secrets or personal data directly in tool input parameters; supply sensitive values via request headers or contact credentials instead
2. Control who can view execution records, granting access only to authorized personnel
3. In private deployments, extend the masked header list via the `TOOL_API_SENSITIVE_HEADERS` environment variable (e.g., `x-functions-key`)

Q: Can I re-execute a failed tool invocation?

A: The system does not currently support directly re-executing historical records. To test a tool, it is recommended to:

1. Ask a question in the chat platform to trigger a tool invocation
2. Use the tool testing feature (if available)
3. Fix the issue and wait for actual users to trigger it

Q: How can I see which tools were invoked in a specific conversation?

A: There are two ways:

1. **From conversation records:** The conversation details show all tools invoked during that conversation
2. **From tool records:** Click the "Associated Conversation" link in the tool execution records

Best Practices

Regular Check Recommendations

Daily check: Review failed records from the past 24 hours to ensure critical tools are operating normally

Weekly review: Count each tool's invocation count and success rate to identify items needing optimization

Monthly analysis: Export a complete report to analyze long-term trends and usage patterns

Alert Configuration Recommendations

Recommended alert settings (if supported by the system):

A tool fails more than 10 times within 1 hour

A tool's success rate drops below 80%

A tool's average execution time exceeds 10 seconds

A tool has not been used for more than 7 days (check if it's still needed)

Documentation Recommendations

Create usage documentation for each tool, recording:

Tool purpose and applicable scenarios

Input parameter descriptions and examples

Expected output format

Common errors and solutions

Optimization history

Add Conversations to FAQ

What problem does this feature solve?

When reviewing an AI assistant's conversation records, one of the most common findings is: “Customers keep asking this question, but the AI does not answer it well.” Previously, fixing this required copying the question and answer, switching to the knowledge base page, and manually creating an FAQ. Switching back and forth could easily interrupt your monitoring workflow.

Now, you can click **Add to FAQ** directly from a conversation record in AI Assistant Monitoring to organize the question and answer as an FAQ and save it to a specified knowledge base for the AI assistant to reference in future responses. You do not need to leave the monitoring page or copy and paste anything separately.

This is a key step in the AgentOps quality management cycle for “feeding issues found in production back into the knowledge base” : **Identify an issue through monitoring → turn it into an FAQ on the spot → continuously improve knowledge content.**

Which conversations should be added to FAQ?

Not every conversation needs to become an FAQ. Prioritize conversations that show both of the following signals:

Signal	What to look for in conversation records	Why it matters
Poor AI response	A low quality score or a user dislike	This indicates that the existing knowledge base cannot answer the question and a standard answer must be provided manually
Frequently asked question	Similar questions appear repeatedly while browsing conversation records	Fixing a frequently asked question once benefits all similar questions in the future

Typical scenarios include:

Questions with fixed answers where mistakes are unacceptable: return and exchange policies, fee calculations, and eligibility requirements. FAQs are suitable for locking in standard answers and preventing the AI from improvising.

New questions not covered by the knowledge base: customers begin asking about a newly launched product or campaign before the corresponding content has been added to the knowledge base.

Answers that are correct but too verbose or inconsistent with company guidelines: use an FAQ to provide a concise, compliant official version.

One-off, case-specific questions, such as the status of a particular customer's order, are **not suitable** for FAQ. The answer applies only to that customer, and saving it to the knowledge base could mislead others.

Entry point

From the left menu, click **AgentOps** → **AI Assistant Monitoring**.



Scroll down the left menu to find the AgentOps section

There are two places where you can launch Add to FAQ:

Entry point 1: Conversation records list

Switch to the **Conversation Records** tab. Each row has an **Add to FAQ** icon in the Actions column. When you spot a poorly answered conversation while browsing the list, click the icon in that row.

搜尋條件

相關紀錄

大小上限

內容限制

語言狀態

針對 90 天

重置篩選

📄 輸出

用戶輸入訊息

輸出訊息

AI 地址

使用者反饋

品質評分

TOKEN 用量

標籤

操作

單一檔案有大小上限嗎？我... 有，單一檔案大小上限依方...

MidAgent 暫...

-

-

4,823

共 1 條資料

1

10 條頁

Conversation records list: the Add to FAQ icon in each row's Actions column

Entry point 2: Conversation details page

Open the **Conversation Details** page for any conversation. An **Add to FAQ** button also appears at the top of the page. Use this option when you want to review the complete input and output, cited data, and quality score before deciding that a correction is needed.

對話詳情

< 上一筆紀錄

下一筆紀錄 >

複製連結

+ 加入 FAQ

↑ 輸出 ↓

基本資訊

用戶	AI 助理	LLM
陳工程師	MaiAgent 智能客服	GPT-5.6 Luna
建立時間	回應時間	
2026-09-05 00:08:09	6.6s	

Conversation details page: the Add to FAQ button at the top

You need edit permission for at least one knowledge base to use this feature. See [What happens if I do not have permission?](#) below.

Create FAQ window

Clicking Add to FAQ opens the **Create FAQ** window, which contains four sections from top to bottom:

[illegible]

Create FAQ window: the question is filled in automatically and the answer is blank by default

Section	Default value	Description
Knowledge Base	The first editable knowledge base is preselected	Required. The drop-down menu only lists knowledge bases you have permission to edit
Question	The customer's original question is filled in automatically	You can rewrite it to make it more general
Answer	Blank	Write the correct answer; click “Insert AI Response” to import the original response and edit it
Source Conversation	Collapsed	Expand it to view the complete conversation

Before saving, review and remove personal information such as order numbers, names, and phone numbers. This question and answer will become searchable knowledge base content and may be cited when the AI assistant responds to other users.

Steps

Select the target knowledge base

Select the knowledge base where you want to save the FAQ from the **Knowledge Base** drop-down menu. The system preselects the first editable knowledge base. If your AI assistant uses multiple knowledge bases, make sure you select **the one this AI assistant actually uses**. An FAQ saved to the wrong knowledge base will not be cited by this assistant.

Rewrite the question to make it more general

The **Question** field is automatically filled with the customer's original wording. Customers may use casual language, omit the subject, or include personal context. Rewrite the question as a general, standard question:

Industry	Customer's original wording (filled in automatically)	Suggested rewrite
Finance	My card won't go through	What should I do if a credit card transaction fails?

Industry	Customer's original wording (filled in automatically)	Suggested rewrite
Technology manufacturing	The machine screen shows E03. Is it broken?	How do I troubleshoot error code E03 on the equipment?
Healthcare	Can I still get an appointment tomorrow afternoon?	How do I check clinic hours and make an appointment?
Hospitality	Is there an extra charge if my child stays with me?	What are the extra-bed options and charges for children?

Rewriting principles: **remove personal context** (“my,” “tomorrow”), **add a complete subject** (“credit card,” “equipment”), and **end with a question**.

Write the answer

The **Answer** field is **blank by default**. This is intentional: you typically use this feature because the AI's original response was poor, so the system does not automatically insert it. Instead, write the correct answer yourself.

The editor works the same way as creating an entry in the knowledge base's **Q&A** section. It supports text formatting such as bold and image insertion, allowing you to clearly highlight key points.

If the AI response is only partially incorrect and its overall structure is reusable, click **Insert AI Response** in the upper-right corner of the answer field to import and edit the original response. The system then displays the conversation's quality score and user feedback (like/dislike) as a reminder:

conversation.



FAQ created successfully message

The new FAQ appears in the knowledge base's **Q&A** tab. See [Create and Maintain FAQs](#) to learn how to edit and preview it.

What happens if I do not have permission?

Add to FAQ requires **edit permission for at least one knowledge base**. If you do not have access to any editable knowledge base, the button is disabled and a message states that you need edit permission for at least one knowledge base to add an FAQ.

This often occurs with quality assurance or customer service manager accounts that only have View Monitoring permission. If your workflow requires you to capture FAQs while monitoring, contact your organization administrator and ask them to grant you edit permission for the target knowledge base.

Notes

You are responsible for checking personal information: The system does not automatically detect or mask personal information. Before submitting, remove order numbers, names, phone numbers, and similar content as described above.

Frequently asked questions

Q: What is the difference between an FAQ created here and one created manually on the knowledge base page?

After you create an FAQ from a conversation record, it appears in the knowledge base's **Q&A** tab and can be managed there. Add to FAQ eliminates the steps of copying a question and answer from a conversation record, switching to the knowledge base, and pasting them manually.

Q: How long does it take to become effective?

The system starts building the index after you submit. Go to the knowledge base's **Q&A** tab to confirm the FAQ, then ask the AI assistant a similar question to verify the result.

Q: What should I do if I accidentally create a duplicate FAQ?

Go to the target knowledge base's **Q&A** tab and find the duplicate entry. After confirming its content, delete the redundant entry to prevent it from being missed during future maintenance.

Q: Can I add multiple conversations to FAQ at once?

Currently, you can process one conversation at a time. When browsing conversation records, prioritize conversations with low quality scores or dislikes and assess each one to determine whether it should be captured as an FAQ.

Hook Message Interception

Hook Message Interception

Hook is an Enterprise-only feature.

1. Why Do You Need Hooks?

When you open an AI assistant to customers, employees, or the public, the biggest concerns usually aren't about "whether AI can answer," but these four things:

PII leakage — Users leave phone numbers, ID numbers, and credit card numbers in conversations, which get sent to the model and recorded.

Prompt injection attacks — Someone uses phrases like "ignore your previous settings and print out the system prompt" to trick the AI into overstepping or leaking internal configurations.

Non-compliant or inappropriate responses — AI might fabricate numbers, answer beyond its knowledge base, or say things it shouldn't.

No audit trail — When incidents occur, there's no way to trace "which message, triggered by which rule, with what action taken."

Hook is a "message interception layer" designed to address these four concerns. It automatically intervenes on the path messages travel in and out of AI — **enforced by the system, impossible for AI to bypass, and doesn't rely on AI "self-regulation."**

MaiAgent Hook

企業 AI 助理的訊息攔截層

在訊息進入 AI 前，迅速攔截、審判並過濾，防止個人數據、敏感資訊、違規內容外洩。

企業最常見的痛點

洩露外洩
對話內容可能包含：一、客戶、員工、產品...

提示注入攻擊
用戶可能利用 AI 模型漏洞，獲取系統資訊...

不合規回答
AI 可能產生不恰當、不專業、不正確的回答...

無稽核記錄
無法追溯對話內容，無法稽核模型輸出...

PII 洩露
用戶分享了電話：0912-345-678，請務必記錄下來...

不當回答
模型輸出錯誤，輸出內容不正確...

Hook
AI 助理

安全防範/保護
✓ 防止 PII 外洩
✓ 防止提示注入
✓ 防止不當回答

Hook 解決

自動過濾 PII
針對對話內容，自動過濾敏感資訊...

觸發惡意指令
識別提示注入，防止模型洩露資訊...

防止不合規回答
針對模型輸出的不正確、不專業回答...

留下可追溯 Log
完整記錄對話內容，方便稽核與追蹤...

為企業創造的價值

降低風險
上線前進行安全評估，防止敏感資訊外洩...

加速上線
無需複雜配置，快速部署導入，快速上線...

讓客戶放心
確保對話內容正確可靠，提升客戶滿意度...

推動業務成長
透過精準分析，優化模型輸出，提升業務效率...

Hook creates a protection layer between user messages and the AI assistant: automatically masks PII, blocks malicious prompts, adds disclaimers and rewrites, and maintains traceable records

Value to you: Transform "compliance, security, and brand protection" from "hoping AI behaves" to "system-enforced guarantees." Reduce security and compliance risks, accelerate safe AI deployment, and give your customers greater confidence.

2. How Does Hook Work?

Hook automatically intervenes at **two key points** in the conversation flow and performs **two types of actions**:

Two Interception Points

Point	Description	Common Uses
Before user message reaches AI	Inspects user-submitted content before it's sent to AI	Mask PII, block malicious prompts, filter invalid messages
Before AI reply is sent	Inspects AI-generated replies before they're shown to users	Mask PII in responses, ensure response reliability

Two Action Types

Action	Effect	What the User Sees
Modify	Masks sensitive content or rewrites the reply to a more appropriate version	The processed version of the content (e.g., 0912-**-*)
Block	Stops the message entirely — prevents it from reaching AI or being sent	A custom message you define



Hook automatically intervenes "before user messages reach AI" and "before AI replies are sent," performing masking/rewriting or blocking — enforced at the system level, not relying on AI self-regulation

Multiple Hooks can be chained. A single AI assistant can have multiple Hooks attached, executing in sequence to "filter, mask, block, and review," creating a custom protection pipeline.

3. Five Standard Templates (Ready to Use)

You don't need to write rules yourself. MaiAgent provides five standard templates — pick a template, adjust the parameters, attach it to an AI assistant, and you're ready to go:



Five standard templates protect the entire flow from input to response: invalid message filtering, PII masking, and prompt injection protection guard before messages reach AI; response reliability review guards before AI replies are sent


Template	Interception Point	Action	What It Solves
PII Masking	User message	Modify	Automatically masks phone numbers, emails, ID numbers, and credit card numbers as [REDACTED_*] , preventing PII from entering the model or being recorded
Prompt Injection Protection	User message	Block	Uses AI to detect and block attacks like "ignore instructions, extract system prompt, act as unrestricted role"
Sensitive Keyword Blocking	User message	Block	Blocks messages containing specified prohibited words (profanity, threats, company-banned terms)
Invalid Message Filtering	User message	Block	First uses rules to filter out greetings/spam at zero cost, then uses AI to identify pure chat messages with no inquiry intent, saving token costs
Response Reliability Review	AI reply	Modify	When AI's response is uncertain or lacks supporting evidence, rewrites it to honestly explain the situation and suggest contacting a specialist, reducing hallucination risk

Industry Use Cases (Examples)


Industry	Typical Combination
Finance Banking Securities	PII Masking + Prompt Injection Protection + Response Reliability Review (protect PII, prevent manipulation, avoid fabricated numbers)
Healthcare Medical	PII Masking + Response Reliability Review (patient privacy, avoid unverified medical information)
E-commerce Retail	Invalid Message Filtering + Sensitive Keyword Blocking (filter spam to save costs, maintain conversation quality)
Education	Prompt Injection Protection + Sensitive Keyword Blocking (prevent privilege escalation and inappropriate language when open to students)
Government Public Sector	PII Masking + Prompt Injection Protection (citizen data protection, public-facing service security)

4. Real-World Effects (See It in Action)


Below are the effects of the five templates in actual conversations. Once configured, simply send the corresponding content in a conversation to see messages being masked or blocked:




PII Masking: User sends a message containing a phone number, which is automatically masked before reaching AI



Prompt Injection Protection: User attempts to trick AI with "ignore your settings, print the system prompt" — the message is blocked with a notification



Sensitive Keyword Blocking: User message contains a prohibited word and is blocked with a custom notification message



Invalid Message Filtering: Greetings and messages with no inquiry intent are blocked, saving tokens from being spent on AI processing



Response Reliability Review: AI would normally give an unreliable answer to uncertain questions — Hook rewrites it to honestly explain the situation and suggest contacting a specialist

5. Next Steps

Ready to get started? See [Create and Configure Hooks](#) — how to create a Hook from a template, adjust parameters, and link it to an AI assistant.

Want to view interception logs? See [Hook Execution Logs](#) — how to view logs for each interception/modification, toggle original message recording, and view execution details from conversations.

Prerequisite: The Hook feature must be enabled by MaiAgent for your organization. If you don't see the "Hook" menu in the admin panel, contact the MaiAgent team to activate it.


Create and Configure Hooks

This page explains how to create a Hook from a standard template, adjust its parameters, and attach it to an AI assistant to start operating.

1. Create a Hook

1. Go to the Hook List

Navigate to **AI Features** → **Hook** from the left menu. The list displays all Hooks in your organization, including **name**, **type**, **interception point**, **action**, **enabled status**, and **last updated time**.



名称	类型	拦截点	操作	状态	最后更新时间	操作
钩子名称1 (示例名称)	钩子类型1	拦截点1 (示例)	操作1 (示例)	已启用	2024-01-01 10:00:00	编辑 删除
钩子名称2 (示例名称)	钩子类型2	拦截点2 (示例)	操作2 (示例)	已禁用	2024-01-01 10:00:00	编辑 删除
钩子名称3 (示例名称)	钩子类型3	拦截点3 (示例)	操作3 (示例)	已启用	2024-01-01 10:00:00	编辑 删除
钩子名称4 (示例名称)	钩子类型4	拦截点4 (示例)	操作4 (示例)	已禁用	2024-01-01 10:00:00	编辑 删除
钩子名称5 (示例名称)	钩子类型5	拦截点5 (示例)	操作5 (示例)	已启用	2024-01-01 10:00:00	编辑 删除
钩子名称6 (示例名称)	钩子类型6	拦截点6 (示例)	操作6 (示例)	已禁用	2024-01-01 10:00:00	编辑 删除

Hook list: View each Hook's type, interception point (when it intervenes), action (modify/block), and enabled status

2. Choose a Template

Click **+ New Hook** in the upper right corner. In the popup window, select a standard template. Each template card indicates its **interception point** and **action**, making it easy to quickly identify its purpose.



Choose a template: Each of the five standard templates displays its interception point (when receiving user message after AI reply is generated) and action (modify block)


The creation flow has three tabs: Template (which type) → Basic Info (name, description) → Apply to AI Assistant (which assistants to attach). Fill them in sequentially after selecting a template.

2. Adjust Parameters

Different templates have different parameters. After creation, go to the Hook settings page to make adjustments. Common parameters include:

Template	Main Parameters	Description
PII Masking	PII types to mask	Select which types to detect (Email Phone ID Number Credit Card)
Sensitive Keyword Blocking	Keyword blocklist, minimum word length	One prohibited word per line; minimum word length filters out overly short words to avoid false positives
Prompt Injection Protection	Judgment model, additional rules	Select an LLM to determine whether it's an attack; add custom rules


Template	Main Parameters	Description
Invalid Message Filtering	Greeting word list, consecutive repeat limit, judgment model	Rules filter first, then an LLM determines whether the message is invalid
Response Reliability Review	Rewrite model, review criteria	Select an LLM to judge whether the response is reliable and rewrite it



Using "PII Masking" as an example: Select which PII types to detect and configure the enabled status

Judgment Model (LLM-based Templates)

Templates that require "semantic judgment" — such as Prompt Injection Protection, Invalid Message Filtering, and Response Reliability Review — have a **judgment model** field. The dropdown lists the models currently available to your organization. Since this judgment runs on every message, **it's recommended to choose a fast, cost-effective model** (such as Haiku, GPT-5 mini, Gemini Flash, or similar). Leave it empty to use the platform default model.



Using "Response Reliability Review" as an example: Select a judgment model from the available LLMs in your organization, and describe the review direction in plain language using "Review Criteria"


Block messages are customizable. "Block" type templates (such as Prompt Injection Protection, Sensitive Keyword Blocking, Invalid Message Filtering) allow you to set the message shown to users when blocked, for example: "Your message contains inappropriate language. Please revise and try again."

3. Link to AI Assistant

A Hook must be **attached to an AI assistant** to take effect. There are two ways:

Attach during creation: In the "Apply to AI Assistant" tab when creating a Hook, directly select the assistants to apply it to.

Attach from assistant settings: Go to an AI assistant's settings page, switch to the **Hook** tab, and select the Hooks to apply.



In the Hook's "Apply to AI Assistant" tab, move Hooks from the "Available" list on the left to the "Selected" list on the right to complete attachment

One assistant can have multiple Hooks. The system executes them in order — for example, first "Invalid Message Filtering" blocks spam, then "PII Masking" masks personal data, then "Prompt Injection Protection" blocks attacks — combining them into a custom protection pipeline.

Enable Hook

After confirming parameters and linked assistants, set the Hook to **Enabled**. Once enabled, whenever users chat with that assistant, the Hook will automatically operate at the corresponding interception points.

4. Next Steps

After setup is complete, you can check [Hook Execution Logs](#) to view logs for each interception or modification and confirm that the Hook is working correctly.

Hook Execution Logs

Every Hook interception, modification, and pass-through is logged. For regulated industries (finance, healthcare, government), **"traceability" is itself a compliance requirement** — when issues arise, you need to trace "which message, triggered by which rule, with what action, at what time."



Execution logs (Audit Log) provide a complete record of each action: matched rule, action taken, processing result, and timestamp, with keyword search, assistant filtering, and report export support

1. View Execution Logs

Navigate to **AgentOps** → **Hook Execution Logs** from the left menu. You can see the following for each execution entry:

- Time** — When it occurred
- Assistant** — Which AI assistant triggered it
- Matched Rule** — Which Hook handled it
- Action** — What was done (mask/rewrite, block, or pass-through)
- Status** — Processing result

You can filter by **date range**, **assistant**, and **keywords**, and **export reports** for further audit analysis. The top of the page also displays statistics for **total executions**, **matches**, and **match rate**.

ID	Hook Name	Trigger Point	Action	Result	Message Content	Log Type
20240610-000001	Hook-001	用户输入消息	敏感信息识别	成功	检测到敏感信息: 138001380000	成功
20240610-000002	Hook-002	AI生成回复	PII替换	成功	将敏感信息替换为 [REDACTED_PHONE]	成功
20240610-000003	Hook-003	用户输入消息	敏感词识别	成功	检测到敏感词: 敏感词	成功
20240610-000004	Hook-004	AI生成回复	敏感词过滤	成功	过滤敏感词: 敏感词	成功
20240610-000005	Hook-005	用户输入消息	敏感词识别	成功	检测到敏感词: 敏感词	成功
20240610-000006	Hook-006	AI生成回复	敏感词过滤	成功	过滤敏感词: 敏感词	成功
20240610-000007	Hook-007	用户输入消息	敏感词识别	成功	检测到敏感词: 敏感词	成功
20240610-000008	Hook-008	AI生成回复	敏感词过滤	成功	过滤敏感词: 敏感词	成功
20240610-000009	Hook-009	用户输入消息	敏感词识别	成功	检测到敏感词: 敏感词	成功
20240610-000010	Hook-010	AI生成回复	敏感词过滤	成功	过滤敏感词: 敏感词	成功

Execution logs list: Each entry shows the AI assistant, Hook name, trigger point, action, and result. You can see actual actions such as PII masking replacing phone numbers with [REDACTED_PHONE], and sensitive keywords and prompt injections being blocked

View Individual Log Details

Click on any log entry to see the complete context of that action: **basic information (Hook name, type, trigger point, action, result), related information, and message content**, giving you a clear picture of what the system did.




Entry details: Using PII masking as an example, you can see the action is "modify" and the message content has been masked as [REDACTED_PHONE]

2. Control Whether Original Messages Are Recorded

By default, execution logs **only record the "action result" without retaining the original sensitive content** — this prevents PII or sensitive messages from being stored again in the logs.

If you need to see "what the original content looked like" during debugging, go to the Hook settings page and change the **Record original values before masking** toggle from **off** to **on**. When enabled, logs will additionally retain the original content before masking/blocking, making it easier to investigate issues.



Control via the "Record original values before masking" field on the Hook settings page: default is off (only records type and count, not original PII), switch to on only when debugging is needed

Privacy trade-off: Enabling "Record original values" means the original PII/message content will be stored in execution logs. For security reasons, it is recommended to keep this off by default and only enable it temporarily when troubleshooting is needed.

3. View Hook Execution Details from Conversations

In addition to centralized execution logs, you can also **trace back from actual conversations**: Navigate to [Customer Service](#) → [All Conversations](#), open any conversation, and you can see which messages in that conversation were processed by Hooks (masked or blocked), along with the corresponding actions taken.

This allows you to understand "why this message was masked/blocked" directly while reviewing customer conversations, without having to cross-reference logs separately.

4. Related Pages


[Hook Message Interception Overview](#) — What Hooks are and the five standard templates

[Create and Configure Hooks](#) — How to create, configure parameters, and link to AI assistants

Meeting Records

Meeting Records Feature Overview

Meeting Records is the voice transcription and AI summary feature of the MaiAgent platform, offering both **live** and **offline** modes. Live mode provides low-latency transcription and translation during meetings; offline mode lets you upload audio recordings for high-accuracy transcription and speaker identification. Both modes support AI auto-summary, Q&A, download/export, and knowledge base integration.



Meeting Records List

Two Modes

	Live Meeting Records	Offline Meeting Records
Input Method	Real-time microphone recording	Upload audio files (mp3, mp4, m4a, aac, wav)
Transcription Features	Low-latency real-time transcription	Multi-language high-accuracy transcription
Translation	Low-latency real-time translation (up to 3 languages)	Post-transcription translation
Speaker Identification	Depends on STT engine support	High-accuracy speaker identification and

	Live Meeting Records	Offline Meeting Records
		management
Timeline	Real-time marking	High-accuracy timeline recognition
Subtitle Playback	—	Dynamic subtitle playback, synced with audio
Backup Audio	Synchronized local backup recording in the browser prevents audio loss during network outages (see Meeting Record Backup Audio)	—

Common Features

Feature	Description
Multi-language Transcription	Supports speech recognition in multiple languages
Multi-language Translation	Translates transcribed content into multiple target languages
Auto Summary	Automatically generates meeting highlights using a custom AI assistant
AI Q&A, Agent Ticket Creation	Ask questions about the transcript; if the selected AI assistant is configured with ticket creation tools, you can ask it to turn meeting action items directly into tasks or tickets
Download & Export	Download MP3 recordings, plain text transcripts, and SRT subtitle files
Knowledge Base Integration	Upload transcripts to the knowledge base for Agent reference

Workflow

Live Mode:

Create Meeting → Configure STT Engine & Translation → Start Recording → Low-latency Real-time Transcription

↓
Meeting Ends → Review Transcript

↓
AI Summary

AI Q&A
Download & Export

Offline Mode:

Upload Audio (mp3/mp4/m4a/aac/wav) → System Transcription & Speaker Identification → Review Transcript

↓
Dynamic Subtitle Playback

AI Summary
AI Q&A
Download & Export

Meeting Record Statuses

Status	Description
Pending (Pending)	Created or uploaded, processing has not started
Recording (Recording)	Transcription in progress in live mode
Processing (Processing)	After recording ends or audio upload, the system is processing
Completed (Completed)	Processing complete, transcript can be viewed and summaries generated
Failed (Failed)	An error occurred during processing

Deployment Options

Meeting Records supports multiple deployment options:

SaaS: Use the MaiAgent cloud service directly

Private Cloud: Deploy in the customer's cloud environment

On-Premises: Install on the customer's local servers

Permissions

Meeting Records requires organization-level activation before use, with three permission levels:

View Permission: Browse the meeting list, view transcripts, and download files

Transcription Permission: Create meetings, start recordings, upload audio files, edit and delete meeting records

Management Permission: Determines whether a user can see all meetings in the organization, or only meetings they participate in

The meeting list **only shows meetings you created or were added to** by default. For details on who can see which meetings and how to share meetings with colleagues, see [Meeting Access Control](#).

Create a Meeting Record

Create a New Meeting

Go to the "Meeting Records" page
Click "New Meeting"
Fill in the following information:


Field	Description	Required
Meeting Name	The meeting title for easy searching and identification	Yes
Participants	Select organization members; added members automatically receive viewing access to this meeting	No
External Participants	For customers, vendors, and other non-organization members. Enter a name and press Enter. Names are recorded only in the meeting information and do not receive viewing access. See Meeting Access Control for details	No
Speech Recognition Engine (STT)	Select the speech-to-text engine	Yes
Recognition Language	Languages used in the meeting (up to 3)	Yes

Select a Speech Recognition Engine

The system provides multiple STT (Speech-to-Text) engines, each with different language support and recognition quality:

Engine	Features
Deepgram	High-quality real-time transcription, supports multiple languages
Azure	Microsoft Speech Services, excellent Chinese and English recognition
AssemblyAI	High accuracy, ideal for English scenarios
WhisperLive	Based on OpenAI Whisper, supports a wide range of languages
Bronci	Customized engine

Different engines support different languages and quantities. After selecting an engine, the system automatically displays the available language list.



Speech Recognition Provider Selection

Configure Real-time Translation (Optional)


If participants speak different languages, enable real-time translation:

Turn on the "Enable Translation" toggle

Select target translation languages (up to 3)

Select the LLM model for translation

Translation occurs simultaneously with transcription, and the transcript displays both the original text and translation results.



Language Selection and Translation Settings

Link Knowledge Base (Optional)

You can specify a knowledge base to link when creating the meeting. After the meeting ends, you can upload the transcript to the knowledge base with one click, allowing the Agent to access meeting content.

Manage Meeting Records

On the meeting records list page, you can:


- Search:** Search by meeting name
- Filter:** Filter by status (Pending, Recording, Processing, Completed, Failed)
- Date Range:** View records within a specific time range
- Edit:** Modify meeting name, language settings, translation settings, and knowledge base links
- Delete:** Remove unwanted meeting records

Live Meeting Transcription

Live mode provides low-latency voice transcription and translation, allowing you to see the real-time transcript while the meeting is in progress.

Start Recording

After creating a meeting record, click "Start Recording" to begin live transcription.



Live Transcription Setup Page

The system will:

Create a real-time communication room (LiveKit)

Start the speech recognition engine

Begin receiving microphone audio and transcribing in real time

Transcription View

During recording, the screen displays in real time:

Recording Timer: Shows the elapsed recording time

- Live Transcript:** Speech recognition results appear on screen in real time
- Speaker Labels:** Indicate the speaker for each text segment
- Language Tags:** Display the original language of each text segment
- Translation Results:** If translation is enabled, translated text appears below the original text
- Transcription Statistics:** Number of completed segments, in-progress segments, and total word count
- Local Backup Status:** "Recording Local Backup" appears below the meeting name, indicating that the browser is simultaneously recording a backup audio file so no audio is lost if the network disconnects. See [Meeting Record Backup Audio](#) for details

Low-latency Real-time Transcription

The core feature of live mode is low latency — transcribed text appears on screen within seconds of speaking, allowing participants to verify the record in real time.

Low-latency Real-time Translation

- When translation is enabled, the system translates content into the specified target languages simultaneously with transcription, also with low latency:
- Supports up to **3 target languages** for simultaneous translation
- Translation is processed in real time by the selected LLM model
- Translation results are attached below each transcript segment
- Ideal for international meetings and multilingual team communication

Controls

Control	Description
Microphone Toggle	Temporarily mute or unmute the microphone
Stop Recording	End the meeting recording; the system begins post-processing
Search Transcript	Search for keywords in the live transcript
Add Remove External Participants	Record or remove customers, vendors, and other non-organization members who join unexpectedly at any time in the "External Participants" section of the Participants card on the right, without interrupting the recording. See Meeting Access Control for details

After Recording Ends

After clicking "Stop Recording," the meeting status changes to "Processing" and the system performs the following in the background:

Audio Conversion: Convert the raw audio to MP3 format

Transcript Organization: Organize live transcription results into a structured transcript file

Cloud Storage: Upload the audio and transcript to S3 storage

Once processing is complete, the status automatically updates to "Completed," and you can access the meeting detail page to view the full content. The local backup audio uploads and merges in the background. When complete, the details page displays "Backup Complete." See [Meeting Record Backup Audio](#) for details.

Meeting Records Quick Access

In the **Office Apps** section on the MaiGPT home page, select **Meeting Records** to start recording immediately or continue viewing an existing meeting under **Recent Records**.



The Meeting Records quick access panel lets you start recording, connect your calendar, and view recent records

Use Case

After a regular project meeting, a project manager wants to check in MaiGPT whether a summary has been generated. She opens **Meeting Records**, selects the meeting under **Recent Records**, and checks the summary status in the panel. When finished, she selects **Back to Entry** to return to the start recording screen. If she needs the transcript or other complete details, she can select the button on the right side of the record to open the full record page. This leaves the MaiGPT chat page; after reviewing the record, she can return to MaiGPT from the sidebar.

View Recent Records

1. Open the Quick Access Panel


On the MaiGPT home page, select **Meeting Records**.

2. Select a Meeting

Under **Recent Records**, select a meeting to view its summary or summary status directly in the panel.

3. Return to the Entry Screen

On the summary screen, select **Back to Entry** to return to the calendar and start recording screen.



After selecting a recent record, you can view its summary status and return to the entry screen

4. Open the Full Record

When you need the transcript or other complete details, hover over the button on the right side of the record. When the **View Full Record** tooltip appears, select the button to open the full record page for that meeting. This leaves the MaiGPT chat page; after reviewing the record, select MaiGPT from the sidebar to return.

Recent Record Statuses

Screen Status	Meaning	Recommended Action
Generating	The meeting is still being transcribed, processed, or summarized	Reopen the quick access panel later to check the latest status
Generation Failed	Meeting transcription or summary generation failed	Open the full meeting record to review the failure status

Screen Status	Meaning	Recommended Action
No Summary Yet	The meeting is complete, but no summary is available to display	Open the full meeting record to generate or view the summary

If Recent Records displays No Meeting Records Yet, complete a meeting recording first, then return to this panel to view it.

For complete instructions on recording, transcription, and summaries, see [Create a Meeting Record](#), [Live Meeting Transcription](#), and [AI Meeting Summary & Q&A](#).

Meeting Record Backup Audio

Live meeting records send microphone audio over the network for cloud transcription. Previously, if a laptop lost its connection during a meeting, audio from the disconnected period could not be sent, leaving a gap in both the cloud recording and transcript. **Backup audio** fills this gap: when recording starts, the browser also records a local backup and uploads it in segments. Even if the network disconnects, local recording continues, and the segments upload automatically when the connection is restored. After the meeting, the system combines all segments into a complete backup audio file—the most complete recording in the system.

Backup audio is created only for live meeting records (recorded with a microphone). Meetings created by uploading an offline audio file already have a complete file and do not generate backup audio.

Use Case

A customer service manager uses a laptop in a conference room to record a product meeting with Live Transcription. Halfway through, the Wi-Fi becomes unstable. The transcript stops updating, but "Recording Local Backup" remains visible below the title, so she knows no audio is being lost and continues as usual. After clicking "Stop Recording," she sees "Backup Complete" on the meeting details page and switches the player to "Backup Audio" to confirm that the discussion during the outage is included. She then clicks "Retranscribe," changes the source to "Backup Audio," reviews the estimated cost, and confirms. A few minutes later, a backup version appears in the Transcript tab, where she can view the complete meeting without asking colleagues to reconstruct the missing section from memory.

Backup Status During Recording

After you click **Start Recording**, the backup status appears below the meeting name:

Display Text	Meaning
Recording Local Backup	Cloud recording and local backup are running simultaneously; everything is working normally
Syncing Backup, N Segments Remaining	The network is unstable or disconnected. Recorded segments are queued locally and upload automatically when the connection is restored; no action is required
Local Backup Unavailable (Reason); Meeting Continues	Backup cannot be enabled (for example, the browser is unsupported, browser storage is unavailable, or another tab is already recording). The meeting and

Display Text	Meaning
Normally	cloud recording are unaffected




The local backup status appears below the meeting name during recording

- Backup recording runs in the browser. Keep the entire meeting open in the same browser tab. Only one tab can record a meeting; a second tab displays "Recording in Another Tab."
- When you click the microphone icon to the left of "Stop Recording" to mute, the corresponding section of the backup audio is also muted, so content you do not want recorded is not captured.
- Local backup data awaiting upload is limited to 1 GB. Once the limit is reached, you cannot start another backup recording.

Backup Status After the Meeting

After you click **Stop Recording**, the meeting status first changes to "Completed." Backup audio has its own sync status, displayed next to the title on the meeting details page:

Status	Meaning
Backup Complete	The meeting is complete and the backup audio is fully synchronized—this backup is the most complete recording in the system
Syncing Backup	The meeting is complete, but the local backup audio is still uploading and merging. The status automatically changes to "Backup Complete" when finished
Backup Incomplete	The backup has not synchronized after more than one day. The system combines the uploaded segments into a backup audio file with the final section missing



Meeting details page: the "Backup Complete" label and audio source switch

In the meeting list, meetings whose backups are not yet finalized display an additional icon beside the status column. Hover over it to view details. Meetings with completed backups have no additional indicator.



Meeting list: the "Syncing Backup" icon beside the status column

What happens if I close the browser immediately after losing the network connection? Segments that have not uploaded remain on the local device. Open the MaiAgent admin console again within one day, and the system automatically resumes uploading while the meeting displays "Syncing Backup." If you do not return within one day, the system combines the uploaded portions and labels the backup "Backup Incomplete."

Switch to Backup Audio for Playback

For a meeting with backup audio, an **Original Audio** **Backup Audio** switch appears above the player in the "Recording" section on the left side of the details page. This switch affects **playback only**: meeting summaries, AI Q&A, knowledge base uploads, and MP3 downloads continue to use the original recording (unless you select another transcript source when generating a summary, as described below).

Retranscribe Using Backup Audio


If the original transcript is missing a section due to a network outage, you can retranscribe the backup audio to create a new transcript version. The existing transcript is retained.

Open Retranscription

Open the meeting details page and click **Retranscribe** in the upper-right corner.

Select "Backup Audio" and Confirm the Cost

Under **Transcription Source**, select **Backup Audio**. Billing information appears at the bottom of the dialog, including the billable item (audio transcription), backup audio duration, and estimated cost. It also explains that "you are charged only if transcription succeeds; failed attempts are free" and "each retranscription is billed separately."



The dialog discloses the estimated cost and billing rules

Submit the Transcription

Click **Confirm Audio Transcription (Approx. X credits)**—the button displays the estimated charge. After submission, the screen indicates that transcription has started and that you will be charged only if it succeeds. A "Processing" entry appears under "Retranscription History."

View the Backup Transcript Version

After transcription is complete, a version labeled "Backup" appears in the version menu on the Transcript tab (for example, "Azure • 09/06 00:42 • Backup"). When selected, the player **automatically switches to the backup audio**, allowing you to click the transcript and jump accurately to the corresponding timestamp.




Version menu on the Transcript tab: original transcript and backup version

- The backup transcript aligns only with the backup audio. If you manually switch to "Original Audio" while viewing the backup version, the screen warns: "The transcript and the audio being played come from different recordings. Clicking the transcript cannot jump precisely to that sentence."
- If the organization wallet has insufficient funds, "Backup Audio" is disabled and displays "Insufficient wallet balance. Please add funds first." No transcription job is created, and no charge is made.
- If the system detects that the end of the original transcript may be missing, a "The end of the transcript may be missing approximately N minutes" notice appears at the top of the Transcript tab, with a [Retranscribe Using Backup Audio](#) shortcut.

Generate a Summary from a Backup Transcript

When you click **Generate Summary** on the Meeting Summary tab, **Transcript Source** defaults to "Original Transcript." If a backup version is available, you can select it so the summary includes the recovered sections. The dialog then states: "The backup transcript timeline corresponds to the backup audio and differs from the original audio."



Select a transcript source when generating a summary

Further Reading

[Live Meeting Transcription](#): Start recording and view the recording screen

[Meeting Record Details](#): Transcript, player, and meeting information

[AI Meeting Summary & Q&A](#): Generate summaries and manage summary versions

Offline Meeting Transcription

Offline mode lets you upload pre-recorded audio files. The system performs high-accuracy transcription and speaker identification, generating a transcript with timeline markers.

Supported Audio Formats

Format	Extension
MP3	<code>.mp3</code>
MP4	<code>.mp4</code>
M4A	<code>.m4a</code>
AAC	<code>.aac</code>
WAV	<code>.wav</code>

Upload Audio

- Go to the "Meeting Records" page
- Select upload audio
- Choose a local audio file (supports mp3, mp4, m4a, aac, wav)
- The system begins background processing
- After uploading, the meeting record status changes to "Processing" and the system automatically starts transcription.

Multi-language High-accuracy Transcription

- Offline mode uses a high-accuracy speech recognition engine (such as WhisperX), which provides higher recognition accuracy compared to live mode:
- Supports multi-language recognition
- Higher recognition accuracy than live transcription
- Ideal for important meetings that require precise records

Speaker Identification and Management

Offline mode supports automatic Speaker Diarization, where the system automatically distinguishes between different speakers:

Auto Identification: The system automatically labels different speakers (e.g., SPEAKER_00, SPEAKER_01)

Custom Names: Replace system labels with real names (e.g., "Manager Zhang", "Team Lead Li")

Management Interface: Manage all speaker label names in one place

After modifying speaker labels, all related segments in the transcript are updated simultaneously.

High-accuracy Timeline Recognition

Each transcribed text segment includes precise start and end times, allowing you to:

Pinpoint the exact time of each utterance

Sync with the audio player for jump navigation

Export as SRT subtitle files with timestamps

Dynamic Subtitle Playback

Offline mode provides dynamic subtitle playback that syncs the audio player with the transcript:

Sync Highlighting: During audio playback, the current transcript segment is automatically highlighted

Click to Jump: Click any transcript segment to jump the audio to the corresponding time point

Search and Navigate: After searching for a keyword, jump directly to that segment's audio position

This feature lets you quickly find specific discussion content in a meeting without listening to the entire recording from the beginning.

Meeting Record Details

After meeting recording processing is complete, you can view the full meeting content on the detail page.

Page Structure


The detail page contains the following sections:

Audio Player: Play the meeting recording MP3. For live meetings with backup audio, you can switch between "Original Audio" and "Backup Audio" above the player. See [Meeting Record Backup Audio](#) for details

Meeting Info Card: Meeting name, participants, recording duration, creation time

Content Tabs: Transcript, Plain Text, Summary, AI Q&A, Download

Transcript Tab



Meeting Details — Transcript Tab


Displays the complete transcript in timeline format:

Timestamps: Start time of each text segment

Speaker Names: Indicate who is speaking

- Language Tags:** Original language labels
- Translation Content:** If translation is enabled, displays translation results
- Keyword Search:** Search for specific content in the transcript

Plain Text Tab



Meeting Details — Plain Text Tab

Presents the transcript in plain text format for easy reading and copying of the complete meeting content. Also supports keyword search.

Meeting Info Sidebar

The right sidebar displays:

- Basic Information:** Meeting date, meeting time, meeting duration, speech recognition provider
- Participants:** The access roster for this meeting (the creator is labeled as host); you can add or remove members — see [Meeting Access Control](#) for details
- External Participants:** Names of attendees who are not organization members. You can add or remove names; this is a text-only roster and does not represent access permissions. See [External Participants](#) for details

Meeting Access Control

Meeting content often involves sensitive information, so the meeting list **does not show all meetings in the organization**. By default, each member can only see meetings they are involved in; to share a meeting with colleagues, grant access through the "Participants" roster within the meeting.

Who Can See Which Meetings

Identity	Visible Scope in Meeting List
Regular Member	Only meetings they created + meetings they were added to as a participant
Organization Owner	All meetings in the organization
Member with "Meeting Record Management Permission"	All meetings in the organization

Meetings without access permission will display a "meeting not found" message even when accessed directly via URL — the system will not reveal its existence.



Meeting list: The "Participants" column shows the access roster for each meeting

The person who created the meeting automatically becomes the host and can always see the meetings they created.

Participants: The Meeting Access Roster

The "Participants" card on the right side of the meeting detail page is the **access roster** for that meeting — only people on the roster can see the meeting.

會議日期

2026/08/17

會議時間

18:29

會議時長

54 分鐘 5 秒

語音識別供應商

-



參與者 (3)



陳

陳文彬

👑 主持人

王

王志明



林

林佳蓉



逐字稿講者 (3)



陳文彬



林佳蓉



王志明

Participants card: Host label, member list, add participant button

Card contents:

Host: The meeting creator, with a host label next to their name

Members: Organization members who have been added

External Participants (bottom of card): Names of attendees who are not organization members; this is a **text-only roster and does not represent access permissions**. See [External Participants](#) for details

The "Participants" column in the meeting list also displays this roster, showing  when there are more than 2 people.

"Participants" and "External Participants" are two different things. External participants are recorded by name only and cannot access the meeting; to grant viewing access, you must add an organization member to the participants roster.

Who Can View and Who Can Edit the Roster

Action	Condition
View participants roster	Anyone who can see the meeting
Add Remove participants	Organization Owner, member with "Meeting Record Management Permission", or any participant already on the roster

Add a Participant

1. Open the meeting detail page

Click on the meeting you want to share in the meeting list.

2. Click the "+" in the upper right corner of the participants card

The card will expand a member selection dropdown.

3. Select an organization member and click "Add"

會議日期

2026/08/17

會議時間

18:29

會議時長

54 分鐘 5 秒

語音識別供應商

-

 參與者 (3) 

選擇成員



黃郁婷 (huang.yut...)

許珮琪 (xu.peiqi@...)

陳

陳文彬

 主持人

王

王志明



林

林佳蓉





Select the person to add from organization members

The dropdown only lists members from the **same organization**, with keyword search available

Members already on the roster will not appear in the dropdown

After successfully adding, the person can **immediately** see the meeting in their meeting list

Remove a Participant

Click the "X" on the right side of the member in the participants roster and confirm to remove them. The person immediately loses viewing access to the meeting.

The host cannot be removed. The meeting creator's row does not have a remove button, ensuring the meeting always has at least one person who can manage it.

External Participants: Record Non-Organization Members

Attendees such as customers, vendors, and consultants who are **not organization members** cannot be added to the "Participants" roster, but you can record them under "**External Participants**." External participants are a **text-only roster**: their names are recorded in the meeting information to provide a complete attendee list, but they **do not receive viewing access or notifications and do not need an account**.

Use Case

A sales manager is holding a product onboarding meeting with two colleagues from her team and two customer contacts. She opens "Live Transcription" and selects the two colleagues under "Participants" so they can view the transcript after the meeting. She then enters the customer contacts' names and company under "External Participants." Halfway through the meeting, the customer's technical manager

joins unexpectedly. She clicks "+" in the "External Participants" section on the right and adds the manager's name without interrupting the recording. After the meeting, the Participants card on the meeting detail page clearly separates the two groups: colleagues appear under "Participants," while customers appear under "External Participants." The full attendee list is easy to review, while the customers still cannot access the meeting content.

Record Participants When Creating a Meeting

In the "Live Transcription" setup form, the "External Participants" field appears below "Participants." Enter a name and press Enter to turn it into a tag. You can add multiple names in succession. When recording starts, the roster is created along with the meeting.

即時轉錄

配置會議轉錄及語音識別和語音轉錄

收起設定

會議名稱

請輸入會議名稱

語音識別供應商

Localai (qwen3-aar-1.7b)

語言

請選擇語言

參與者

選擇成員

加入的參與者將自動獲得此會議的權限

外部參與者

王大爺 (ABC 公司) 李小姐 (聯興)

外部參與者僅用於記錄會議資訊，不會獲得權限

啟用翻譯

將即時轉錄內容翻譯成多種語言

"Participants" only lets you select organization members and grants them viewing access; "External Participants" lets you freely enter names that are recorded only in the meeting information

The "External Participants" field is currently available in the Live Transcription setup form. It is not available in the offline audio upload form or the edit dialog in the meeting list. For offline meetings or names added afterward, use the sidebar on the meeting detail page (see the next section).

Add or Remove Participants in the Sidebar During or After Recording

The Participants card on the right side of both the Live Transcription page (while recording) and the meeting detail page (after the meeting) includes an "External Participants" section. You can add or remove names at any time:

1. Add

Click "+" next to "External Participants," enter a name, and press Enter or click "Add." The roster updates immediately after it is saved and remains current after you leave, return to, or refresh the page.

2. Remove

Click the "X" next to a name to remove it. No confirmation dialog appears; if you remove someone by mistake, you can add the name again immediately.



The sidebar has two sections: "Participants" are organization members with viewing access, while "External Participants" are a text-only roster

Anyone who can add or remove external participants has the same permissions as someone who can edit the meeting: members with "Transcription Permission" who can view the meeting. Other users can only view the roster

If saving fails, the page displays "Failed to update external participants." The roster remains unchanged, and recording is not affected

Roster Rules

Rule	Description
Maximum number	Up to 20 people per meeting. Once the limit is reached, no more names can be added and a notification appears
Name length	Up to 50 characters per person. If the name is too long, a notification appears and the name is not automatically truncated
Duplicate names	Identical names are treated as the same person and merged automatically
Whitespace	Empty input or input containing only spaces is not added to the roster

These rules apply consistently to the setup form, the sidebar during recording, and the sidebar after the meeting.

Meetings started from a calendar invitation (for example, starting a recording after connecting a calendar in the MaiGPT meeting record shortcut panel) automatically import the invitee list. Invitees who are organization members become "Participants" and receive viewing access, while everyone else is recorded as an "External Participant." You can still adjust the roster afterward in the sidebar.

Grant "See All Meetings" Access

For scenarios like auditing, handover, or department managers who need an overview of all meetings in the organization, grant access through role permissions:

1. Go to the Role Permission page in Organization Settings

See [Role Permission Management](#).

2. Edit role permissions and check "Meeting Record Management Permission"



Check "Meeting Record Management Permission" to allow role members to see all meetings in the organization

"Meeting Record Management Permission" is located at the bottom of the permission list and is an independent main permission:

Unchecked: Role members can only see meetings they participate in (created or added to)

Checked: Role members can see all meetings in the organization

When checking "Meeting Record Permission", the system will **automatically check** this permission as well; if you only want the role to see their own meetings, manually uncheck it

Members with "Meeting Record Management Permission" who only want to view their own meetings can filter using the API parameter `?scope=mine`.

FAQ

A colleague says they can't see a meeting I created?

Add them to the meeting's participants roster — it takes effect immediately.

How do I record customers or vendors who attend a meeting?

Add their names under "External Participants." This is only a roster and does not let them access the meeting content. If the person is an organization member, add them to the "Participants" roster instead to grant viewing access.

After unchecking "Meeting Record Management Permission", can I still see old meetings?

You can only see meetings you created and meetings where you were added to the participants roster. Meetings created by others that you were not added to will disappear from the list — if needed, ask a participant of that meeting to add you.

AI Meeting Summary & Q&A

Meeting summaries are processed in the background, so you do not need to remain on the summary page while a long meeting is being processed. You can leave the page during generation and return later to check the progress. When processing is complete, the summary is created as a new version.

Use Case

A project manager has just finished a long quarterly meeting and wants to organize the decisions before leaving work. She opens the **Meeting Summary** tab for the meeting and starts generating a summary with the default template. After seeing that the system is analyzing the transcript, she leaves the page to work on other tasks. When she returns to the same meeting later, the page continues to display the progress. Once processing is complete, she can review the decisions and action items in the newly created summary version.

Custom Summary Agent

The meeting records summary feature lets you select a custom AI assistant (Agent) to generate summaries. Different Agents can have different role instructions to produce summaries in different styles:

General Summary Agent: Organize meeting highlights, decisions, and action items

Technical Meeting Agent: Focus on technical decisions, architecture discussions, and technical debt items

Business Meeting Agent: Focus on client requirements, sales progress, and opportunity tracking

Generate a Summary

Go to the meeting detail page

Switch to the **Meeting Summary** tab

To customize the summary format, select the AI assistant to generate the summary; if none is selected, the default template is used

Click **Generate Summary**

The AI uses the complete transcript to organize the meeting highlights, decisions, and action items. When you use the default template, the page displays the **Analyzing transcript N/M** status followed by the summary compilation status. N represents the number of completed segments, and M represents the

total number of segments. When you select a custom AI assistant, the assistant processes the transcript directly and the segmented N/M progress is not displayed.

Whether you use the default template or a custom AI assistant, you can leave the page while processing continues in the background. Refresh the page or return to the meeting later to see the current status.



Leave the page and return later to continue viewing the default template's segmented progress



When the summary is complete, the new version appears in the version list

When the summary is complete, a new version is created automatically; you do not need to save it separately. If generation fails, the page prompts you to try again. Confirm that the transcript is available, then click **Regenerate**.

Summary Version Management

Each time a summary is generated or edited, the system retains a complete version history:

Add Summary: Each summary generation creates a new version

Edit Summary: Modify existing summary content

Version History: View all summary versions and their creation times

Delete Summary: Remove unwanted summary versions

Copy a Summary

When viewing a saved summary, a **Copy Summary** button appears in the toolbar above it. Use the button to copy the entire summary to the clipboard without manually selecting all the text, making it easy to paste into email, documents, or messaging apps.

Go to the meeting detail page and switch to the "Meeting Summary" tab

Select a saved summary

Click **Copy Summary**

After the "Copied successfully" notification appears, paste it into another application



The toolbar in the "Meeting Summary" tab

The summary's original text is copied with its paragraph structure preserved.

Copy Summary is available only when viewing a saved summary. It does not appear while editing a summary or previewing an unsaved draft. The button is disabled when the summary is empty.

If "Copy failed" appears, the browser has usually blocked clipboard access. Select and copy the text manually instead.

Read a Summary Aloud

When reading is inconvenient, have the system read the summary aloud. When viewing a saved summary, a **Read Aloud** button appears in the toolbar. Click it to hear the summary.

Go to the meeting detail page and switch to the "Meeting Summary" tab

Select a saved summary

Click **Read Aloud**. The button displays "Preparing..." before playback begins

During playback, click **Pause** to pause, **Resume** to continue, or **Stop** to end playback



During playback, "Read Aloud" changes to "Pause" and "Stop"

Before playback, the system automatically removes Markdown symbols such as `#`, `*`, and `-` from the summary so they are not read aloud.

Voice and Supported Languages

The voice is determined by **the meeting's spoken language or target translation language**, not by the actual language of the summary. The following seven languages are currently supported:

Language	Language Code
Traditional Chinese (Taiwan)	zh-TW
English (United States)	en-US
Japanese	ja-JP
Korean	ko-KR
Thai	th-TH
Indonesian	id-ID
Vietnamese	vi-VN

If the meeting language is not on the list above, such as Simplified Chinese zh-CN or Cantonese zh-HK, the toolbar displays a message saying that the language is not currently supported and the default language will be used. You can still use Read Aloud, but it uses a Traditional Chinese voice.

Limitations

- Read Aloud** is available only when viewing a saved summary; it does not appear while editing a summary or previewing an unsaved draft
- The button is disabled when the summary is empty
- Playback stops automatically when you switch summary versions or leave the "Meeting Summary" tab, so audio does not continue in the background
- The button is not displayed if the browser does not support audio playback
- If playback fails, "Text-to-speech failed. Please try again later" appears; try again later

AI Q&A

In addition to the summary feature, you can ask questions directly about the transcript content:


Switch to the "AI Q&A" tab

Select the AI assistant to use

Enter your question (e.g., "What decisions were made in this meeting?" "What are the next action items?")

The AI finds relevant content from the transcript and provides an answer

This feature is especially useful for quickly reviewing specific content from long meetings without re-reading the entire transcript.



AI Q&A Tab

Agent Ticket Creation

After a meeting ends, action items and decisions usually need to be manually entered into a project management or ticketing system. You can let the AI assistant handle this step: in the **AI Q&A** tab, ask it to turn the items discussed in the meeting directly into tasks or tickets.

Meeting Records does not have a dedicated "Create Ticket" button. **AI Q&A** uses the full AI assistant conversation, so as long as the selected AI assistant is configured with tools that can create tickets, you can ask it to create tickets in the conversation — the ticket creation capability comes from the AI assistant, not from Meeting Records.

How It Works

Step	Details
① Meeting Transcript	Meeting transcription is complete, and the transcript has been uploaded to the knowledge base and parsed
② Select AI Assistant	In the AI Q&A tab, select an AI assistant configured with ticket creation tools
③ Issue Ticket Creation Command	For example, "Compile the action items from this meeting and create a ticket for each one"
④ AI Assistant Calls Tools	The assistant calls API tools or MCP tools based on the field mappings and ticket creation rules in its role instructions
⑤ Create Tickets	Tickets are written to your company's ticketing or project management system
⑥ Report Results	The assistant reports each ticket's number and link in the conversation

Prerequisites

Condition	Description
Meeting is linked to a knowledge base, and transcript parsing is complete	AI Q&A requires the transcript to be in the knowledge base before it can be used. You can link a knowledge base when creating a meeting ; if not set at that time, you can also select a knowledge base in the AI Q&A tab and click Upload Transcript . See Download and Knowledge Base Integration for details
An AI assistant configured with ticket creation tools	The tool, not Meeting Records, is what actually calls the ticketing system. See the setup steps below
Ticketing system provides a callable API	MaiAgent communicates with external systems via API tools or MCP tools, without being limited to any specific vendor; as long as the other party provides an API for creating tickets, it can be integrated

Setup Steps

Create a ticket creation tool: On the [Tools](#) page, click [Add API Tool](#) and configure your company's ticketing system's "Create Ticket" API (you can paste a cURL command directly and the system will automatically parse the URL, headers, and HTTP method). See [Create API Tool](#) for details. If the other party provides an MCP server, you can also use [MCP Tool](#) instead.

Assign the tool to the AI assistant: Add the newly created tool to the AI assistant's available tools list. See [Configure Tools for AI Assistant](#).

Define ticket creation rules in the role instructions: Whether the AI calls the tool at the right time and how fields are filled depends on the role instructions. Refer to the example below, and see the [Role Instructions Design Guide](#) for writing principles.

Use in meetings: Go to the **AI Q&A** tab on the meeting detail page, select this AI assistant, and issue a command directly, such as "Compile the action items from this meeting and create a ticket for each one."

Ticket creation writes to an external system. It is recommended to verify field mappings with a test project or test environment before switching to your production ticketing system.

Role Instructions Example

```
# Role
You are a meeting action item ticketing assistant.

# Task
Users will ask you to create tickets based on meeting transcripts. Follow this process:

1. Identify clear action items and decisions from the transcript; ignore pure discussions or items that haven't been finalized
2. Organize each action item into a ticket with the following field mappings:
  - Title: One sentence describing what needs to be done, under 30 characters
  - Content: Add the background and rationale mentioned in the meeting
  - Assignee: The person explicitly assigned in the transcript
  - Due date: The date explicitly mentioned in the transcript
3. Before calling the tool, present the compiled ticket list to the user for confirmation; only create after getting approval
4. After creation, report each ticket's number and link

# Constraints
- Only create tickets based on transcript content; do not add tasks that were not discussed
- If the assignee or due date cannot be determined from the transcript, leave blank and note it in the response; do not guess
- When the same item is mentioned multiple times in the meeting, merge into a single ticket
```

Verification and Troubleshooting

The parameters actually received by the tool and the returned results can be viewed in AgentOps' [Tool Execution Records](#) — this is the fastest way to troubleshoot ticket creation issues.

Issue	Cause and Resolution
AI Q&A is unavailable	The meeting is not linked to a knowledge base, or the transcript is still being parsed. Select a knowledge base in the tab and upload the transcript, then try

Issue	Cause and Resolution
	again after parsing is complete
AI listed the action items but did not actually create tickets	The tool was not added to the assistant's tool list, or the role instructions do not explicitly require calling the tool. Check steps 2 and 3
Ticket creation failed	Check the error message in the tool execution records. Common causes include missing required fields, expired authentication, or insufficient permissions
Ticket field content is incorrect	Make the field mapping rules more specific in the role instructions, and explicitly instruct it to leave fields blank rather than guess when information cannot be determined
AI's response is unrelated to this meeting	The currently selected AI assistant may not be linked to this meeting's knowledge base; the interface will show a reminder. Switch to a linked assistant


Download & Knowledge Base Integration

Download Files

After meeting processing is complete, you can download the following files:

Format	Description	Use Case
MP3 Audio	Meeting recording in MP3 format	Archiving, offline playback
Plain Text Transcript	Complete transcript in plain text format	Copy and paste, document archiving
SRT Subtitle File	Transcript in subtitle format with timestamps	Video playback, subtitle import

On the meeting detail page, go to the "Download" tab and click the corresponding download button to get the file. The plain text transcript can also be copied with one click from the "Plain Text" tab.



Download Tab

Upload to Knowledge Base

After uploading the meeting transcript to the knowledge base, the Agent can answer questions based on meeting content:

On the meeting detail page or list page, select "Upload to Knowledge Base"

Select the target knowledge base (if a knowledge base was linked when creating the meeting, it will be pre-filled)

Confirm the upload

After uploading, the transcript becomes a document in the knowledge base, and the Agent can retrieve meeting discussion content through RAG.

Use Cases

Project Meetings: Import each project meeting transcript into the knowledge base so the Agent can answer "What was decided in the last meeting?"

Client Interviews: Import client interview records for easy reference of client requirements

Training Records: Import internal training content to build an organizational knowledge base

Customer Service Conversations & Platform Integration

Choose an Integration Platform

MaiAgent's AI Assistant can be connected to different types of chat platforms based on your needs, whether for external customer service or internal team support. Below are the definitions and use cases for each connection type:



System Default Platform Interfaces

When an AI Assistant is created, the system automatically creates the following three platform interfaces:

API: Provides a programmatic interface for developer integration

Web: Built-in web chat interface that can be embedded directly on websites

Internal: A Q&A platform for internal company use

External Public Deployment

Suitable for external services, making the AI Assistant the frontline communication channel for your business:



Connect Chat Platform: Website

Function: Embed the AI Assistant on your company website

Use case: Instantly respond to visitor inquiries, provide product information, service descriptions, FAQ answers, etc.

Advantage: 24/7 online customer service to improve website conversion rates



Connect Chat Platform: LINE

Function: Integrate with your LINE Official Account

Use case: Provide customer service and user interaction through the widely used LINE messaging interface

Advantage: Reach the large LINE user base and lower the barrier for customer engagement



Connect Chat Platform: FB Messenger

Function: Connect to your Facebook Page's Messenger

Use case: Instantly reply to customer questions through social media, handle order inquiries, event Q&A, etc.

Advantage: Integrate with social marketing to boost fan engagement



Connect Chat Platform: Telegram

Function: Create a Telegram Bot for customer service

Use case: Supports individual conversations and group discussions; suitable for tech communities and international customer service

Advantage: Simple connection setup, smart group responses, conversation reset support, and other practical features

Internal Use Only

Suitable for internal knowledge queries and process support to improve team efficiency:

Internal Q&A

Function: AI Assistant is available only to internal staff

Use cases:

Connect to internal documents or policy databases

Help employees quickly look up company policies and process guides

HR policy and SOP reference

New employee onboarding support

Advantage: Serves as an intelligent knowledge hub for the enterprise, reducing repetitive inquiry workload

API Integration

Function: Integrate with internal enterprise systems via API

Use cases:

Integrate with ERP and CRM systems

Connect to internal workflow platforms

Develop custom applications

Advantage: Deep integration with existing systems for a seamless user experience

Internal Q&A Features

Introduction to Internal Q&A

💬 Conversation Area

Engage in real-time Q&A interactions with the AI Assistant

📋 Conversation History

View and manage all past conversation records

🤖 AI Assistant Selection

Switch between different AI Assistants for conversations

⚙️ AI Assistant Settings

Adjust the current conversation's AI Assistant settings in real time, including permissions and tool configurations

📖 Knowledge Base Document Management

Precisely control which documents the AI Assistant can reference for each conversation



How to Use Internal Q&A?

1. Navigate to the Internal Q&A Page



Continue a Past Conversation (Select Conversation)



Start a New Conversation (New Conversation)

4. Customize Document Reference Scope (Optional)

Use the Knowledge Base management feature on the right side to select specific documents for the AI Assistant to reference during this conversation.



You can freely select reference documents and FAQs for each message

You can reconfigure settings before each message is sent

Document Access Scope

When sending the first message, all documents are accessible, including "Ambient-Light-Series.md"

The AI Assistant references the document and responds: "Document content: The official website has been closed and no longer sells any products"



When sending the second message, uncheck "Ambient-Light-Series.md"

The AI Assistant cannot reference it and responds: "I cannot find the specific content of the '**Ambient-Light-Series.md**' file."



FAQs Access Scope

When sending the first message, all FAQs are accessible, including "What is the working principle of LLM?"

The AI Assistant references the FAQ and provides a relevant answer



When sending the second message, uncheck "What is the working principle of LLM?"

The AI Assistant cannot reference "What is the working principle of LLM?" and responds "I cannot find information about the working principle of LLM in the FAQ"



Tag Filtering

You can use tags configured in the Knowledge Base to filter the document content you want to make available to the AI Assistant:

For tag settings, refer to: [Document Management: Tags and Metadata](#)

a. Click "Show Tag Filter" in the upper right corner



b. Select AND or OR filter conditions by clicking the "AND" text to toggle



c. Add tags

Using "AND" filtering

Documents: Only "Advanced-Camping.pdf" meets all conditions and is available for the AI Assistant to reference

FAQ: No FAQs meet all conditions simultaneously, so none are displayed

Tag filtering applies to both documents and FAQ content simultaneously






5. Adjust AI Assistant Settings in Real Time (Optional)

Use "AI Assistant Settings" in the upper right to modify permissions, tools, and other configurations in real time.



Click to open a side panel on the right, where you can freely edit the required permissions, tools, and more



Usage Tips

Conversation Management

Use conversation history to quickly return to previous discussions

Start a new Q&A for entirely new topics

Document Control Precision

Restrict reference documents for specific questions to get more precise answers

Avoid irrelevant documents interfering with the AI Assistant's judgment

Flexible Settings Adjustment

Adjust AI Assistant capabilities in real time for different conversation scenarios

Test the impact of different settings on response quality

Web Chat Overview

1. What is WebChat?

You can embed AI assistants created on MaiAgent into your website to provide real-time customer service or other conversations. Web Chat features a **responsive design that adapts perfectly to desktop, tablet, and mobile devices**.

MaiAgent currently offers the following three embedding methods:

1. Bottom-Right Corner Widget

Embed into your website as a built-in smart assistant

Non-intrusive to normal browsing, available whenever needed

Ideal for general customer service and consultation



Bottom-right corner embed (draggable window)

2. Full-Screen Embed

Provides a complete conversation experience

Suitable for complex consultation and service workflows

Customizable design to meet enterprise needs




Full-screen embed demo

3. MaiGPT Window Mode Embed

A full ChatGPT-style conversation interface (conversation history list on the left + main conversation area on the right) that can be embedded directly into a page as the primary UI

Can be permanently displayed in a specified container, or opened full-screen by clicking a floating button in the bottom-right corner

Ideal for AI portal sites, internal knowledge portals, or as the primary conversation interface for a product



MaiGPT Window Mode: Full ChatGPT-style interface

MaiGPT window mode can be selected in the "**Embed**" window of the chat platform settings. For complete setup steps, see [MaiGPT Window Mode Embed](#).

This design delivers:

Immediacy: Visitors get professional answers instantly without waiting, enhancing user experience

Convenience: The floating bottom-right design doesn't interfere with normal browsing, available whenever needed

Cost Efficiency: AI assistant provides 24/7 service, significantly reducing labor costs

Professional Image: Demonstrates the enterprise's technological capabilities and service quality

2. Build a Brand-Styled Web Chat Service

MaiAgent offers customization features including custom logos, avatars, and Q&A theme colors to help you build a brand-styled Web Chat service

You can:

Choose the assistant display name

Upload your company logo

Choose the AI assistant display avatar
Select the Web Chat theme color
Select the message background color for AI assistant replies
to create your own Web Chat appearance:



Default style



Custom style

This way, you can create a Web Chat service that matches your brand colors and corporate identity



Dark Mode Optimization

MaiAgent's Web Chat dark mode has been fully optimized to provide a more comfortable visual experience:

Dynamic Color Adjustment

Automatically adjusts theme color brightness to ensure visibility on dark backgrounds

Intelligently adjusts text contrast for optimal readability

Optimizes icon and button colors for better visibility in dark mode

Contrast Optimization

Message bubble background colors are carefully designed to prevent visual fatigue

Links and important information use high-contrast colors for clear identification

Input field and button border colors are optimized to enhance the user experience

Seamless Switching Experience

Users can switch between light and dark modes at any time

Conversation content and operation states are preserved during switching

The system remembers the user's theme preference

Dark mode is particularly suitable for:

Nighttime use or low-light environments

Extended conversations to reduce eye strain

Enterprise brand designs with dark color schemes

3. Message Customization

When entering the Web Chat service for the first time, you can freely configure your AI assistant's opening message based on the following scenarios.

1. Greeting Customization

Branded Greeting

Set: "Welcome to our website! I'm your dedicated customer service assistant" to showcase your enterprise's professional image and friendly service attitude

Guided Opening

Set: "I can help you look up product information, answer usage questions, and assist with order processing" to clearly inform users of the available service scope and build trust in the service capabilities



Custom greeting

2. Guided Starter Questions

Set "**conversation starter questions**" as quick options

Users can directly click on common questions without typing

Example questions:

"How do I check my ordered items?"

"How do I place an order?"

"Membership points rules"

"Business hours"

Users can tap starter questions to begin a conversation, speeding up the Q&A process:



Starter question setup example



User taps to start conversation

Stop a Response While It Is Being Generated

If you notice an error in your question after sending it, you do not need to wait for the AI assistant to finish the entire response. While the response is being generated, the send button to the right of the input field changes to a square **Stop** button. Click it to stop the current response, then enter and send your corrected question.

For example, suppose a customer service representative asks the AI assistant to plan an annual product launch but realizes after sending the request that they forgot to specify the event budget. They can click **Stop** while the AI assistant is generating the response, then add the budget and number of attendees and ask again. This avoids waiting for content that does not meet their needs to finish generating.

Stop the Current Response

1. Find the Stop Button

After sending a question, wait for the AI assistant to start generating a response. A square **Stop** button appears to the right of the input field.



Click the Stop button to the right of the input field while a response is being generated

2. Stop and Ask Again

Click **Stop**. Once the "Typing a response for you" message disappears, enter the corrected question in the input field and send it again.

Stopping only interrupts the response currently being generated. It does not delete the question you sent or any previous conversation history.

4. Diverse Q&A Content

1. Multi-Language Support

Responds in the user's language, eliminating communication barriers.

Auto-detection: Selects the assistant's reply language based on the user's input language



Manual switching:

Use the Web Chat interface to select the language. See: [Multi-Language Support: Web Chat Chat Settings](#).

Use the Web Chat initialization script (JavaScript) to specify the AI assistant's reply language. See: [Technical Manual - Web Chat Initialization](#).

For the list of supported languages, see: [Multi-Language Support: Supported Language List](#)

2. Multimodal Q&A

In addition to text-based Q&A, the AI assistant also supports uploading different types of files:

Spreadsheets: .xls, .xlsx, .csv, .ods

Word processing documents: .doc, .docx, .odt, .pdf, .md, .txt

Presentation files: .ppt, .pptx, .odp

Web files: .html, .htm

Data formats: .json, .jsonl

Audio files: .wav, .mp3, .m4a, .aac

Video files: .mp4

You can disable file uploads for users, restricting conversations to text only to ensure conversation security.

Image Content Processing

Image Recognition and Analysis

Upload image analysis: Users can upload product images, and the AI automatically identifies and provides relevant information

Screenshot processing: Supports screenshot analysis to help resolve usage issues

Image descriptions: Automatically generates image content descriptions and related suggestions

Image Response Feature


Product image display: Replies include actual product images

Operation guide images: Provides step-by-step operation screenshots

Examples:

A user asks where to configure the MaiAgent assistant, and the assistant responds with an image of the settings screen

A user sends a screenshot to the assistant to confirm whether the path is correct, and the assistant correctly interprets and responds



Assistant responds with images



Assistant interprets user screenshot content

Document Processing Capabilities

Supported Document Formats

PDF files: Automatically parses PDF content and answers related questions

Word documents: Processes .doc/.docx format documents

Excel spreadsheets: Analyzes spreadsheet data and provides insights

PowerPoint presentations: Extracts presentation content and answers questions

Document Interaction Features

Document content queries: Answers specific questions based on uploaded documents

Data analysis: Automatically analyzes table data and provides statistical information in table format

Document summaries: Generates key content summaries for long documents

Example 1: Ask the AI assistant to present data content in table format and provide reference links

☰

預設助理樣式

⋮

請用表格幫我整理 LLM 現有的服務，並給予參考連結，表格不要超過三格

2025/08/15 16:46:01

 預設助理樣式

服務提供商	主要LLM服務
OpenAI	GPT-4, GPT-3.5, ChatGPT, AI
Google	Gemini (原Bard), PaLM API, V
Anthropic	Claude系列模型, Claude API

這個表格涵蓋了目前最主要的三個LLM服務提供商，包含了消費者和開發者都能使用的服務選項。

2025/08/15 16:46:04

 請輸入您想問的問題



Powered by



✕

Information organized in table format



Providing external links

Example 2: Ask the assistant to analyze sales content and provide table analysis



Example 3: Upload a docx file and ask the assistant to help analyze CV writing issues



The AI assistant replies in Markdown format, organizing response content in a well-structured manner.

5. Conversation Memory and Sharing

1. Conversation History Memory


To provide a continuous service experience, WebChat automatically remembers users' conversation history, allowing each conversation to build upon previous ones.

Memory Scope

Conversation content: Complete Q&A records, including text, images, documents, etc.

Service status: Incomplete queries, pending issues

Users can view past conversation records in the conversation list on the left side of Web Chat, or start a new conversation





The conversation memory feature can be used for:

Continuous consultation: When a user asks "What tent model did you recommend last time?", the AI can immediately recall and answer

Progress tracking: When a user says "Continue the order inquiry from last time", the AI automatically retrieves the relevant information

Personalized service: Provides more accurate recommendations and suggestions based on conversation history

2. Conversation Sharing

When you want to share a conversation with friends or colleagues after consulting the AI assistant, or when internal teams need to collaborate on customer issues, you can use MaiAgent's service to generate a conversation sharing link:





根據文件 `cv_revised_01082025.docx`，我發現以下幾個問題：

格式與結構問題

分享對話的公開連結



<https://chat.maiagent.ai/web-chats/1ccd7cae-bfaa-47a2-a242-8b062f0a5a53/conversations/f4646896-07f2-42c0-8ebf-d70e0a29cf60/share>

複製連結

1. 時間邏輯錯誤：


- 目前工作經驗顯示"Aug. 2025 -- Now"，但這是未來的日期
- 學業預計2026年畢業，但已經有2025年的工作經驗



請輸入您想問的問題



Powered by





Shared Content

Complete conversation record: Includes all text, images, documents, and other content

Conversation summary: Automatically generated conversation highlights summary

Related resources: Product links, documents, etc. mentioned in the conversation

Conversation record sharing can be used for:

Customer sharing: "This camping gear recommendation is great, share it with the hiking club"

Internal collaboration: Customer service staff share complex issues with the technical support team

Training purposes: Share typical conversation cases as employee training materials

3. Conversation Timeout Management

To ensure conversation quality and system performance, Web Chat features an intelligent conversation timeout management mechanism:

Intelligent Filtering Mechanism

Automatic greeting filtering: The system identifies and filters opening greetings, excluding them from conversation timeout calculations

Correct message count reset: When a conversation restarts, the message counter resets correctly

Optimized timeout determination: Timeout calculations only apply to substantive conversation content, providing more accurate conversation management

Conversation Continuity Assurance

Ensures users do not unexpectedly time out during normal conversations

Conversations can continue after extended periods of thinking or pauses

The system intelligently determines whether a conversation is still in progress

How It Works

Start conversation: The user sends a greeting (e.g., "Hello"), which the system does not count toward the timeout

Enter conversation: The user begins asking substantive questions, and the timeout timer starts

Conversation pause: If the user temporarily leaves, the system maintains the conversation state


Resume: The user returns to continue the conversation, and the message count updates correctly

This optimization makes the conversation experience smoother and more natural, so users don't need to worry about interruptions due to extended thinking time.

6. Choose Whether to Display Tool Usage and Citation Lists

1. Display Tool Usage List

You can choose whether to let users see the list of tools the AI assistant uses during reasoning. If enabled, the following screen appears when the AI assistant generates a response:



You can use this list to verify whether the AI assistant is using the correct tools.

2. Display Citation List

Each AI response automatically generates a citation list, clearly indicating information sources to enhance the credibility and traceability of answers.

Display Settings

Enterprises can configure whether to let users see the tool usage list and reference document names. If display is enabled, the following list appears at the bottom of the AI assistant's response. Clicking on a "citation node" reveals the referenced document excerpt:



Depending on different usage scenarios, you can choose whether to display citation information:

Public customer service: Display all citation information to build trust

Internal support: Hide sensitive documents to protect confidential information

Educational institutions: Display learning resources for convenient student reference

7. Conversation Access Control - A Secure and Controllable Conversation Environment

1. Login Configuration

Mandatory Login Mechanism

To ensure the security and controllability of conversation content, enterprises can enable the login feature, requiring users to log in before using the AI assistant service. Once configured, visitors clicking on WebChat will first be directed to a login page, and only successfully authenticated users can begin conversations with the AI assistant

Security Advantages

Identity verification: Ensures every conversation user has a verified identity

Access control: Provides corresponding service content based on user identity

Conversation tracking: Maintains complete conversation history records for each user

Data protection: Prevents unauthorized users from accessing sensitive information

2. Three Login Methods

AD (Active Directory) Login

Use case: For internal employees, integrated with existing AD systems

Features: Single Sign-On (SSO), no additional account management needed

Setup: Configure AD server connection information and synchronization rules

Practical application: Employees log in directly with company accounts and automatically receive corresponding permissions

Keycloak Login

Use case: For enterprises requiring complex identity management, supporting multiple identity providers

Features: Supports standard protocols such as OAuth2, SAML, and OpenID Connect

Setup: Configure Keycloak server, user pools, and identity providers

Practical application: Integrate user identities across multiple systems with unified permission management

MaiAgent Login

Use case: For simple user management needs with quick deployment

Features: Built-in user management system, simple and intuitive operation

Setup: Create user accounts and permissions on the MaiAgent platform

Practical application: Create a dedicated login system for specific user groups

Once enabled, the following screen will appear before starting a conversation, requiring login:



登入

請登入以繼續使用 預設助理模式

電子郵件

請輸入您的電子郵件

密碼

請輸入密碼



[忘記密碼?](#)

登入


[還沒有帳號? 立即註冊](#)

8. Data Collection (Pre-chat Form)

After data collection is enabled, visitors see a form before starting a conversation and must complete it to continue. Their responses are saved to the visitor's contact, so customer service agents can view them directly on the conversation page. This is useful when you need to identify visitors in advance. For example, distributor support can collect a company name and business registration number, while event inquiries can collect a name and email address.

1. Configure the Data Collection Form

Path: **Customer Service** → **Conversation Platforms** → select the Web Chat conversation platform → **Data Collection** tab.



Data Collection tab: configure the toggle, title, and fields on the left; the Live Preview on the right updates the form's appearance in real time

Enable the Form

Turn on **Enable Pre-chat Form**. When it is disabled, visitors enter the conversation directly, following the original flow.

Set the Title and Subtitle (Optional)

Form title: Appears at the top of the form. If left blank, the system uses the default text "A few questions before we start."

Form subtitle: Descriptive text for the form. Use `{name}` to insert the conversation platform name. If left blank, the default text is "Leave your information so {name} can provide more personalized service."


Both fields support multiple languages. Switch language tabs to enter each translation separately.

Add Fields

Click **Add Field**. You can configure the following options for each field:

Setting	Description
Field type	Text, Email, Phone, or Dropdown
Field label	The field name visitors see; enter a separate label for each language
Placeholder	Gray hint text displayed inside the field
Required	When selected, visitors cannot submit the form without completing the field
Contact field mapping	Determines where the data is stored in the contact: Name (Contact.name) , Email (Contact.email) , Phone (Contact.phoneNumber) , or Custom Field (Metadata)

Dropdown fields also have an **Options** list. Select **Allow custom "Other" input** to let visitors enter their own response after selecting "Other."



Configure each field's type, localized label, placeholder, Contact field mapping, and required status

Drag fields to reorder them. A form can contain up to 10 fields. The preview area on the right updates the form's appearance in real time. Switch the **Preview Language** to check each language version.

Save

After confirming that the preview is correct, save the form. The changes take effect immediately in Web Chat.

Contact field mapping restrictions Name can only map to a Text field, Email can only map to an Email field, and Phone can only map to a Phone field. Incompatible field types are disabled with an explanatory message. Name, Email, and Phone can each be mapped by only one form field. For all other fields, such as business registration number or company name, select Custom Field (Metadata).

2. Visitor Form Flow

When visitors open Web Chat, they first see the form with its title, subtitle, configured fields, and a language selector in the upper-right corner. After completing the form, they click [Start Conversation](#) to enter the conversation.

開始對話前的小提問



留下您的資料以便 测试PreForm 提供更個人化的服務

姓名

Email

電話

從哪裏聽說我們？

請選擇



開始對話

Visitors see the form when they open Web Chat

開始對話前的小提問



留下您的資料以便 测试PreForm 提供更個人化的服務

姓名

王小明

Email

xiaoming@example.com

電話

0912345678

從哪裏聽說我們?

搜索引擎



開始對話

Click Start Conversation after completing the form

If a required field is empty or an email address or phone number has an invalid format, an error appears below the field and the form cannot be submitted.

After completing the form once in a browser, the visitor will not be asked to complete it again when opening the same Web Chat.


To correct their information, visitors can open the  menu in the upper-right corner of Web Chat and click [Refill Information](#). The form appears again.



Refill Information in the upper-right  menu

3. Use with Login: Associate Data with a Login Account

If both the [login feature](#) and data collection are enabled for the same Web Chat, the flow is: **Log in → Complete form → Start conversation**. The information entered by a user is associated with their login account rather than their browser.



Login Settings tab: enable the login feature and select a login source; these settings are independent of the Data Collection tab



MAIAGENT

登入

請登入以繼續使用 手冊示範－登入＋資料收集

電子郵件

請輸入您的電子郵件

密碼

請輸入密碼



[忘記密碼](#)

登入

[還沒有帳號？立即註冊](#)

① The login screen appears first when Web Chat opens

開始對話前的小提問



留下您的資料以便 手冊示範－登入＋資料收集 提供更個人化的服務

姓名 *

統一編號 *

Email

開始對話

② The data collection form appears only after login

Situation	Web Chat behavior
The account is being used for the first time and has not completed the form	The form appears after login. Submitting it creates a contact associated with the account.
The same account logs in on another device or browser	The form does not appear again; the conversation starts directly with the account's existing data.
The device previously submitted the form without logging in, but the account has no data	A "Is this information yours?" confirmation screen appears (see below).
The account already has data, and the device has separate data submitted while logged out	The account data takes precedence. The system does not ask for confirmation or merge the records.
A logged-in user clicks Refill Information	The existing contact associated with the account is updated; no additional contact is created.

Confirmation screen: Is this information yours?

If someone previously completed the form on the device without logging in and the account currently logging in does not have any data, the form does not appear immediately after login. Instead, a confirmation screen first displays the information previously entered on the device, partially masked, and provides two buttons:

Yes, this is mine: Associates that information with the current account. The same account will not need to complete the form again when logging in on any device.

No: Does not merge the records. The original information remains unchanged, and the form then appears for the user to complete.

這筆資料是您的嗎？

我們發現這台裝置先前填過的資料。這是您本人留下的嗎？選「是」會把它掛到您的帳號，之後換裝置登入就不必再填。

您在這台裝置留下的資料（部分顯示）：

姓名	陳大****
統一編號	12345****
Email	dahua****

是，這是我的

不是

The confirmation screen masks some of the content, showing only enough for users to determine whether they provided the information

Web Chats without the login feature are unaffected and continue to use the browser to determine whether the form has already been completed. The system does not automatically populate the form from account data. The first time an account is used, the user must complete the form even if their information already exists in the admin console.

After the conversation begins, both the AI assistant and customer service agents in the admin console see the same contact associated with the account:



The AI assistant addresses the user by the name entered in the form

林美玲

手冊示範 - 登入 + 資料收集 關閉詳情

AI Agent 處理中

我來接手

Step 3：服務加購

視需求加購 AI 顧問服務或客製化開發

後續聯繫方式

由於經銷商合作涉及商務條款、報價與合作細節，我建議您直接聯繫我們的業務團隊：

Email: sales@maiagent.ai

LINE: <https://lin.ee/RWjv8vc>

地址: 台北市大安區光復南路 438 號 11 樓

我們的業務團隊會為您詳細說明經銷商方案、定價與合作流程。有其他產品功能或應用場景的問題，也歡迎繼續提問！

參考

加入我們-開啟新視...

Agent-開發平台.md 平台方案.md

AI-KM.md 公司介紹.md

MaiAgent-企業級生... 參考節點

6 分鐘前

ID 對話 ID

e49c6c2c-4976-...

對話滿意度

尚未評分

Credit 消耗彙總

279.67 Credits

向量化

79.23

LLM 輸入

147.37

LLM 輸出

53.06

輪數

平均

標準差

1

279.67

0.00

中位數

最高

最低

279.67

279.67

279.67

> 進階分析

對話時間軸

對話開始

2026/08/26 14:05

客戶從 手冊示範 - 登入 + 資料收集 發起對話

聯絡人資料

電子郵件

meiling@example.com

統一編號

87654321

在此輸入訊息以回覆...

傳送

The conversation page in the admin console displays the contact and the business registration number collected by the form

4. View Collected Data in the Admin Console

View in a conversation: Go to [Customer Service](#) → [All Conversations](#) → open a conversation. The [Contact Information](#) section on the right lists each response collected by the form. Fields mapped to Email and Phone appear as "Email" and "Phone Number," while custom fields use the field labels you configured.

匿名

測試PreForm

關閉詳情

AI Agent 處理中

我來接手

• 查找相關信息和上下文

3. 數據整理與報告

• 生成可下載的報告和文檔

• 協助數據分析和整理

4. 一般諮詢與協助

• 回答問題和提供建議

• 協助解決各類問題

如果你有具體的需求或問題，歡迎告訴我！例如：

• 需要分析某份文件？

• 想查詢之前的對話內容？

• 需要生成報告或文檔？

• 有其他問題需要協助？

請告訴我你具體需要什麼幫助，我會盡力為你服務！😊

測試PreForm | 2026-08-26 13:09:16

在此輸入訊息以回覆...

傳送

對話滿意度

尚未評分

Credit 消耗彙總

109.91

Credits

向量化

0.04

LLM 輸入

74.33

LLM 輸出

35.54

輪數

平均

標準差

1

109.91

0.00

中位數

最高

最低

109.91

109.91

109.91

> 進階分析

對話時間軸

對話開始

2026/08/26 13:07

客戶從 測試PreForm 發起對話

聯絡人資料

電子郵件

xiaoming@example.com

電話號碼

0912345678

姓名

王小明

從哪裏聽說我們?

搜索引擎


The Contact Information section on the right side of the conversation page lists each response collected by the form

The Contact column in the contact list displays the Contact's name. If the form's name field is mapped to Custom Field (Metadata) instead of Name (Contact.name), logged-out visitors appear as Anonymous, and their names are available only in custom attributes. To display visitor names directly in the list, map the name field to Name (Contact.name).

View on the Contacts page: Go to **Customer Service** → **Contacts** → open a contact. The **Basic Information** tab displays the name and phone number, while the **Custom Attributes** tab displays the other form fields.



Search the contact list by name, email address, or phone number



The Custom Attributes tab lists the value of every form field


Information entered by a logged-in user is associated with their account rather than an unclaimed anonymous contact. The contact name depends on the form configuration: if the name field is mapped to Name (Contact.name), the name entered in the form is displayed; if the form has no name field, the login account name is displayed.

How Do I Handle a Connection Failure?

For example, a website administrator may receive a report from a visitor that Web Chat displays a connection error when opened. The visitor should first wait for the system to retry automatically. If the error screen remains, they should capture the entire screen—including the failed stage, error category, and time of occurrence—and send it to the administrator so the issue can be diagnosed more quickly.

If the network connection is interrupted while Web Chat is loading the AI assistant configuration, it retries automatically and remains on the loading screen during this period. If it still cannot connect after retrying, the error screen displays:

- The loading stage that failed
- The error category, together with the HTTP status code when available
- The local time when the error occurred



After automatic retries fail, the screen lists the failed stage, error category, and time of occurrence

If the system has received an HTTP error status code, the server has already responded. In this case, Web Chat does not retry automatically and displays the diagnostic error screen immediately.



First confirm that the device can access the internet, then click **Retry** on the error screen. If the issue persists, preserve all diagnostic information in the screenshot and ensure that no conversation content or personal data is visible.

Integrate with Platform: Website

I. Access the Settings Page


1. Path 1: Add a New Website Chat

Click "**Customer Service > Chat Platforms**" in the left menu, click "**+ Connect Chat Platform**" in the upper right corner, and select "**Website**"



2. Path 2: AI Assistant Associated Chat Platform

In the left sidebar, go to **AI Features** → **AI Assistants**, select the assistant you want to configure, and switch to the **Chat Platforms** tab. Find the website chat platform you want to adjust, then click **Platform Settings**.



Click "Platform Settings" on the AI Assistant's Chat Platforms tab.

After completing the settings, click **Back to AI Assistant Settings** at the top of the page to return to the original assistant's **Chat Platforms** tab. This return option remains available when you switch between **Basic**, **Quota**, and **Auto-Reply Settings**.



After completing the settings, return directly to the original AI Assistant settings page.


Back to AI Assistant Settings is displayed only when you access the page from the AI Assistant's **Chat Platforms** tab. If you access it directly from **Chat Platforms** in the left sidebar, the page retains its original navigation.

Use Case: Return to the Assistant After Updating Website Customer Service

A customer service manager needs to update the website customer service settings for the "Product Consultation Assistant." She clicks **Platform Settings** on this AI Assistant's **Chat Platforms** tab. After making the changes, she clicks **Back to AI Assistant Settings** to return to the same assistant's list of chat platforms and continue checking other customer service channels, without having to find the assistant again in the AI Assistant list.

3. Path 3: Chat Platform Operations

Go to "**Customer Service > Chat Platforms**" in the left menu, select the AI Assistant you want to configure, and click the "**Operations**" icon to enter the chat platform configuration.



II. Basic Settings

1. Basic Settings



Set Platform Name

Enter the AI Assistant name you want to display on the website in the "Platform Name" field

For example: "Customer Service Helper", "Smart Assistant", etc.



Platform name display

Select AI Assistant

Choose the AI Assistant to use from the "AI Assistant" dropdown menu

This determines which AI Assistant answers questions in the chat room

Set Return URL

Enter your company website URL in the "Return URL" field

When users click the back button in the chat room, they will be redirected to this URL

Obtain and Use the Public URL

The system automatically generates a "Public Access URL"

You can:

Click the "Copy" button to copy the URL

Click the "Visit" button to open the chat room in a new window

Click the "Embed" button to get the embed code

Click the "Embed Preview" button to preview the embed effect

Embed Script Optimization

When you choose the embed method, the system provides an optimized embed.js script with the following features:

Duplicate execution prevention: Even if the script is loaded multiple times, it only initializes once to avoid conflicts

Better error handling: Clear error messages are provided when loading fails, making debugging easier

Optimized loading performance: The script is optimized to not affect website loading speed


Improved compatibility: More compatible with various website environments and frameworks

Embed Best Practices:


Place the embed script before the `<body>` tag to ensure main page content loads first

Only include the embed script once on the page to avoid duplication

For SPA (Single Page Application) usage, refer to the technical documentation for advanced settings



Embed method selection



Embed preview



After adding a return URL, it will be displayed in the background when navigating back

Select Embed Method

After clicking "**Embed**", you can switch between three embed methods at the top of the window:

JavaScript: Loaded via